

CSDS 451: Designing High Performant Systems for AI

Lecture 1

8/26/2025

Sanmukh Kuppannagari

sanmukh.kuppannagari@case.edu

<https://sanmukh.research.st/>

Case Western Reserve University

Outline

- Course Information
- Course Overview
- Introduction to Heterogeneous Computing
- Introduction to Key AI/ML Kernels

Outline

- Course Information
- Course Overview
- Introduction to Heterogeneous Computing
- Introduction to Key AI/ML Kernels

Introduction



Assistant Professor, Case – Fall 2022 - present
Senior Research Associate/Post Doc, USC – Fall 2018 to Summer 2022
PhD, Fall 2018, USC
B.Tech, CSE, 2011, IIT Guwahati

Research Focus: Accelerating AI Algorithms

Team

3 PhDs, 3 Masters, 2 undergraduates

<https://sanmukh.research.st/people>

Research

- Deploy real-time, low-power, low-latency CNN models on Processing-in-Memory (PIM) architectures
- Scaling Transformer Models to Extremely Long Sequences by Exploiting Sparsity
- Scalable Training of Geometric Graph Neural Networks
- Scalable Training of Fourier Domain Models

Basic Information

- Lecture: Tuesday, Thursday 4:00 – 5:15 pm, Olin 305
- Office Hours
 - Location: Olin 506
 - Times: TBD

Course Goals

- **Objective:** Provide a broad overview of the challenges and opportunities that exist in designing high performance AI systems
 - Key Focus: Algorithmic Innovations
- Course Audience
 - Data Scientists who want to look under the hood and speed up their models
 - Parallel Algorithms or Hardware Acceleration researchers

Course Pre-requisites

- Algorithms
 - Complexity Analysis
 - Problem solving
 - ...
- Mathematical/Computer Science maturity
- Familiarity with Programming (preferably C/C++)
 - Ability to write programs given an algorithm
 - Ability to debug your programs or find answers online if you are stuck
 - In this world of Generative AI, programming language should not be a barrier
- Most importantly, ability to quickly learn new concepts
 - Most critical skills for a computer scientist given how fast our field evolves

Course Objectives

- Theory
 - Parallel Computing Basics
 - Algorithms to Accelerate Convolutional Neural Networks (CNN), Large Language Models (LLMs), and Distributed Training on a Cluster
- Practice
 - BLAS libraries for High Performant AI/ML kernel implementations
 - Heterogeneous Programming using AMD HIP
 - Pytorch Lightning for distributed training on a cluster

Grading Policy (1)

- Grading
 - Written Assignments (30%)
 - Programming Assignments (20%)
 - Mid-term exam (15%)
 - Final exam (15%)
 - Project (20%)
- Watch out for extra credit assignments or bonus points in your assignments

Grading Policy (2)

- Written Assignments (30%)
 - 3 assignments to test the theoretical concepts covered in the class
 - Fairly involved
- Programming Assignments (20%)
 - 2 assignments on programming frameworks covered in the class
- Due date typically within 2 week from the assigned date

Grading Policy (3)

- Mid-term (15%)
 - 9th October
 - In class
- Final (15%)
 - 25th November
 - In class
- Similar to assignments but in a timed manner
- Will be open books/slides
- No exam during the final exam day

Grading Policy (4)

- Project (20%)
 - Implement one or more AI model
 - Using one or more acceleration technique
 - Using one or more heterogeneous programming frameworks
- Groups of 2
- Source code and project report due during the finals week
- I will provide more details during the semester on expectations, topics, team formation and submission guidelines

Grading Policy – Late Assignments

- 10% reduction in achieved score per day with a total grace of 3 days.
0 points after 3rd day
- Only exceptions for medical reasons

Pre-Requisite Knowledge Assignment

- Programming Assignment 0 has been uploaded
- Objectives:
 - Warm up your C/C++ skills
 - Illustrative of the complexity of the programming assignments in the course
- Counts as 2% extra credits

Reading Materials

- No textbook as the topics discussed are based on cutting edge research
- Course slides are intended to explain topics in sufficient detail
- List of papers will be provided for students to review

Course Guidelines (1)

- Please start assignments early
- Please take advantage of the office hours
- Asking questions in class is encouraged
- Please check the course canvas regularly

Course Guidelines (2)

- Policy on using AI tools
- Use of AI tools encouraged for background research
 - Caveat: beware of hallucinations or misinformation
 - Could be a good tool to understand a topic covered in class
 - Could be a good tool to debug your programs if you are stuck
- Please refrain from directly asking the solutions for written or programming assignments
- Please keep me updated on your use of these tools. We as faculty would like to learn how students can best benefit from these tools while maintaining academic integrity

Course Slack Channel

- To Be Created
- Will put an announcement on canvas


Outline

- Course Information
- **Course Overview**
- Introduction to Heterogeneous Computing
- Introduction to Key AI/ML Kernels

Why This Course? (1)

AI Increasingly Integral to Real-World

Autonomous driving



Surveillance



AI

Social Media




Power Grids




Why This Course? (2)


- Models are increasing in size and complexity



Why This Course? (3)



Why This Course? (4)




- Gap between training compute requirements and available computation performance keeps on increasing
- Source: <https://epoch.ai/trends>

Why This Course? (5)

- <https://claude.ai/public/artifacts/c4d13724-3130-4d23-89b6-63b08514b299>
- (Not sure how much CO2 got emitted when I created this interactive graph in Claude)

AI Model Training Carbon Emissions



Key Insight: Training emissions have grown exponentially with model size. GPT-3 training produced equivalent emissions to 626,000 pounds of CO₂- more than 100 times the annual footprint of an average person.


Why This Course? (6)

- Rise of Domain Specific Architectures
 - End of Moore's Law: Power becomes key constraint
 - Limitations and inefficiencies in exploiting instruction level parallelism: reaching Amdahl's law's limits
 - Shift in Application focus: individual devices, IoT, data-oriented, AI/ML...



CPU

+



NPU



GPU



TPU




Figure from Presentation: The End of Moore's Law, CPUs (as we know them), and the Rise of Domain Specific Architectures, John Hennessy, September 2019

Why This Course? (7)

Performance


Latency, throughput, ...



...

Privacy


Security Level, Scheme, ...



...

Data

image, graph, ...



Systems Person




Power

Datacenter, edge, ...



Programming

General Purpose, vector, spatial, ...



Technology

Optical, ASICs, Reconfigurable

Why This Course? (8)



Data Scientist

PyTorch



TensorFlow

Why This Course? (9)

- Building accurate AI/ML models is not sufficient
- You should be able to deploy them efficiently in real world for them to be impactful
 - Time is \$\$\$
- Deep learning frameworks are optimized for widely used use case scenarios
 - What if your problem does not fit into this?

Why This Course? (10)

- Hypothetical Scenarios (inspired by real world anecdotes)
- You are a data scientist at a security company that performs intrusion detection using images. The client wants the model to be hosted on their premises on a low cost server.
- You are a research scientist at a biomedical imaging company. A new dataset you want to explore has 10x the number of pixels and cannot be simply handled using your current GPU server.

Why This Course? (11)

- Hypothetical Scenarios (inspired by real world anecdotes)
- You are a data scientist at a security company that performs intrusion detection using their premises. Will you ask your client or company to upgrade server?
- You are a research scientist at a biomedical imaging company. A new dataset you want to explore has 10x the number of pixels and cannot be simply handled using your current GPU server.

Why This Course? (12)

- Hypothetical Scenarios (inspired by real world anecdotes)
- You are a data scientist at a security company that performs intrusion detection using GPU parallel processing. Your company has been asked to host a new system on their premises. Will you ask your client or company to find a systems/parallel computing engineer?
- You are a research scientist at a biomedical imaging company. A new dataset you want to explore has 10x the number of pixels and cannot be simply handled using your current GPU server.

Topics

- Parallel Computing Basics
 - Learn how to model three popular computing platforms – processor-memory architecture (GPUs), Systolic arrays (TPU), cluster
- Heterogeneous Computing Framework
 - Learn how to program CPU+accelerator platforms
- AI/ML Acceleration – Apply the above two concepts to accelerate AI models
 - Convolutional Neural Network Acceleration
 - Transformer Model Acceleration
 - Distributed Training of Neural Networks


Outline

- Course Information
- Course Overview
- **Introduction to Heterogeneous Computing**
- Introduction to Key AI/ML Kernels

Acknowledgements

- Slides on Heterogeneous Computing are borrowed from the TOUCH program
 - Faculty Training Workshop for Teaching Heterogeneous Computing
- Heterogeneous Computing Resources:
https://touch.cs.txstate.edu/training_ws22_resources.php

Von Neumann Architecture What's a Computer?




Evolution of Parallelism in Processors



Single Core
Computing

- Pipelining
- Out of order execution
- ...



Hidden from Programmer



Multi-Cores/ Multi-
Processor Computing

- Multi-threading
- Parallel For
- Parallel Tasks
- Tasks to core mapping
- ...




Heterogeneous
Computing

- Parallel For
- Parallel Tasks
- Capability based Task to Device Mapping
- ...



+

Single Core Computing



Intel 4004



Intel Pentium

Single Core Computing



code

```
int main() {
    int x, y, result;
    x = 17;
    y = 13;
    result = x + y;
    return result;
}
```


compile

```
.text
.globl _main
_main:
LFB2:
        pushq   %rbp
LCFI0:
        movq    %rsp, %rbp
LCFI1:
        movl    $17, -4(%rbp)
        movl    $13, -8(%rbp)
        movl    -8(%rbp), %eax
        addl    -4(%rbp), %eax
        movl    %eax, -12(%rbp)
        movl    -12(%rbp), %eax
        leave
        ret
```



execute

...11001010100...





Program Execution



*Instruction execution follows
program order*


*Processor executes one
instruction at a time**



Multi-Core/Multi-Processor Computing



Intel i7



AMD Ryzen




Parallel Program Execution



code

```
int main() {  
    int x, y, result;  
    x = 17;  
    y = 13;  
    result = x + y;  
    return result;  
}
```

compile




The two assignment statements

$x = 17;$ and $y = 13;$ will execute in parallel




Parallel Program Execution




```
int main() {
    int x, y, result;
    x = 17;
    y = 13;
    result = x + y;
    return result;
}
```

*Cannot arbitrarily assign
instructions to processors*



Dependencies in Parallel Code

- If statement s_i needs the value produced by statement s_j then, s_i is said to be dependent on s_j



A diagram illustrating a dependency between two statements. A red arrow points from the label s_j to the label s_i , indicating that s_i depends on the value produced by s_j . The code block contains:
 $s_j \quad y = 17;$
...
...
 $s_i \quad x = y + \text{foo}();$


- All dependencies in the program must be preserved
- This means that if s_i is dependent on s_j then we need to ensure that s_j completes execution before s_i

*responsibility lies with
programmer (and software)*

Heterogeneous Computing

Heterogeneity


- CPU + XPU
 - XPU = GPU, FPGA, NPU, custom accelerator
- CPU
 - Runs OS; performs general bookkeeping
 - Offloads Task to XPU
- XPU
 - Processors specialized to perform a particular type of task
 - Trades off generalizability with performance



For more info on XPU:
<https://www.intel.com/content/www/us/en/architecture-and-technology/xpu.html>




CPU + GPU Heterogeneity




CPU + GPU Heterogeneity

- CPU
 - Runs the main program
- GPU
 - Runs the “kernel” program
 - Specializes in Data Parallel Programming
- Excels at
 - Dense, parallel, structured, computations
 - Image based; Dense Matrix based
- Limited for
 - Sparse irregular workloads
 - Graph based; ...



CPU+FPGA Heterogeneity

- CPU
 - Runs the main program
- FPGA
 - Specializes in “User Programmed” Task
- Excels at
 - Dense parallel structured computations
 - Sparse Irregular computations
- Caveats
 - FPGA computation speeds 10x slower than GPUs



Crash Course in FPGAs (1)

- Field Programmable Gate Array
 - Programmable Logic Blocks +
 - On-Chip Memory +
 - Programmable Interconnects
- Logic + Interconnect can be programmed to represent ANY program/circuit/algorithm

Crash Course in FPGAs (2)

- Why FPGAs?
 - “Programmable” Hardware
 - Processors/Microcontroller
 - Inherently constrained by Instruction Set
 - Limited parallelism
 - FPGAs
 - No software, no instruction set constraints, a digital circuit that implements the exact functionality
 - High fine grained parallelism


We will not discuss in the class but if you are interested to do projects on FPGA as opposed to GPUs, contact me directly




Heterogeneous Computing

Heterogeneity

```
int dostuff(int z) {  
    int x, y, result;  
    x = 17;  
    y = Complicated_F(z);  
    result = x + y;  
    return result;  
}
```



*Offload Complicated Function to
XPU*



Example Frameworks (1)

- Pytorch/Tensorflow: CPU+GPU
 - python based ML framework
- Deep Graph Library: CPU+GPU
 - Python based graph learning framework
- OpenMP: CPU + XPU (currently XPU = GPU)
 - Pragma based programming for heterogeneous computing

Example Frameworks (2)

- OneAPI: CPU+XPU (XPU = GPU or FPGA)
 - C++ based device agnostic programming
- CUDA/HIP: CPU+GPU
 - C/C++ based programming for GPU
- High Level Synthesis (HLS): CPU + FPGA
 - Pragma based programming for FPGA
 - Typically, used in conjunction with OpenCL


Pytorch

- Open Source ML Framework
- Explore
 - Basics of ML Model development
 - Distributed ML (later in the course)
- Pytorch: <https://pytorch.org/>



Heterogeneous-compute Interface for Portability (HIP)

- C++ based runtime API and kernel language
- Enables developers to design platform-independent GPU programs that can run on both AMD and NVIDIA GPUs
- Syntax very similar to CUDA



Our Focus in this Course

- Target Platform – CPU+GPU Heterogeneous Platforms
- Programming Frameworks
 - C++ - CPU-only
 - HIP – CPU+GPU
 - Pytorch – CPU+GPU

Outline

- Course Information
- Course Overview
- Introduction to Heterogeneous Computing
- **Introduction to Key AI/ML Kernels**

Key Computational Kernels in Machine Learning

- Majority of Computations in Machine Learning occur using a small set of kernels (operations that are a bit more complex than addition/multiplication)
- Vector – Vector Operations
- Matrix – Vector, Matrix – Matrix Operations
- Sparse Matrix Operations
- Hashing
- Message Passing on Graphs

Ungraded HW Assignment

- Use your favorite generative AI tool to learn about each of the kernels
- We will introduce them in this class as and when needed

Next Class

- 8/28 Lecture 2 – Processor-Memory Architecture – Modeling, Analysis, Challenges, Opportunities for Optimization

Thank You

- Questions?
- Email: sanmukh.kuppannagari@case.edu