

CSDS 451: Designing High Performant Systems for AI

Lecture 2

8/28/2025

Sanmukh Kuppannagari

sanmukh.kuppannagari@case.edu

<https://sanmukh.research.st/>

Case Western Reserve University

Outline

- Why Performance Modeling?
- Processor Memory Architectures
 - Basic Execution Model
 - Memory System Modeling

Outline

- Why Performance Modeling?
- Processor Memory Architectures
 - Basic Execution Model
 - Memory System Modeling

Performance Modeling

- Identifying important parameters that determine the performance of an **algorithm** mapped to the **platform**
 - Algorithm – Parallel algorithm that we have designed
 - Platform – The target hardware on which we wish to run the parallel algorithm
- Identifying equations that govern the performance

Why Performance Modeling?

- Serial Code
- For $i = 1$ to N
 - $j = 2^i$
 - If $j \geq M$ break
 - Else: Do work()
- What are the important parameters?
- What are the equations governing the performance of the algorithm?

Why Performance Modeling?

- Serial Code
- For $i = 1$ to N
 - $j = 2^i$
 - If $j \geq M$ break
 - Else: Do work()
- What are the important parameters?

 M, N
- What are the equations governing the performance of the algorithm?

Min{log M,N}

Why Performance Modeling?


- Complexity Analysis of Serial Algorithms
 - Bubble sort (n^3) vs Merge sort ($n \log n$)
 - Parameter that determine performance: input size n
 - Equations governing performance – complexity of the algorithm
- Provide understanding of bottlenecks and what to optimize
- Do not provide actual implementation estimates
 - Quick sort vs Merge sort ???

Parallel Algorithm Performance Modeling

- Performance is platform dependent
- Different platforms -> different parameters
- Different platforms -> different equations governing performance

Targeted Platforms

- Processor-Memory Architecture
 - CPUs, GPUs
- Systolic Array Architecture
 - TPUs, Tensor Cores
- Clusters
 - HPC



Why Performance Modeling?

- Most importantly, to get a mental model of how the algorithm you are designing runs on the hardware, without going into specifics of the hardware.
- Parallel programs can be notoriously hard to reason about
- Simple performance models, even if not accurate, can help us get clarity on execution and aid in debugging


Outline

- Why Performance Modeling?
- Processor Memory Architectures
 - Basic Execution Model
 - Memory System Modeling

Basic Execution Model


Fetch-Execute

A simple view



Parameters:

Memory access = 1 cycle for any location
Execution = 1 cycle



Basic Execution Model

Calculating Algorithm Performance

- Consider the following computation
- $C[i] = \sum_k A[i][k] * B[k]$
- $i = 0$ to M
- $k = 0$ to N
- What is the above operation?

Basic Execution Model


Calculating Algorithm Performance

- Consider the following computation
- $C[i] = \sum_k A[i][k] * B[k]$
- $i = 0$ to M
- $k = 0$ to N
- What is the above operation? **Matrix-Vector Multiplication**



Matrix Vector Multiplication

$$C[i] = \sum_k A[i][k] * B[k]$$



Basic Execution Model


Calculating Algorithm Performance

- $C[i] = \sum_k A[i][k] * B[k]$

Decompose the equation into two types of instructions

Fetch: Read from memory or Write to memory

Execute: Arithmetic/Logical Operation



Basic Execution Model

Calculating Algorithm Performance

- $C[i] = \sum_k A[i][k] * B[k]$

RA: Read A[i][k]


RB: Read B[k]

M: Res <- Mult A[i][k]*B[k]

RC: Read C[i]

A: Add C[i] + Res

S: Store C[i]



Basic Execution Model

Calculating Algorithm Performance

- $C[i] = \sum_k A[i][k] * B[k]$

RA: Read $A[i][k]$

RB: Read $B[k]$

M: Res \leftarrow Mult $A[i][k] * B[k]$


RC: Read $C[i]$

A: Add $C[i] + \text{Res}$

S: Store $C[i]$

Red arrows point from the first four lines to the fifth line.

Assume we are storing these
in a local memory (such as
registers) in the processor



Basic Execution Model

- Assumptions
- No limit on the local memory
 - Putting a limit may require us to use a different set of instructions
- Addition and Multiplications are each single cycle operations
 - Not necessarily true in modern architectures

Basic Execution Model

Calculating Algorithm Performance

- $C[i] = \sum_k A[i][k] * B[k]$

RA: Read A[i][k]


RB: Read B[k]

M: Res <- Mult A[i][k]*B[k]

RC: Read C[i]

A: Add C[i] + Res

S: Store C[i]



6 cycles per entry of A

Total Time for the entire matrix = ??

Basic Execution Model

Calculating Algorithm Performance

- $C[i] = \sum_k A[i][k] * B[k]$

RA: Read A[i][k]


RB: Read B[k]

M: Res <- Mult A[i][k]*B[k]

RC: Read C[i]

A: Add C[i] + Res

S: Store C[i]




6 cycles per entry of A

Total Time for the entire matrix = **6*M*N** cycles


Basic Execution Model

- Assume Processor is running at 2 GHz
- Total time (in seconds) to execute Matrix-Vector Multiplication as per Basic Execution Model
- ???



Basic Execution Model

- Assume Processor is running at 2 GHz
- Total time to execute Matrix-Vector Multiplication as per Basic Execution Model
- $3MN$ seconds



Basic Execution Model: Limitations

- Does not account for memory transfer latencies
- Does not account for differences in the number of cycles in various arithmetic operations
 - Multiplication usually requires more cycles than addition

Basic Execution Model: Uses


- It will be a part of more complicated models that we will study next
- With modern pipelined architectures, 1 cycle per execute operation is not a bad assumption

Outline

- Why Performance Modeling?
- Processor Memory Architectures
 - Basic Execution Model
 - Memory System Modeling

Memory System

- System performance (execution time) depends on the “rate” at which data can be accessed by the program



Updated Execution Model

Latency: Time to receive the first word since the fetch instruction


Bandwidth: Rate of data transfer = bitwidth of the interconnect x frequency

Execute – 1 cycle

Fetch – Latency + Data size/bandwidth cycles

Memory System

- System performance (execution time) depends on the “rate” at which data can be accessed by the program



Updated Execution Model

Latency: Time to receive the first word since the fetch instruction


Bandwidth: Rate of data transfer =
bitwidth of the interconnect x frequency

Execute – 1 cycle

Fetch – Latency + Data size/bandwidth cycles

Memory System


- **Latency:** Every fetch request requires some setup operation in memory before it can actually start transferring the data
- **Bandwidth:** Once the data starts transferring, it can continuously transfer without any interruptions depending upon the speed of the interconnect, as long as it is contiguous



Fetch request: Given me 10 data elements starting at location 0


Memory System

- Read A[0-31]
- Latency: 10 cycles
- Bandwidth: 8 data elements per cycle
(assuming 8 bits/data element)
- Number of cycles to finish the read operation?




Memory System

- Read A[0-31]
- Latency: 10 cycles
- Bandwidth: 8 data elements per cycle
(assuming 8 bits/data element)
- Number of cycles to finish the read operation? **10 + 4 = 14 cycles**




Memory System

- Read A[0-31], A[32-63], A[64-95]
- Latency: 10 cycles
- Bandwidth: 8 data elements per cycle
(assuming 8 bits/data element)
- Number of cycles to finish the read operation?




Memory System

- Read A[0-31], A[32-63], A[64-95]
- Latency: 10 cycles
- Bandwidth: 8 data elements per cycle
(assuming 8 bits/data element)
- Number of cycles to finish the read operation? $10 + 4 + 10 + 4 + 10 + 4 = 42$ cycles




Memory System

- Read A[0-95]
- Latency: 10 cycles
- Bandwidth: 8 data elements per cycle
(assuming 8 bits/data element)
- Number of cycles to finish the read operation?



Memory System

- Read A[0-95]
- Latency: 10 cycles
- Bandwidth: 8 data elements per cycle
(assuming 8 bits/data element)
- Number of cycles to finish the read operation? $10 + 12 = 22$ cycles




Memory System

- By fetching the entire 96 words in a single go, we saved 20 cycles
- Prefetching:
 - If you know what data your algorithm needs to access beforehand (not as easy as it sounds)
 - Prefetch it to save on latency
- (Note: above 20 cycles is based on a simple model, in reality, it is difficult to exactly calculate the savings, but we do see substantial savings)

Memory System

- System performance (execution time) depends on the “rate” at which data can be accessed by the program



Updated Execution Model

Latency: Time to receive the first word since the fetch instruction


Bandwidth: Rate of data transfer =
bitwidth of the interconnect x frequency

Execute – 1 cycle

Fetch – Latency + Data size/bandwidth cycles

Memory System

- **Latency:** Every fetch request requires some setup operation in memory before it can actually start transferring the data
- **Bandwidth:** Once the data starts transferring, it can continuously transfer without any interruptions depending upon the speed of the interconnect, as long as it is contiguous



Fetch request: Given me 10 data elements starting at location 0

Memory System: Modeling

- Performance will be measured in terms of “System Performance Metric”
- Number of “useful computations” performed per second
- Useful computation – arithmetic computations such as Add, Mult, etc.
 - Does not include memory fetching time in calculations

System Performance Metric

- FLOPs – Floating Point Operations Per Second
- Number of useful “floating point” operations in a second
- GFLOPs – Giga (1024) Floating Point Operations Per Second
- Note: Unless explicitly specified, we will use 64 bit floating point as data type for performance modeling in this class.
- We will call each data element as a **word**

System Performance Metric

- 1024 read operations
- For every two read operations, we perform one add operation.
- All the operations mentioned above take a total of 10 seconds.
- What is the system performance?

System Performance Metric

- 1024 read operations
- For every two read operations, we perform one add operation.
- All the operations mentioned above take a total of 10 seconds.
- What is the system performance? **512 Floating point operations in 10 seconds**


System Performance Metric

- 1024 read operations
- For every two read operations, we perform one add operation.
- All the operations mentioned above take a total of 10 seconds.
- What is the system performance? **51.2 FLOPs**

System Performance Metric

- 1024 read operations
- For every two read operations, we perform one add operation.
- All the operations mentioned above take a total of 10 seconds.
- What is the system performance? **0.05 GFLOPs**

Memory System Modeling



Processor

2 Ghz (0.5 ns per cycle)

Memory

Memory Latency = 10 ns (20 cycles)

Bandwidth = 64 bits at 1 GHz

(64 Gbits/sec or 8 GB/sec) or

(64 bits in every 2 processor cycles)

Memory System Modeling

- We will consider three metrics:
 - Peak Performance
 - Sustained Performance (best case)
 - Sustained Performance (worst case)
- Actual obtained performance on the hardware should be somewhere in between Sustained Performance (worst case) and Peak Performance
 - If not, there was some error in modeling or the implementation of the algorithm
 - It can be higher or lower than the Sustained Performance (best case), this information governs what optimizations to use

Memory System Modeling: Peak Performance

- Captures the Raw Computer power of the processor (memory not considered)
- Peak Performance = Clock Rate X Num ops per cycle
 - = 2 (Ghz) x 1 (ops/cycle)
 - = 2 Giga Ops per Second
- If Floating point ops = 2 GFLOPs (Giga Floating Point Operations per second)

Memory System Modeling: Peak Performance

- Captures the Raw Computer power of the processor (memory not considered)
- Peak Performance = Clock Rate X Num ops per cycle
 - = 2 (Ghz) x 1 (ops/cycle)
 - For multiple processors/GPUs, this number will change
 - = 2 Giga Ops per Second
- If Floating point ops = 2 GFLOPs (Giga Floating Point Operations per second)

Memory System Modeling: Sustained Performance (Worst Case)


- Each fetch takes: Latency + (data size = 1)/bandwidth \approx latency number of cycles
- $C[i] = \sum_k A[i][k] * B[k]$



Memory System Modeling: Sustained Performance (Worst Case)

- **Each fetch takes: Latency + (data size = 1)/bandwidth \approx latency number of cycles**
- Example: When Prefetching is not used
 - In reality, the data size will be cache line length, but we are simplifying the model here
 - Also, the sophisticated caching mechanisms in modern architectures ensure the performance is not this bad

Memory System Modeling: Sustained Performance (Worst Case)



Processor

2 Ghz (0.5 ns per cycle)

Memory

Memory Latency = 10 ns (20 cycles)

Bandwidth = 64 bits at 1 GHz

(64 Gbits/sec or 8 GB/sec) or

(64 bits in every 2 processor cycles)

Can you calculate the Sustained Performance (worst case) in GFLOPs for the entire matrix?

Memory System Modeling: Sustained Performance (Worst Case)


- Fetch = ??
- Execute = ??
- Total time (per each loop over i) = ??
- Number of Operations (per each loop over i) = ??
- Performance: ?? GFLOPs

Memory System Modeling: Sustained Performance (Worst Case)

- Fetch = 10 ns
- Execute = 0.5 ns
- Total time (per each loop over i) = $4 \times 10 + 2 \times 0.5 = 41$ ns
- Number of Operations (per each loop over i) = 2
- Performance: $2/41$ ns = 0.048 GFLOPs

Memory System Modeling: Sustained Performance (Best Case)

- In Sustained Performance (worst case), we were fetching one word at a time.
 - Separate fetch instructions
 - Each instruction incurring latency
- Sustained Performance (best case): Assume data is contiguously stored
 - Single fetch can be performed
 - Assume latency is 0




Memory System Modeling: Sustained Performance (Best Case)

- How can the latency be 0?
- Data prefetching: Processor initiates data transfer request way before the data is actually needed
- Predicting data prefetching parameters using ML is a hot research topic:
https://scholar.google.com/scholar?as_ylo=2018&q=data+prefetching+machine+learning&hl=en&as_sdt=0,36
- We will not cover in this class


Memory System Modeling: Sustained Performance (Best Case)

- Each fetch takes: Latency (=0) + data size/bandwidth \approx data size/bandwidth
- $C[i] = \sum_k A[i][k] * B[k]$




Memory System Modeling: Sustained Performance (Best Case)

- Each fetch takes: Latency (=0) + data size/bandwidth \approx data size/bandwidth
- $C[i] = \sum_k A[i][k] * B[k]$



Memory System Modeling: Sustained Performance (Best Case)



Processor

2 Ghz (0.5 ns per cycle)

Memory

Memory Latency = 10 ns (20 cycles)

Bandwidth = 64 bits at 1 GHz


(64 Gbits/sec or 8 GB/sec) or

(64 bits in every 2 processor cycles)

Can you calculate the Sustained Performance (best case) in GFLOPs for the entire matrix?

Memory System Modeling: Sustained Performance (Worst Case)

- Each fetch takes: Latency (=0) + data size/bandwidth \approx data size/bandwidth
- $C[i] = \sum_k A[i][k] * B[k]$



Memory System Modeling: Sustained Performance (Best Case)

- Total time (per loop over i) = 5 ns
- Number of Operations (per loop over i) = 2
- Performance: $2/5 \text{ ns} = 0.4 \text{ GFLOPs}$
- Notes
 - Again we assumed no limit on local memory of processor
 - Also, note how we reordered the read instructions to maximize bandwidth use

System Performance Metrics



Actual System Performance

Latency Hiding Mechanisms

Efficient Bandwidth Utilization

0.048 Gflops/s

0.4 Gflops/s

2 Gflops/s

Sustained Worst Case

- Latency dominant

Sustained Best Case

- Latency Hidden
- Memory bandwidth dominant

Peak Performance

- Compute Only

Memory System: Conclusion

- A simple model to perform “pen and paper” calculations to analyze an algorithm
- Compare various algorithms in a similar spirit as complexity analysis
- Get a mental model of the algorithm execution to potentially find better optimizations

Memory System: Conclusion

- Ungraded HW assignment 1: Can you create a simple code to perform these calculations? (May help you in your assignments later)

Next Class

- 9/2 Lecture 3 – Processor-Memory Architecture II
 - Memory Optimizations,
 - Processor Memory Architectures with Cache
 - Blocked Matrix Multiplication

Thank You

- Questions?
- Email: sanmukh.kuppannagari@case.edu