

# CSDS 451: Designing High Performant Systems for AI

Lecture 7

9/16/2025

Sanmukh Kuppannagari

[sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)

<https://sanmukh.research.st/>

Case Western Reserve University

# Outline

- Parallel Program Analysis (Quick Review)
- Data Parallel Algorithms
  - Sorting
- Task Parallelism

# Outline

- Parallel Program Analysis (Quick Review)
- Data Parallel Algorithms
  - Sorting
- Task Parallelism

# Speedup

- Improvement in time due to parallelization

$$\text{Speedup} = \frac{\text{Serial time (on a uniprocessor system)}}{\text{Time after parallelization}}$$

- Can be empirical – gives us actual speedup of implementation
- Can be in order notation – helps us in analyze whether a particular parallel algorithm is good - scalable or not

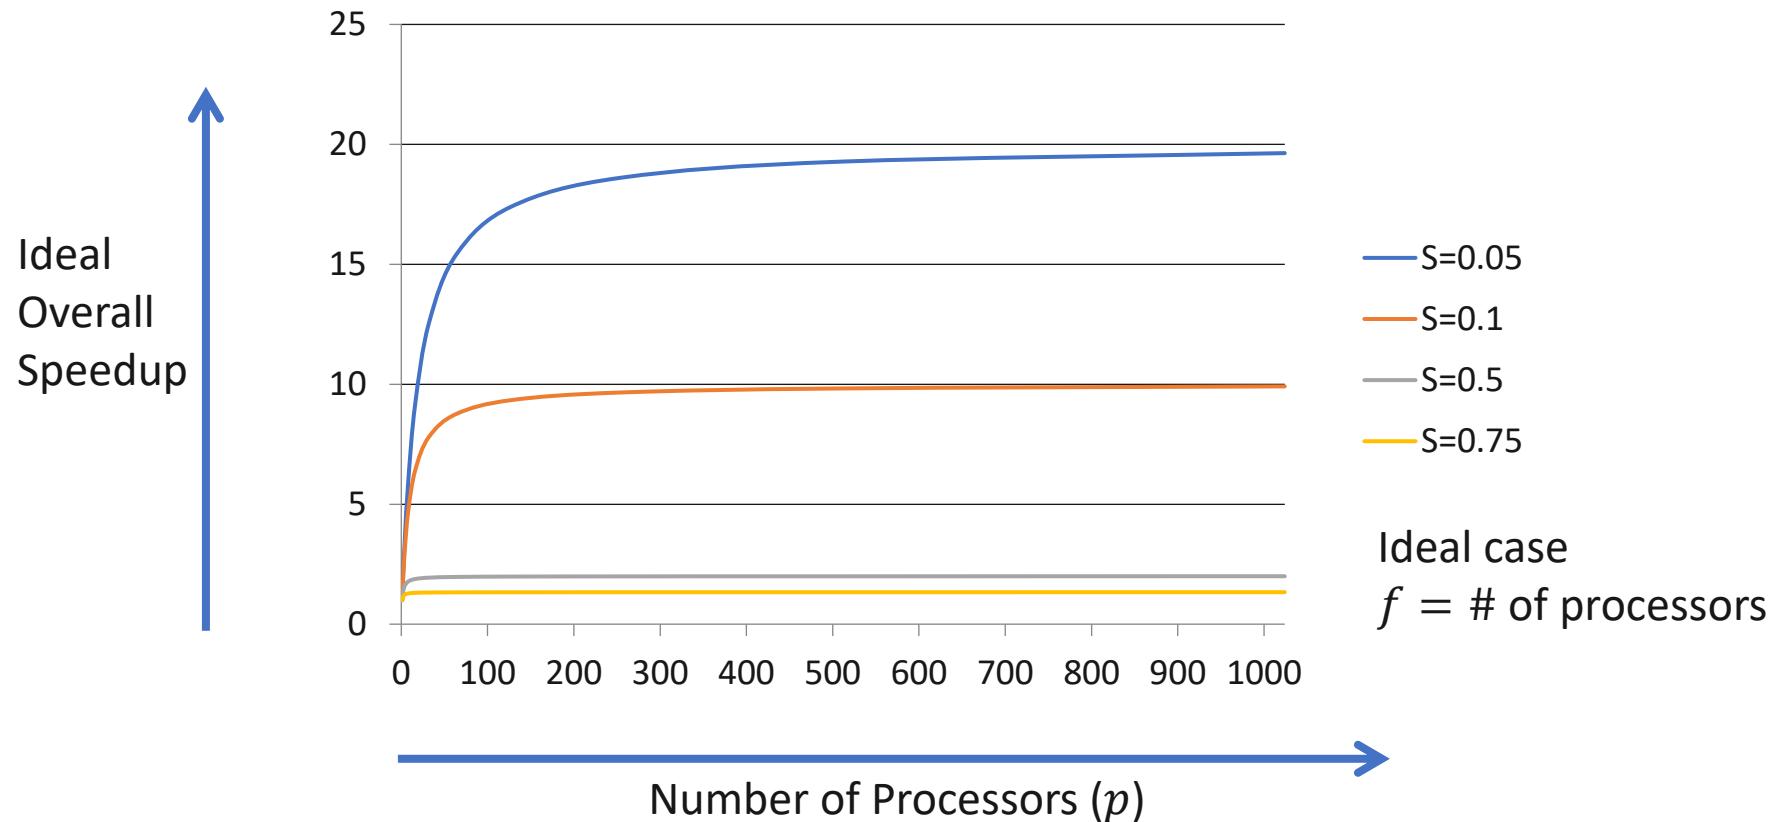
# Scalability

$$\text{Speedup} = \frac{\text{Serial time (on a uniprocessor system)}}{\text{Parallel time using } p \text{ processors}}$$

If speedup =  $O(p)$ , then it is a **scalable** solution

# Amdahl's Law

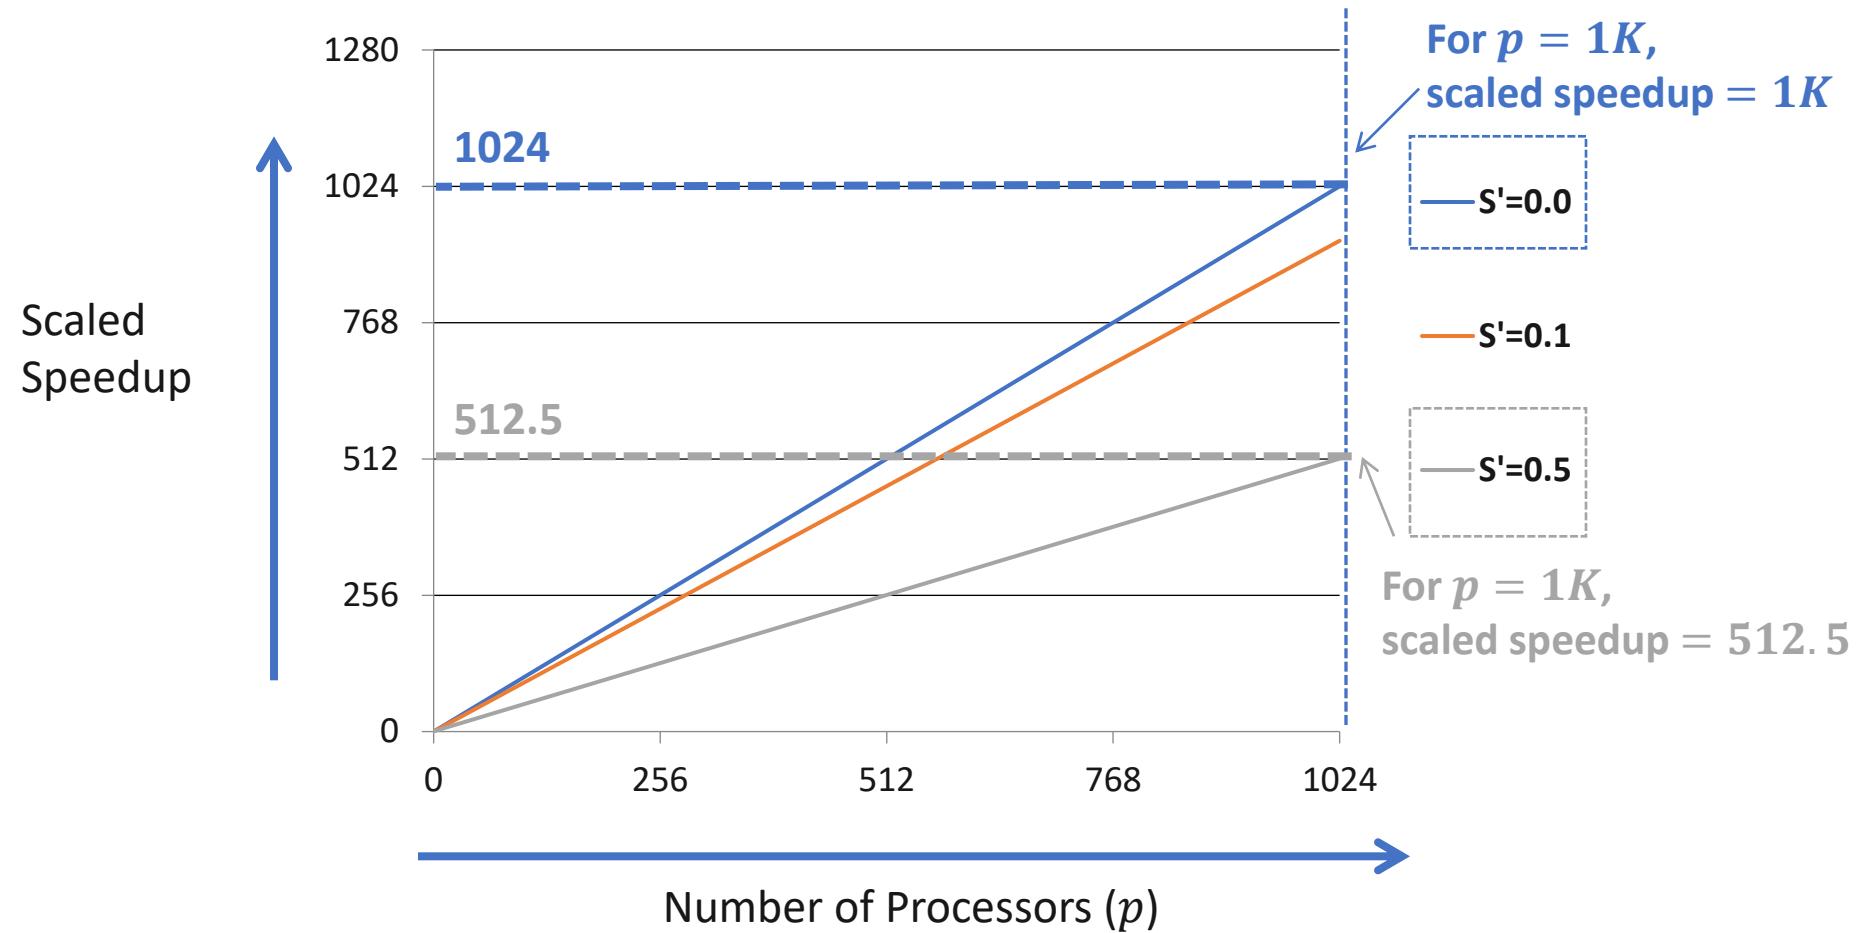
$$\text{Overall Speedup} = \frac{1}{S + P/f} = \frac{1}{S + (1 - S)/f} \rightarrow \frac{1}{S}$$



# Scaled Speedup (Gustafson's Law) (1)

- Amdahl's Law — Serial portion of code limits performance  
(As we use more processors)
- As we use more processors
  - we use more data
  - e.g., more fine grained model
- E.g. Processing  $N \times N$  image
  - Using  $p$  processors
  - As we increase  $p$  we usually increase image size

# Scaled Speedup (Gustafson's Law) (2)



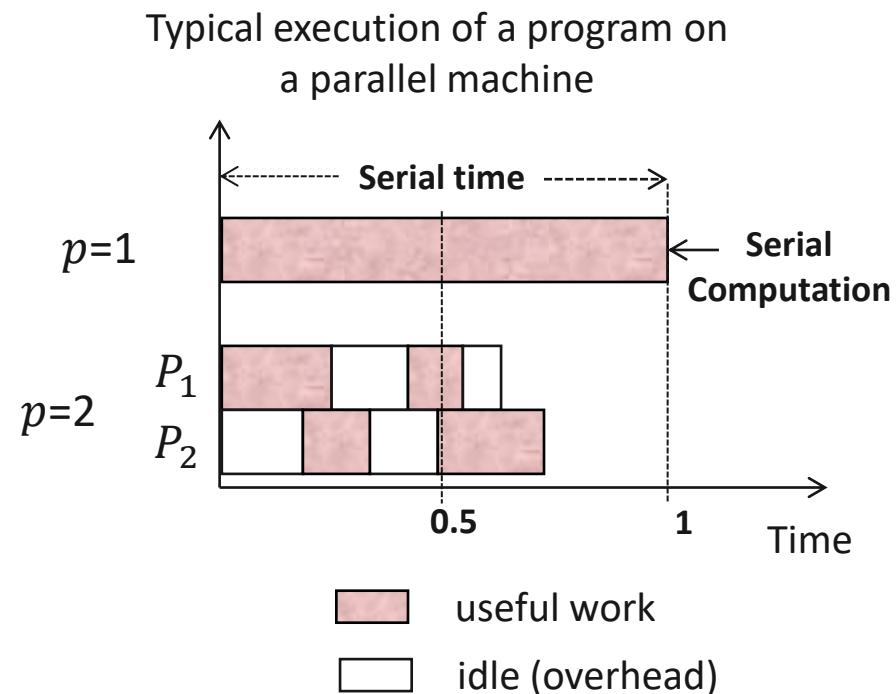
# Strong or Weak Scaling

- Strong Scaling
  - Governed by Amdahl's Law
  - The number of processors is **increased** while the problem size **remains constant**
  - Results in a **reduced** workload per processor
- Weak Scaling
  - Governed by Gustafson's Law
  - **Both** the number of processors and the problem size are **increased**
  - Results in a **constant** workload per processor

# Performance (1)

## Efficiency

- Question: If we use  $p$  processors, is speedup =  $p$  ?
- Efficiency  $\triangleq$  Fraction of time a processor is usefully employed during the computation
- $E = \text{Speedup} / \# \text{ of processors used}$ 
  - $E$  is the average efficiency over all the processors
  - Efficiency of each processor can be different from the average value

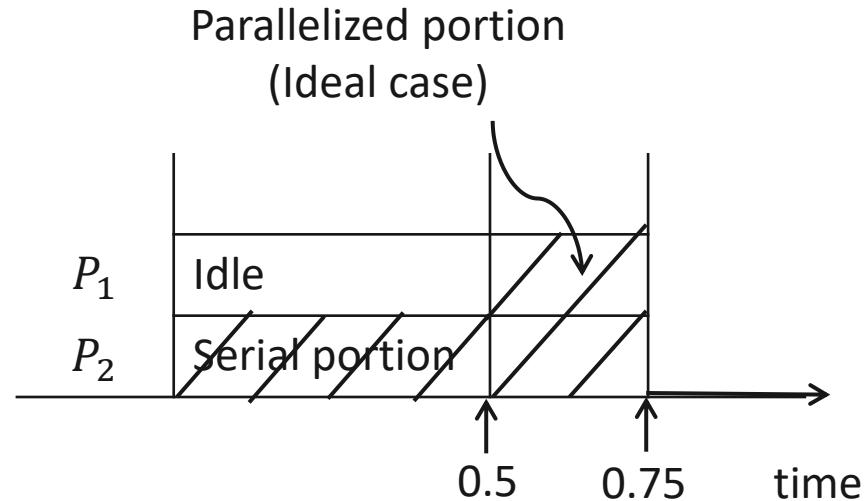


# Performance (2)

Ex.

$$\begin{aligned} S &= 0.5 \\ P &= 0.5 \end{aligned}$$

2 processor system



$$\text{Speedup} = \frac{1}{0.75} = 4/3$$

$$(\text{Average}) \text{ Efficiency} = \frac{4/3}{2} = 2/3$$

$$\text{Efficiency of } P_1 = \frac{0.25}{0.75} = 1/3$$

$$\text{Efficiency of } P_2 = \frac{0.75}{0.75} = 1$$

$$(\text{Average}) \text{ Efficiency} = \frac{1/3+1}{2} = 2/3$$

# Performance (3)

- Cost = Total amount of work done by a parallel system
  - = Parallel Execution Time  $\times$  Number of Processors
  - =  $T_p \times p$
- Cost is also called **Processor Time Product**
- COST OPTIMAL (or WORK OPTIMAL) Parallel Algorithm
  - Total work done = Serial Complexity of the problem

# Blocked Matrix Multiplication

- Serial Algorithm -  $O(n^3)$

- Parallel Algorithm:
- Number of steps  $n/b$
- Work done by each processor in each step:  $O(b)$

Cost of the algorithm:  $O\left(n^2 \times \left[\frac{n}{b} \times b\right]\right) = O(n^2 \times n) = O(n^3)$

Total Number  
of Processors

Work done in  
each processor

- Cost of the algorithm  $\sim$  serial complexity. So work optimal.

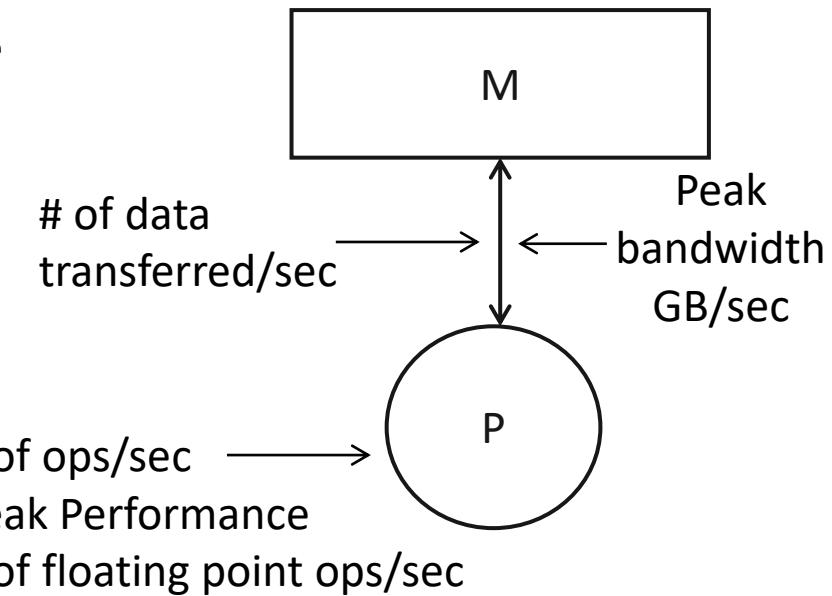
# Sum of Elements in an Array

- Serial Algorithm:  $O(p)$
- Parallel Algorithm:
- Number of steps -  $O(\log_2 p)$
- Number of processors in each step -  $O(p)$ 
  - Note: We can also consider only active processors to come up with a different analysis.
- Cost of the Algorithm -  $O(p \log_2 p)$
- Cost of the algorithm  $\sim$  serial complexity. So not work optimal (assuming GPU model). It is work optimal if we only consider active processor in the analysis.

# Roofline Model

- **In Practice:** Processor Performance  $\gg$  Memory Performance
- Processor may idle waiting for data

$$\text{Machine Balance} = \frac{\text{Peak Performance (Flops/sec)}}{\text{Peak Bandwidth (GB/sec)}}$$



The Roofline Model  
 Samuel Williams, Lawrence Berkeley National Lab  
<https://escholarship.org/content/qt0qs3s709/qt0qs3s709.pdf>

# Roofline Model

- Arithmetic Intensity (AI) = For every fetched data (from memory) how many ops make use of them

$$= \frac{\# \text{ of ops performed}}{\text{Amount of data moved from external memory}}$$

- Algorithm Dependent

- Example

$$\text{AI of vector inner product} = \frac{2n}{2n} = O(1)$$

$$\text{AI of Dense matrix multiplication} = \frac{2n^3}{2n^2} = O(n)$$

# Roofline Model

- Understand system bottlenecks – Computation vs Memory
- High level model using few parameters (ignores ISA etc.)
- Can be used to understand and evaluate an application performance on a target
- Can be used to make architectural decisions for custom accelerators

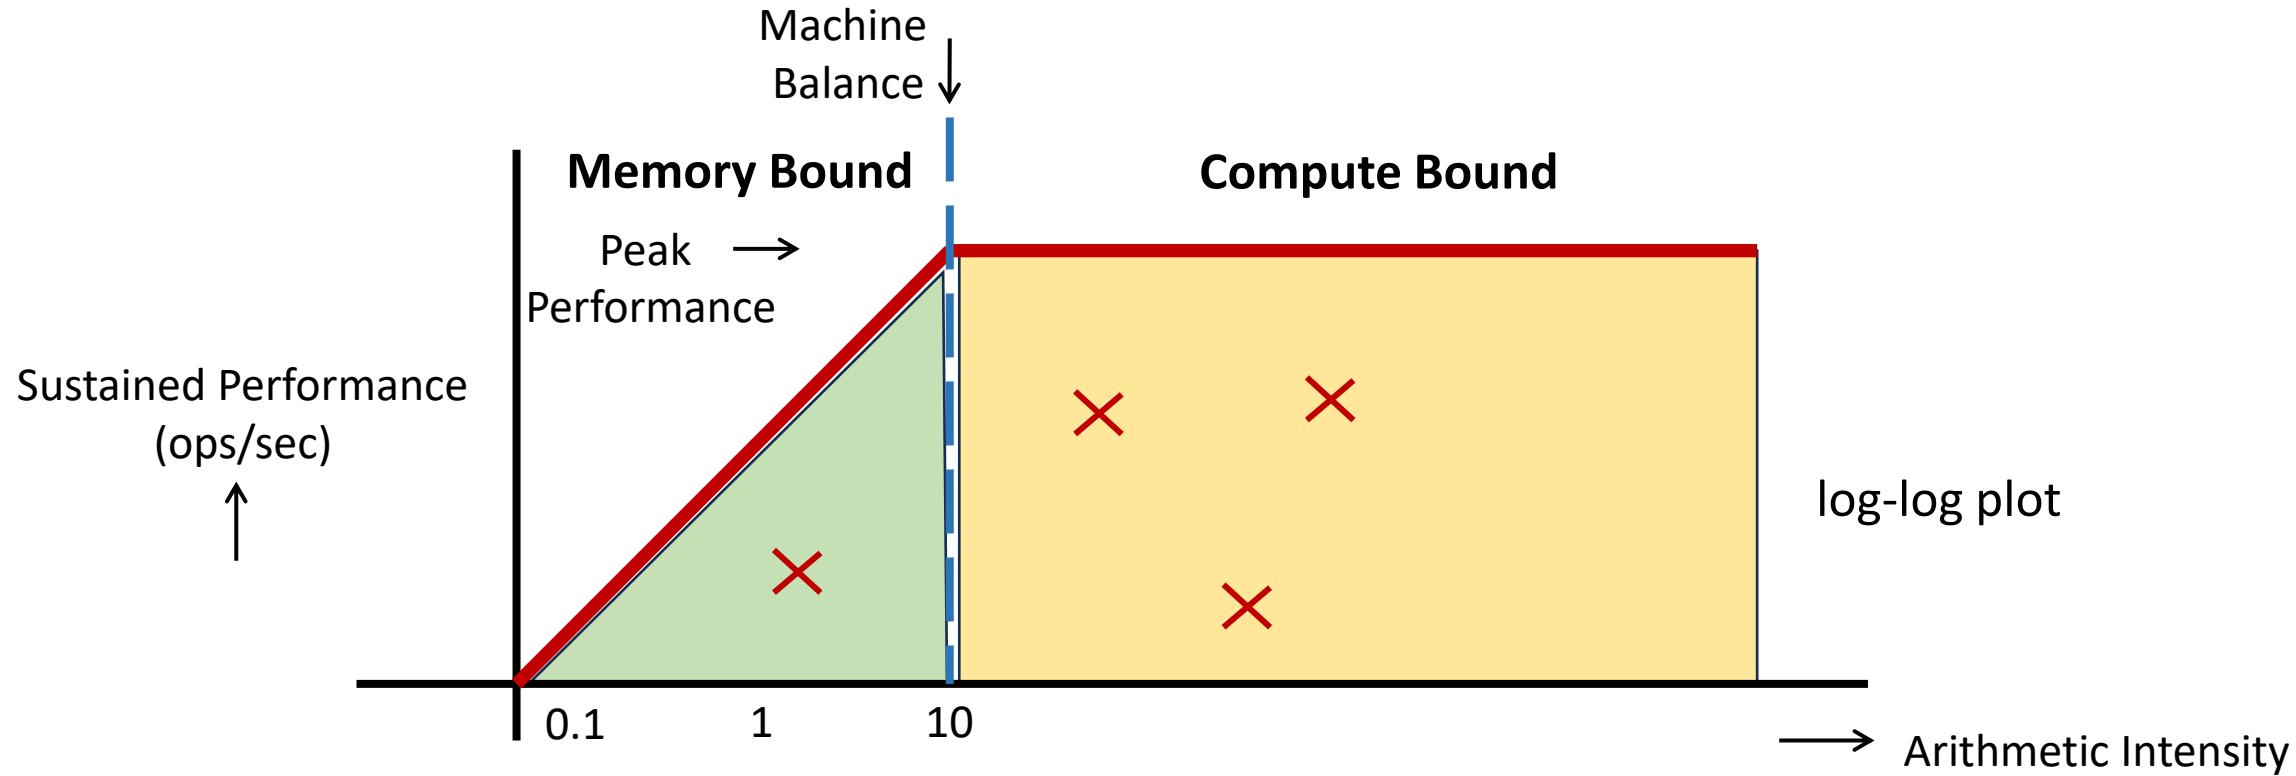
The Roofline Model

Samuel Williams, Lawrence Berkeley National Lab

<https://escholarship.org/content/qt0qs3s709/qt0qs3s709.pdf>

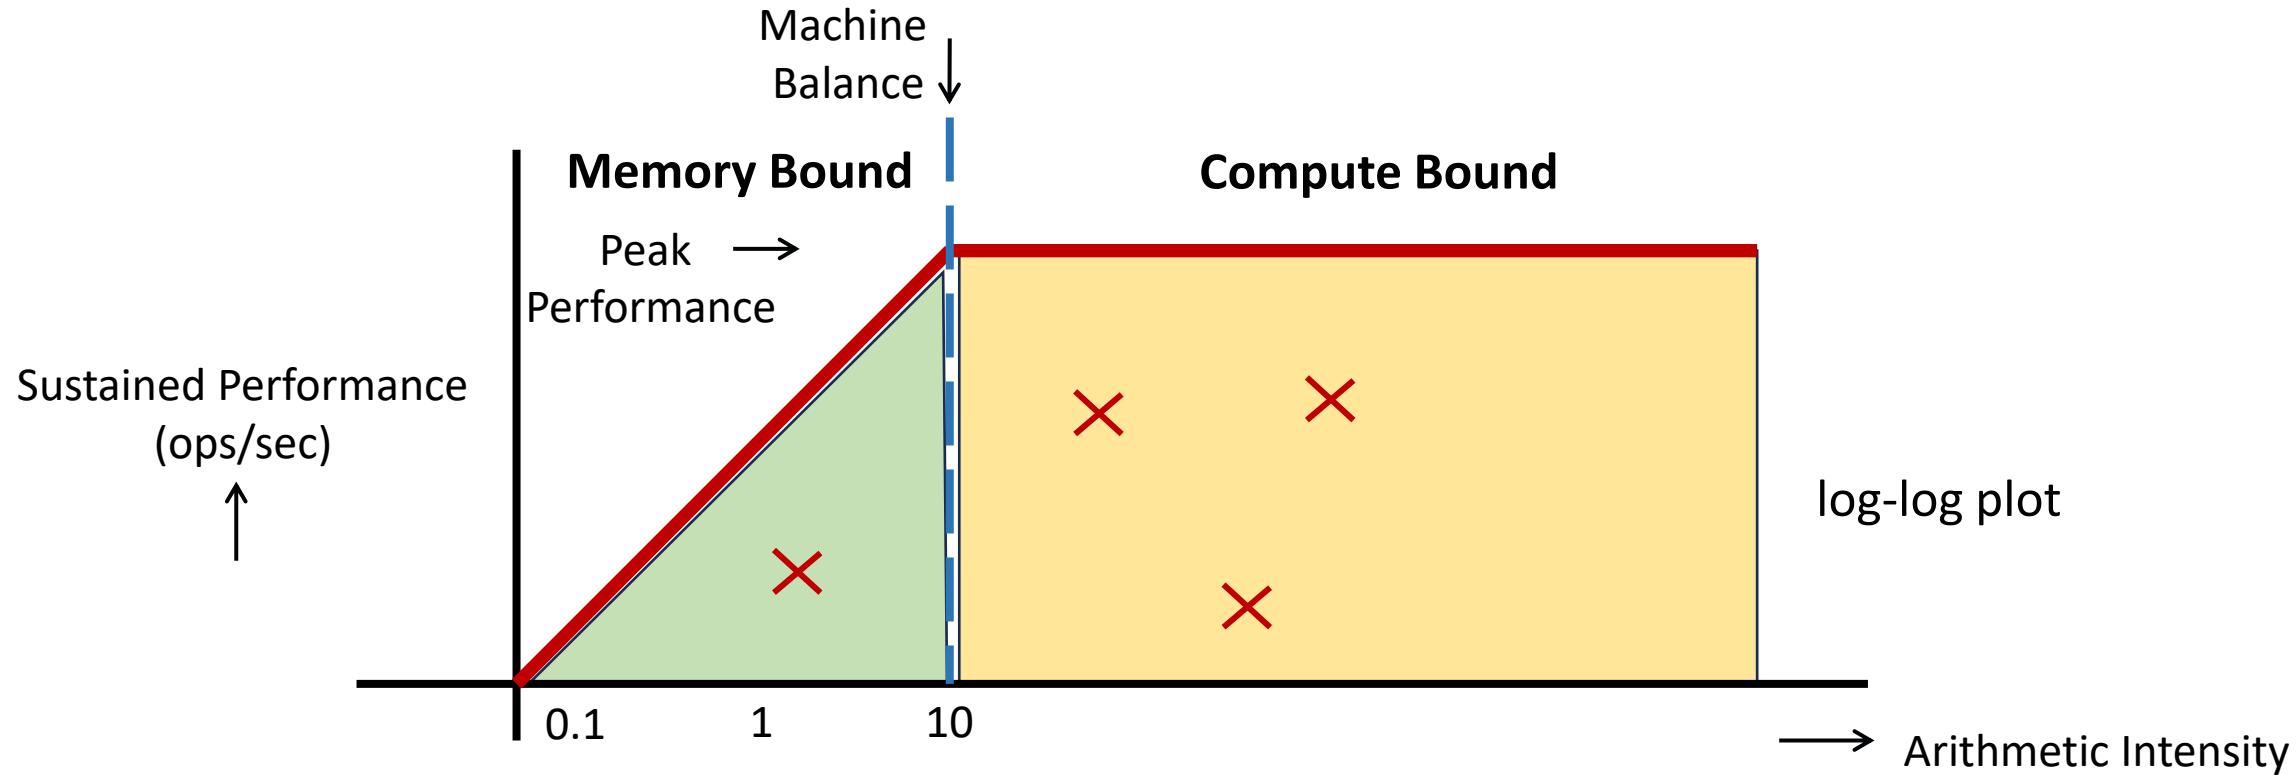
# Compute Bound vs Memory Bound

## Roofline Model (for a given platform)



# Compute Bound vs Memory Bound

## Roofline Model (for a given platform)

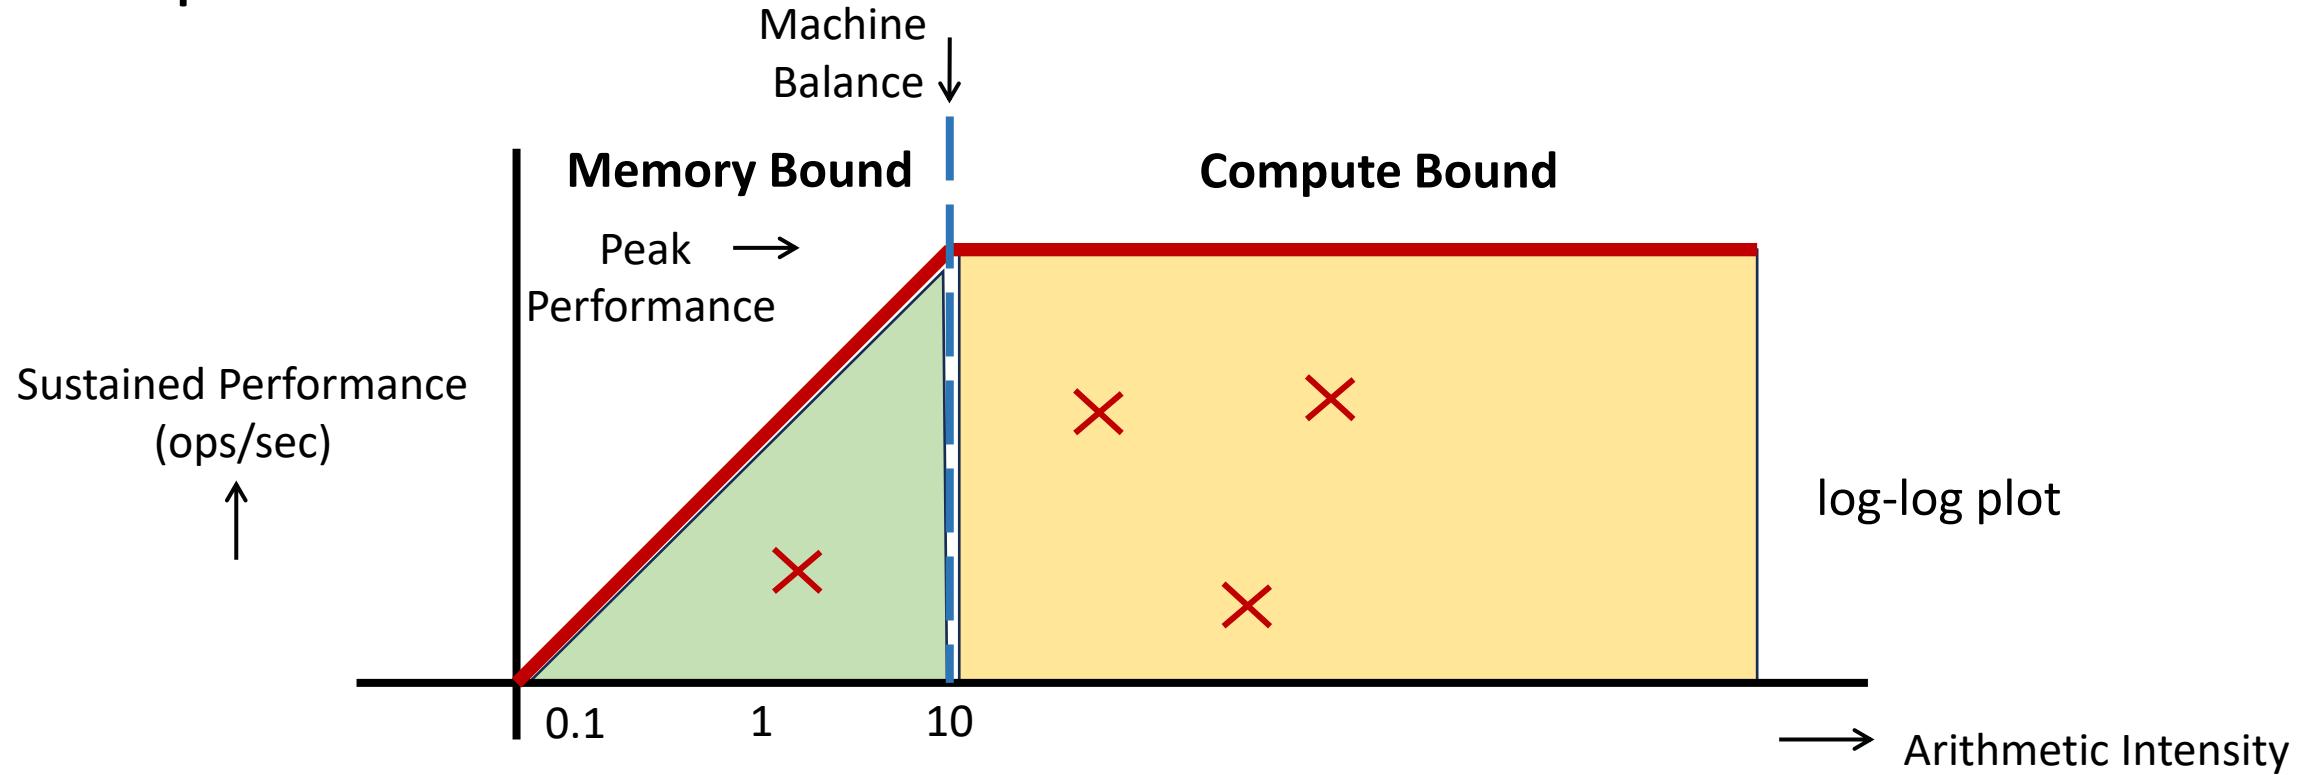


- Bold Red line denotes the maximum achievable performance for a given arithmetic intensity
- Memory bound region: maximum achievable performance bounded by memory bandwidth
- Compute bound region: maximum achievable performance bounded by compute capability

# Compute Bound vs Memory Bound Roofline Model (for a given platform)

- Compute Bound: Maximum Performance is limited by available compute resources
  - if we add more compute resources, the performance may improve (for a given bandwidth)
  - Useful in designing accelerators. Not a focus of this class
- Memory Bound: Maximum Performance is limited by memory bandwidth
  - Reorder computations to more efficiently utilize bandwidth
  - In other words, improve arithmetic intensity to move right on x axis, which also improves y axis

# Roofline Model – What about actual performance?

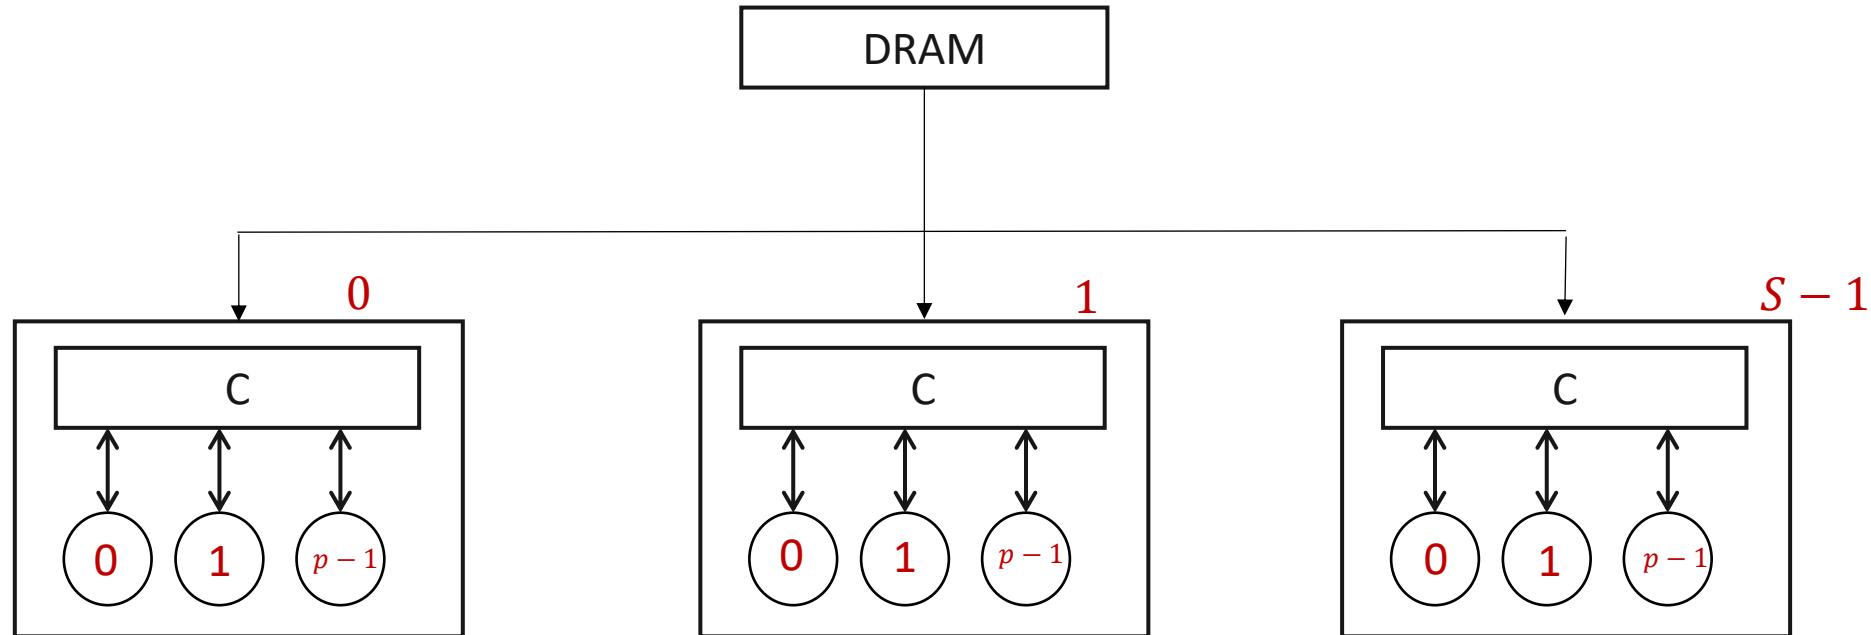


- For a fixed AI: improving performance (going up) requires better parallelism – Data and Task Parallel Techniques
- Once the Red bold line is reached
  - Green Region: Write better algorithm to improve arithmetic intensity
  - Yellow Region: No algorithmic solution, need to add more computation power

# Outline

- Parallel Program Analysis (Quick Review)
- Data Parallel Algorithms
  - Sorting
- Task Parallelism

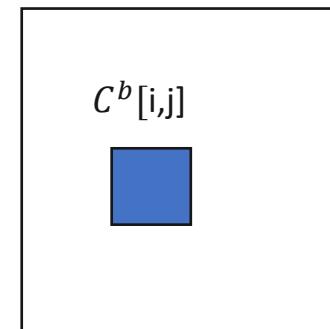
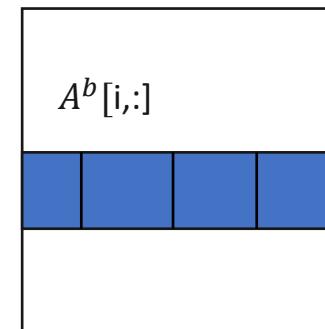
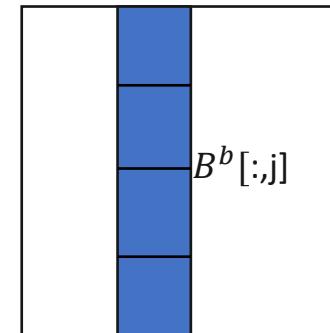
# Modeling GPU Architectures



- Number of Blocks can be (in fact should be) greater than the number of SMPs
- Number of threads per block need not be equal to the number of compute cores per SMP
  - It is good to have it as a multiple though
- We will discuss these considerations when we discuss GPU programming later
- For now, assume the simple model above.
  - $S$  Blocks
  - $p$  Threads per block

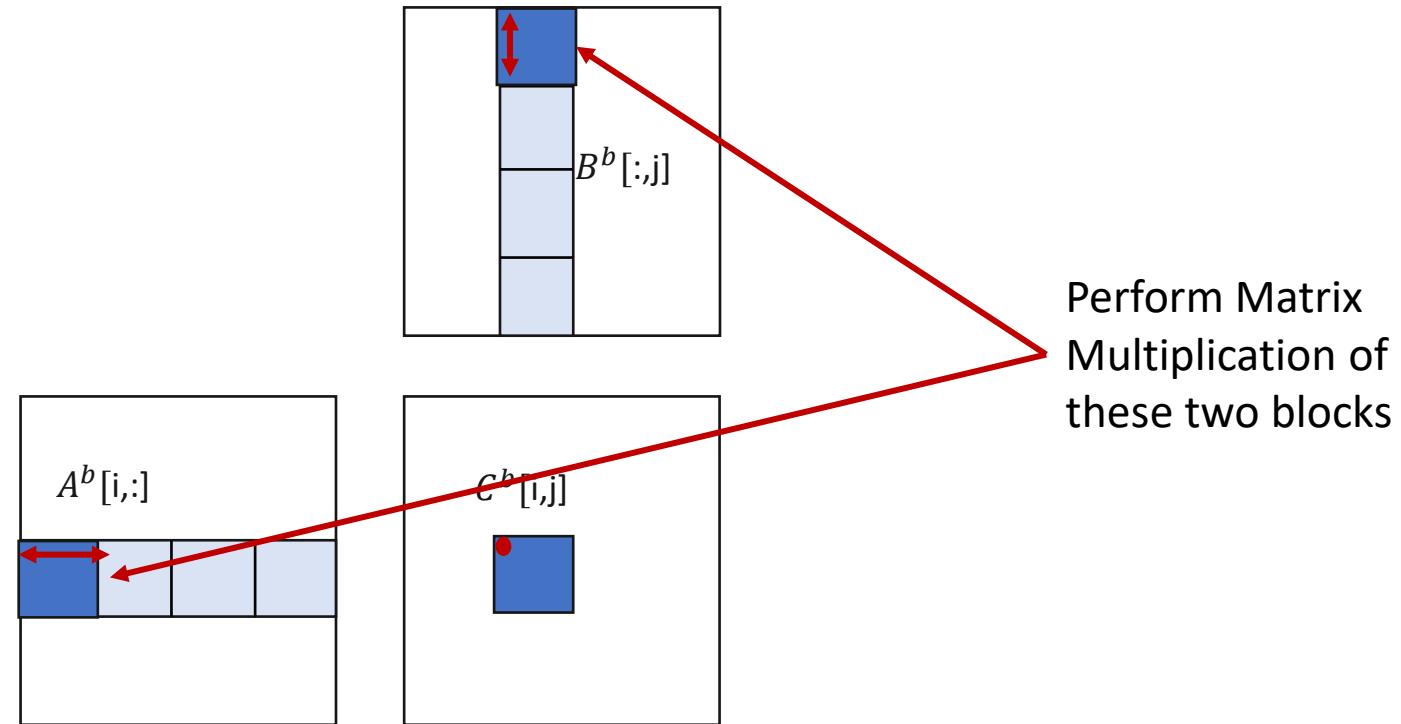
# Block Matrix Multiplication

- Key Idea: Partition matrices A/B/C into  $b \times b$  size blocks
- Blocks denoted as  $A^b[1:\frac{N}{b}][1:\frac{N}{b}]$ /  $B^b[1:\frac{N}{b}][1:\frac{N}{b}]$ /  $C^b[1:\frac{N}{b}][1:\frac{N}{b}]$
- $C^b[i][j] = \text{BlockedMM}(A^b[i][:], B^b[:,j])$
- For  $k = 0$  to  $N/b$ 
  - Fetch  $A^b[i][k]$ ,  $B^b[k][j]$
  - Matrix Multiply  $A^b[i][k]$ ,  $B^b[k][j]$  and Accumulate  $C^b[i][j]$



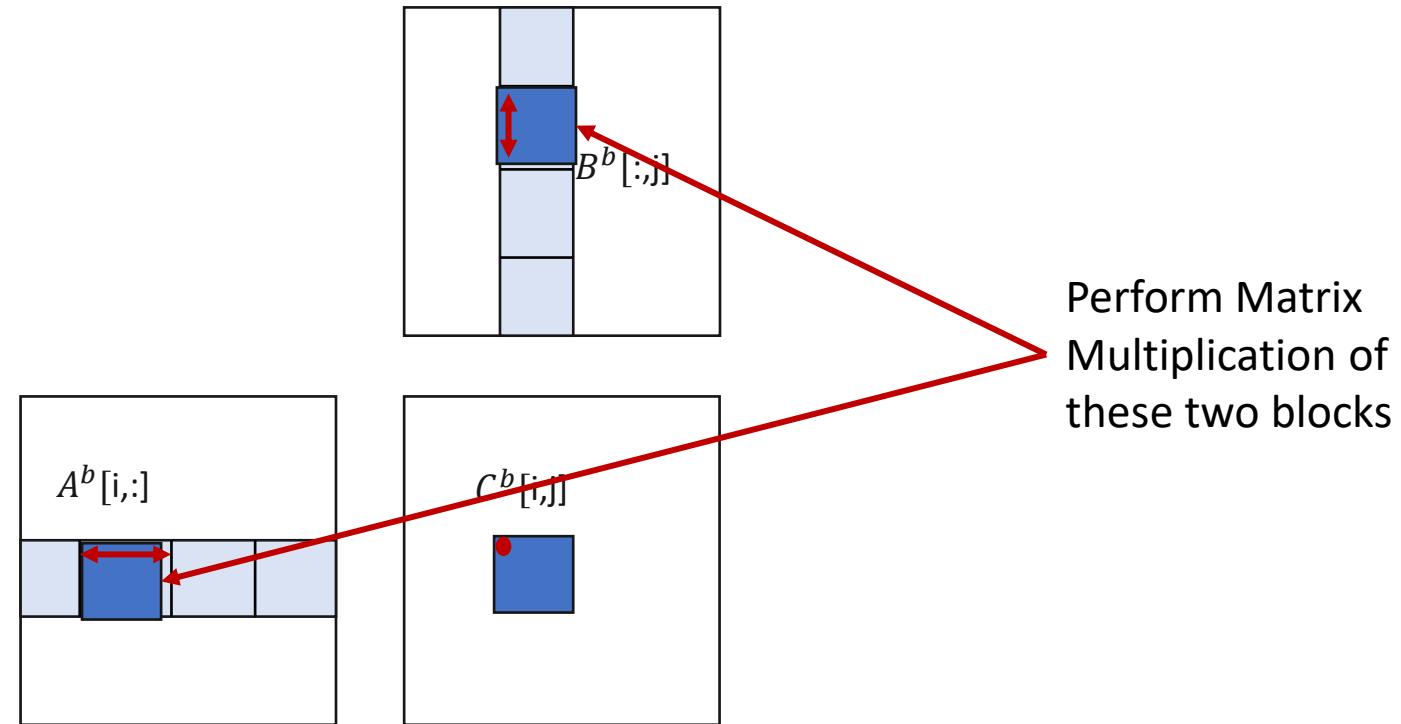
# Block Matrix Multiplication

SMP:  $BID_i, BID_j$



# Block Matrix Multiplication

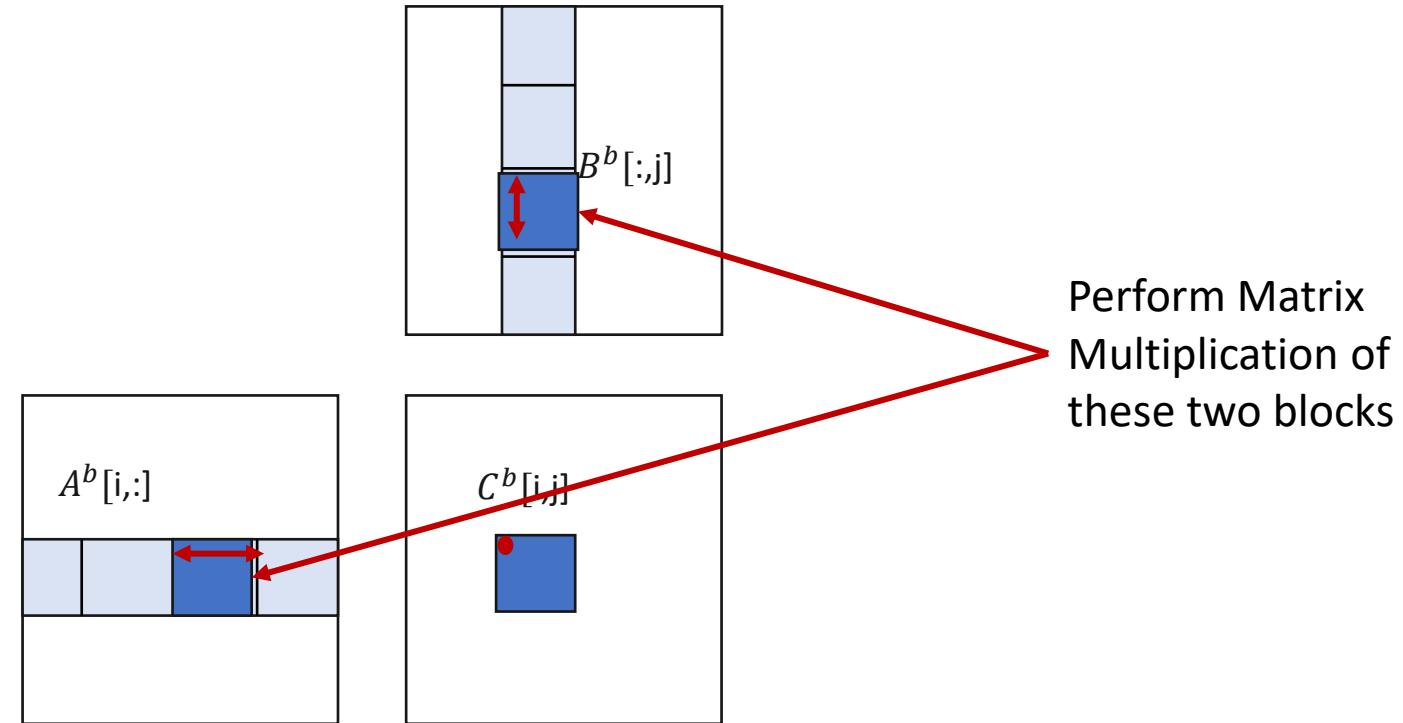
SMP:  $BID_i, BID_j$



Perform Matrix  
Multiplication of  
these two blocks

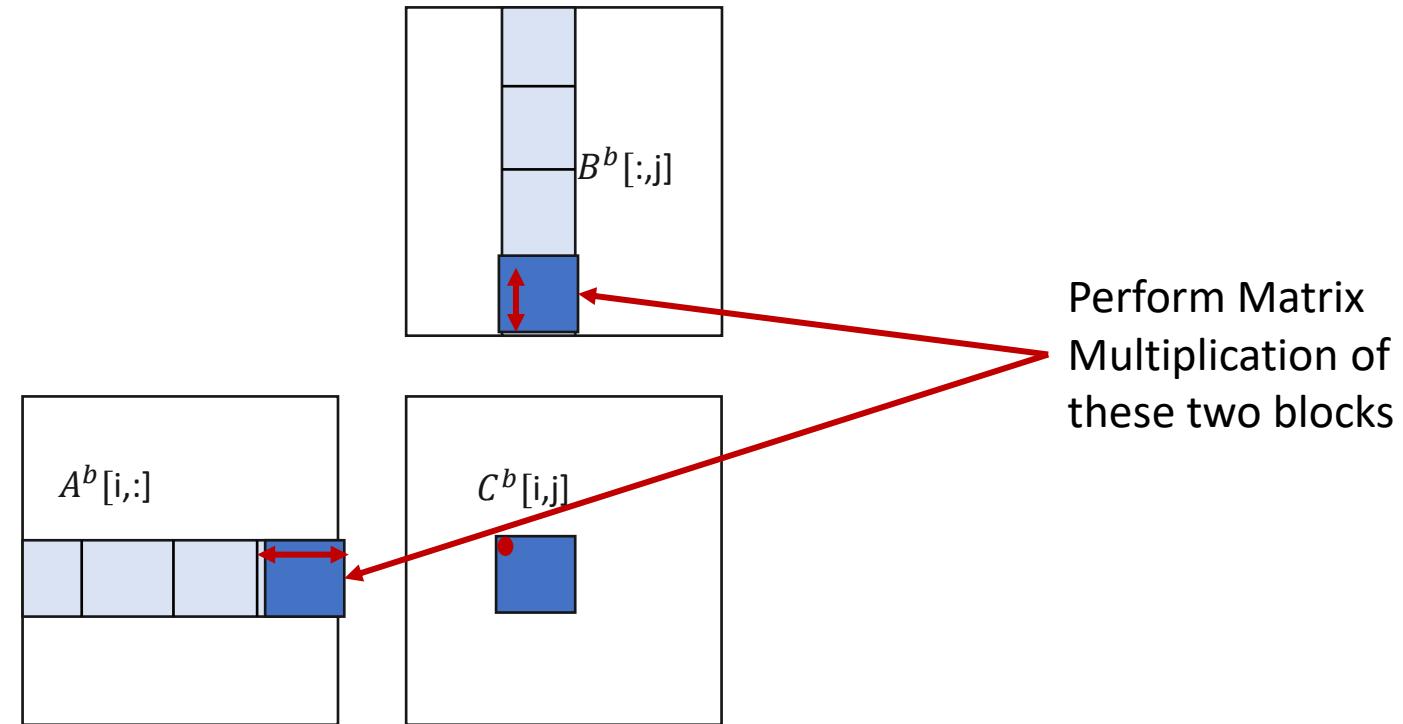
# Block Matrix Multiplication

SMP:  $BID_i, BID_j$



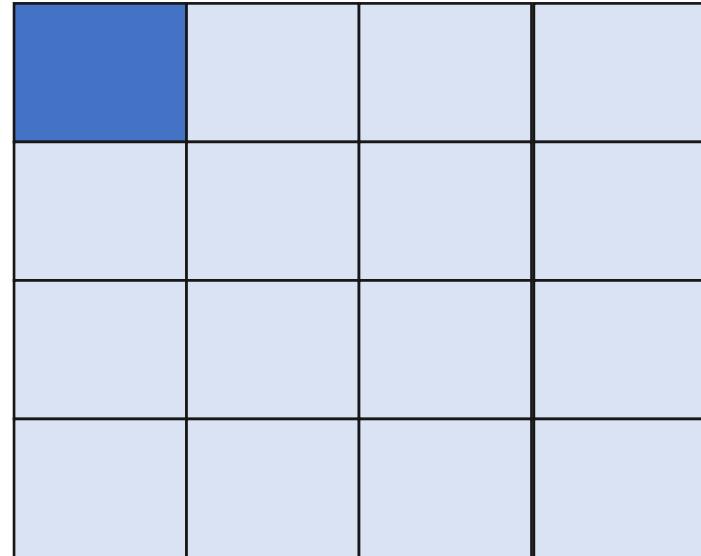
# Block Matrix Multiplication

SMP:  $BID_i, BID_j$



# Block Matrix Multiplication

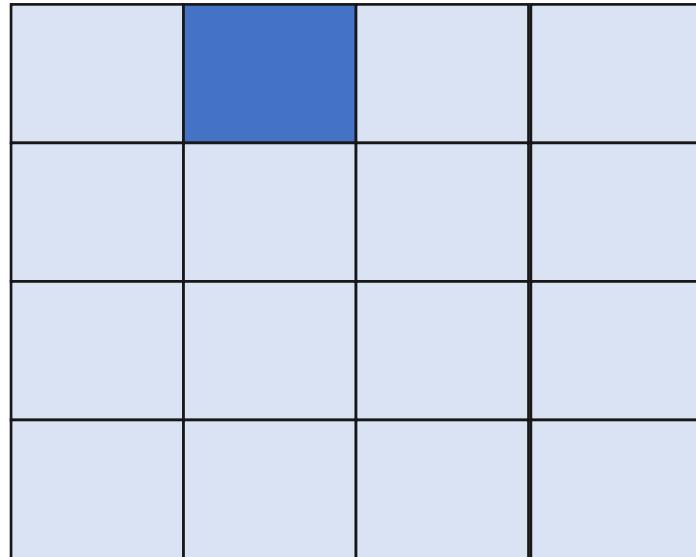
SMP: 0,0



$C$

# Block Matrix Multiplication

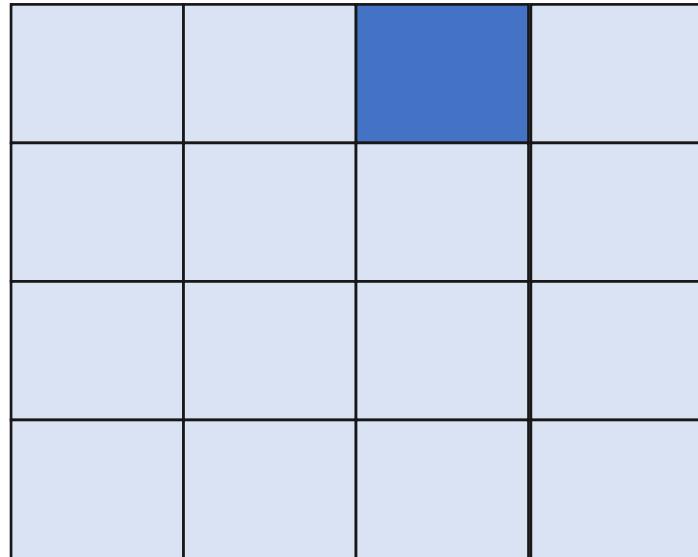
SMP: 0,1



$C$

# Block Matrix Multiplication

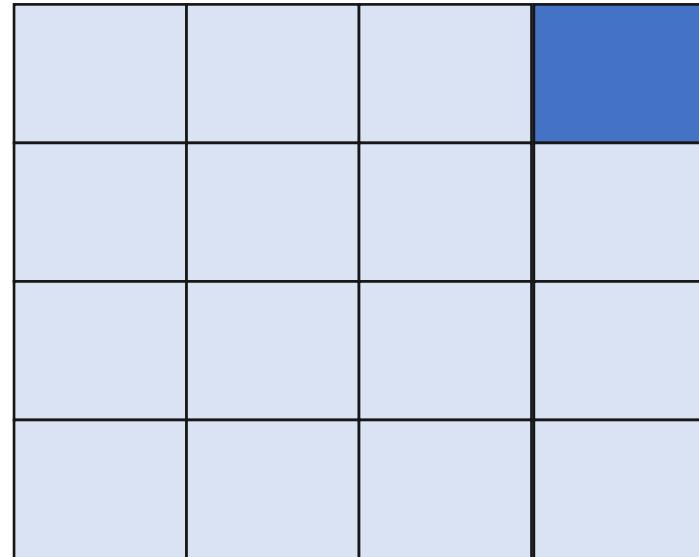
SMP: 0,2



$C$

# Block Matrix Multiplication

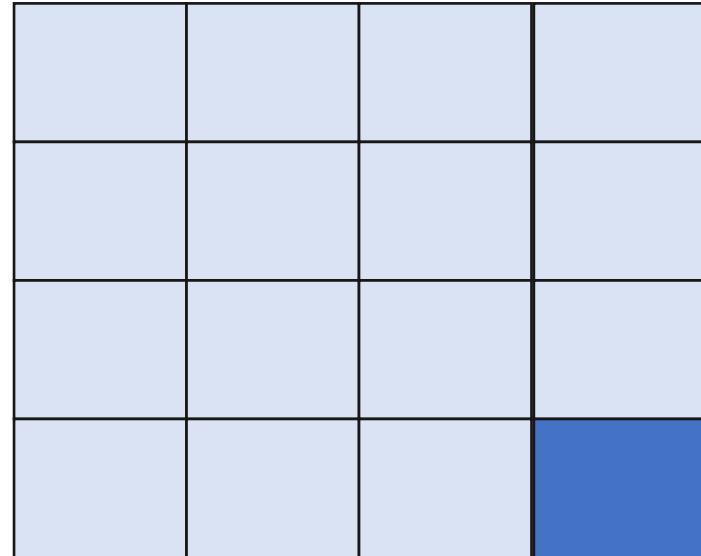
SMP: 0,3



$C$

# Block Matrix Multiplication

SMP: ?, ?



$C$

# Block Matrix Multiplication

SMP: 0,0

SMP: 0,1

SMP: 0,2

SMP: 0,3

SMP: 1,0

SMP: 1,1

SMP: 1,2

SMP: 1,3

SMP: 2,0

SMP: 2,1

SMP: 2,2

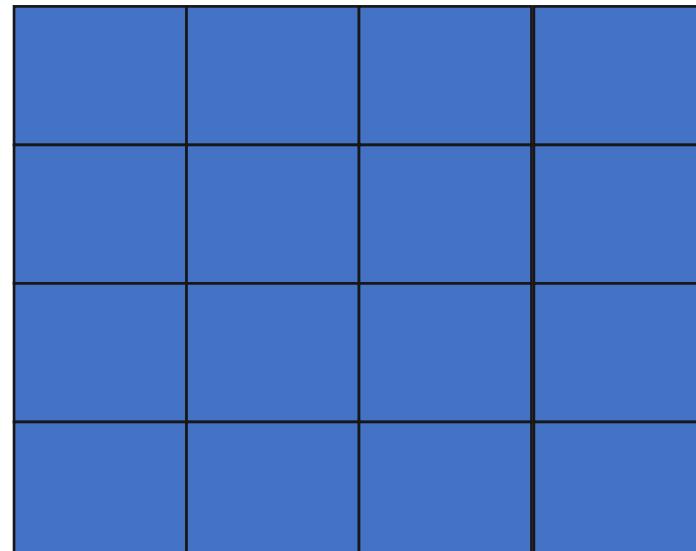
SMP: 2,3

SMP: 3,0

SMP: 3,1

SMP: 3,2

SMP: 3,3



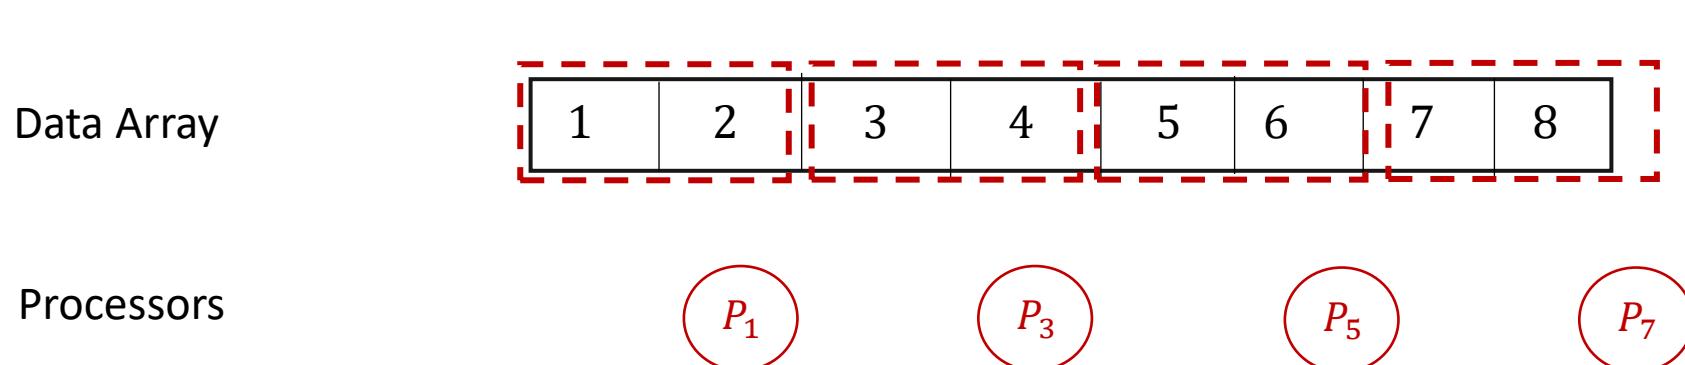
*C*

In parallel produce all the blocks of *C*

- No conflicts. Each SMP can work independently
- In practice, number of blocks can be greater than # SMPs available on the hardware
  - Hardware performs time multiplexing, hidden from programmer,
  - We will discuss in our HIP programming class

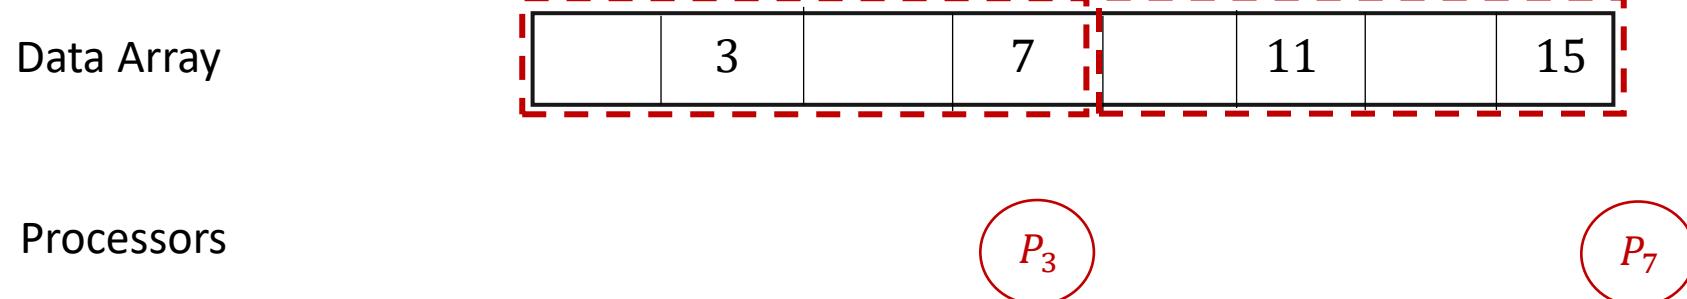
# Sum of an Array

- Block MM – Output Partitioning
- Sum of an Array – Input Partitioning



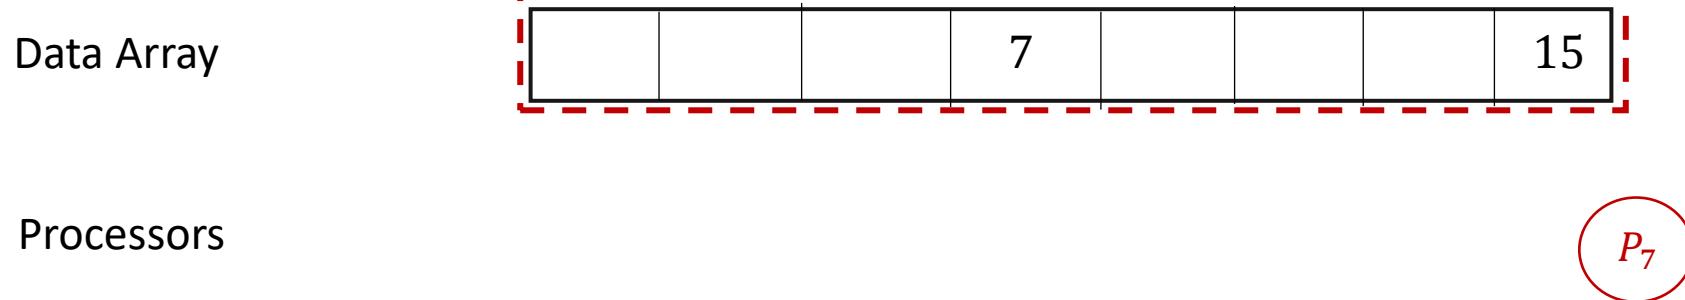
# Sum of an Array

- Block MM – Output Partitioning
- Sum of an Array – Input Partitioning



# Sum of an Array

- Block MM – Output Partitioning
- Sum of an Array – Input Partitioning



# Parallel Sorting

- Sort a list of elements in ascending (or descending) order
- Input: [8,7,6,5,4,3,2,1]
- Sorting Order: Ascending
- Output: [1,2,3,4,5,6,7,8]

# Parallel Sorting

Input:

8	7	6	5	4	3	2	1
---	---	---	---	---	---	---	---

Output:

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Can I do conflict-free output partitioning based data parallelism?

# Parallel Sorting

Input:

8	7	6	5	4	3	2	1
---	---	---	---	---	---	---	---

Output:

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Can I do conflict-free output partitioning based data parallelism?

Maybe, for example if we do bucket or radix sort

# Parallel Sorting

Input:

8	7	6	5	4	3	2	1
---	---	---	---	---	---	---	---

Output:

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Can I do conflict-free output partitioning based data parallelism?  
**Think about how that could work. (May include as a question in Exam)**

# Parallel Sorting

Input:

8	7	6	5	4	3	2	1
---	---	---	---	---	---	---	---

Output:

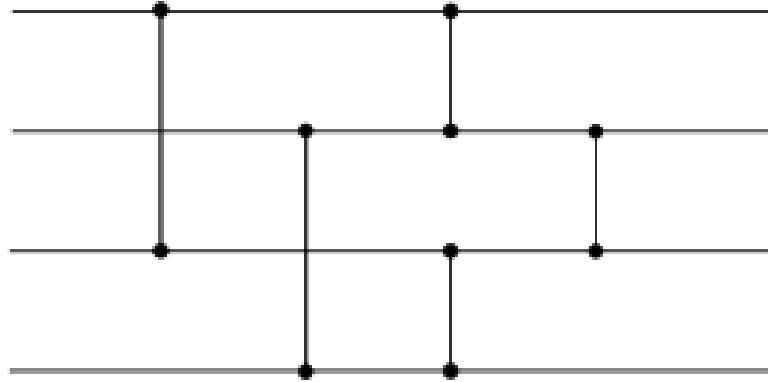
1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Can I do conflict-free output partitioning based data parallelism?

We will discuss an input partitioning based approach in this lecture  
with a focus on comparison based sort

# Parallel Sort

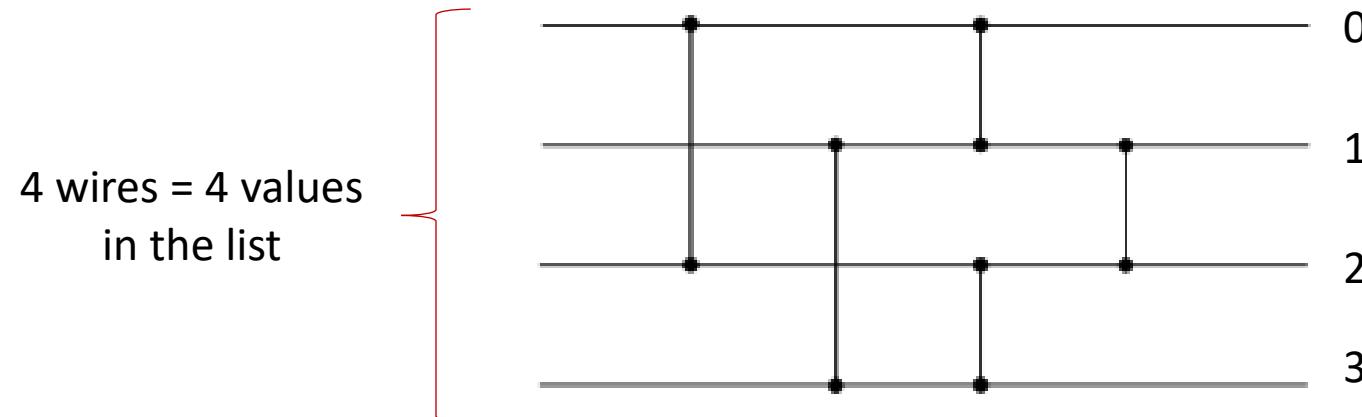
- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



[https://en.wikipedia.org/wiki/Sorting\\_network](https://en.wikipedia.org/wiki/Sorting_network)

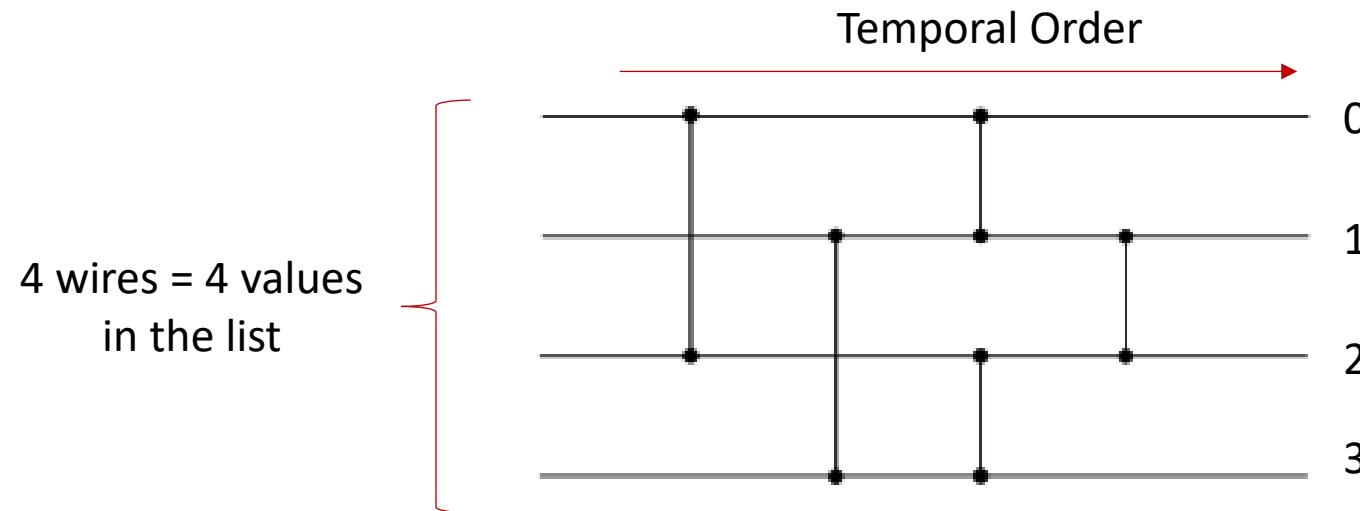
# Parallel Sort

- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



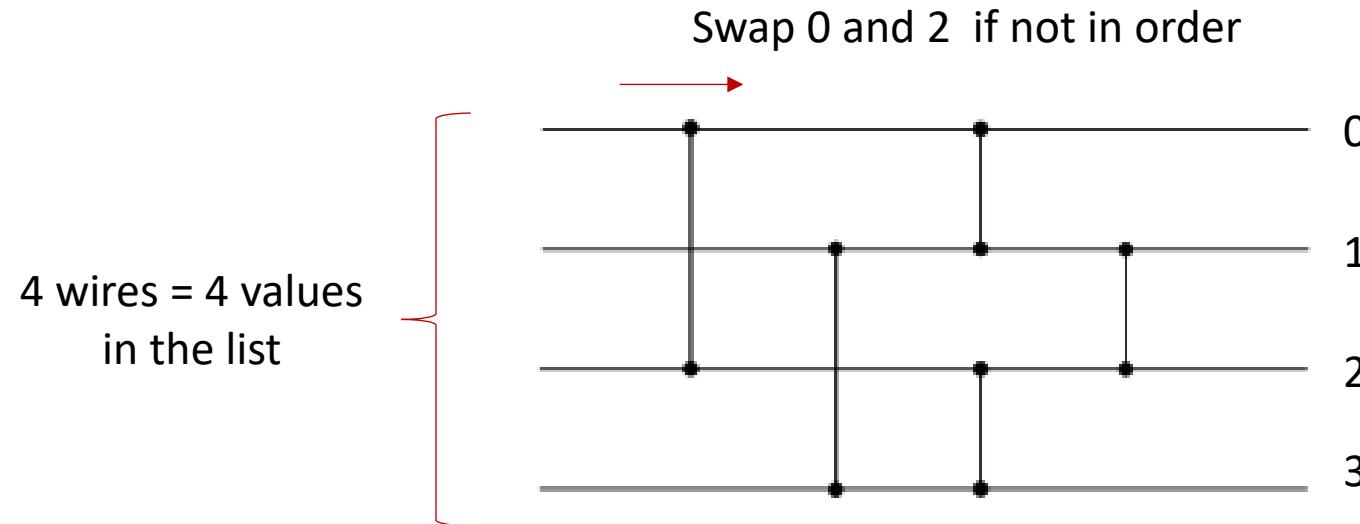
# Parallel Sort

- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



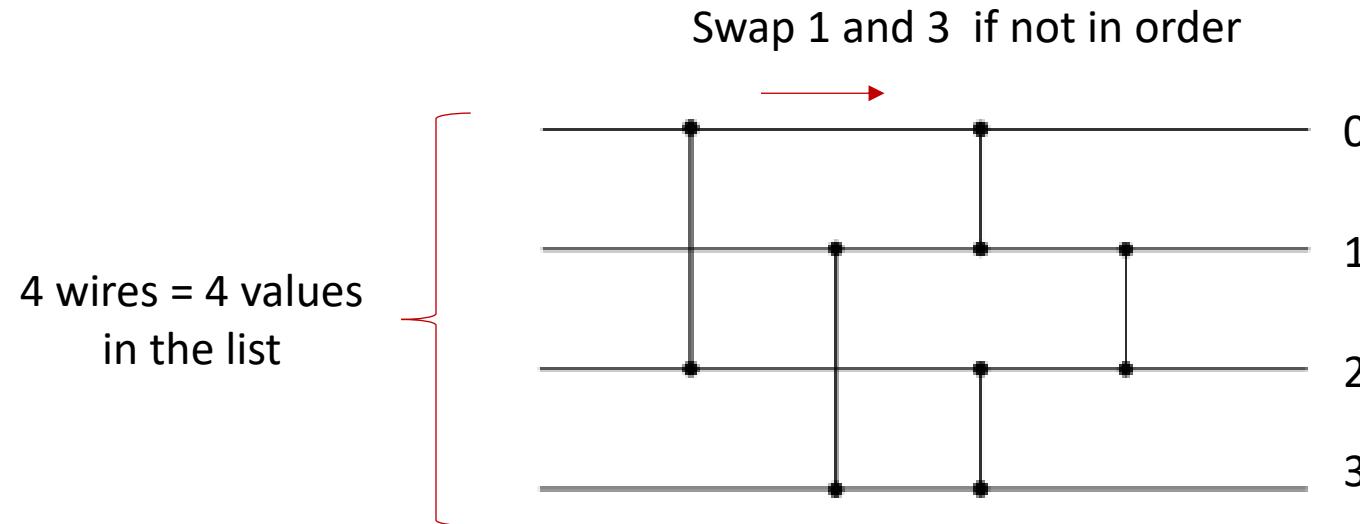
# Parallel Sort

- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



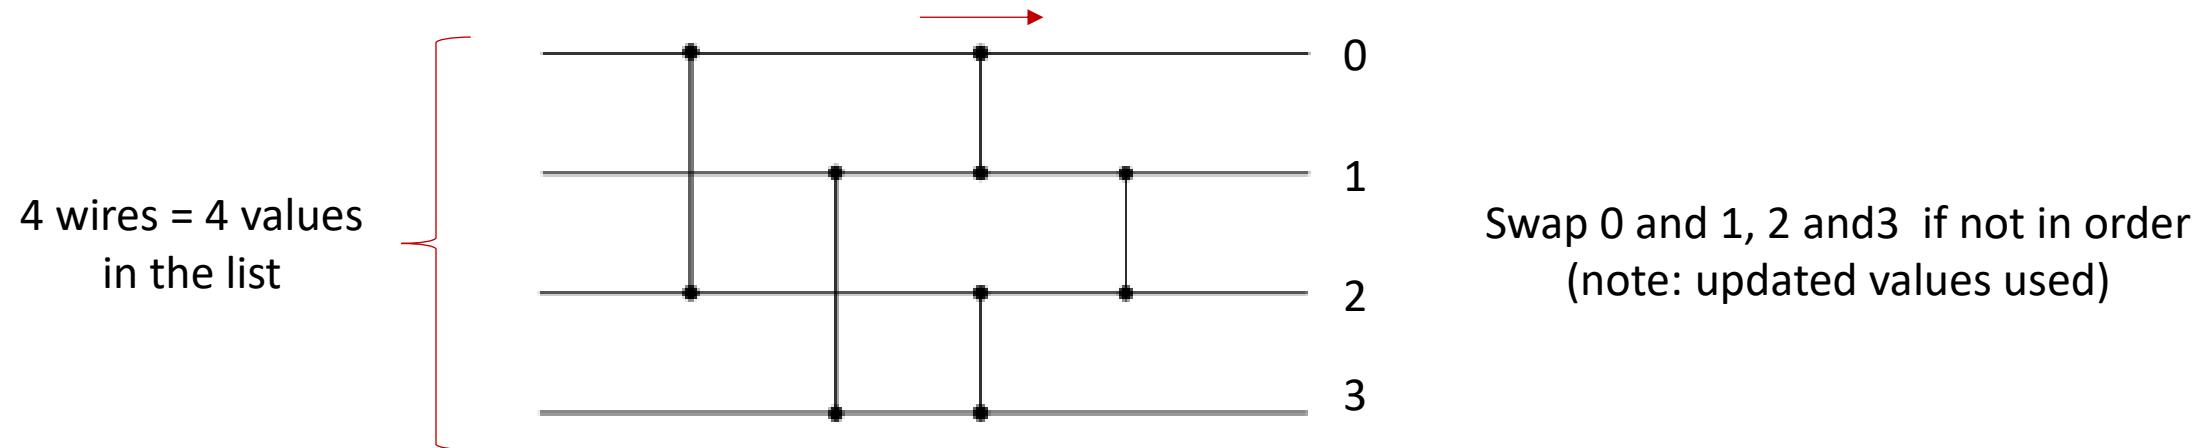
# Parallel Sort

- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



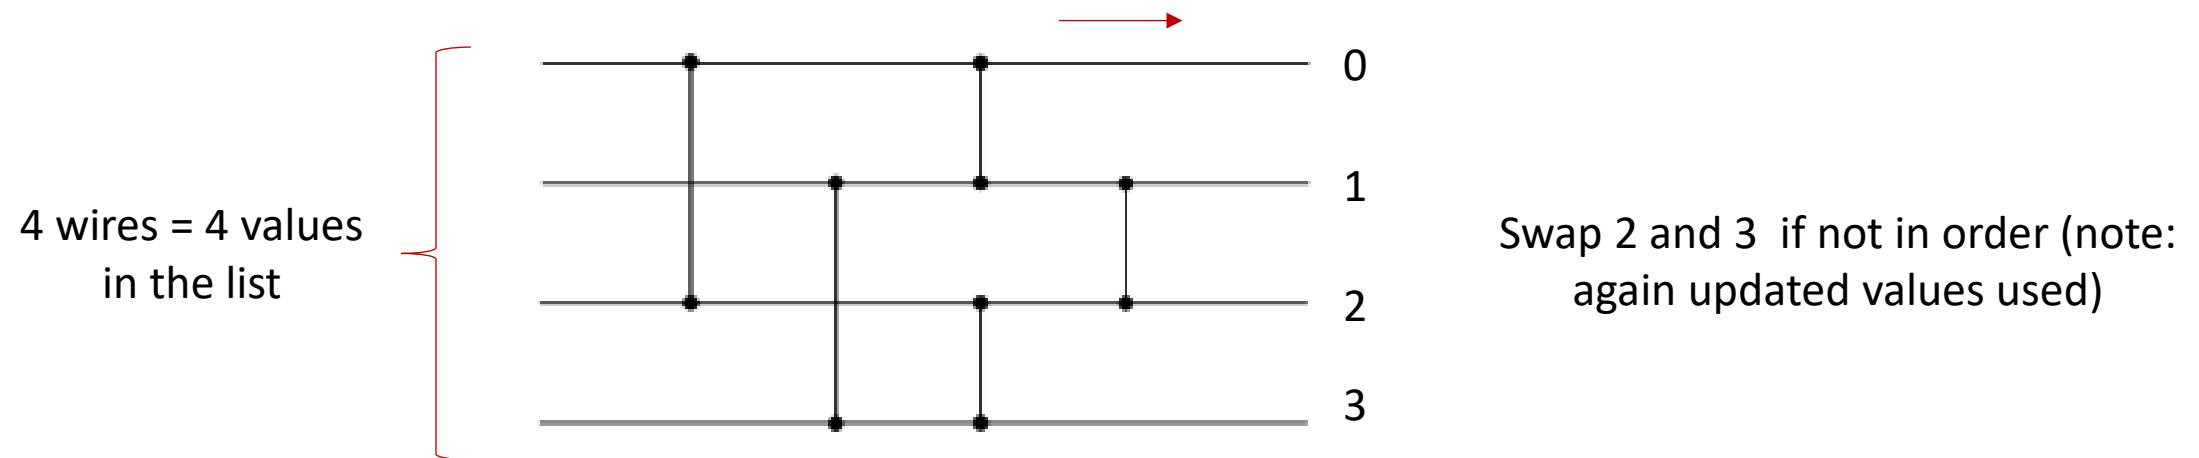
# Parallel Sort

- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



# Parallel Sort

- Definition: Sorting Networks
  - An abstraction built of a fixed number of “wires” that carry value, and comparators between pairs of “wires” that swap values if they are not in the desired order



# Parallel Sorting

- Key Idea: Build a sorting network to visualize the computations
- Use the visualization to understand parallelism and design a parallel program
- A number of sorting algorithms exist: bitonic sort, odd-even sort, ....
  - Bitton, Dina, et al. "A taxonomy of parallel sorting." *ACM Computing Surveys (CSUR)* 16.3 (1984): 287-318.
  - Singh, Dhirendra Pratap, Ishan Joshi, and Jaytrilok Choudhary. "Survey of GPU based sorting algorithms." *International Journal of Parallel Programming* 46.6 (2018): 1017-1034.

# Parallel Sorting

- We will discuss Odd-Even Sort
- Source: <https://developer.nvidia.com/gpugems/gpugems2/part-vi-simulation-and-numerical-algorithms/chapter-46-improved-gpu-sorting>
- (Note the page uses an old GPU programming model)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively
- Two types of iterations
  - Odd: Compare pairs starting at odd index – (1,2); (3,4); (5,6)
  - Even: Compare pairs starting at even index – (0,1); (2,3); (4,5)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

Input Array

80	●	??
70	●	??
60	●	??
50	●	??
40	●	??
30	●	??
20	●	??
10	●	??

Iteration 1  
(even)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

Input Array

80	70
70	80
60	50
50	60
40	30
30	40
20	10
10	20

Iteration 1  
(even)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

Input Array

80	●	70	●	??
70	●	80	●	??
60	●	50	●	??
50	●	60	●	??
40	●	30	●	??
30	●	40	●	??
20	●	10	●	??
10	●	20	●	??

Iteration 2  
(odd)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

Input Array

80	70	70	70
70	80	50	
60	50	80	
50	60	30	
40	30	60	
30	40	10	
20	10	40	
10	20	20	

Iteration 2  
(odd)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

Input Array

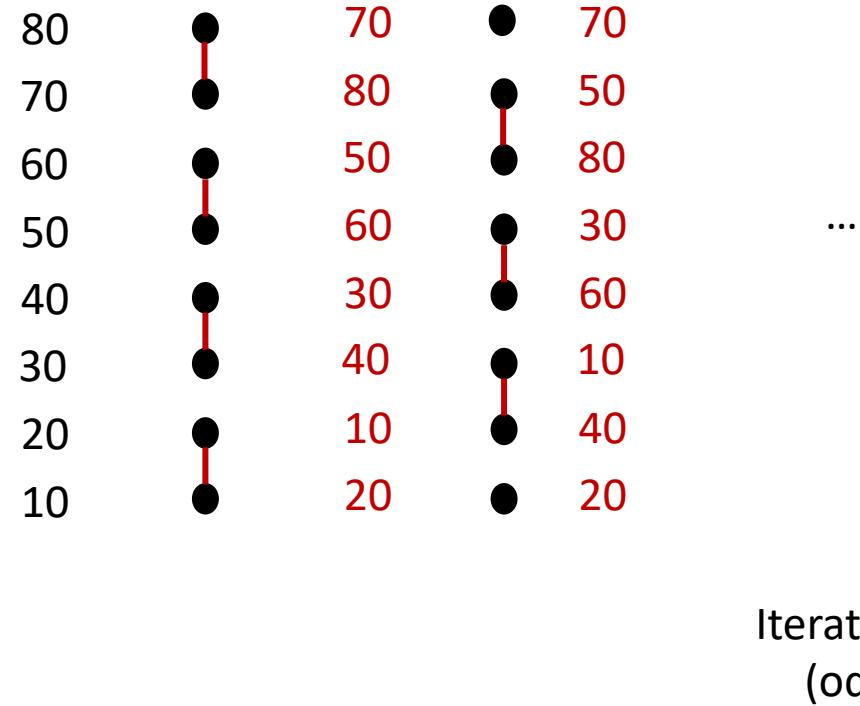
80	70	70	70
70	80	50	
60	50	80	
50	60	30	...
40	30	60	
30	40	10	
20	10	40	
10	20	20	

Iteration 3  
(even)

# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

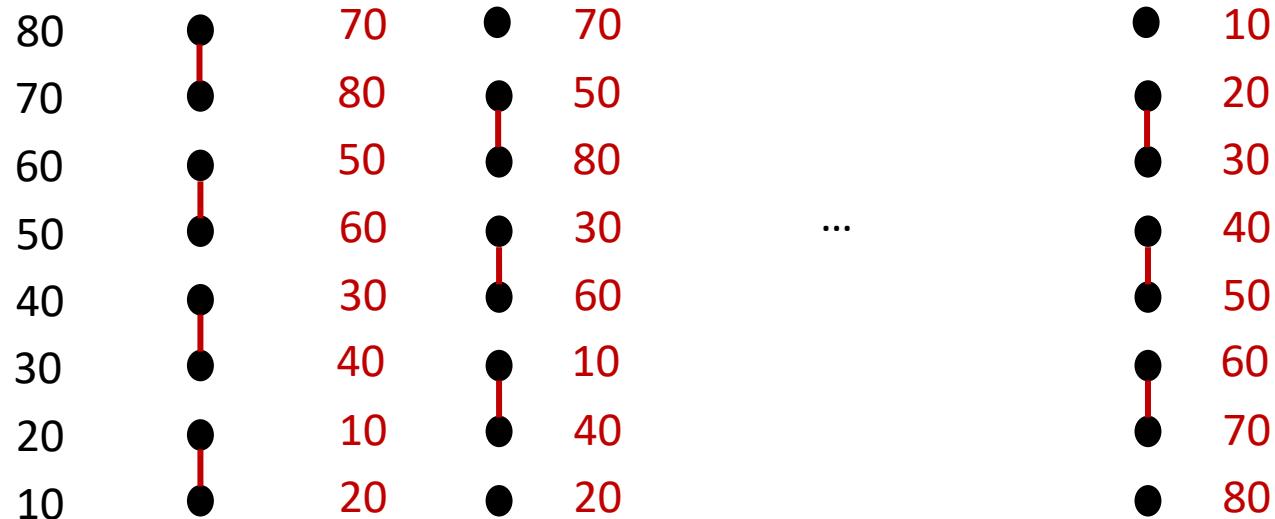
Input Array



# Parallel Sorting

- Odd-Even Sort – Similar to bubble sort,
  - In each iteration, move small keys upwards (and large key downwards) iteratively

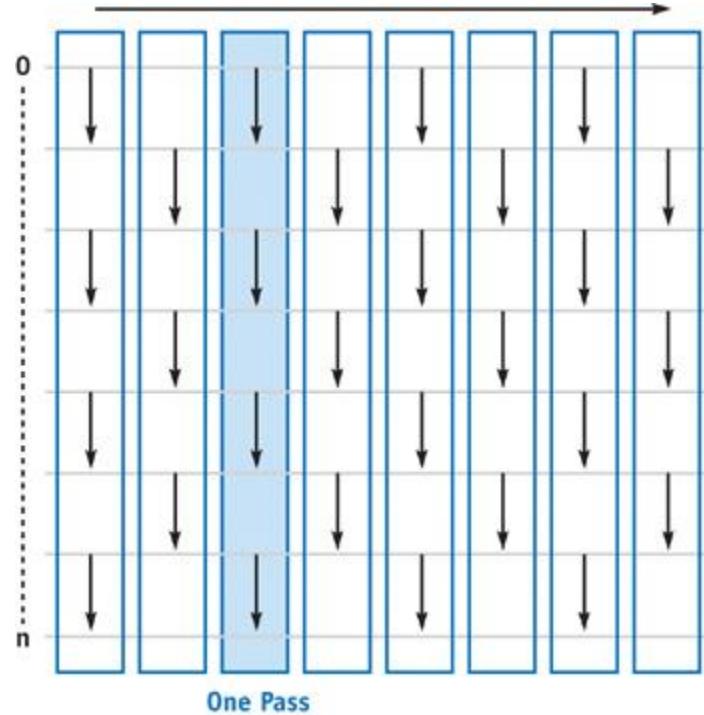
Input Array



Iteration 8  
(odd)

Ungraded HW Assignment: Can you  
simulate the steps of all 8 iterations

# Parallel Sorting



Entire Sorting Network – Arrow can be used to represent the direction of comparison

<https://developer.nvidia.com/gpugems/gpugems2/part-vi-simulation-and-numerical-algorithms/chapter-46-improved-gpu-sorting>

# Parallel Sorting

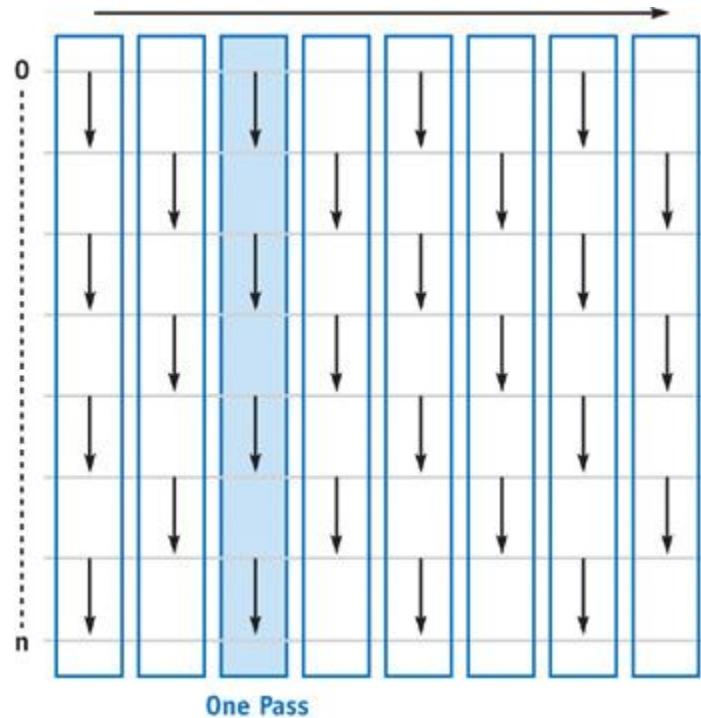
- Odd-Even Sort
- For an array of  $n$  elements, terminates in at most  $2n$  iterations
- Serial Complexity: ??

# Parallel Sorting

- Odd-Even Sort
- For an array of  $n$  elements, terminates in at most  $2n$  iterations
- Serial Complexity:  $O(n^2)$

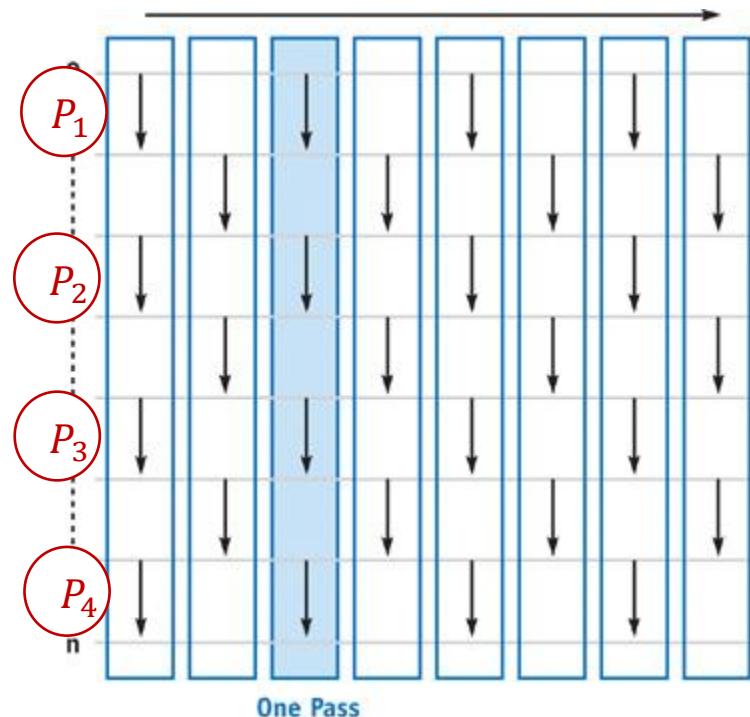
# Parallel Sorting

- How do we parallelize?



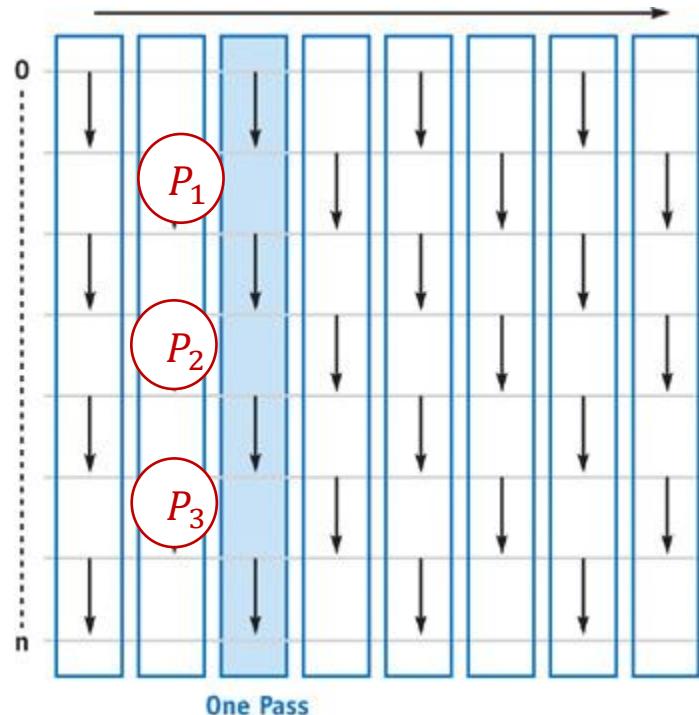
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



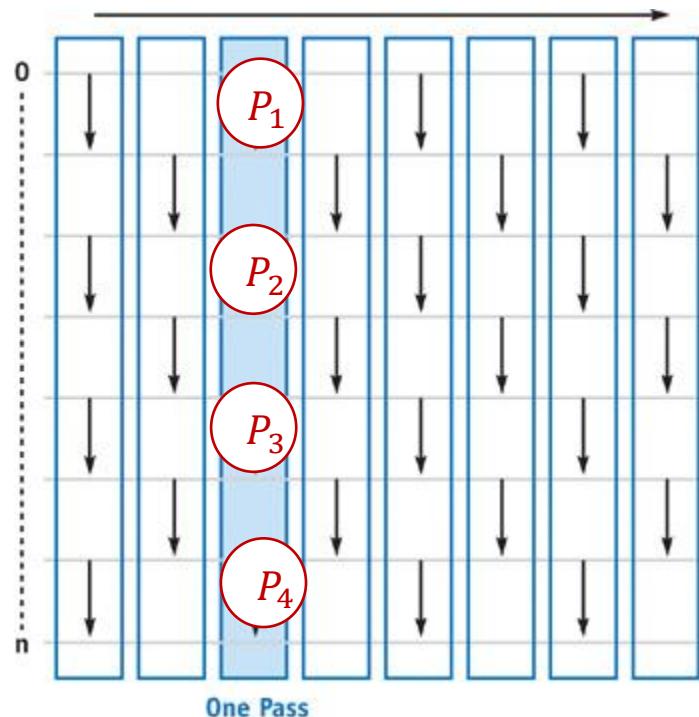
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



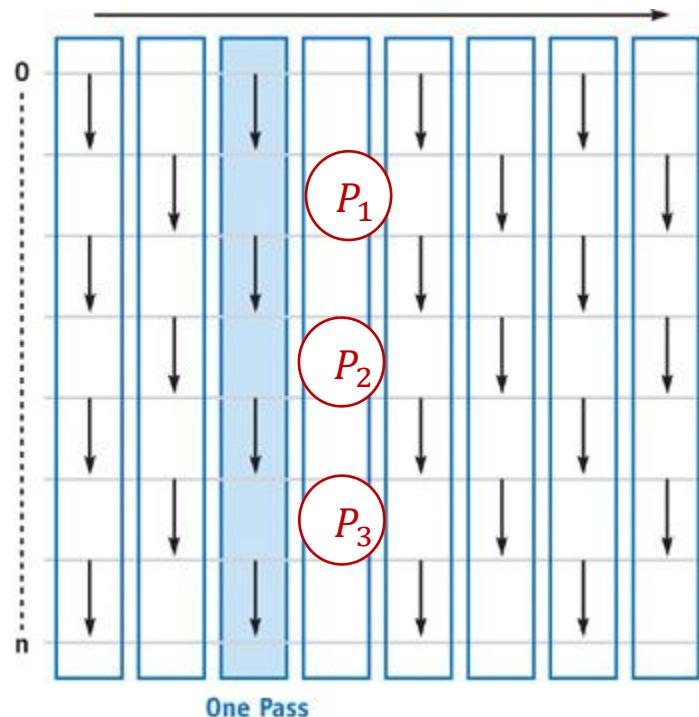
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



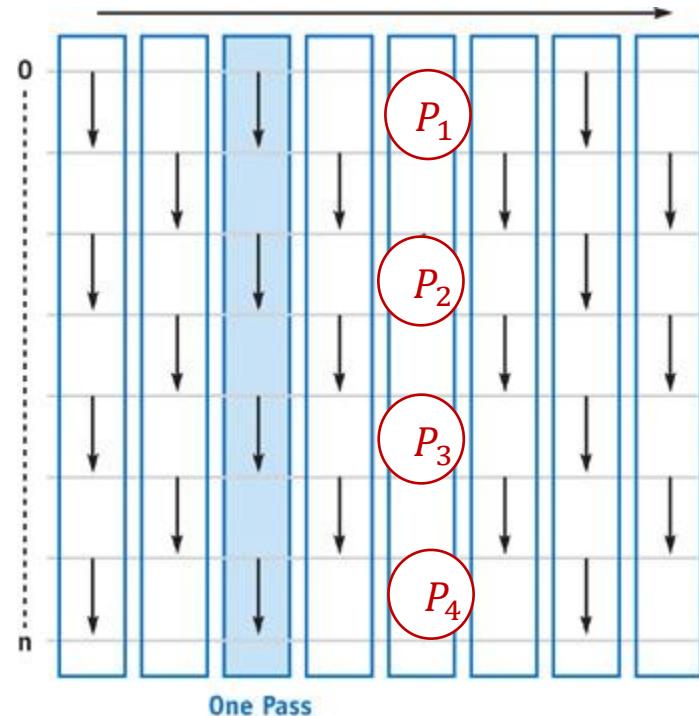
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



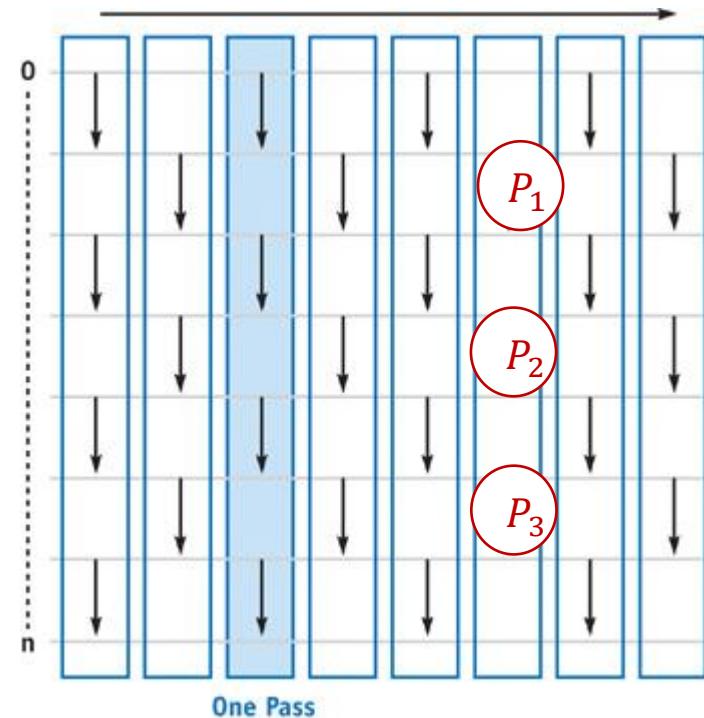
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



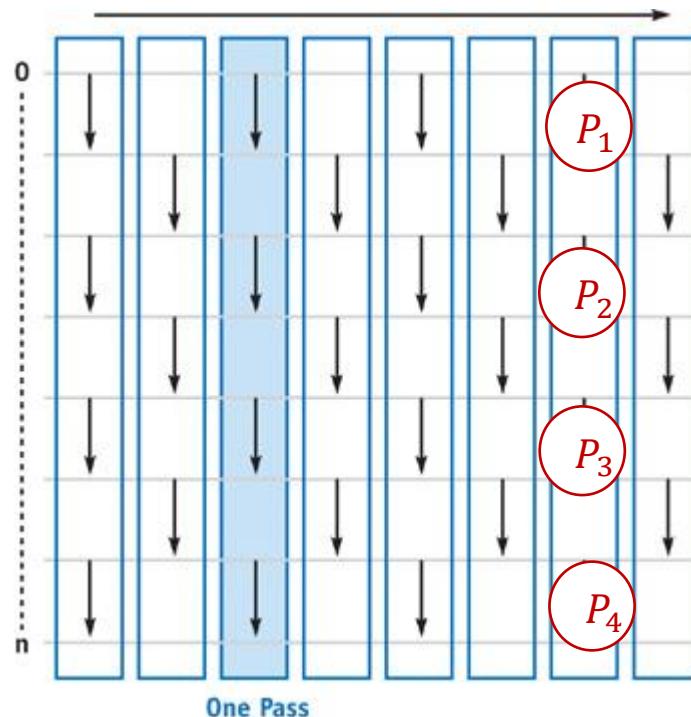
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



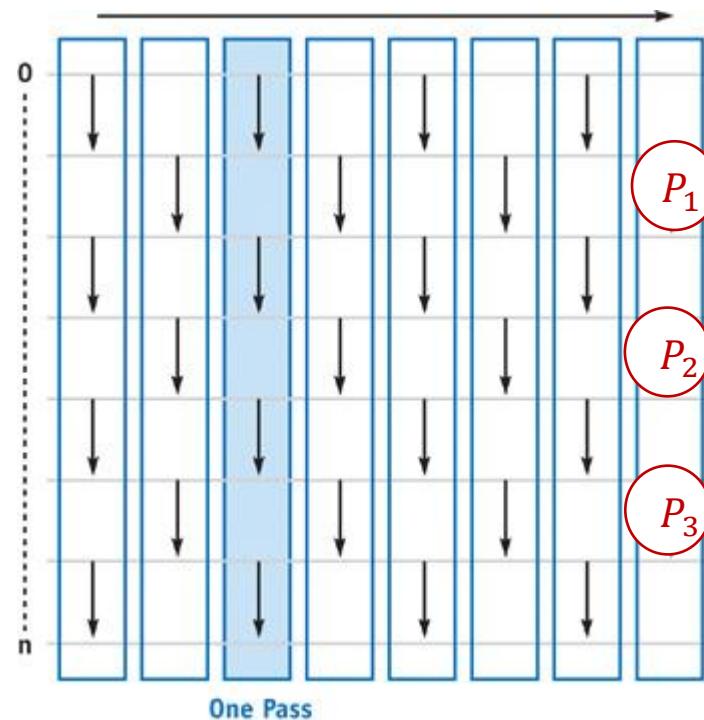
# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



# Parallel Sorting

- How do we parallelize? In each iteration, use a maximum of  $\frac{n}{2}$  processors in parallel to do the compare and swap operation
  - Iterations still need to occur sequentially



# Parallel Sorting

- Can you calculate speedup. Efficiency, cost, and whether this parallel algorithm is scalable or cost optimal or not? A question on WA 1
- Can you create a GPU code of the algorithm?
- Assume  $N = 2 \times S \times p$
- Where do you see a challenge? (We will explore this more in WA 2)

# Parallel Sorting

- Even for comparison based sort, algorithms with better cost/work exist
- Check out these papers:
  - Bitton, Dina, et al. "A taxonomy of parallel sorting." *ACM Computing Surveys (CSUR)* 16.3 (1984): 287-318.
  - Singh, Dhirendra Pratap, Ishan Joshi, and Jaytrilok Choudhary. "Survey of GPU based sorting algorithms." *International Journal of Parallel Programming* 46.6 (2018): 1017-1034.
- You are not expected to know these algorithms for exams
  - But I may give you an algorithm and ask you to do analysis or convert into a GPU code.

# Outline

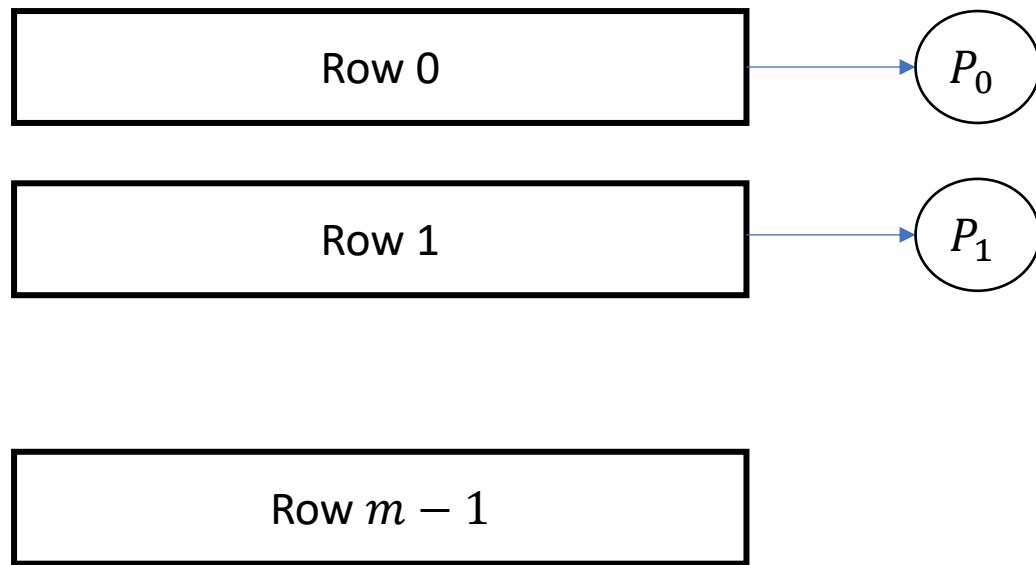
- Parallel Program Analysis (Quick Review)
- Data Parallel Algorithms
  - Sorting
- Task Parallelism

# Writing Parallel Programs

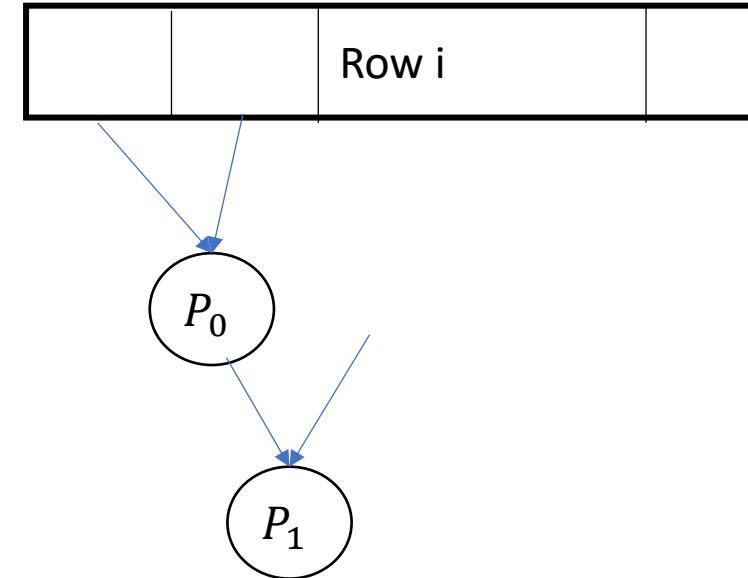
- Key Idea: Need to decompose an algorithm into chunks of independent tasks.
- Such that it can get best performance
- Two approaches
  - Data Parallelism (SPMD): decompose algorithm into tasks that perform “same” work on different data
  - Task Parallelism: decompose algorithm into tasks that perform “different” work

# Task vs Data Parallelism

- Consider Matrix Vector Multiplication



Data Parallelism



Task Parallelism

Note: Data Parallelism is a  
special case of task parallelism

# Task Parallelism Techniques

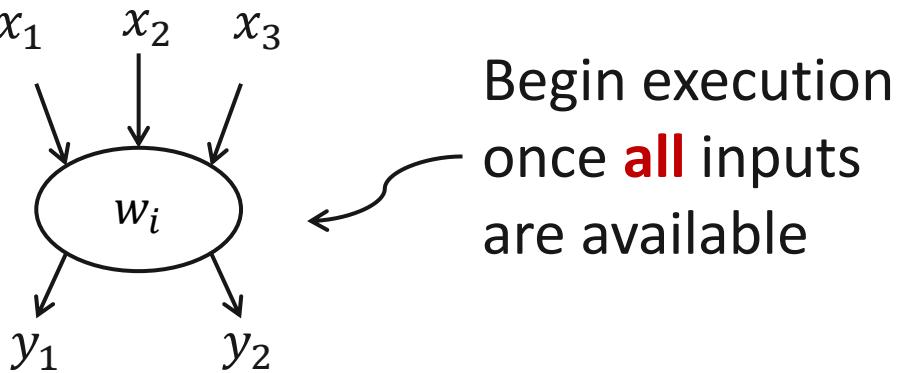
- How to write parallel program using task parallelism paradigm?
- Can be tricky, needs careful analysis of tasks (operations to be performed) and their dependencies
- Careful selection of the granularity of tasks
  - Lowest granularity – operation level (e.g., Task parallel MV)
  - Coarser granularity – example, matrix row level in MV
  - Coarser granularity leads to reduced overheads, but may lose on opportunities of parallelism

# Tasks and Dependencies

Computation = decompose into tasks

Task = Set of instructions (program segment)

## Inputs & outputs



Task size?

= weight of the node  
(e.g., # of instructions executed)

Fine grain

Coarse grain

Note: tasks need not be of the same size

# Task Dependency Graph

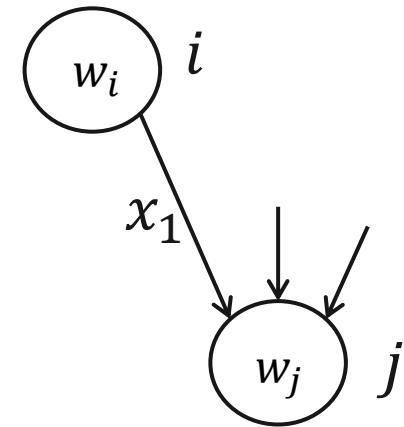
## Directed Acyclic Graph

Task<sub>j</sub> cannot start until Task<sub>i</sub> completes

Data  $x_1$ : output of Task<sub>i</sub>

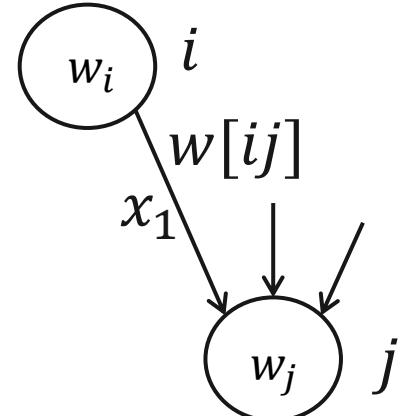
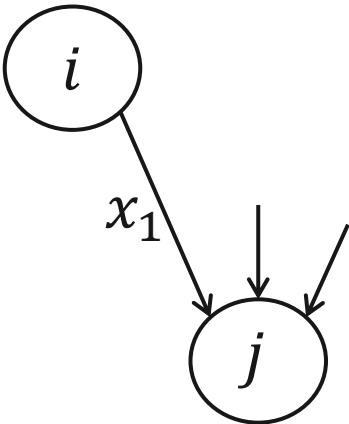
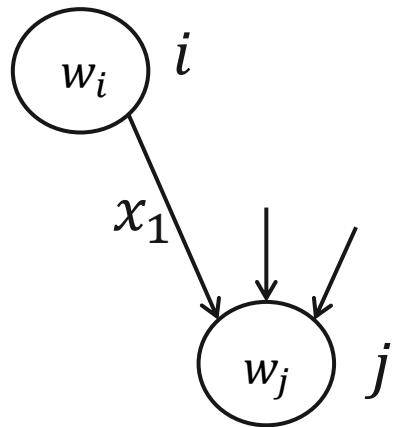
Input to Task<sub>j</sub>

Weight of a node: task size



Task dependency graph need not be connected  
Graph is acyclic

# Task Dependency Graph



# Example (1)

## Code

$A[2] = A[0] + 1$

$B[0] = A[2] + 1$



## Instructions

Load  $R0 \leftarrow A[0]$

Add  $R1 \leftarrow R0 + 1$

Add  $R2 \leftarrow R1 + 1$

Store  $A[2] \leftarrow R1$

Store  $B[0] \leftarrow R2$

## Tasks

$T_0$

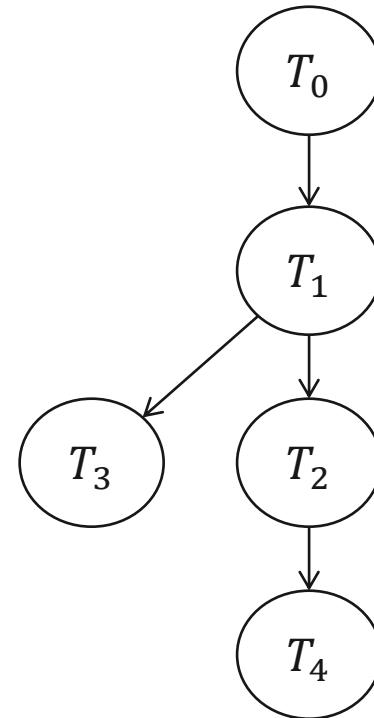
$T_1$

$T_2$

$T_3$

$T_4$

## Task Dependency Graph



# Example (2)

## Matrix Vector Multiplication

$$C[i] = \sum A[i][k] * B[k]$$

Task<sub>i</sub> – Compute C[i]

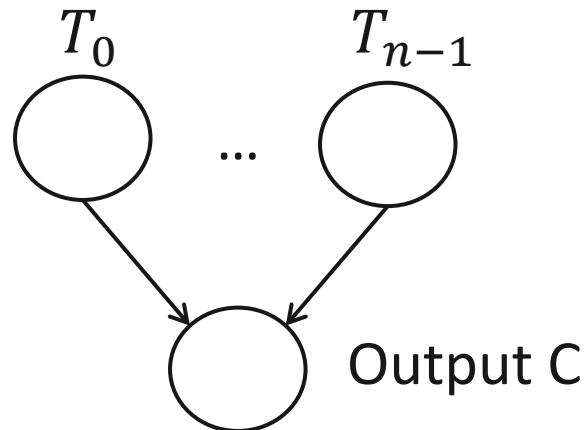
$n \times n$

Parallel for i = 1 to n

Compute C[i]

End

Output C



# Maximum Degree of Concurrency (1)

Given a task dependency graph,  
maximum number of tasks that can be executed  
concurrently

Maximum Parallelism Achievable at any Point in  
the code -> Decide number of processors

# Maximum Degree of Concurrency (2)

Example: **level by level ordering (topological sort)**

Task dependency graph

Order tasks level by level

- Any level  $i$  task has dependency with some task in level  $i - 1$  (and possibly with other lower levels) but no dependency with level  $i$
- All tasks in any level  $i$  are independent

Execute level  $i$  tasks (in parallel), and then level  $i + 1$  tasks

# Maximum Degree of Concurrency (3)

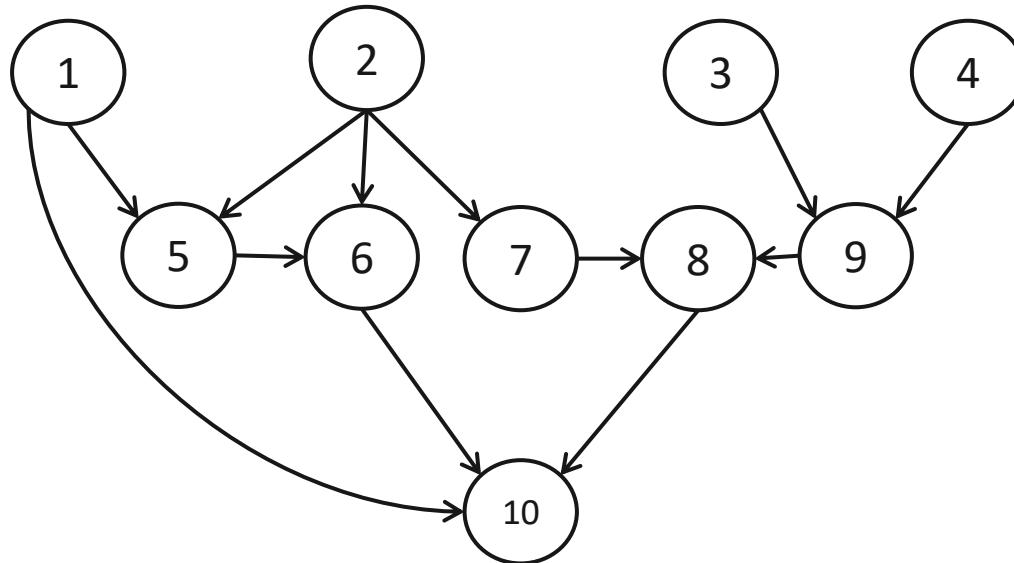
Example: **level by level ordering (cont.)**

Find the number of tasks in each level

Take maximum over all levels = maximum degree of concurrency  
(if scheduled using level by level ordering)

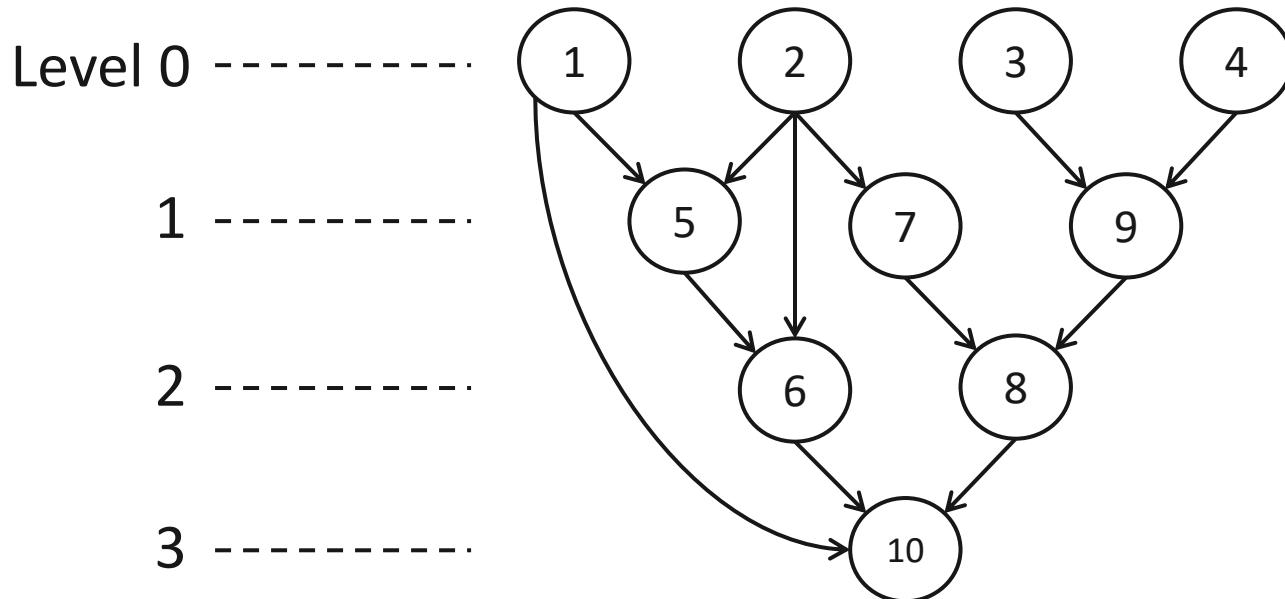
# Maximum Degree of Concurrency (4)

Example



# Maximum Degree of Concurrency (5)

Example (cont.)



Maximum degree of concurrency = 4

# Critical Path (1)

Dependency graph

Start nodes (indeg = 0)

Finish nodes (outdeg = 0)

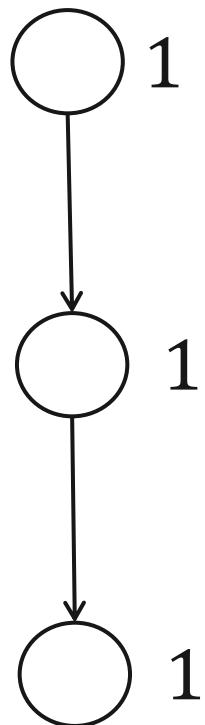
Critical path = A longest path from a start node  
to a finish node (# of edges)

Critical path length = Sum of the task weights of the  
nodes along the critical path

# Critical Path (2)

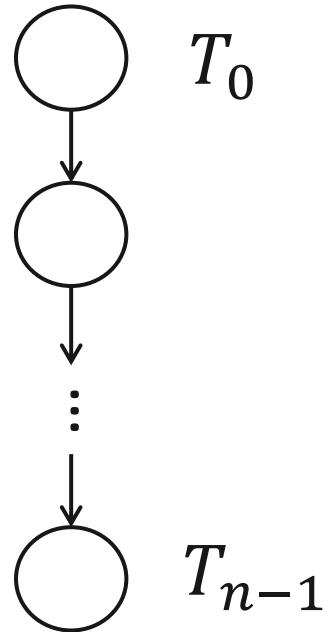
Note: Critical path length considers critical path only (long paths)

For a given number of tasks,

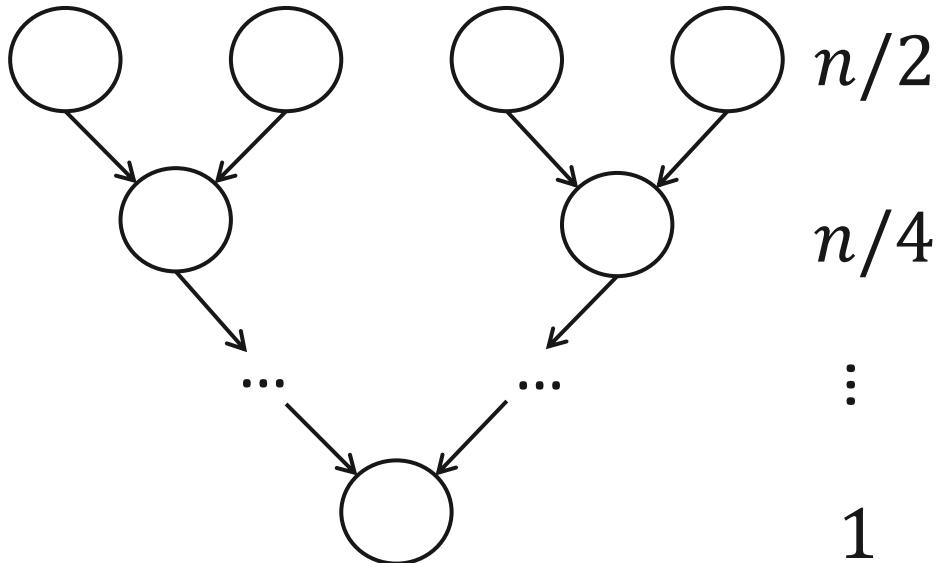


Longer critical path  $\Rightarrow$  Longer execution time  
(may also mean less concurrency)

# Critical Path (3)

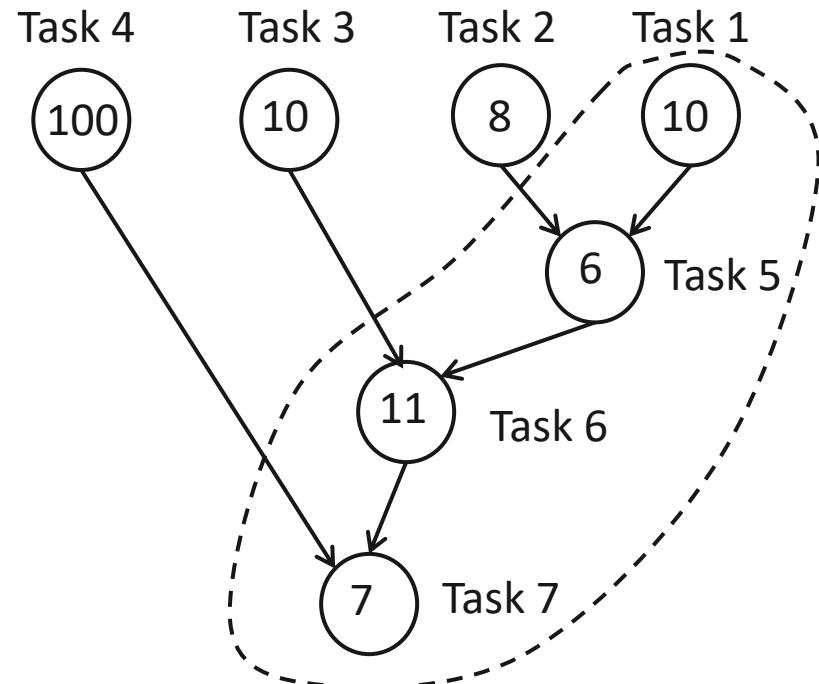


Total no. of tasks =  $n$   
Critical path length =  $n$



Total no. of tasks =  $n - 1$   
Critical path length =  $\log_2 n$

# Critical Path (4)



Critical path length =  $10+6+11+7 = 34$

# Task dependency graph to Parallel Program (1)

Given a task dependency graph

Assume weight of each node = 1

Maximum degree of concurrency =  $c$

Critical path length =  $l$

# Task dependency graph to Parallel Program (2)

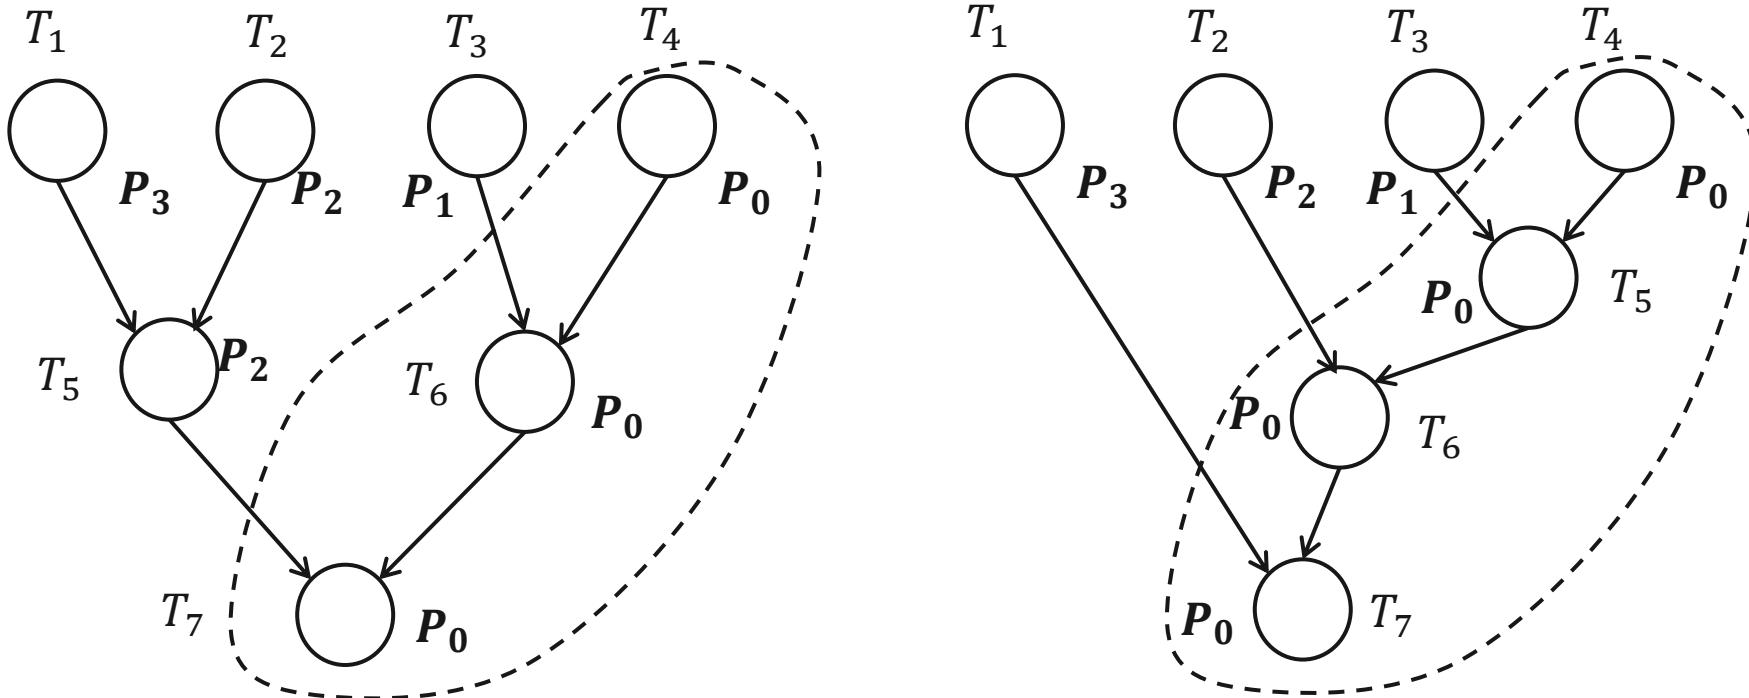
## Idea

DAG → Organize into levels  $0, 1, \dots, l$  (level by level ordering)  
 $(l + 1)$  levels

Execute level by level,      0 to  $l$

Total number of processors needed  $\leq c$

# Mapping Tasks to Processes



# Next Class

- 9/18 Lecture 8 – Systolic Array Architecture – Modeling, Matrix Multiplication;

# Thank You

- Questions?
- Email: [sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)