

# CSDS 451: Designing High Performant Systems for AI

Lecture 9

9/23/2025

Sanmukh Kuppannagari

[sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)

<https://sanmukh.research.st/>

Case Western Reserve University

# Outline

- Accelerating Convolutional Neural Network Models: Basics

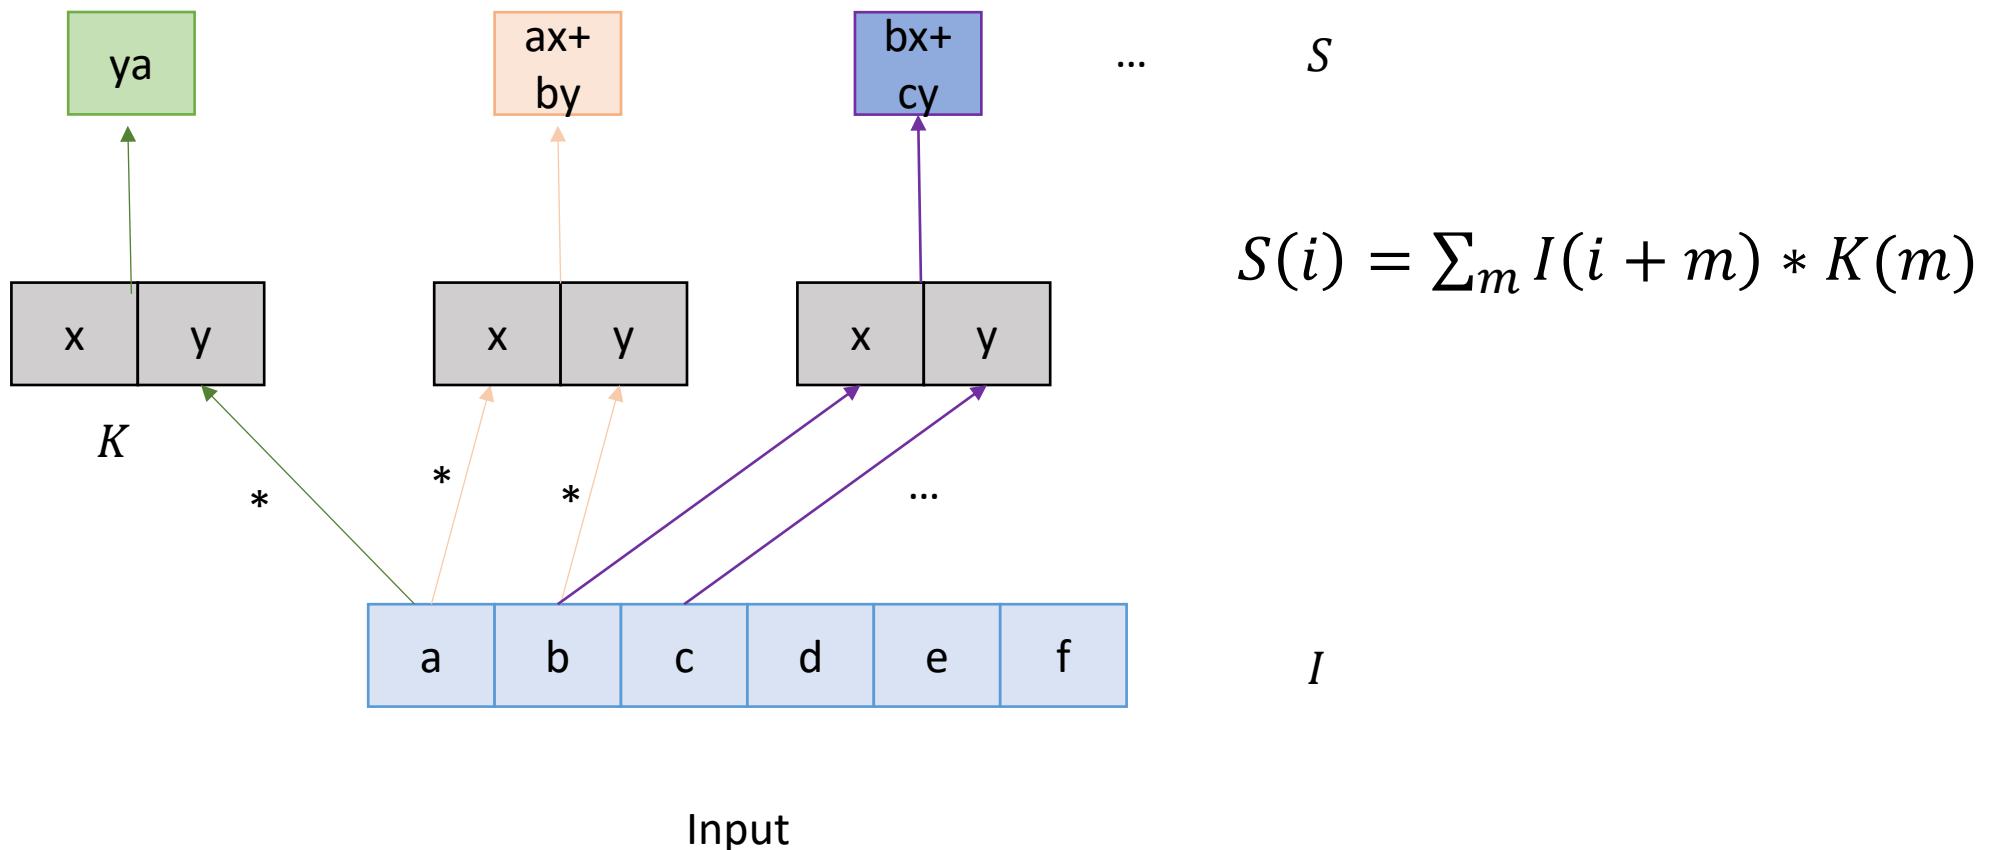
# Reading Materials

- Deep Learning, by Ian Goodfellow, Yoshua Bengio, Aaron Courville
  - Available Online: <https://www.deeplearningbook.org/>
  - CNNs: Chapter 9
- A Guide to Convolution Arithmetic for Deep Learning
  - Available Online: <https://arxiv.org/pdf/1603.07285.pdf>

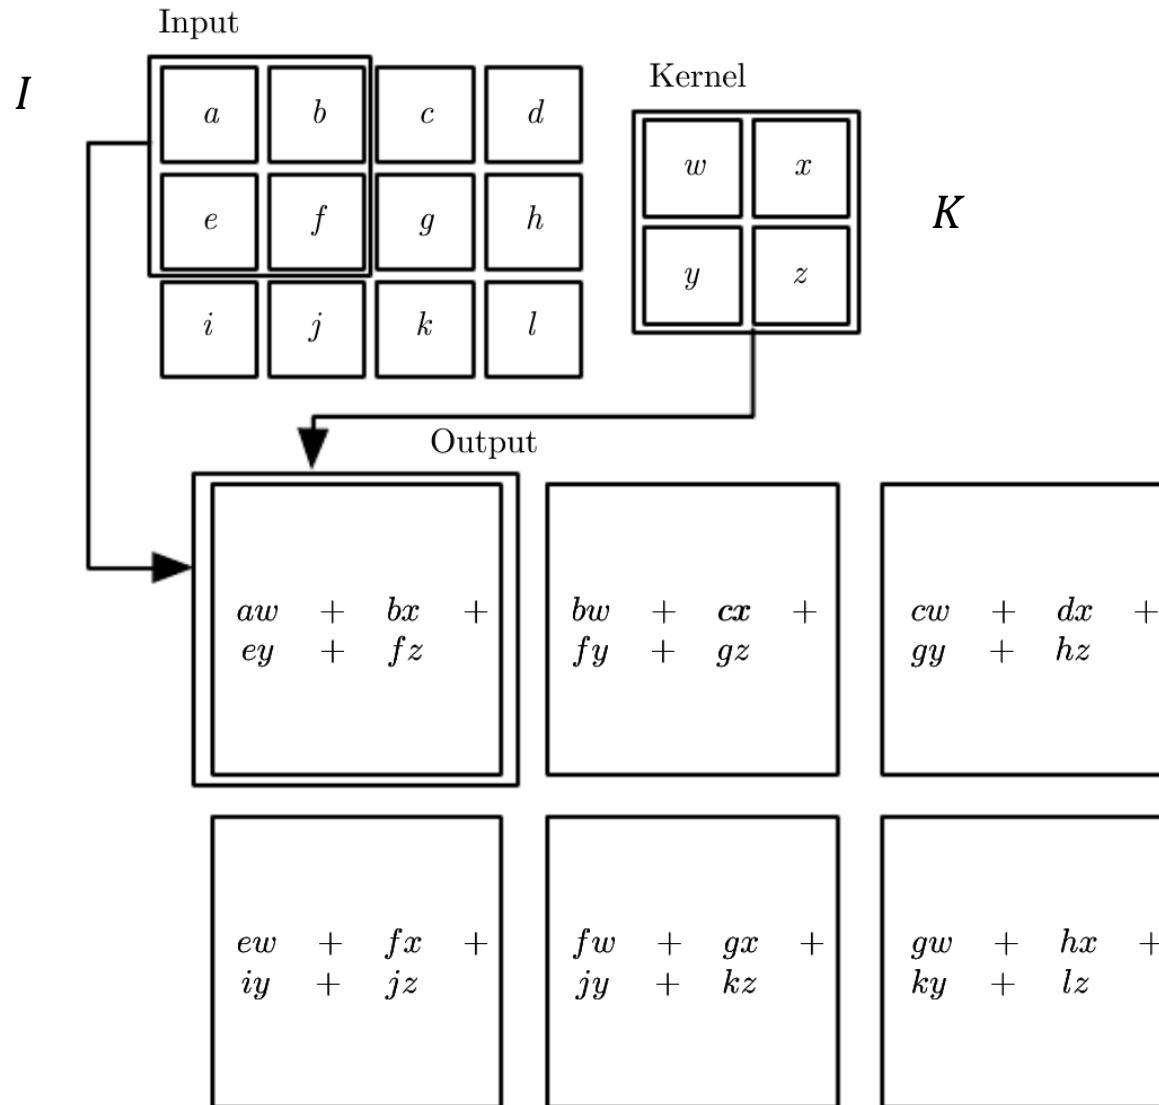
# Convolutional Neural Networks

- Neural Network specialized for a grid-like topology
  - 1D: timeseries data, e.g., (12 pm, 63), (1pm, 65), (2 pm, 68),...
  - 2D: Images
  - 3D: Images with multiple channels, e.g., RGB images
- Name derives from “Convolution” operations performed in the “Convolution” Layers

# 1D Convolution



# 2D Convolution



Source: Deep Learning Book

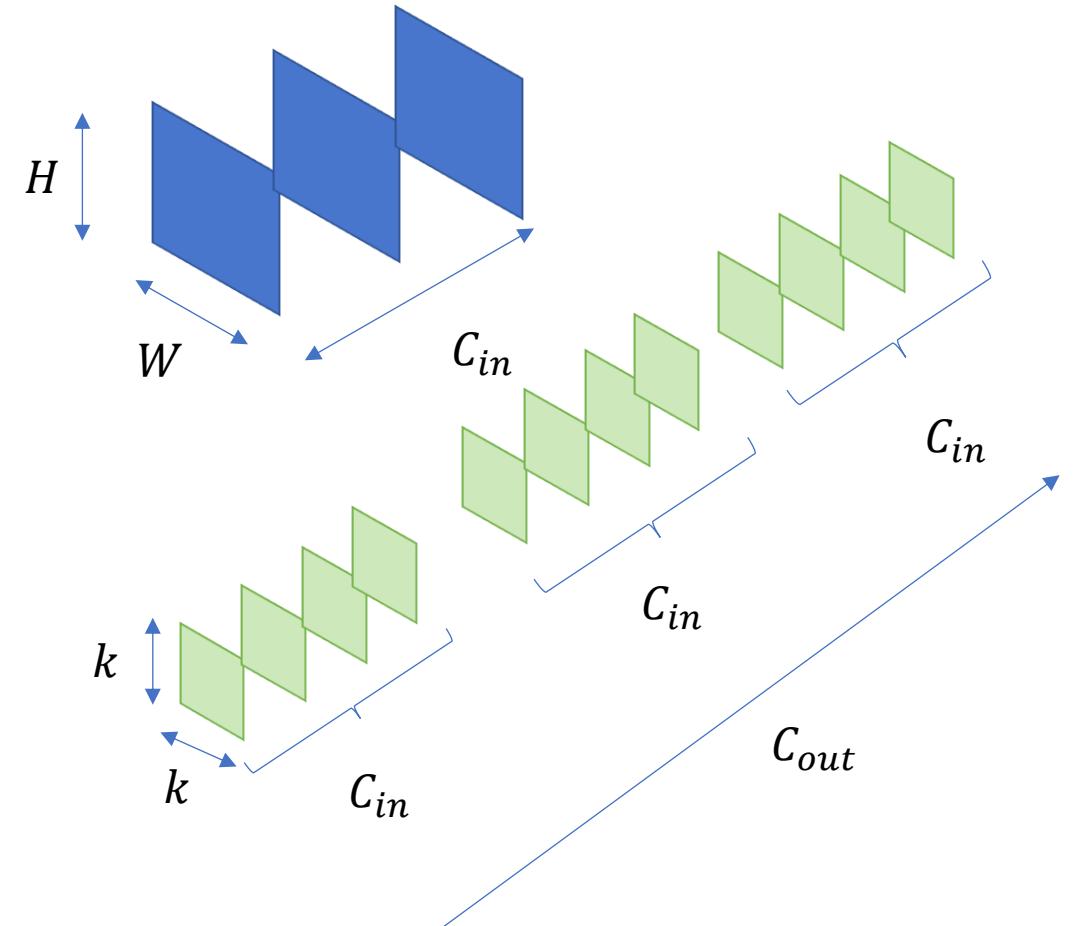
What happens for corner cases???

# Convolutional Neural Network

- A Neural Network with:
  - Convolution Layers
  - Activations
  - Pooling Layers
  - ...

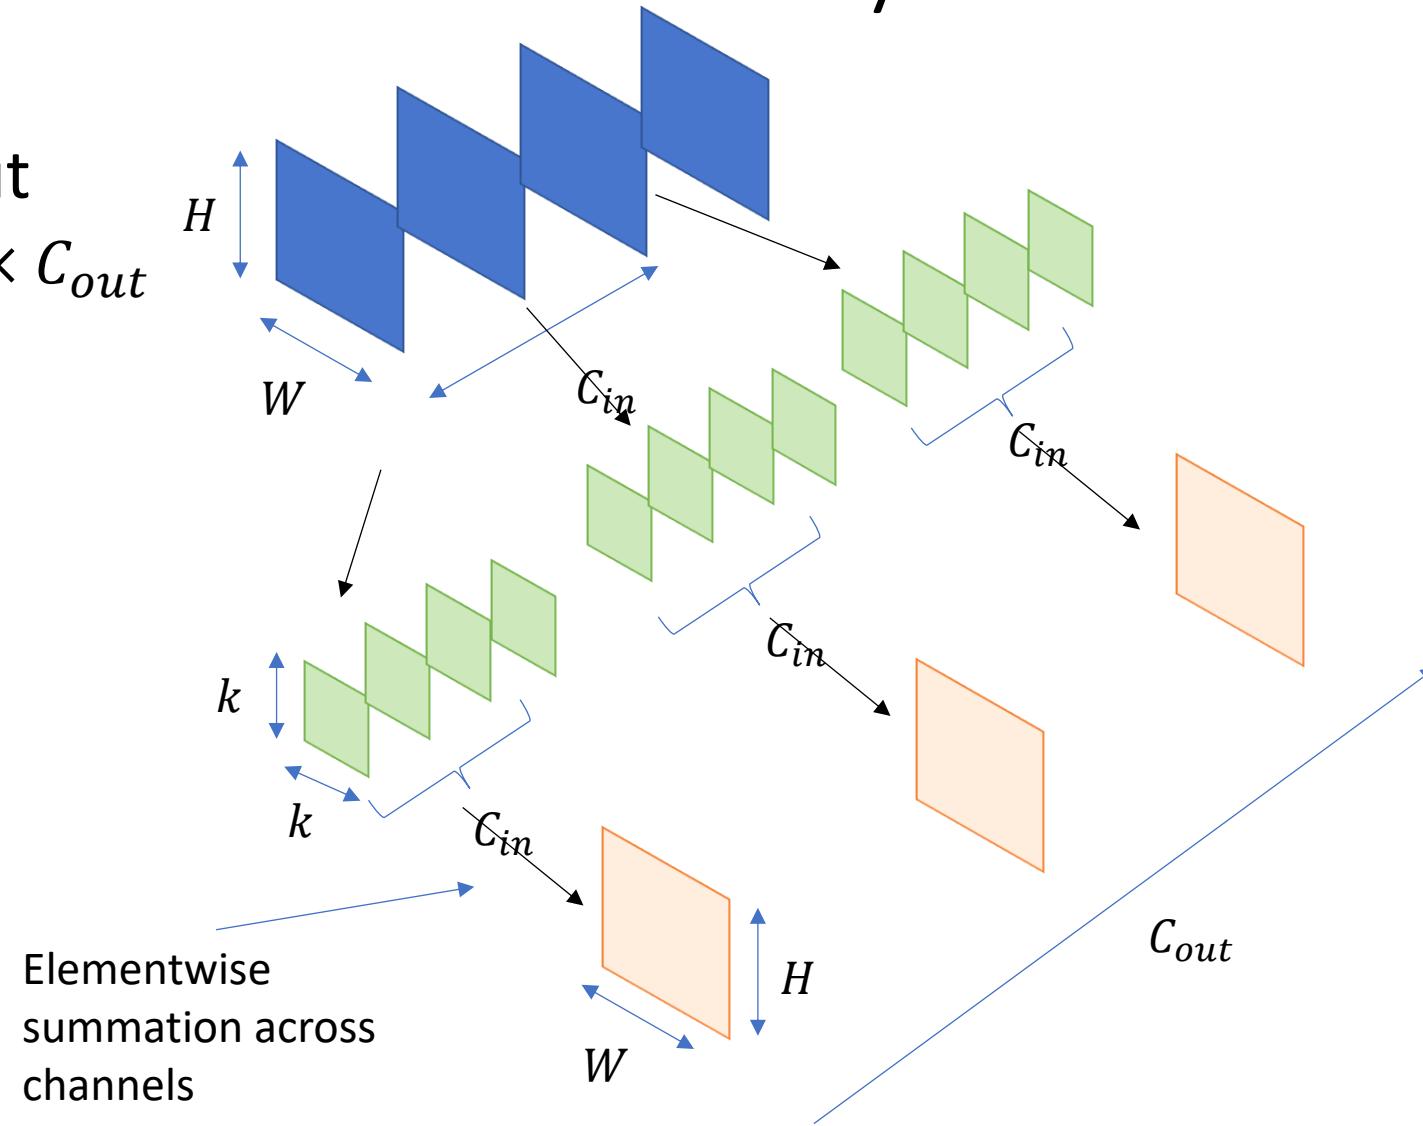
# Convolution Layer

- Given an Input Feature Map
  - Size:  $H \times W \times C_{in}$
  - $H$ : Height of the input feature map
  - $W$ : Width of the input feature map
  - $C_{in}$ : Number of Channels in the feature map
- $C_{in} \times C_{out}$  Kernels, each of size  $k \times k$



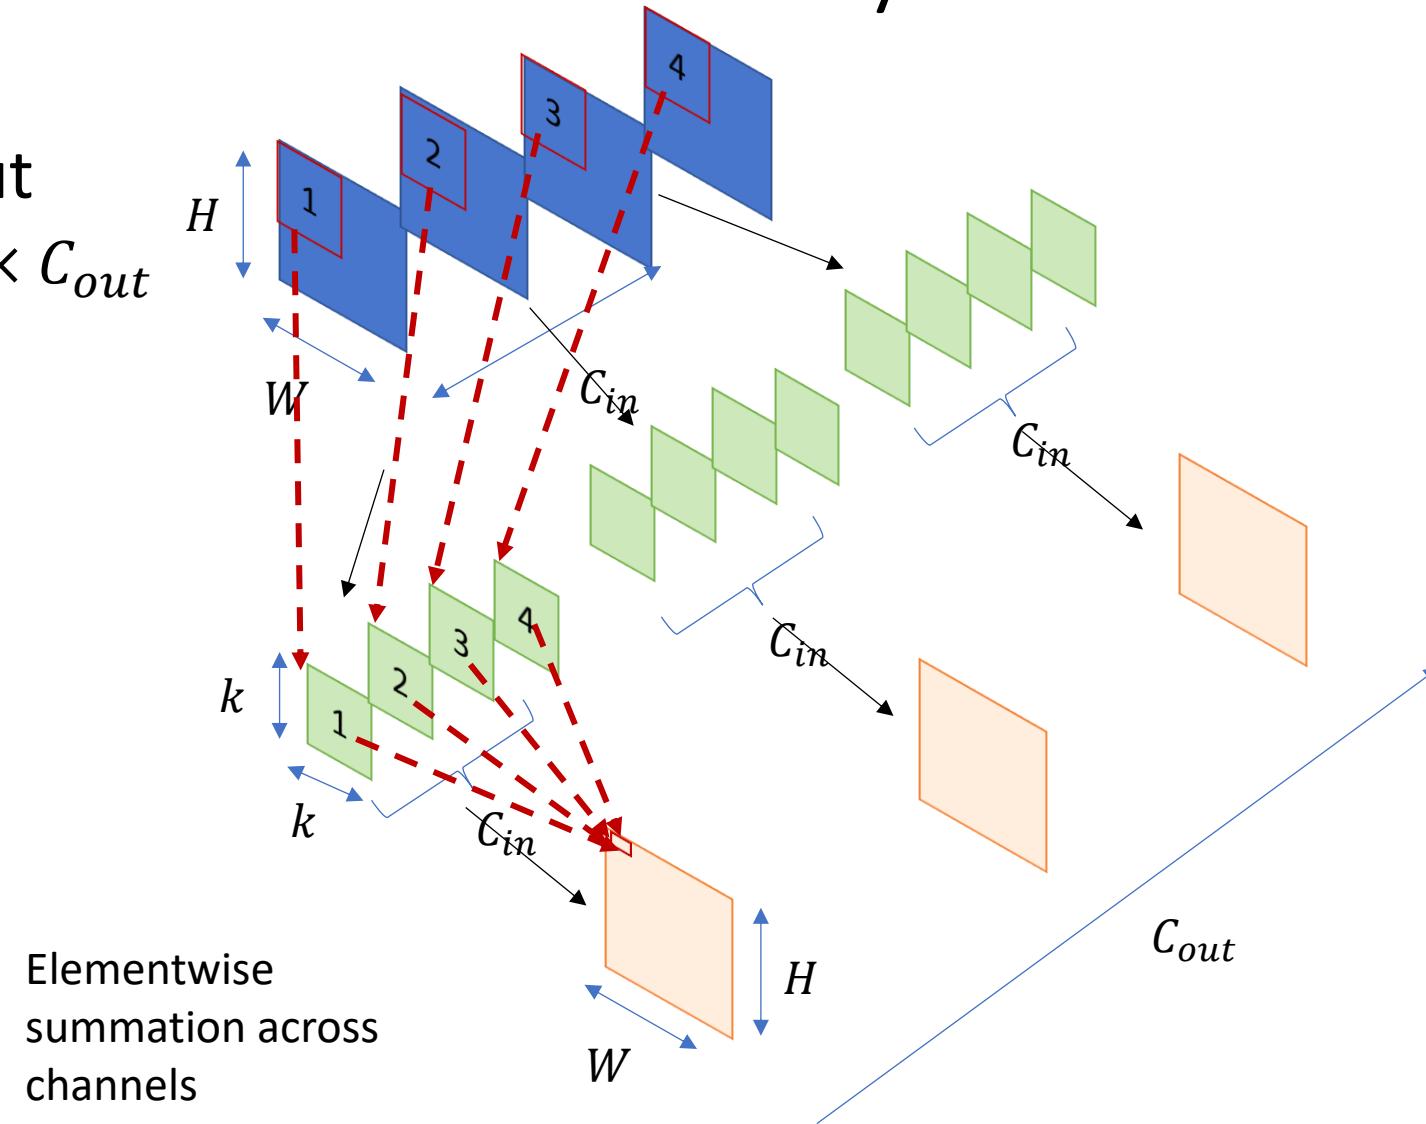
# Convolution Layer

- Produce Output
  - Size:  $H \times W \times C_{out}$



# Convolution Layer

- Produce Output
  - Size:  $H \times W \times C_{out}$



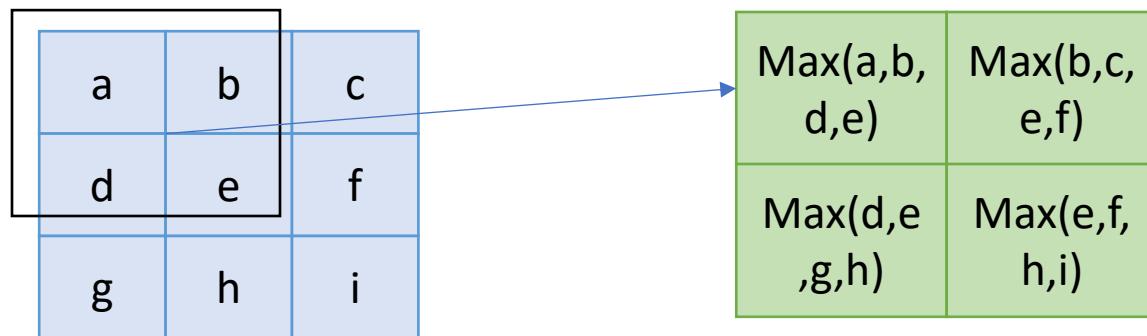
# Activation Layer



$\sigma$  : Activation Function, Sigmoid, ReLU, Softmax

# Pooling Layer

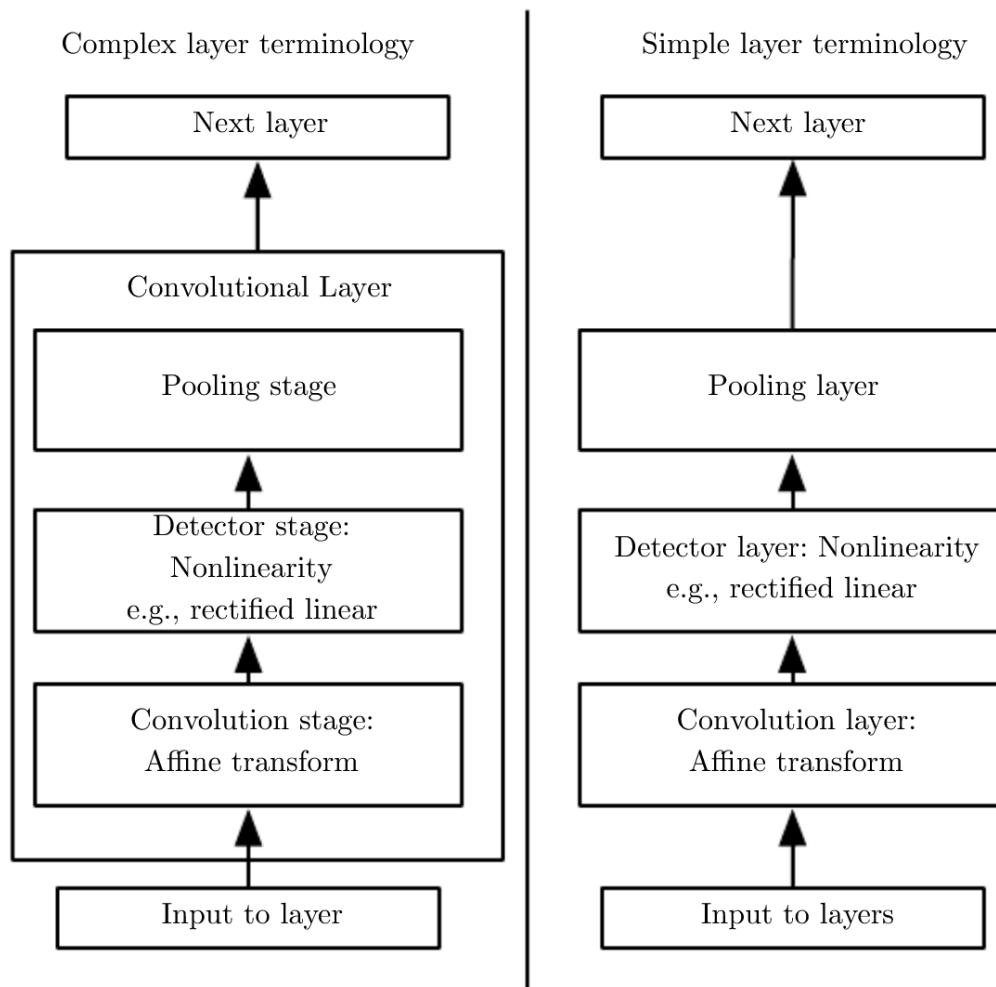
- Replaces the output of a grid with a summary of the grid



Max Pooling: Size  
2 X 2

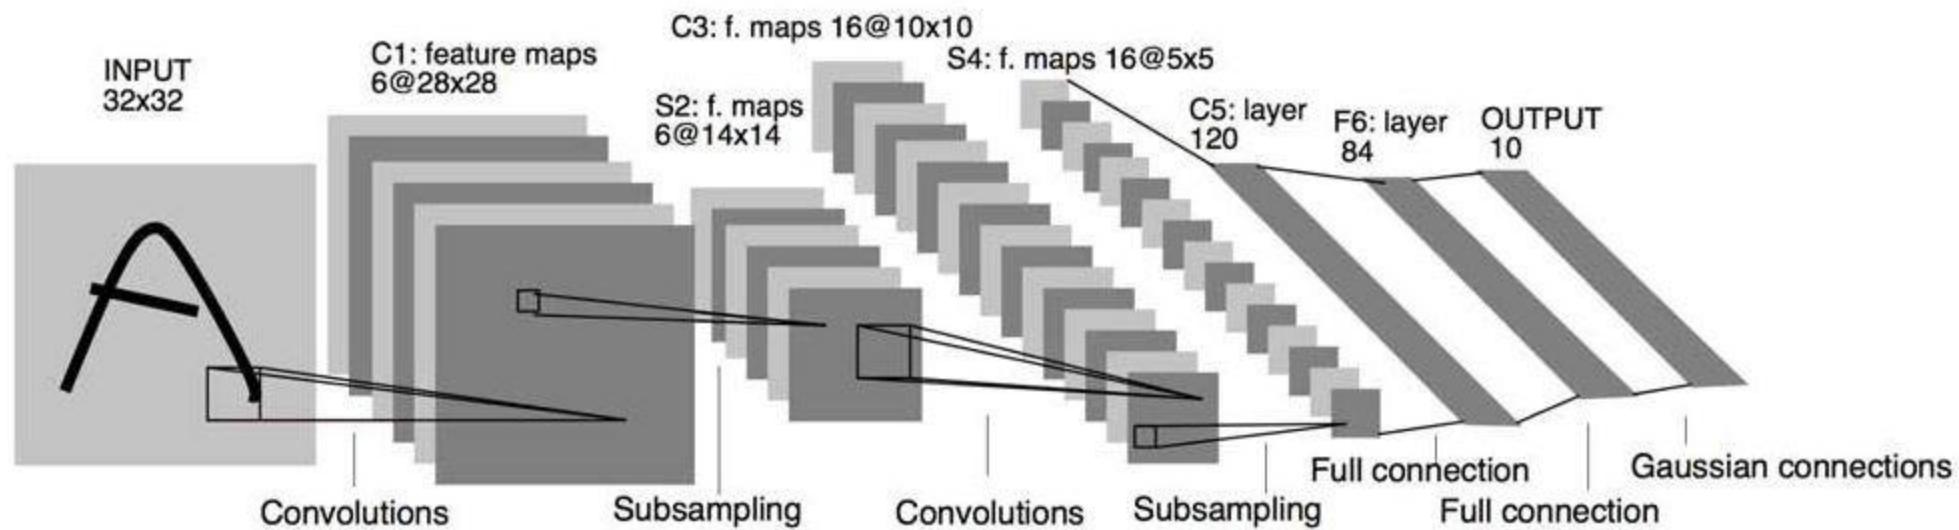
Pooling Types:  
Average Pooling  
Max Pooling

# Convolution Neural Network



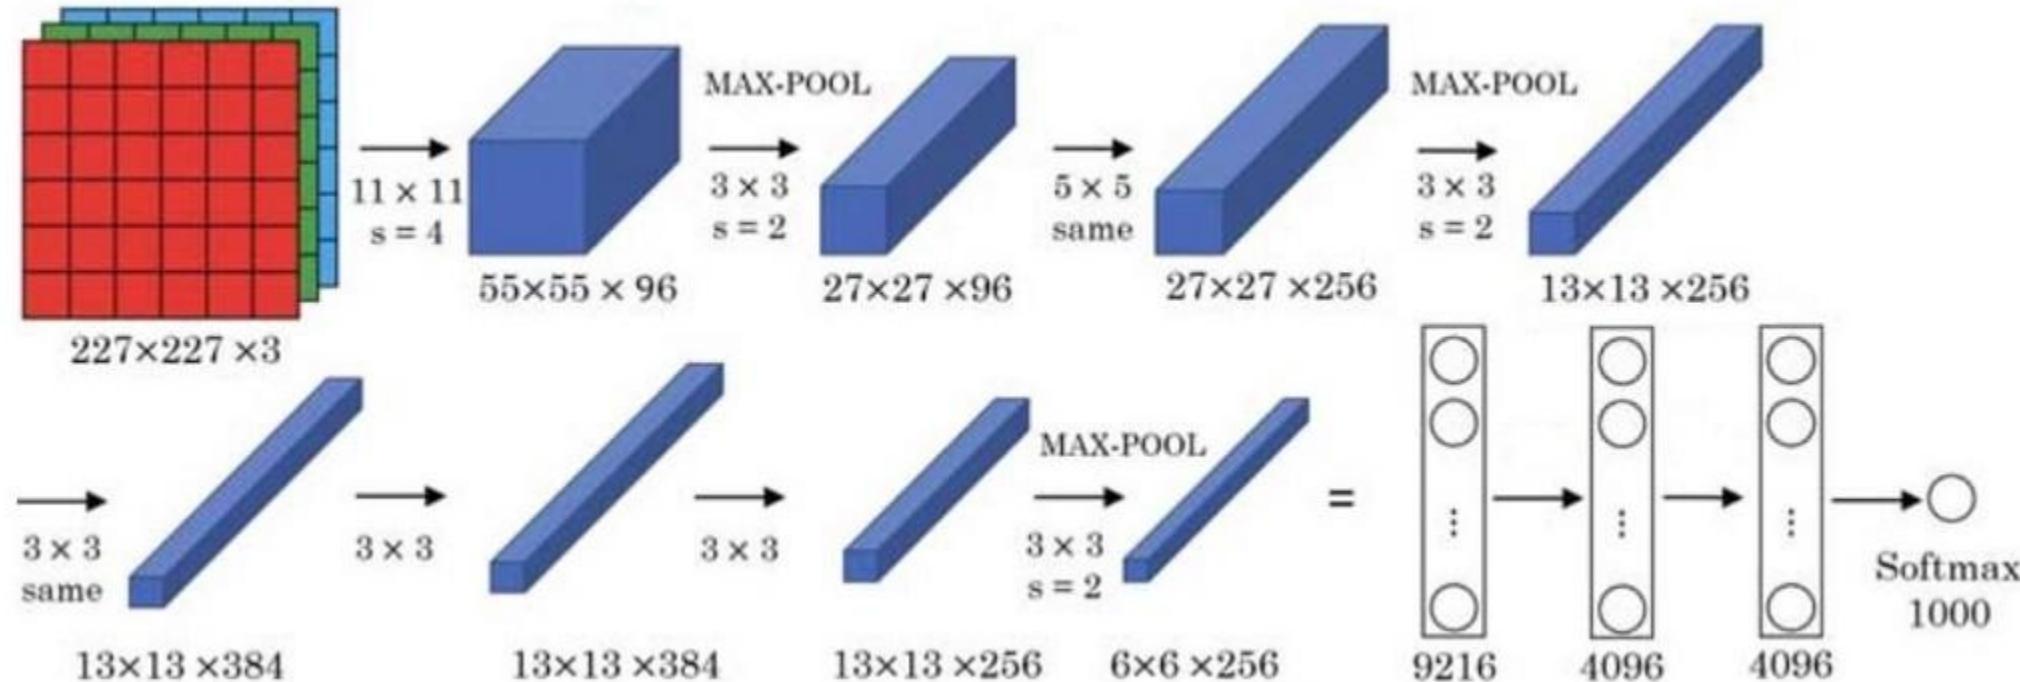
Source: Deep Learning Book

# Example CNN Networks



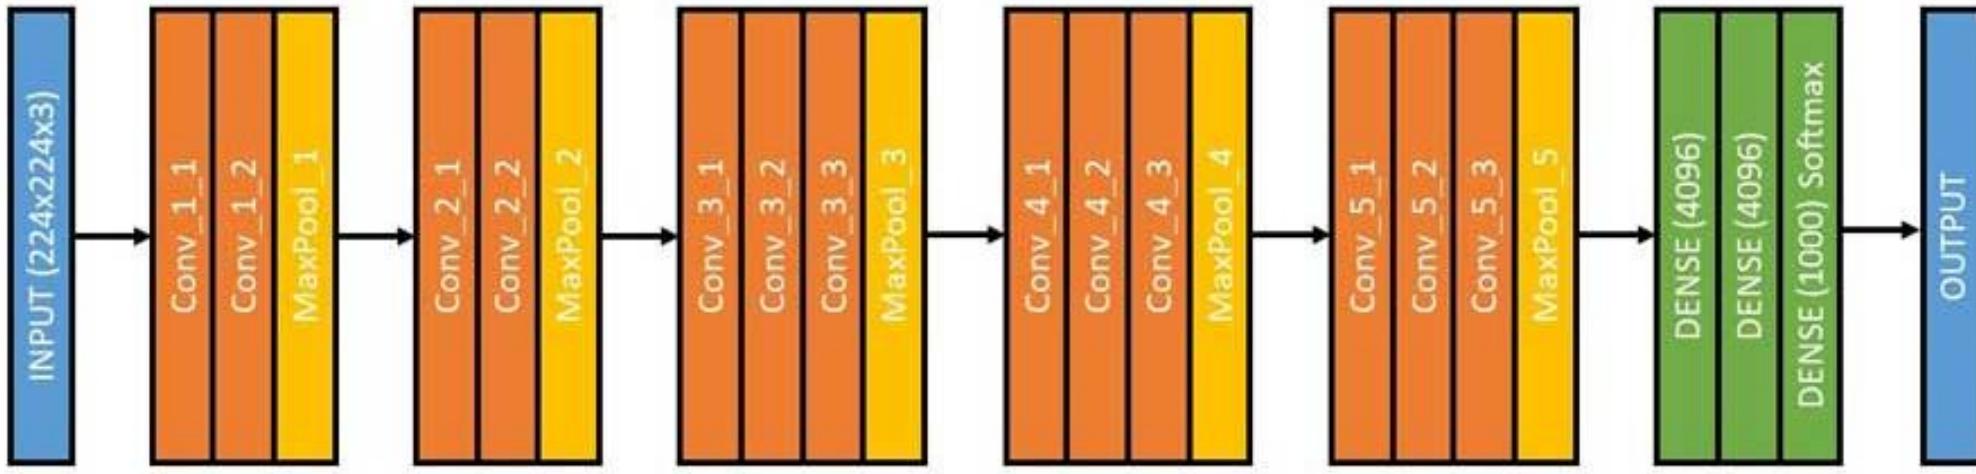
LeNet (1998): Gradient Based Learning Applied to Document Recognitions,  
<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>

# Example CNN Networks



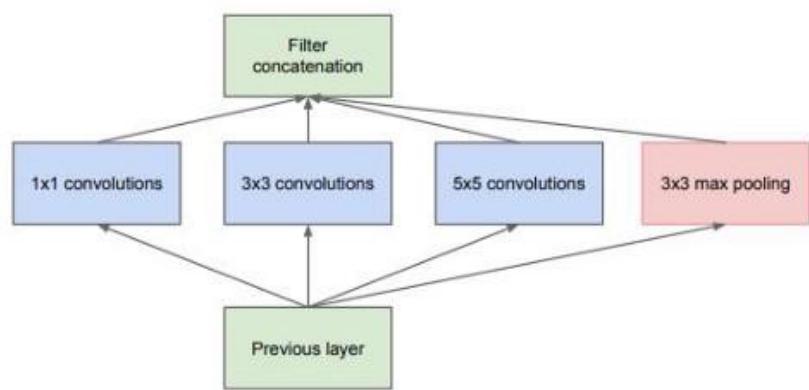
AlexNet (2012): ImageNet Classification with Deep Convolutional Neural Networks,  
<https://proceedings.neurips.cc/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf>

# Example CNN Networks

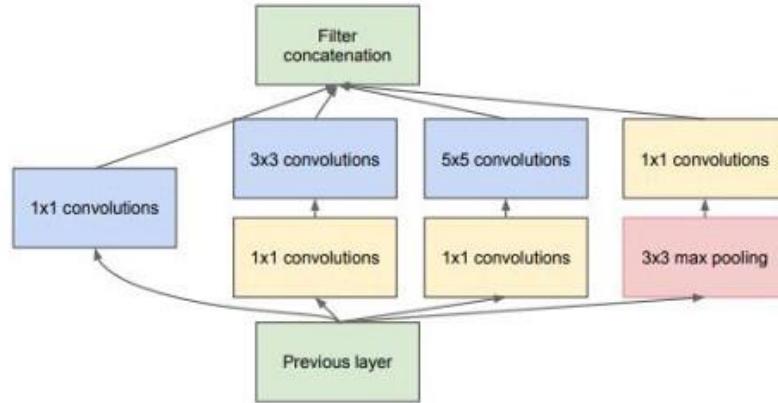


VGG16 (2014): VERY DEEP CONVOLUTIONAL NETWORKS  
FOR LARGE-SCALE IMAGE RECOGNITION,  
<https://arxiv.org/pdf/1409.1556.pdf>

# Example CNN Networks

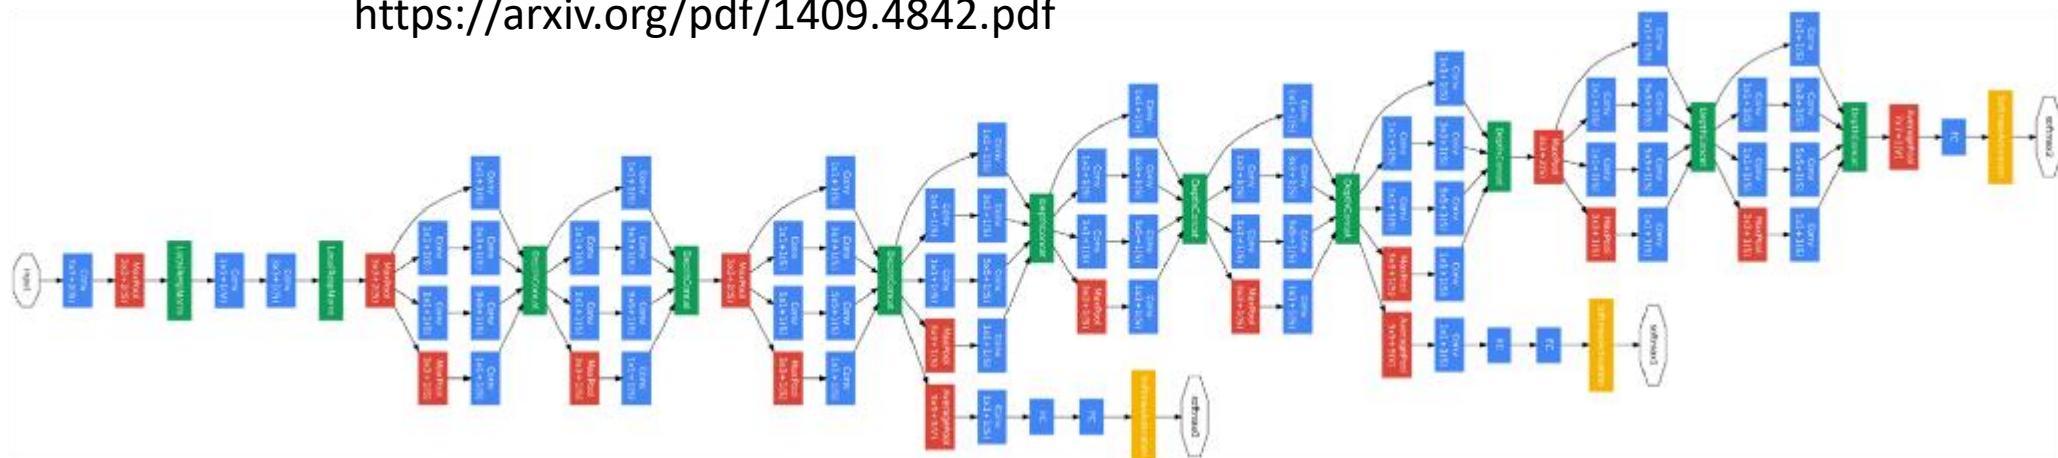


(a) Inception module, naïve version

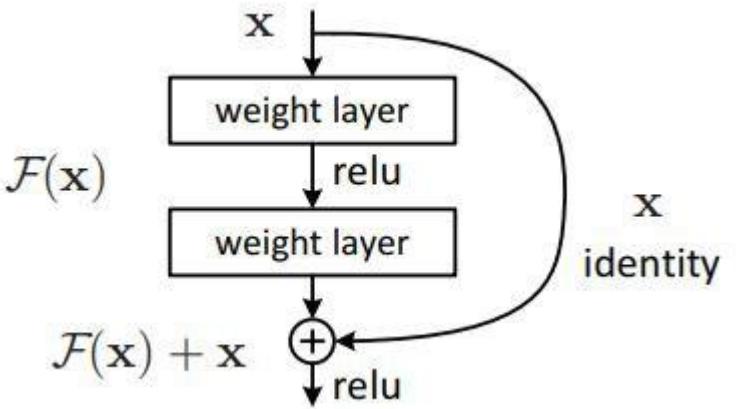


(b) Inception module with dimension reductions

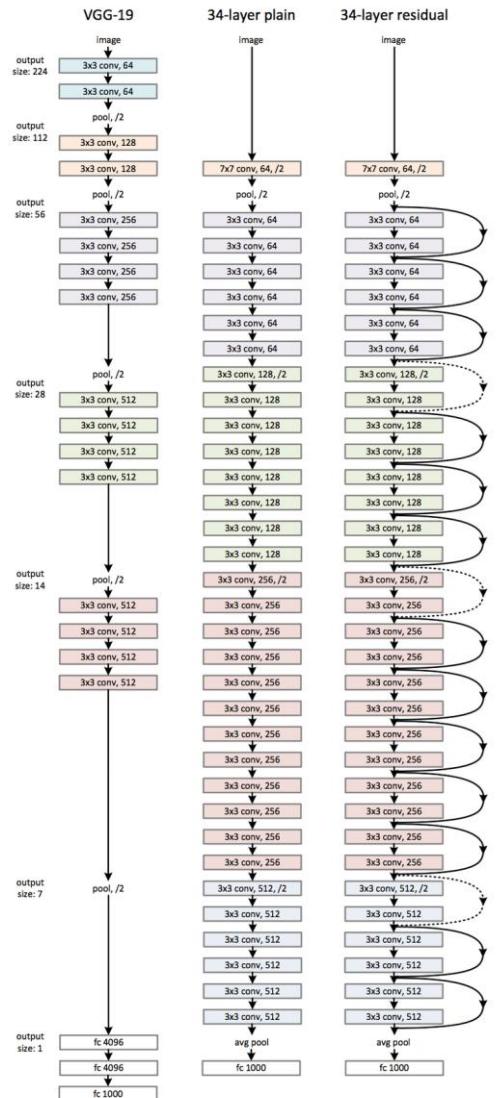
GoogLeNet/InceptionNet (2014): Going deeper with convolutions,  
<https://arxiv.org/pdf/1409.4842.pdf>



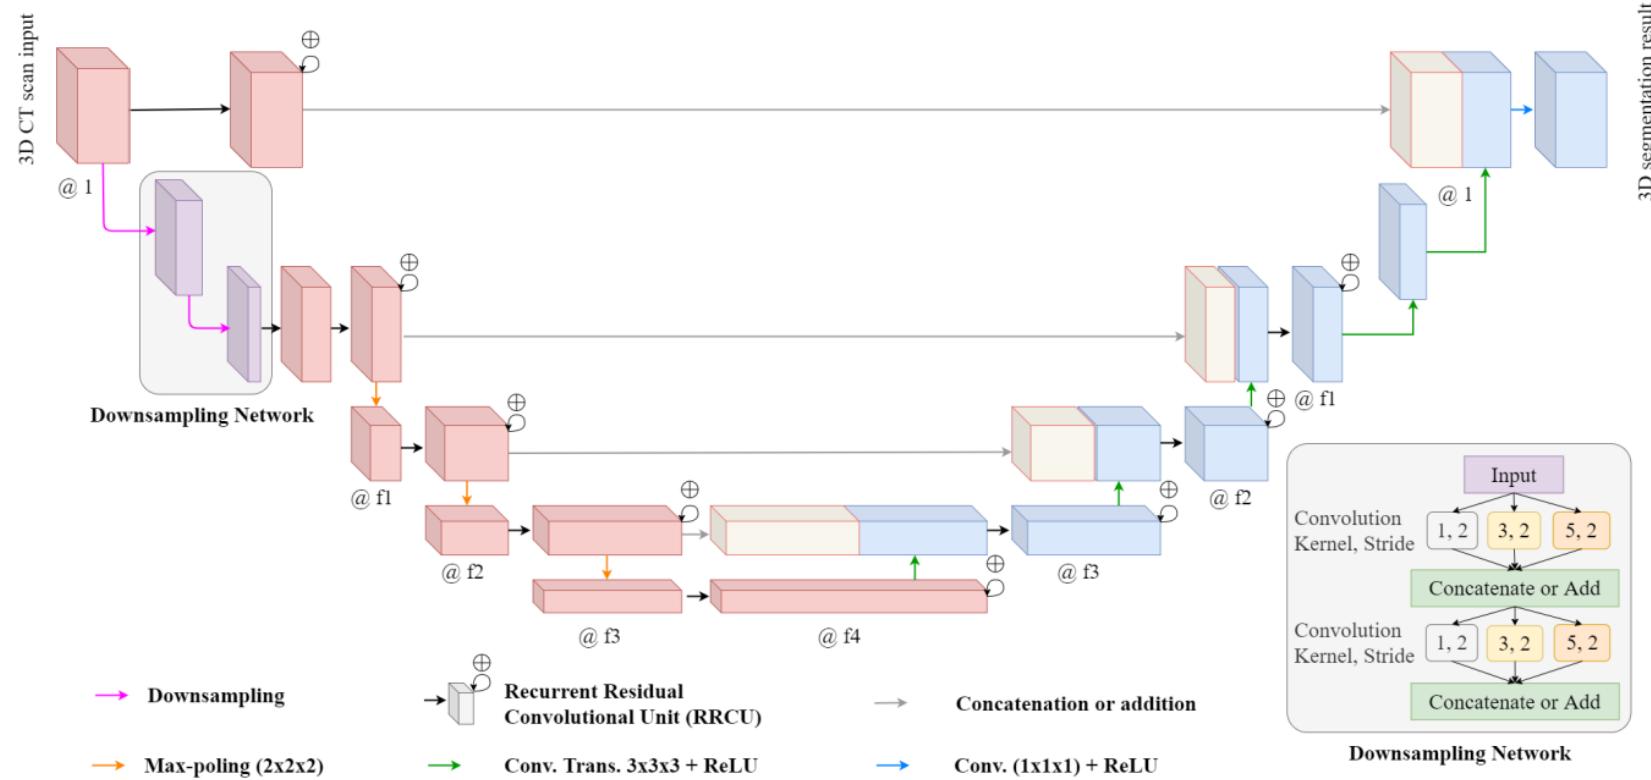
# Example CNN Networks



ResNet (2015): Deep Residual Learning for Image Recognition,  
<https://arxiv.org/pdf/1512.03385.pdf>



# Example CNN Networks



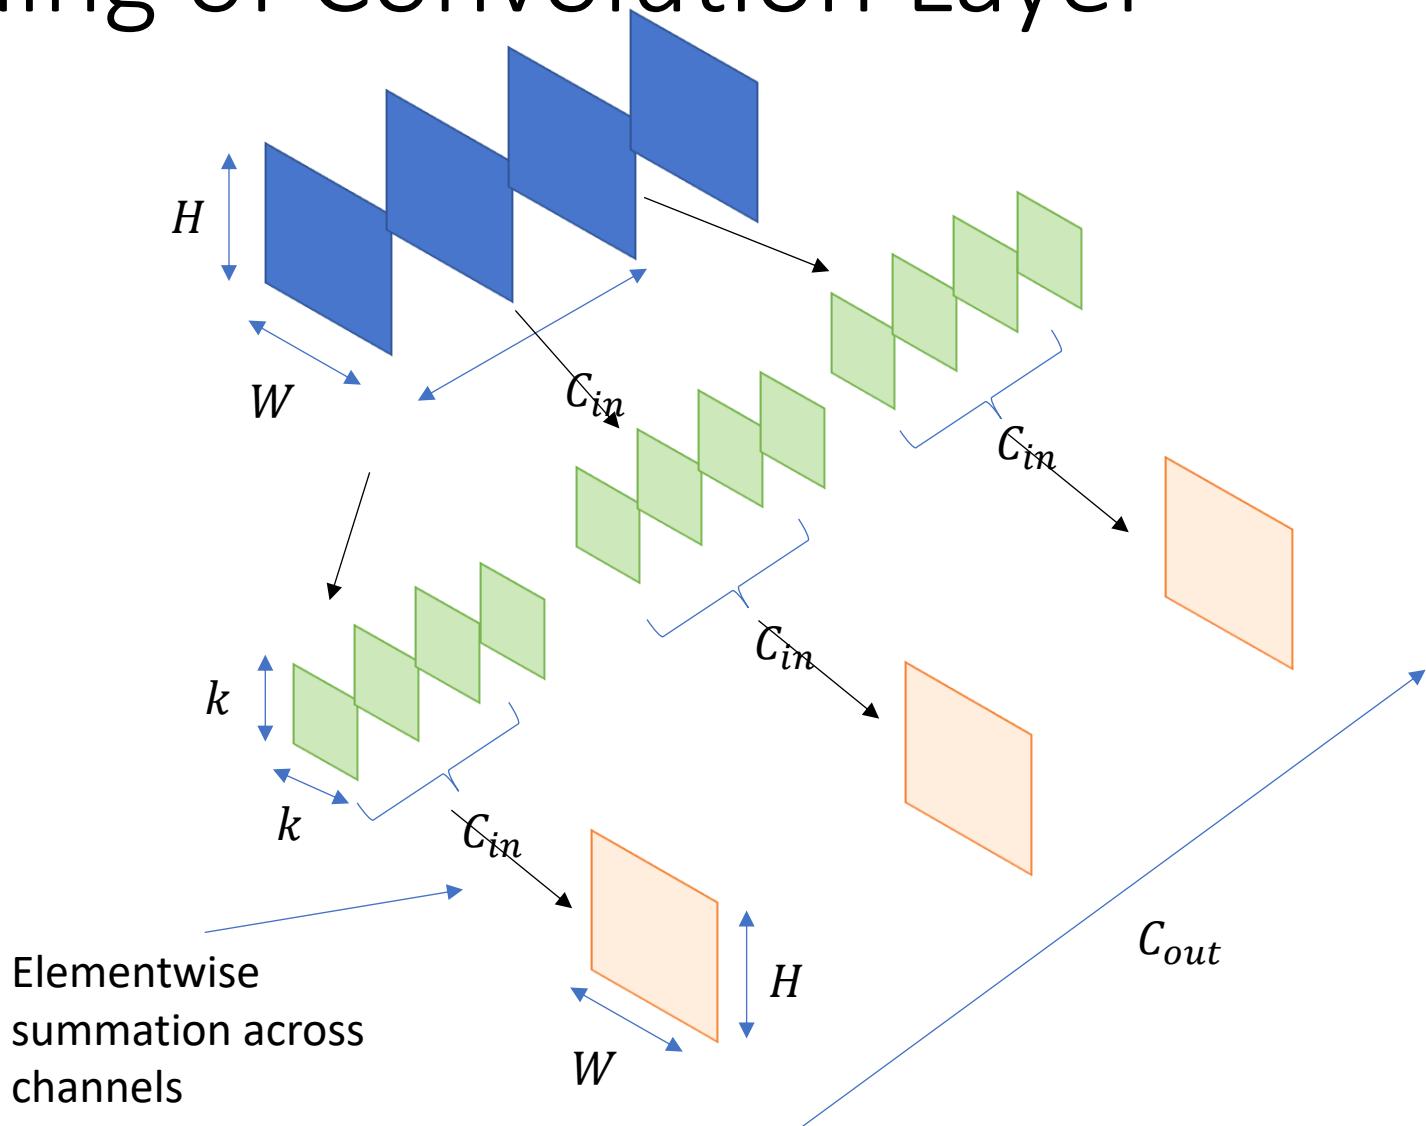
U-net: State of the art in image segmentation, especially in biomedical domain

# Example CNN Networks

- EfficientNet (2019): Tan, M., & Le, Q. V. (2020). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. *ArXiv:1905.11946 [Cs, Stat]*. <http://arxiv.org/abs/1905.11946>
- FixEfficientNets (2022): Touvron, H., Vedaldi, A., Douze, M., & Jégou, H. (2020b). Fixing the train-test resolution discrepancy: FixEfficientNet. *ArXiv:2003.08237 [Cs]*. <http://arxiv.org/abs/2003.08237>
- AmoebaNets (2019): Regularized Evolution for Image Classifier Architecture Search, <https://ojs.aaai.org/index.php/AAAI/article/view/4405>

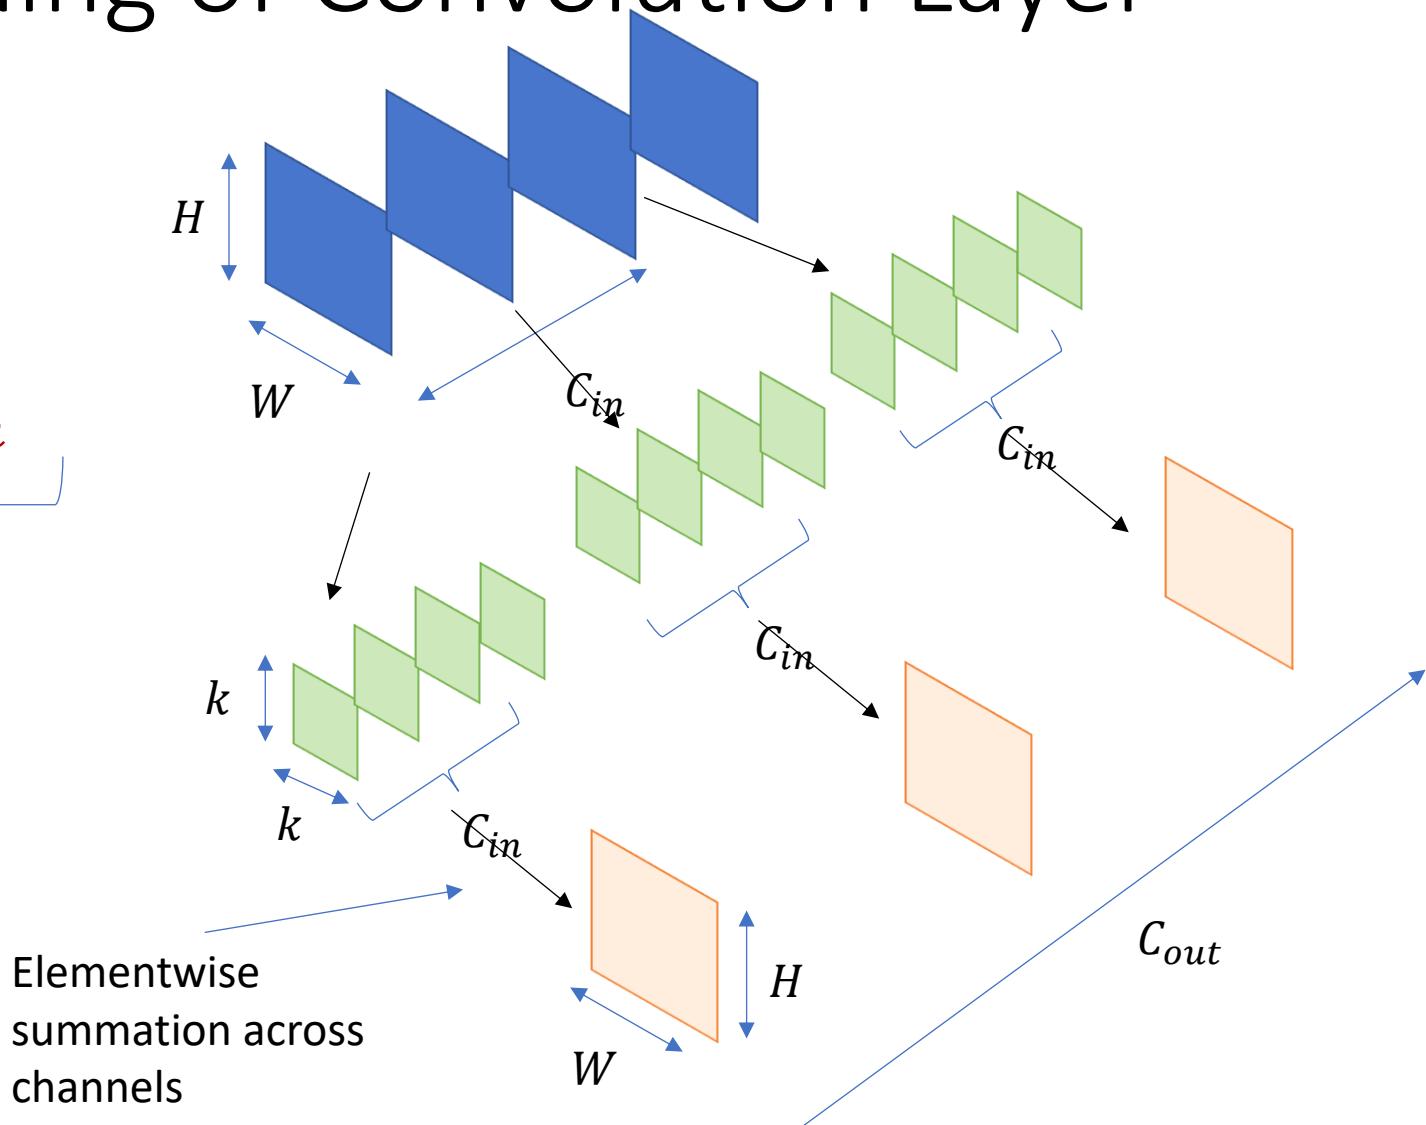
# Performance Modeling of Convolution Layer

- Number of Operations: ??



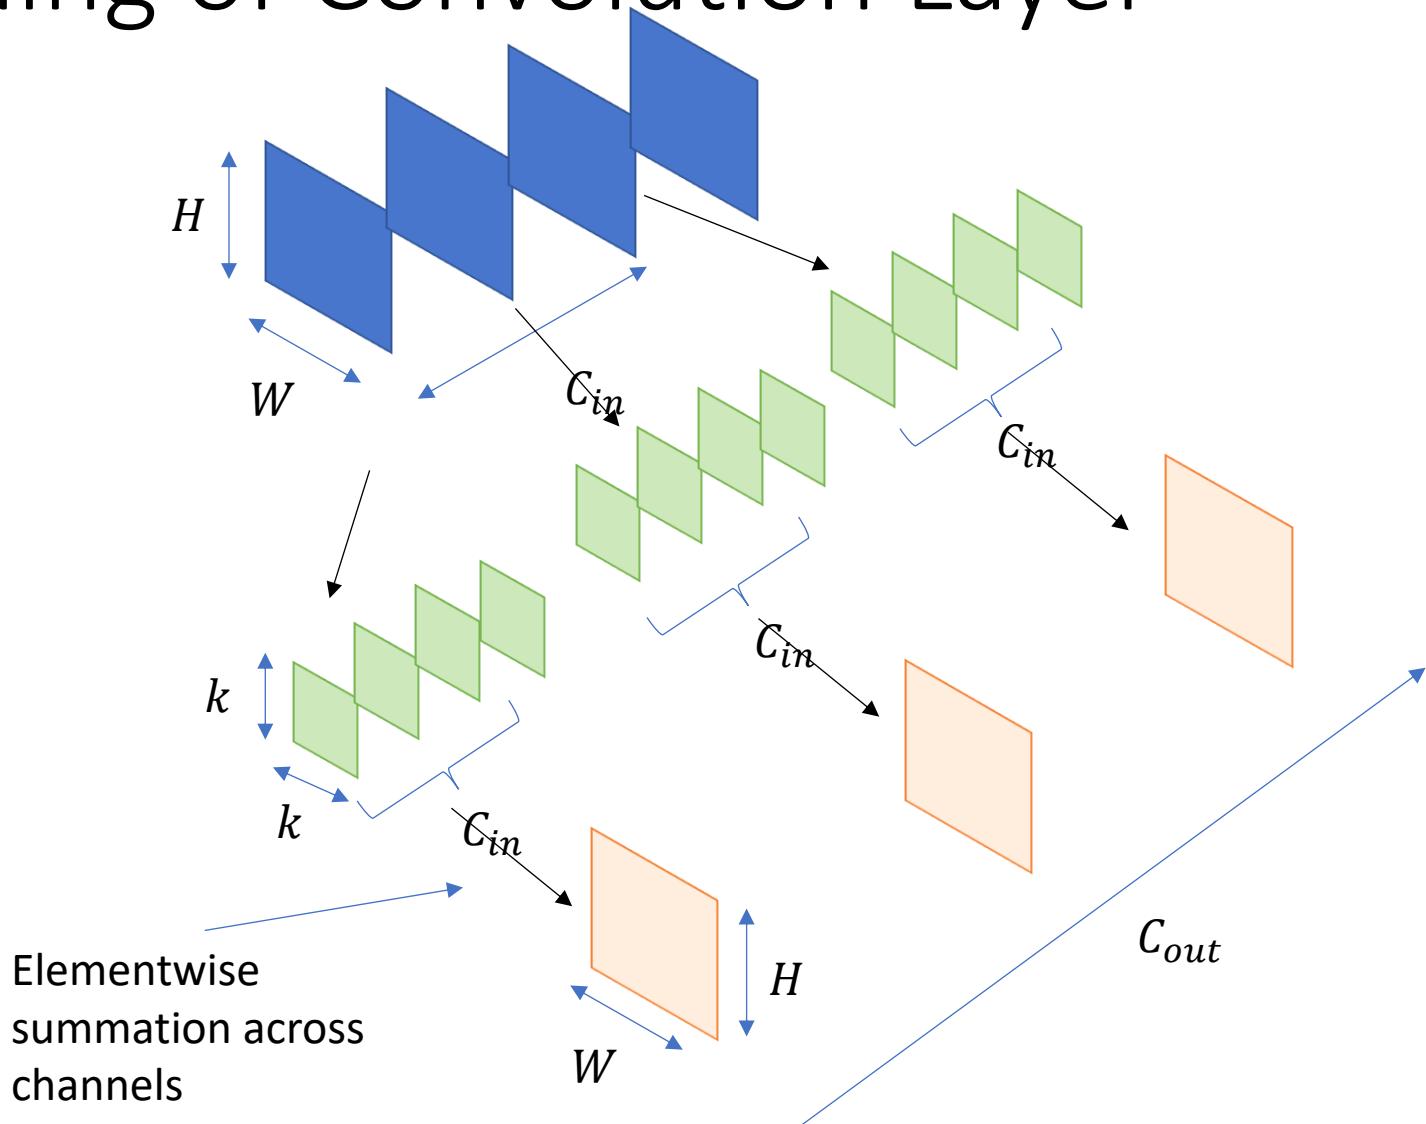
# Performance Modeling of Convolution Layer

- Number of Operations:
  - $H \times W \times C_{out} \times k \times k \times C_{in}$
- { Output Pixels }      { Operations per Pixels }



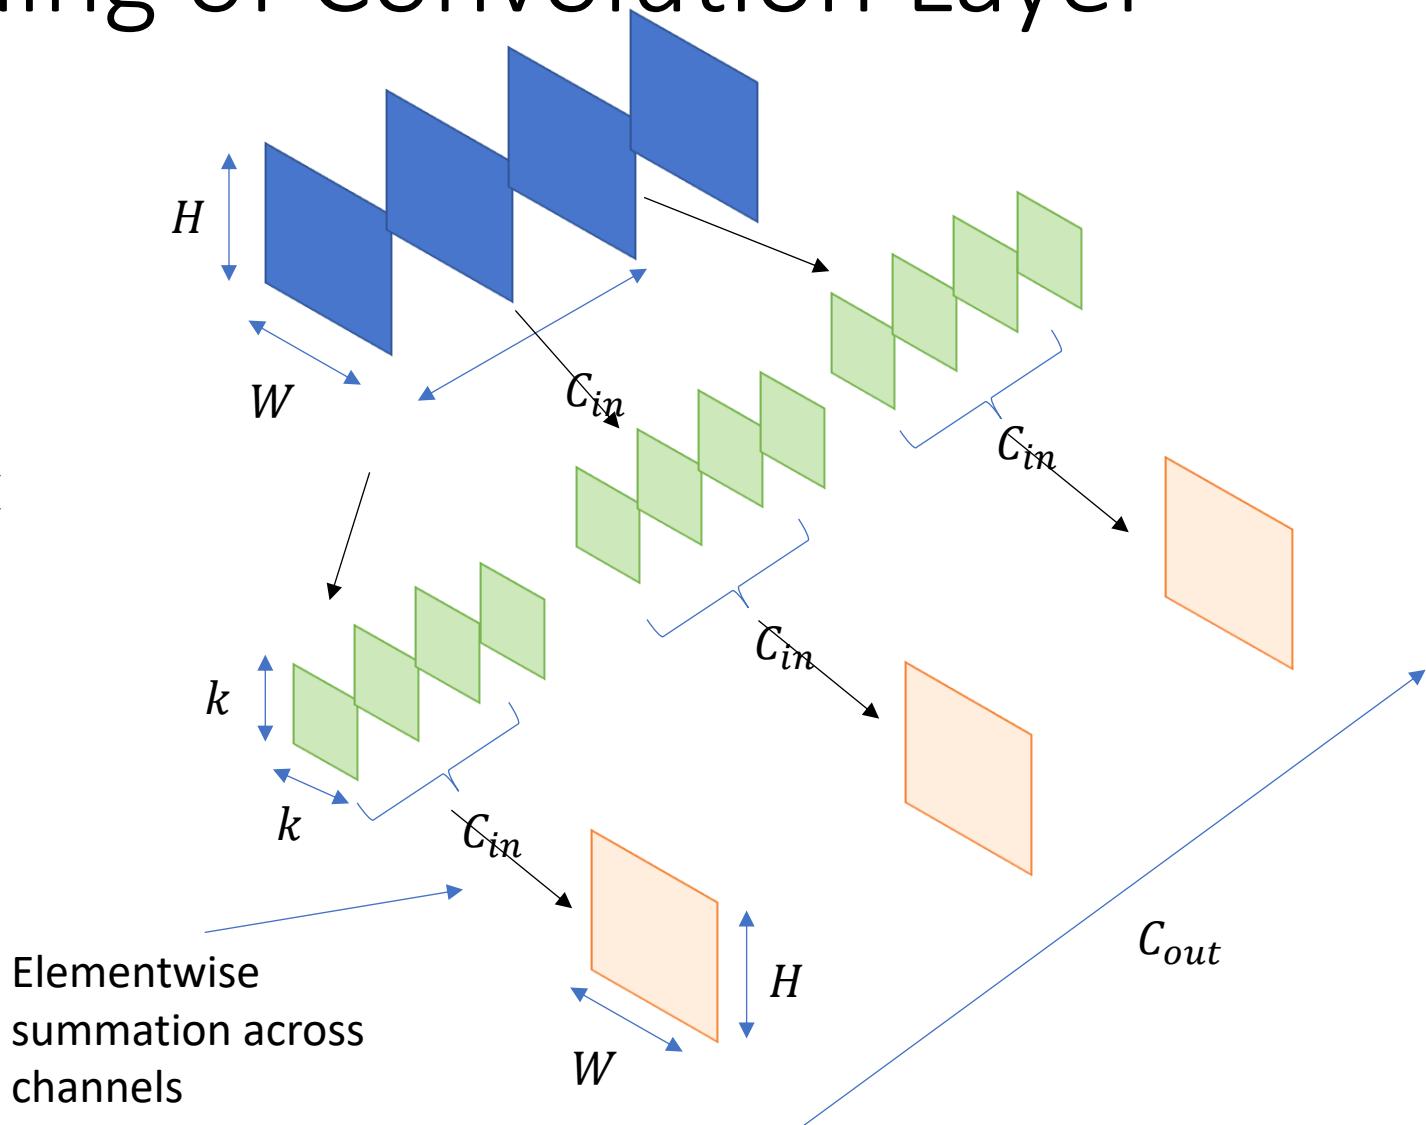
# Performance Modeling of Convolution Layer

- Memory Transfers (infinite cache): ??



# Performance Modeling of Convolution Layer

- Memory Transfers (infinite cache): ??
- $H \times W \times C_{in} + C_{in} \times C_{out} \times k \times k + H \times W \times C_{out}$
- $\approx H \times W \times (C_{in} + C_{out})$



# Performance Modeling of Convolution Layer

- Theoretical Arithmetic Intensity

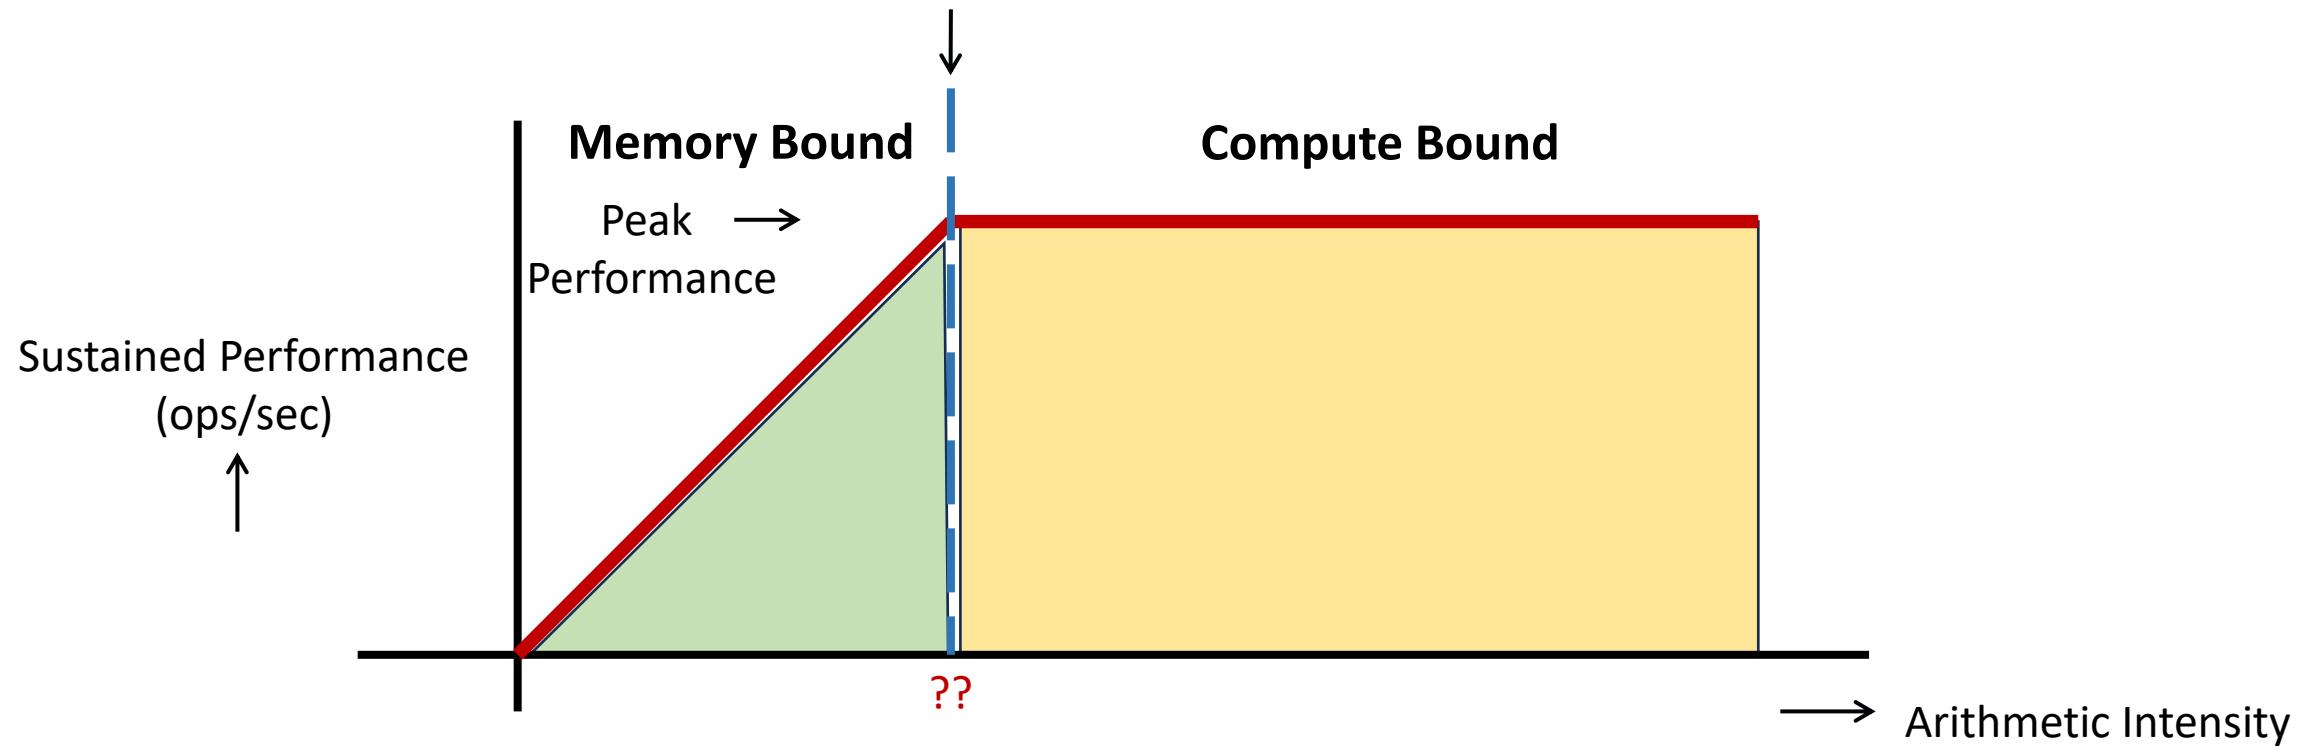
$$\frac{k^2 C_{in} C_{out}}{C_{in} + C_{out}}$$

Example value for  $k = 9, C_{in} = 3, C_{out} = 64$

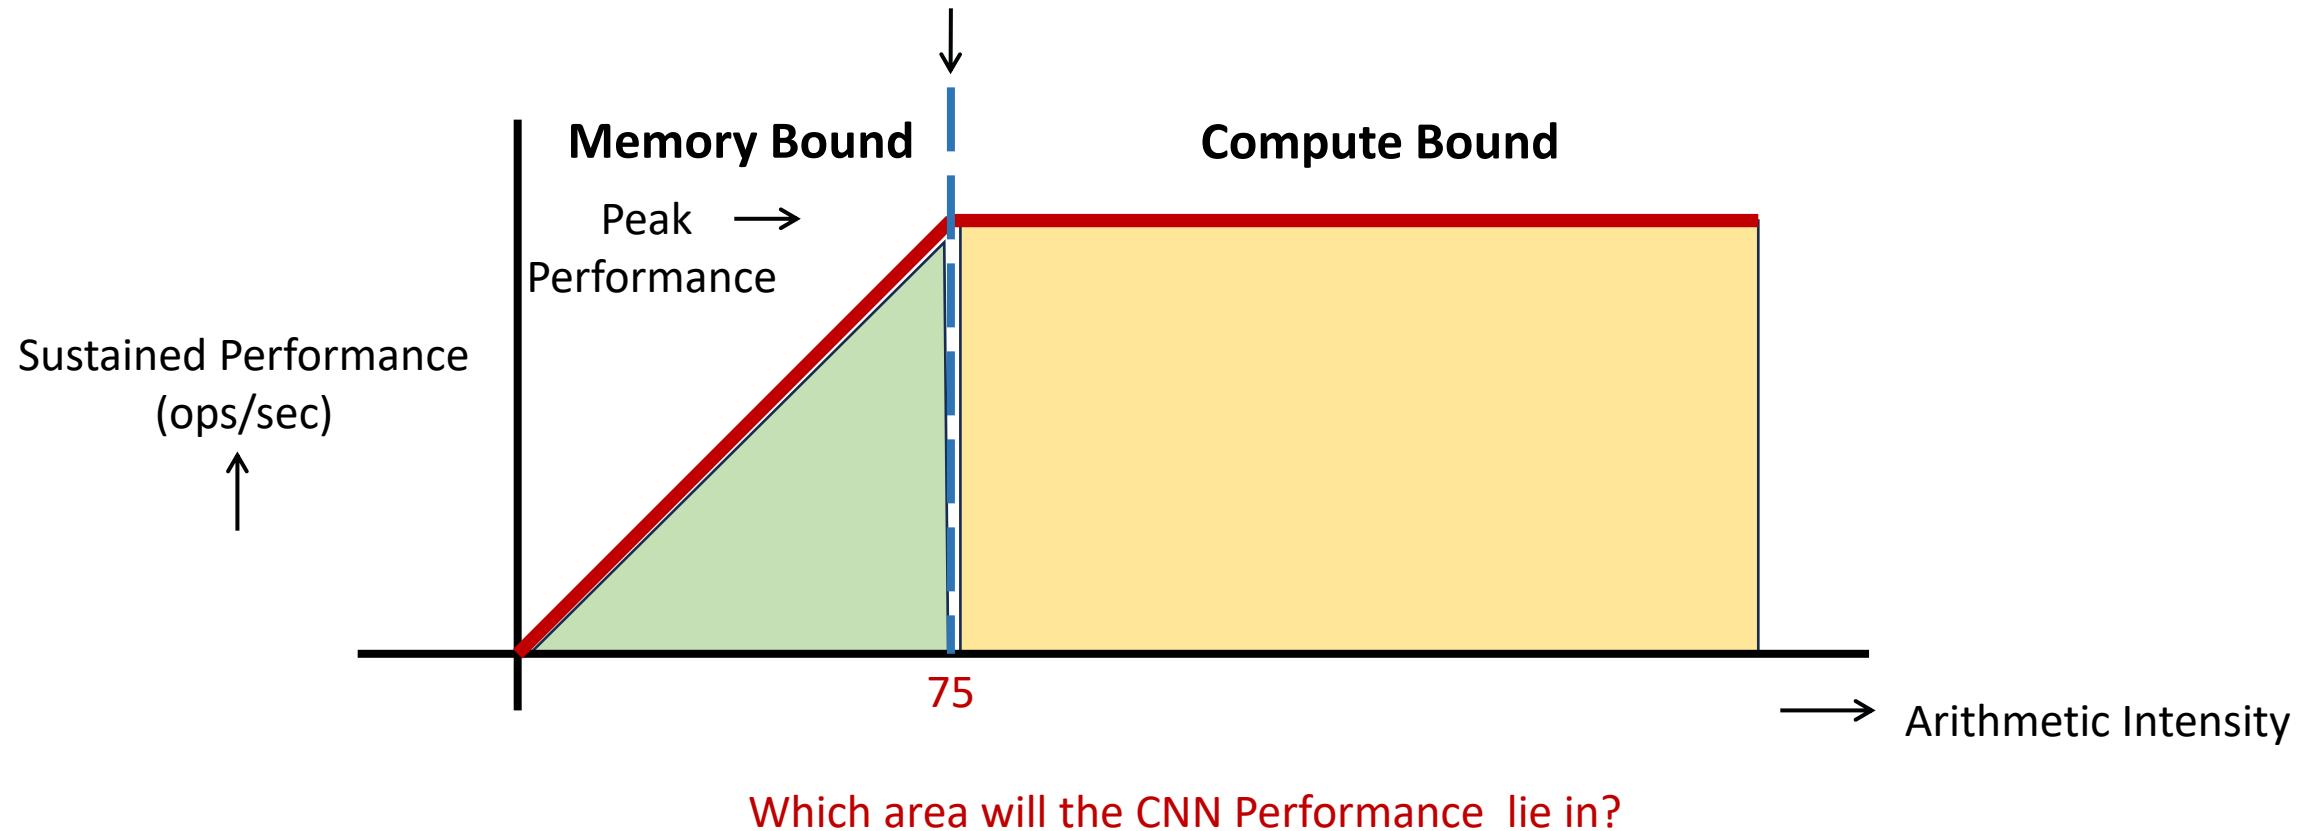
232

Example Machine Intensity = 75 (Nvidia A100\*) **how???**

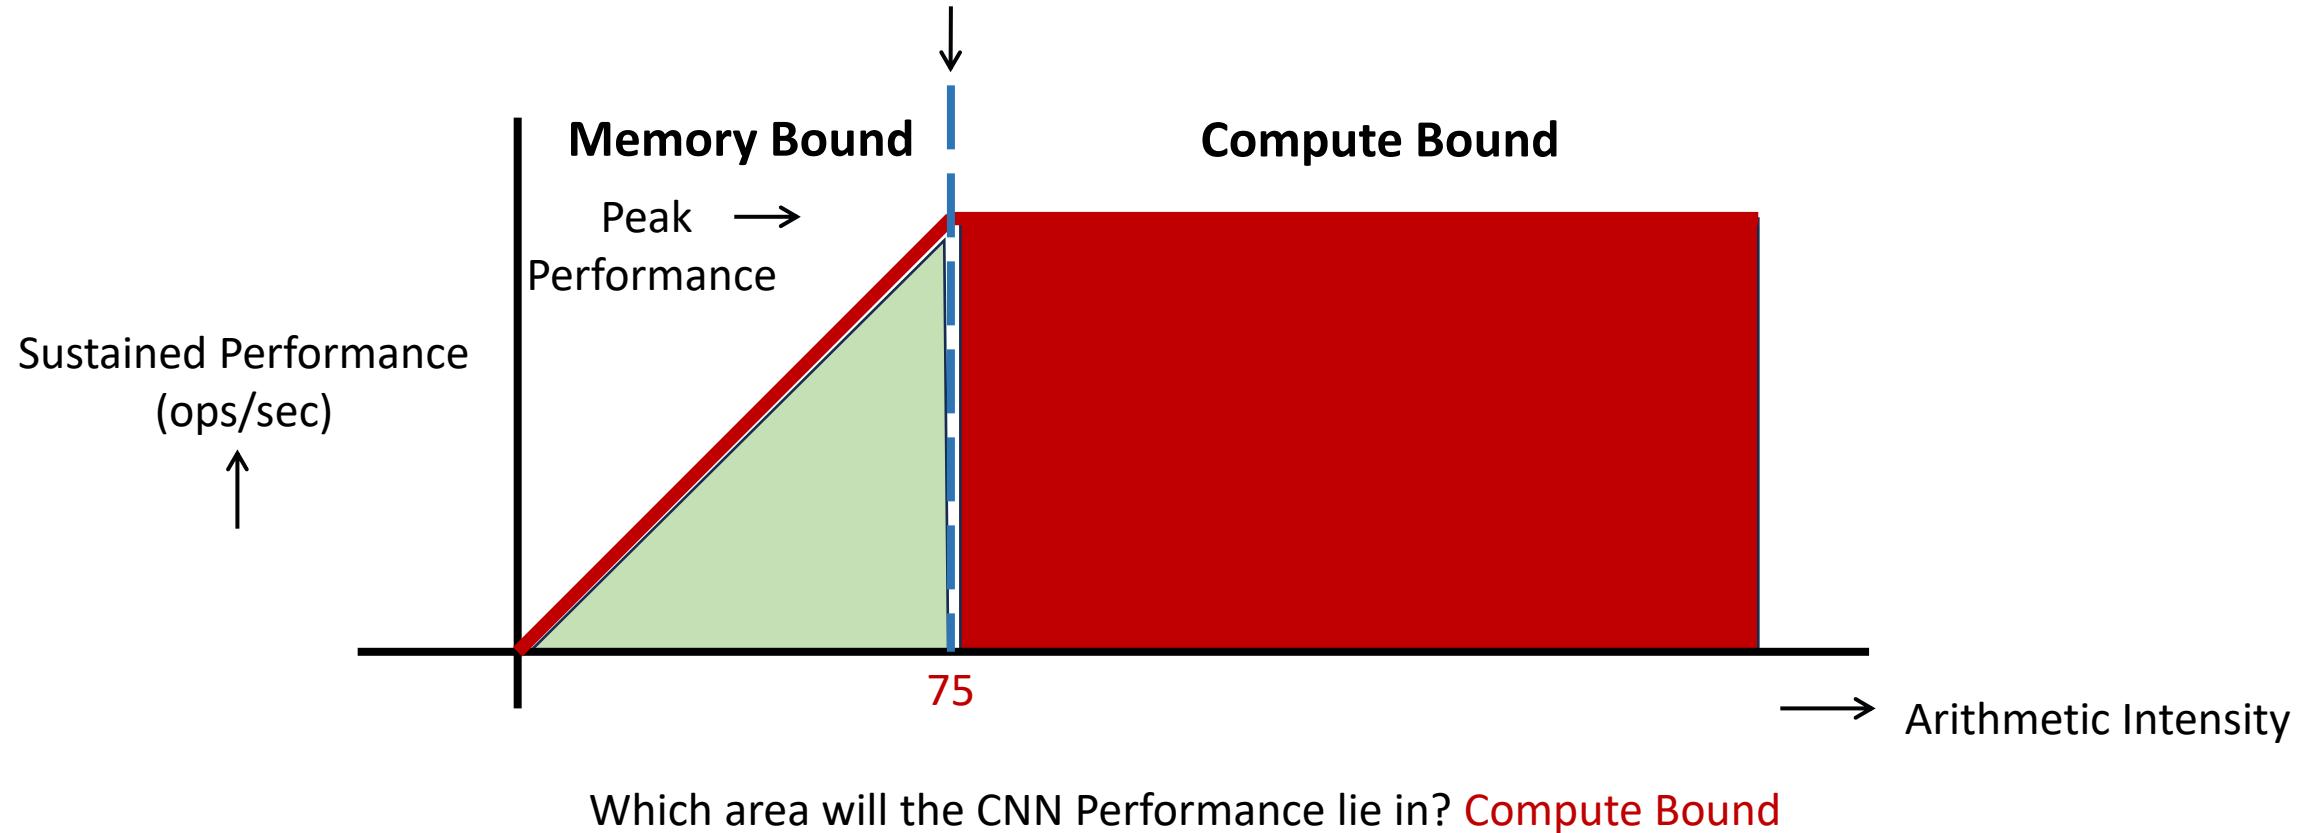
# CNN on A100 Roofline Plot (assuming infinite cache)



# CNN on A100 Roofline Plot (assuming infinite cache)



# CNN on A100 Roofline Plot (assuming infinite cache)



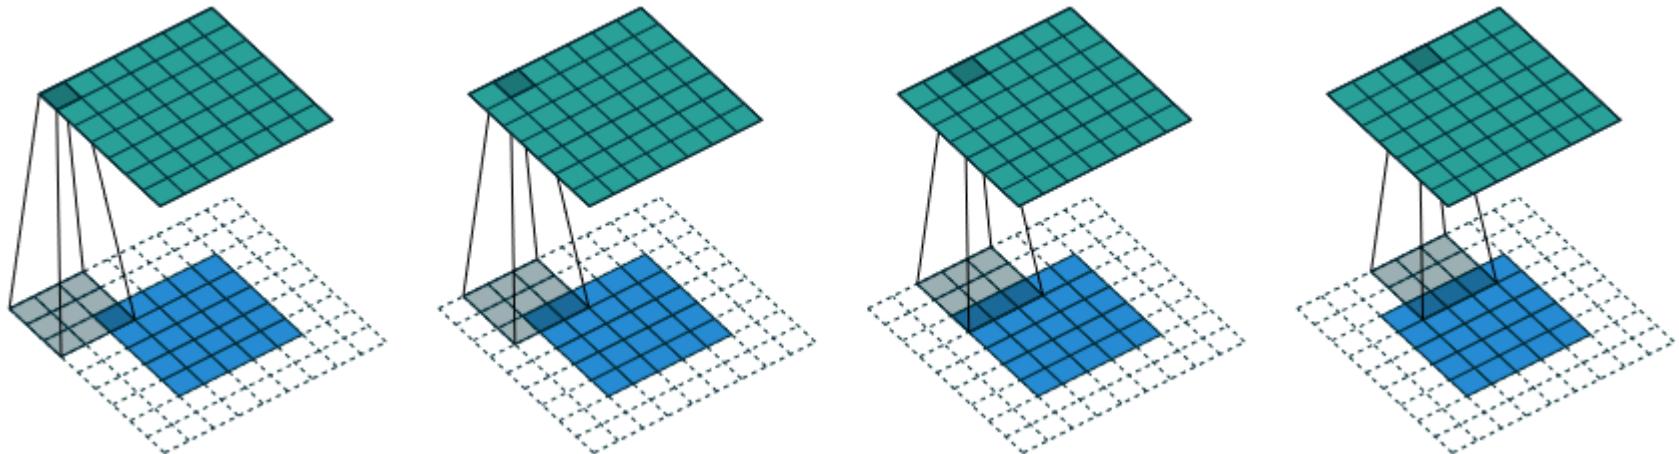
# Implications

- Essentially a Compute Bound Problem, assuming large cache
- Objective: How to ensure efficiency is  $O(1)$ , i.e., all compute units are fully utilized
- If cache is limited
  - Objective 1: How to improve data reuse so that compute units have enough to work on?
  - Objective 2: How to ensure efficiency is  $O(1)$ , i.e., all compute units are fully utilized
- Another approach: increase Number of Devices
  - + Larger On-chip memory
  - + Higher bandwidth
  - - Communication/Synchronization overheads

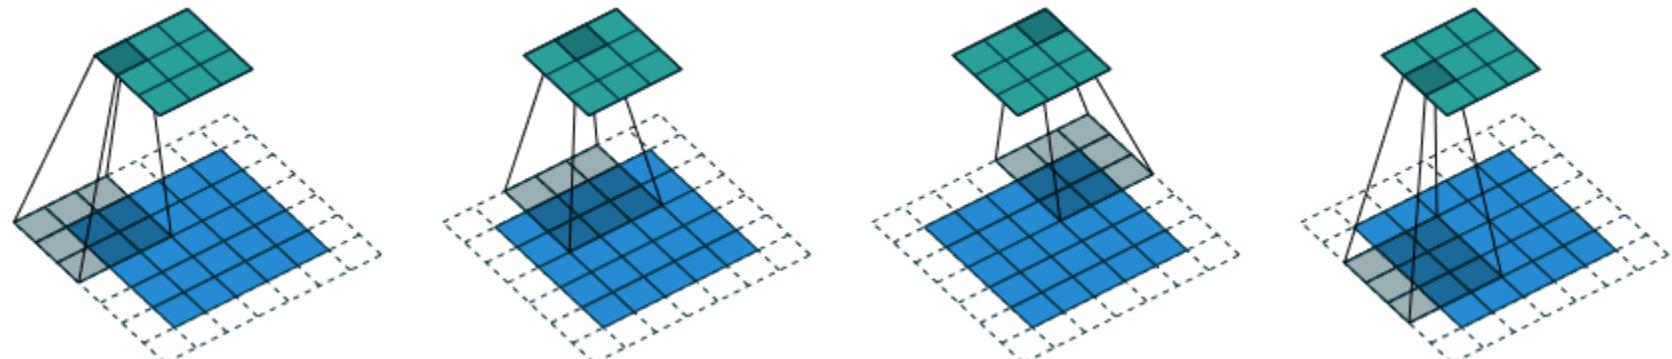
# Variations

## Reading Homework:

A guide to convolutional  
arithmetic for deep learning



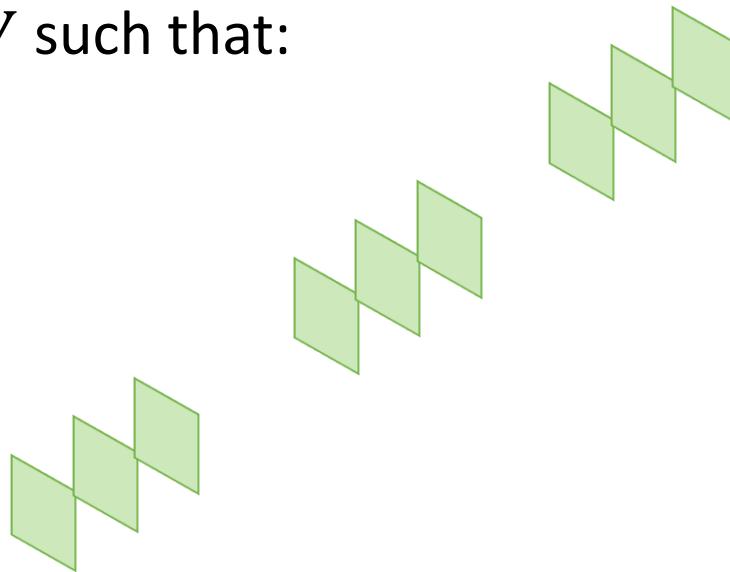
- Padding



- Strides

# Convolutional Neural Network

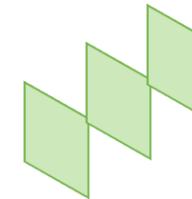
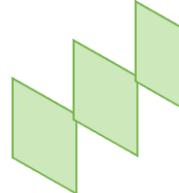
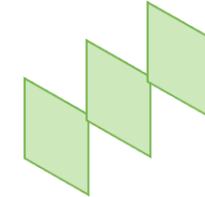
- Training
  - Given  $x_i \in X$  (image from a database),  $y_i$ : Label for the image
  - “Learn” a model  $F$ , defined by the filters  $W$  such that:
    - $\sum(y_i - F(x_i; W))^2$  is minimized
  
- Inference
  - Given a new image  $x'$
  - Predict a label  $y' = W(x')$
  - No change in the filters



# Convolutional Neural Network

- **Training**

- Given  $x_i \in X$  (image from a database),  $y_i$ : Label for the image
- “Learn” a model  $F$ , defined by the filters  $W$  such that:
  - $\sum(y_i - F(x_i; W))^2$  is minimized



- **Inference**

- Given a new image  $x'$
- Predict a label  $y' = W(x')$
- No change in the filters

# Training

$$E(W): \min_W \sum (y_i - F(x_i: W))^2$$

$$W_{t+1} \leftarrow W_t - \alpha \frac{\delta E(W)}{\delta W}$$

Iteratively

$$\frac{\delta E(W)}{\delta W} = 2 \sum_i (y_i - F(x_i: W)) \times \frac{\delta F(x_i: W)}{\delta W}$$

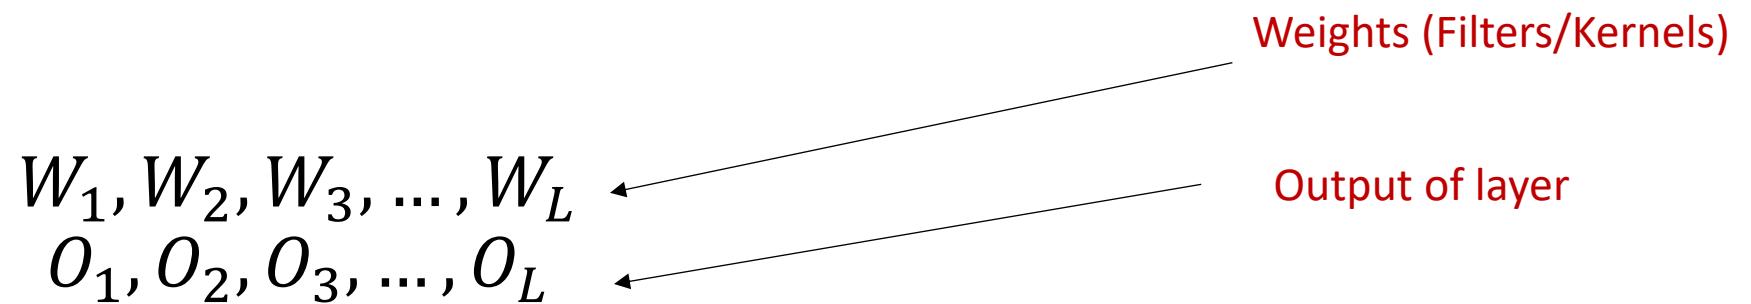
For all samples, in each iteration



Error

# Training

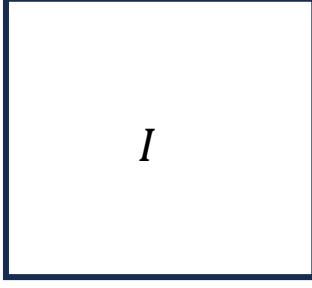
$L$  layers in CNN



$$2 * \text{Error} * \frac{\delta F(x_i: W_{1:L})}{\delta W_1}, \frac{\delta F(x_i: W_{1:L})}{\delta W_2}, \dots, \frac{\delta F(x_i: W_{1:L})}{\delta W_L}$$

Compute these for  
all the samples

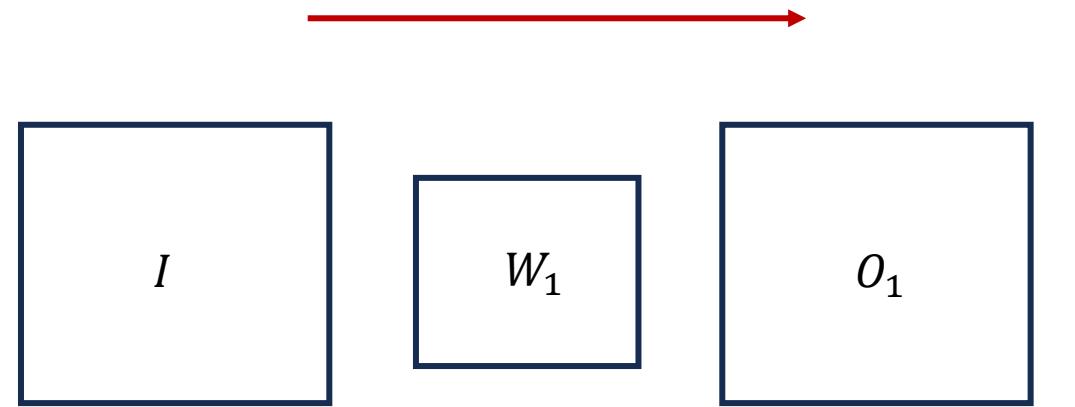
# CNN Training: Forward Propagation



*I*

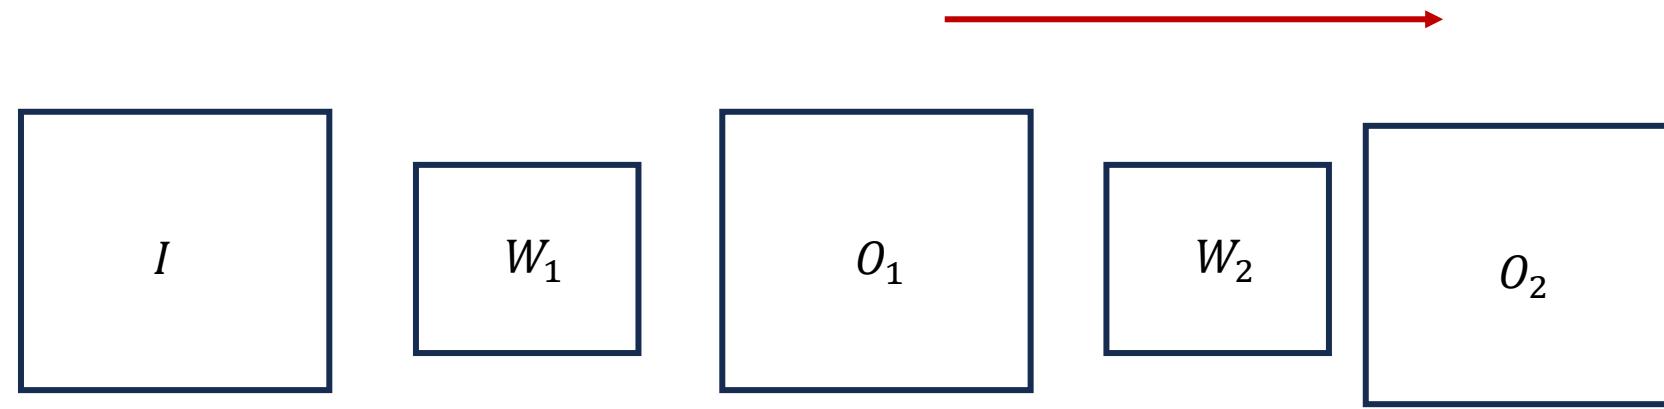
For a single image

# CNN Training: Forward Propagation



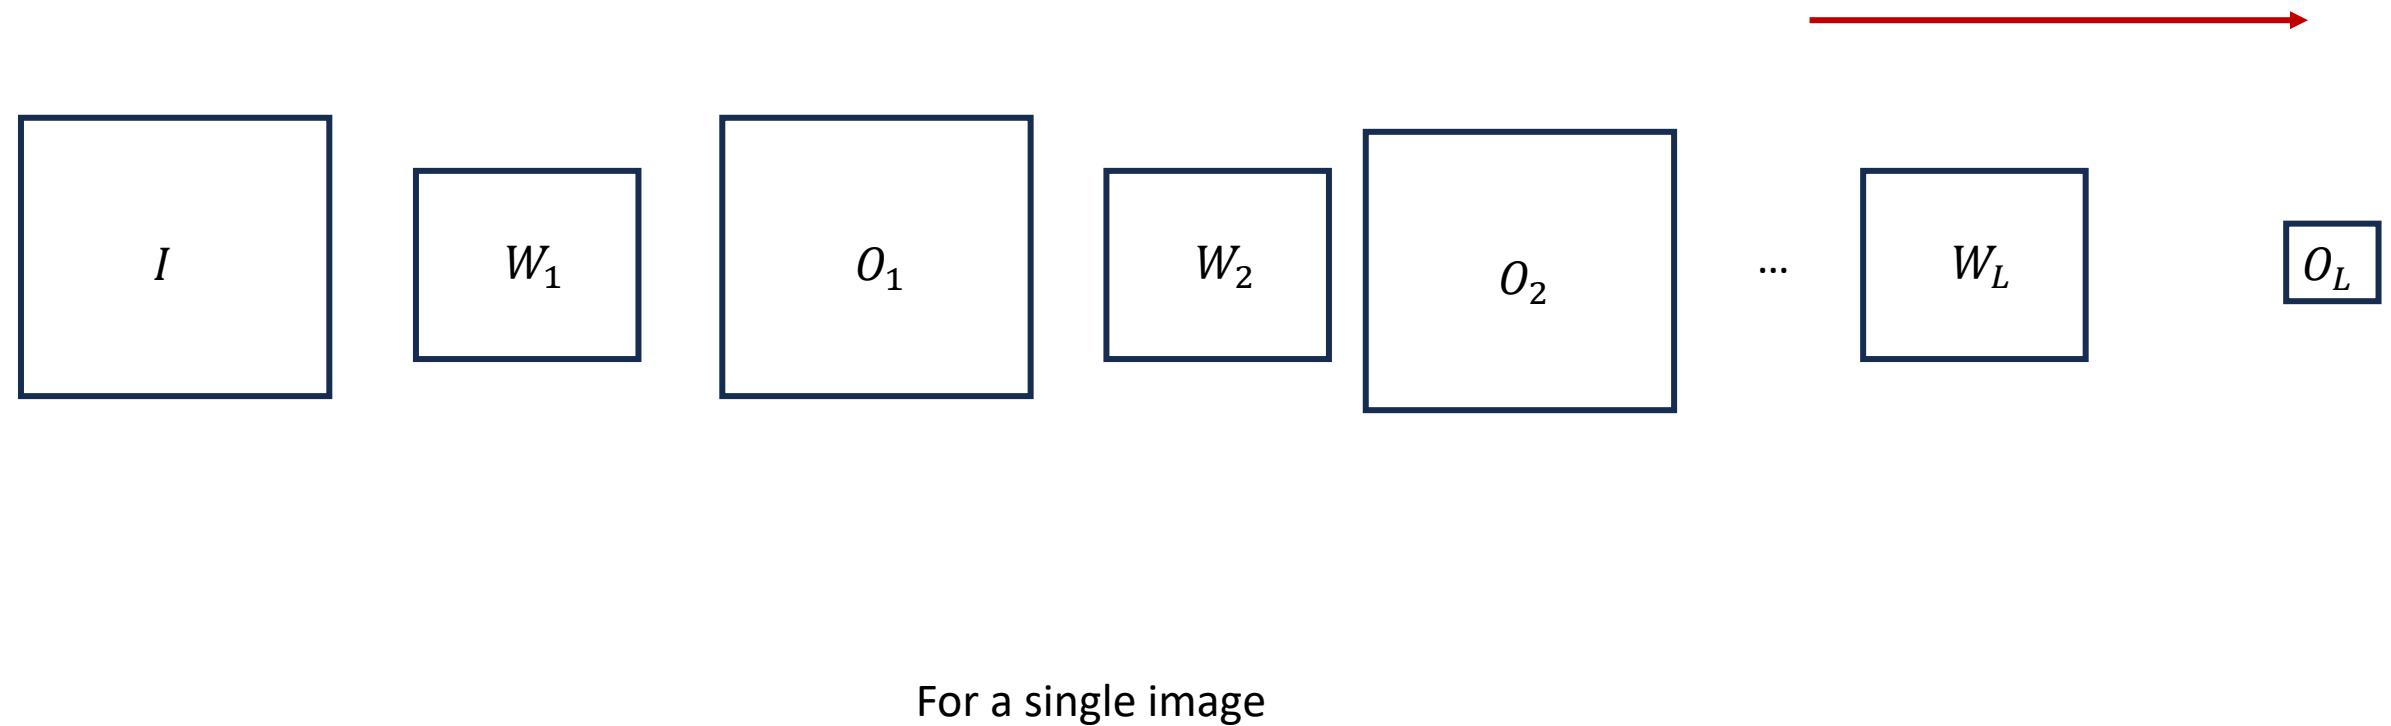
For a single image

# CNN Training: Forward Propagation



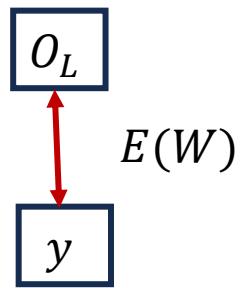
For a single image

# CNN Training: Forward Propagation



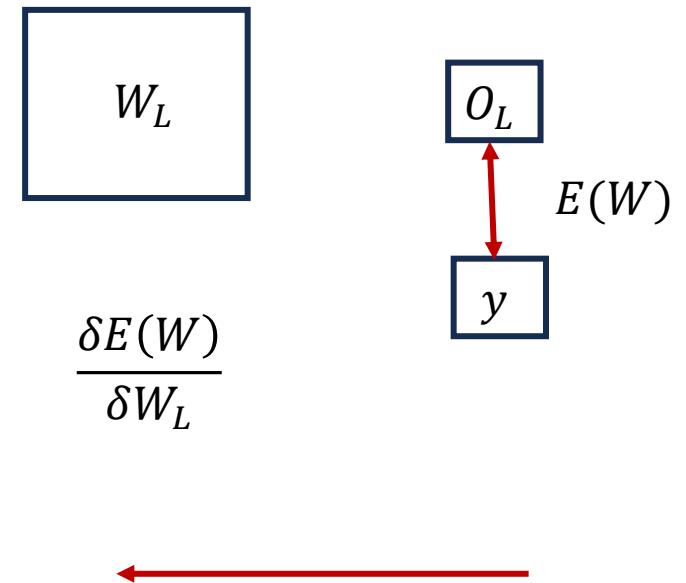
# CNN Training: Error Calculation

For a single image

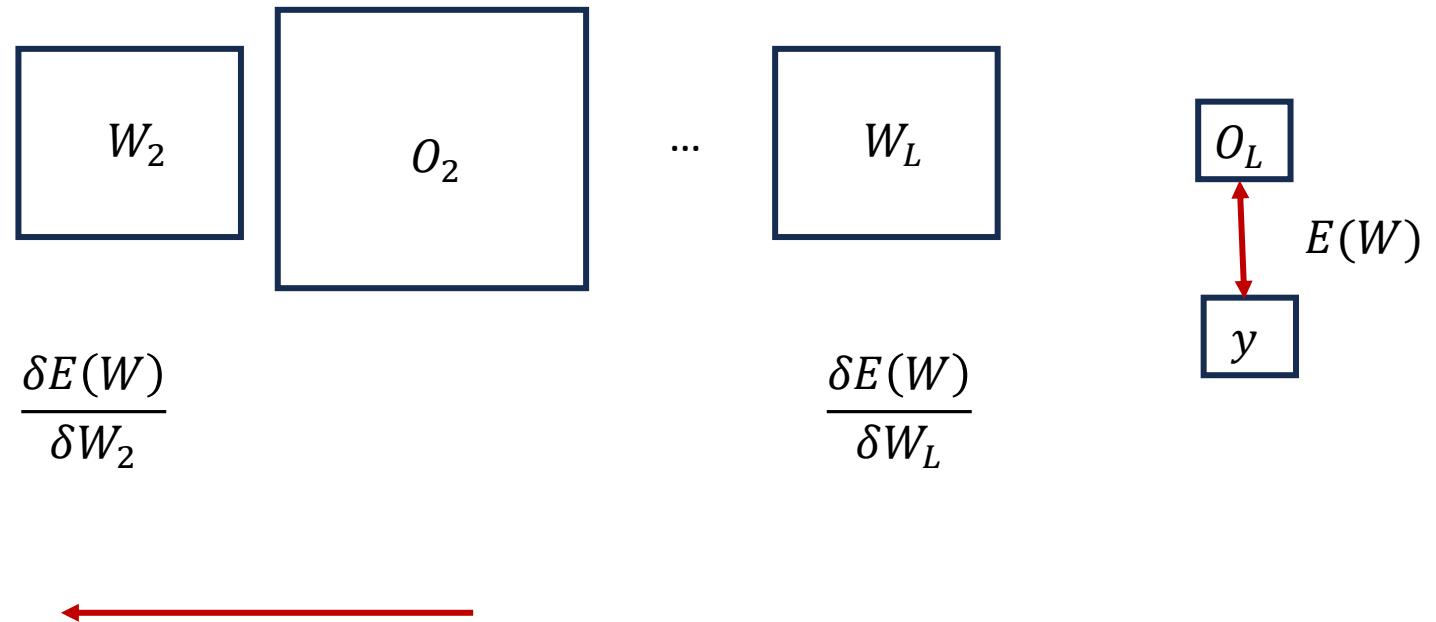


# CNN Training: Backward Propagation

For a single image

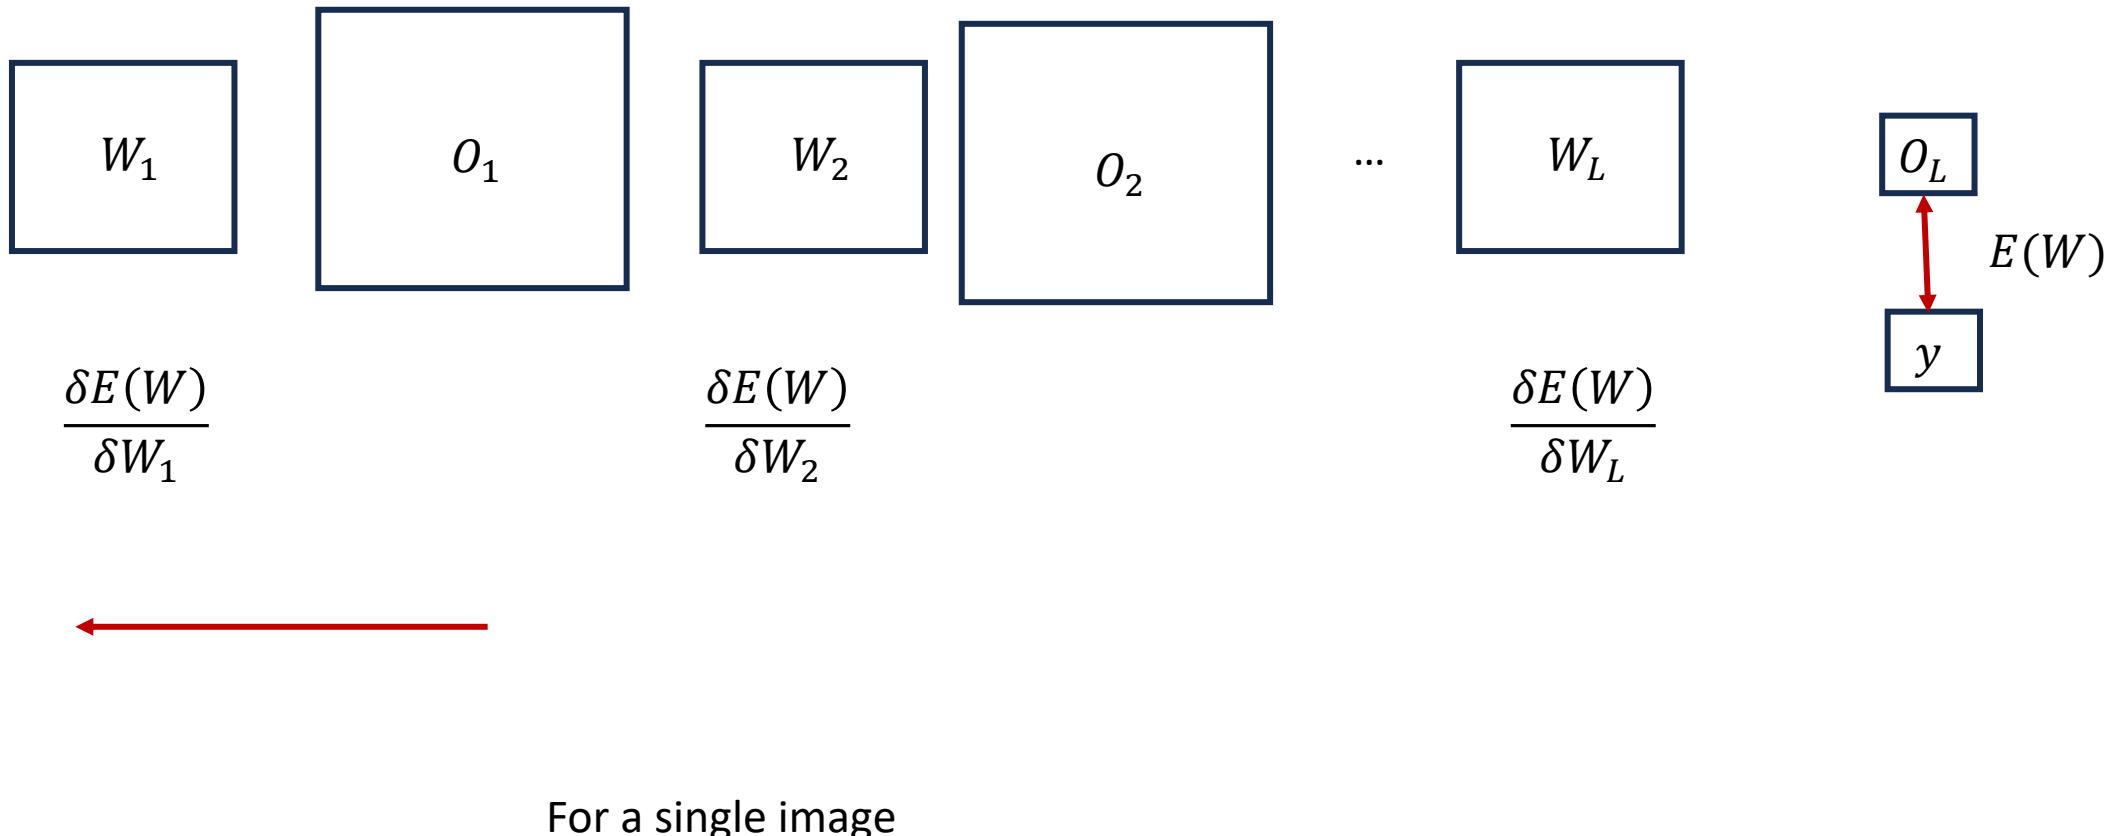


# CNN Training: Backward Propagation

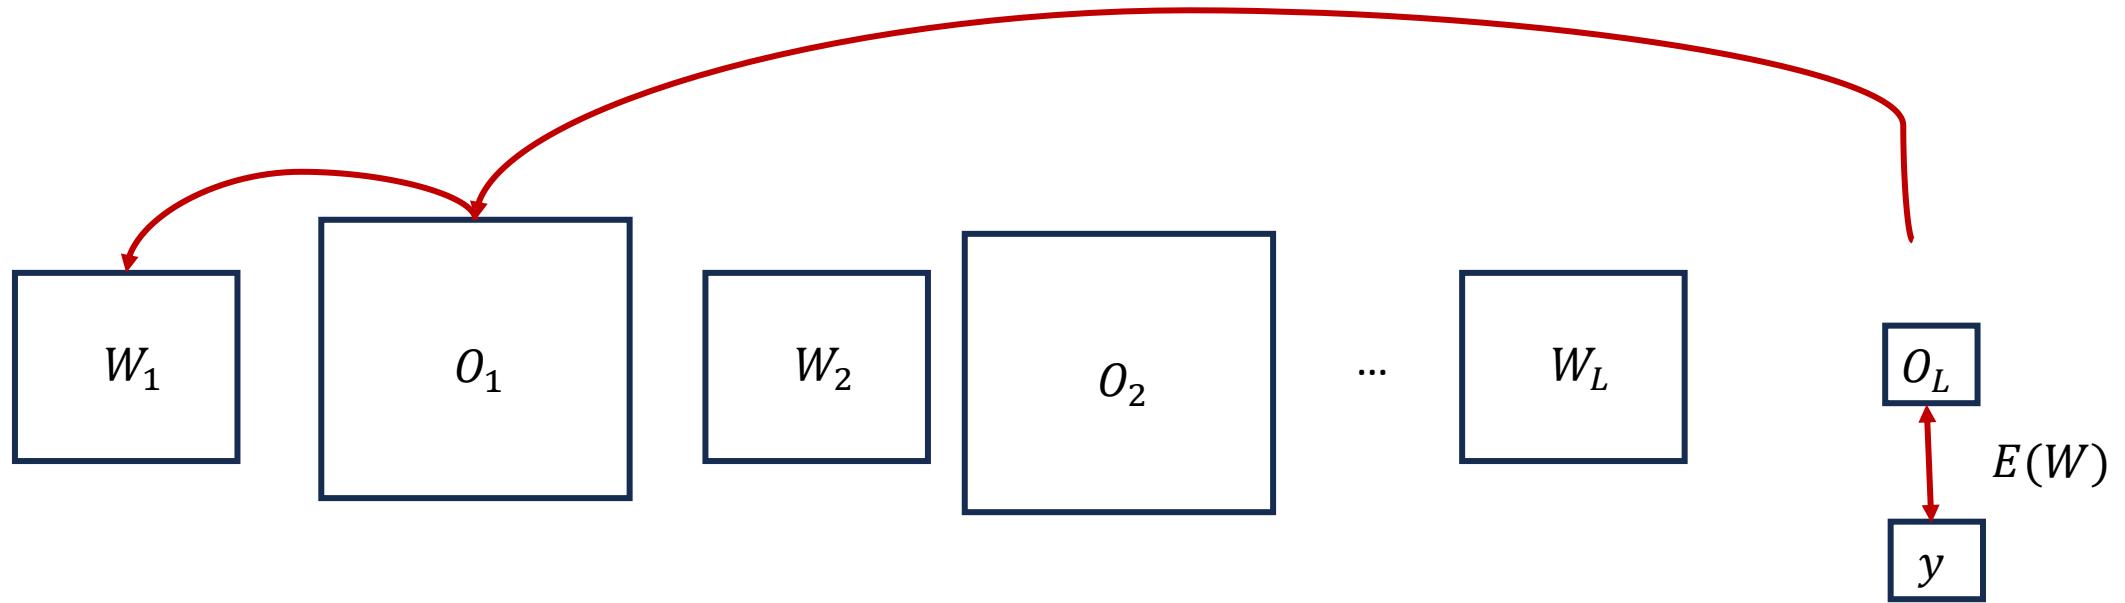


For a single image

# CNN Training: Backward Propagation



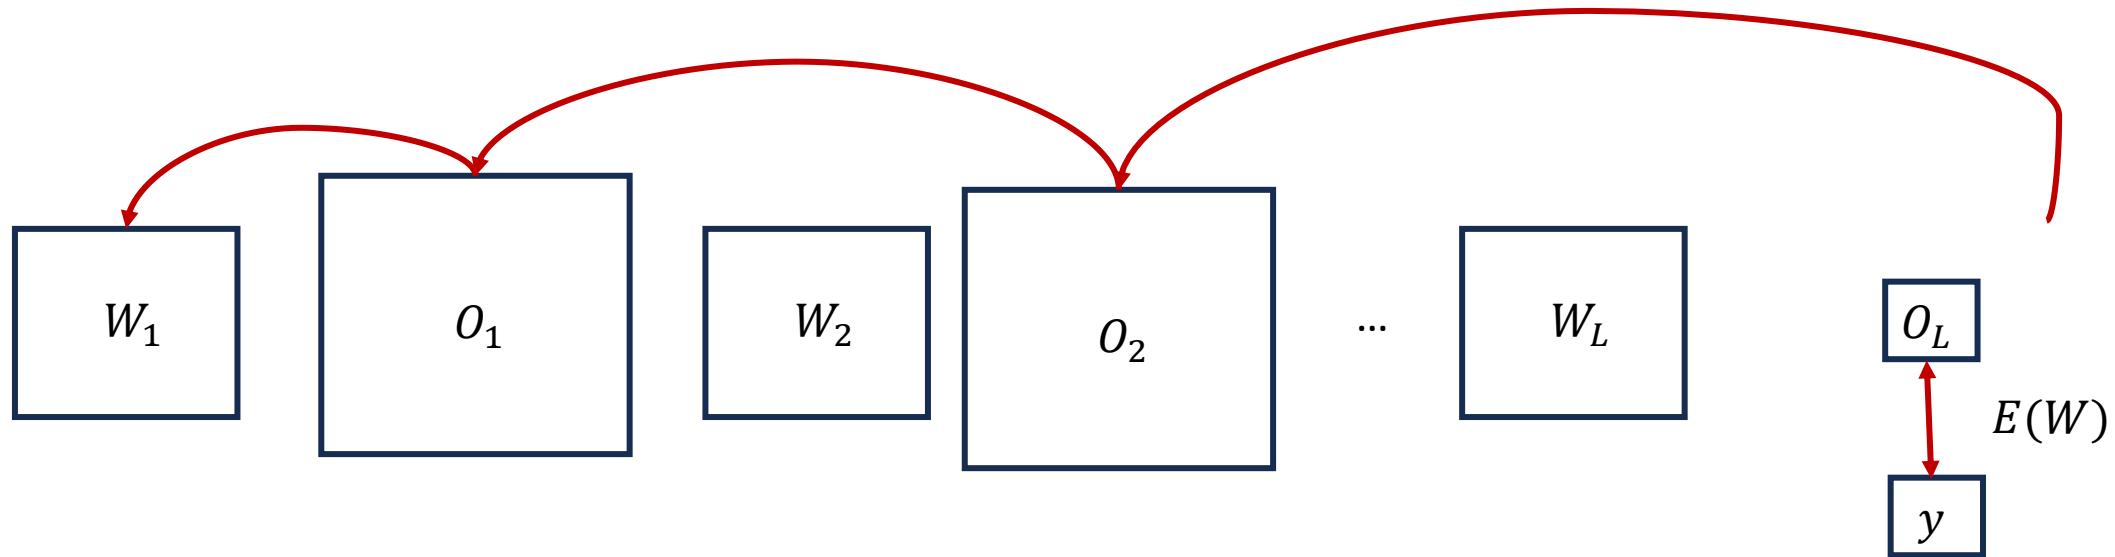
# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta E(W)}{\delta O_1}$$

For a single image

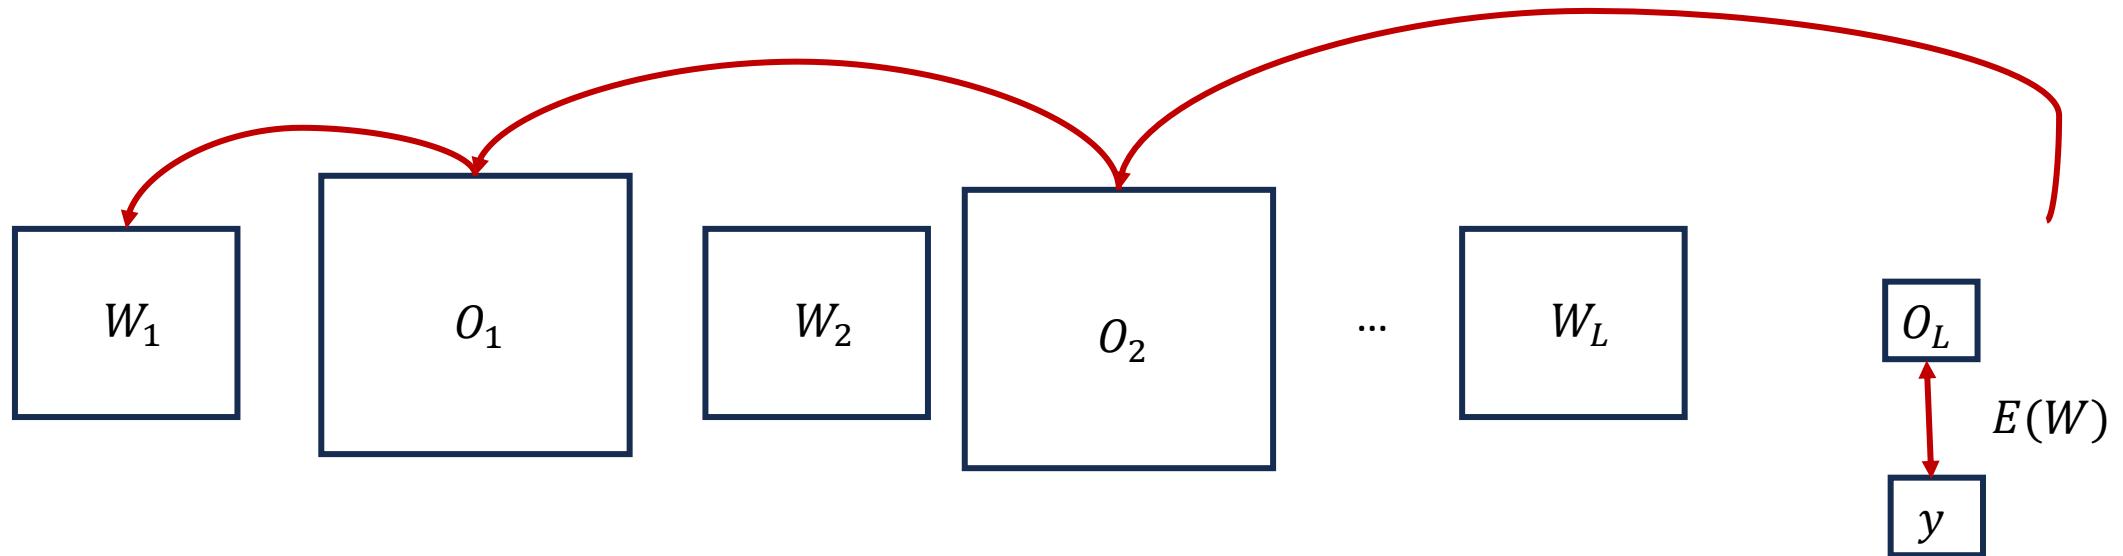
# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta O_2}{\delta O_1} \times \frac{\delta E(W)}{\delta O_2}$$

For a single image

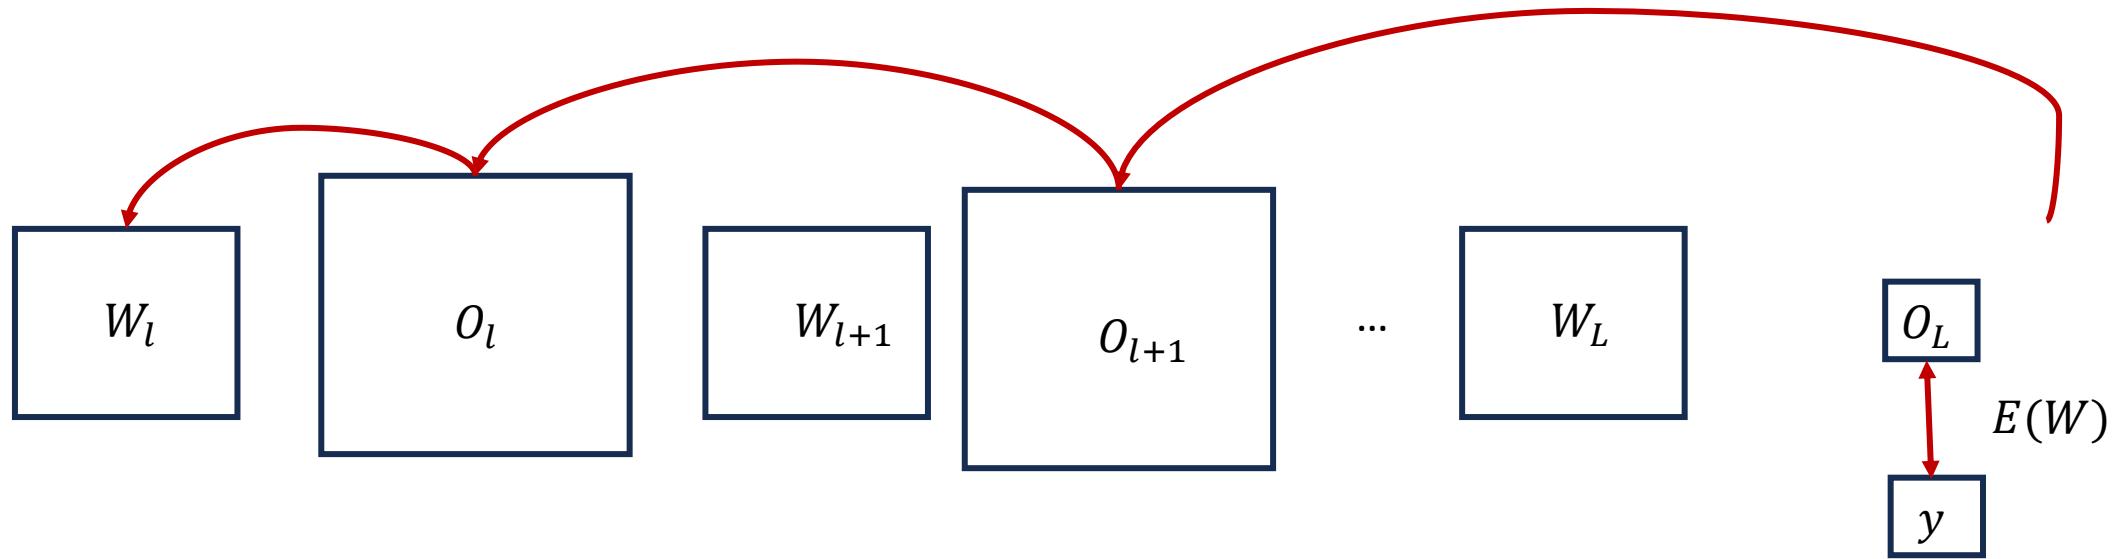
# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta O_2}{\delta O_1} \times \frac{\delta O_3}{\delta O_2} \times \cdots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

For a single image

# CNN Training: Backward Propagation

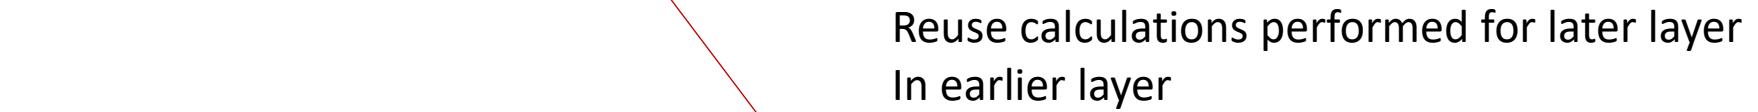


$$\frac{\delta E(W)}{\delta W_l} = \frac{\delta O_l}{\delta W_l} \times \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

For a single image

# CNN Training: Backward Propagation

$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta O_2}{\delta O_1} \times \frac{\delta O_3}{\delta O_2} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

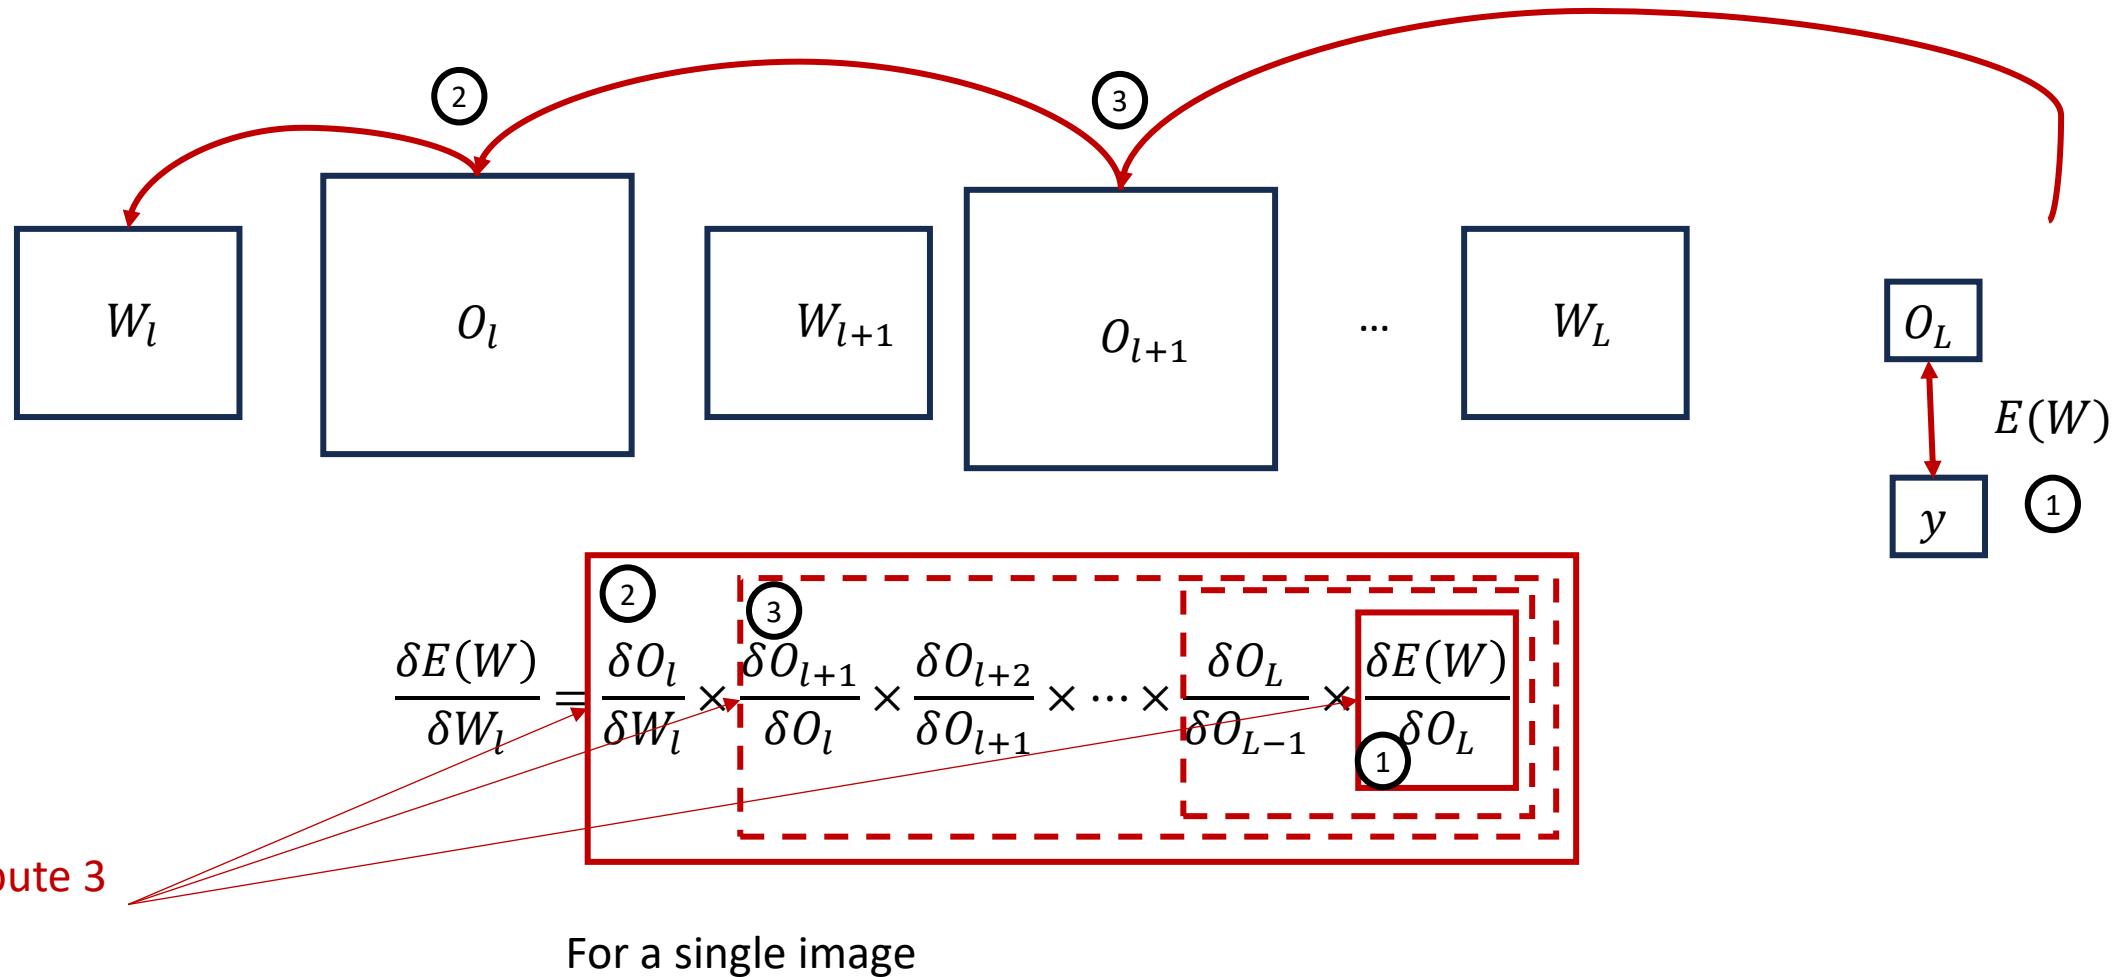


A red arrow points from the term  $\frac{\delta O_L}{\delta O_{L-1}}$  in the first equation to a bracket below, indicating that this calculation is reused from the previous layer's error term.

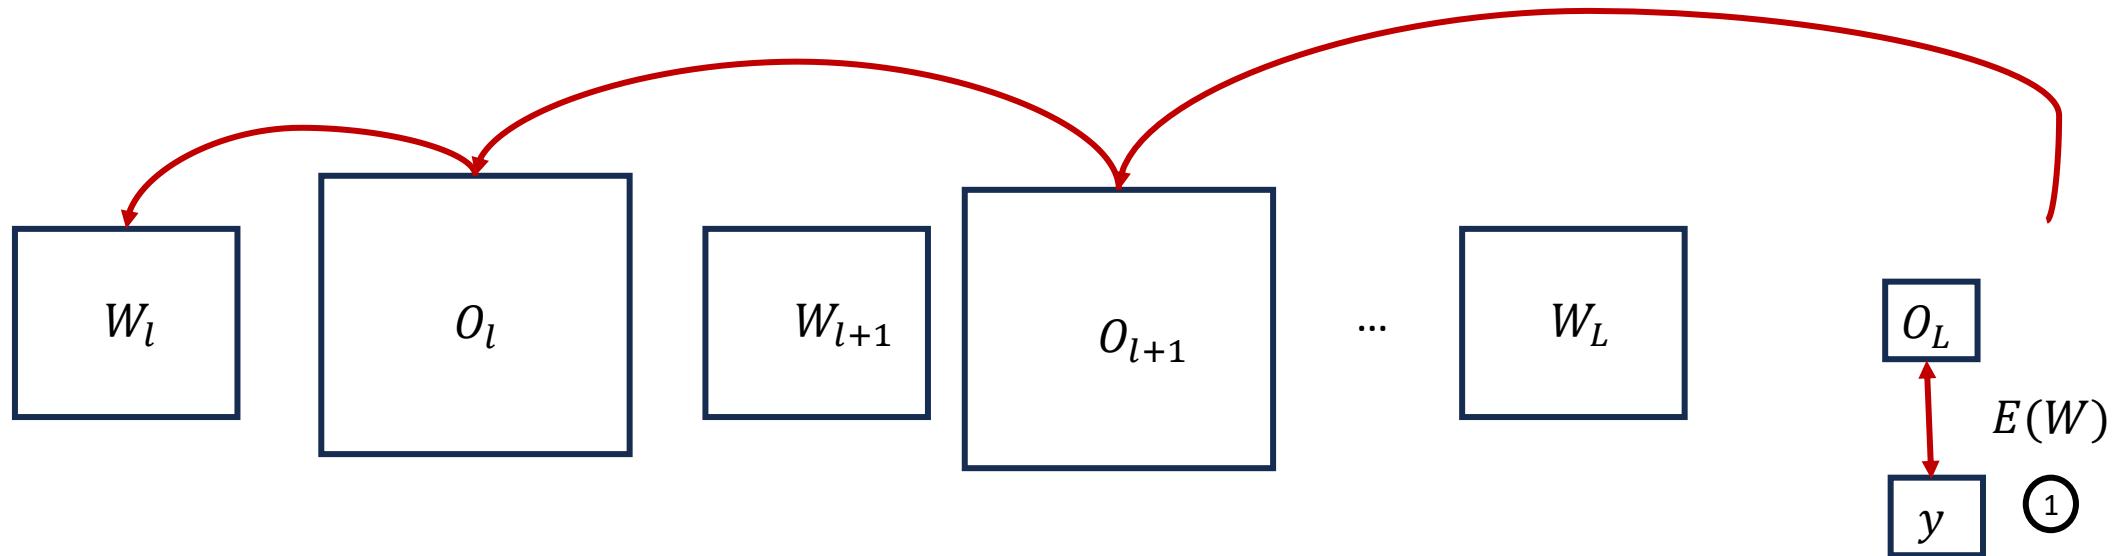
Reuse calculations performed for later layer  
In earlier layer

$$\frac{\delta E(W)}{\delta W_l} = \frac{\delta O_l}{\delta W_l} \times \underbrace{\frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}}_{\text{Reuse calculations performed for later layer in earlier layer}}$$

# CNN Training: Backward Propagation



# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_l} = \frac{\delta O_l}{\delta W_l} \times \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

We need to compute 3 types of terms

For a single image

# CNN Training: Backward Propagation

- $\frac{\delta E(W)}{\delta O_L}$  ???

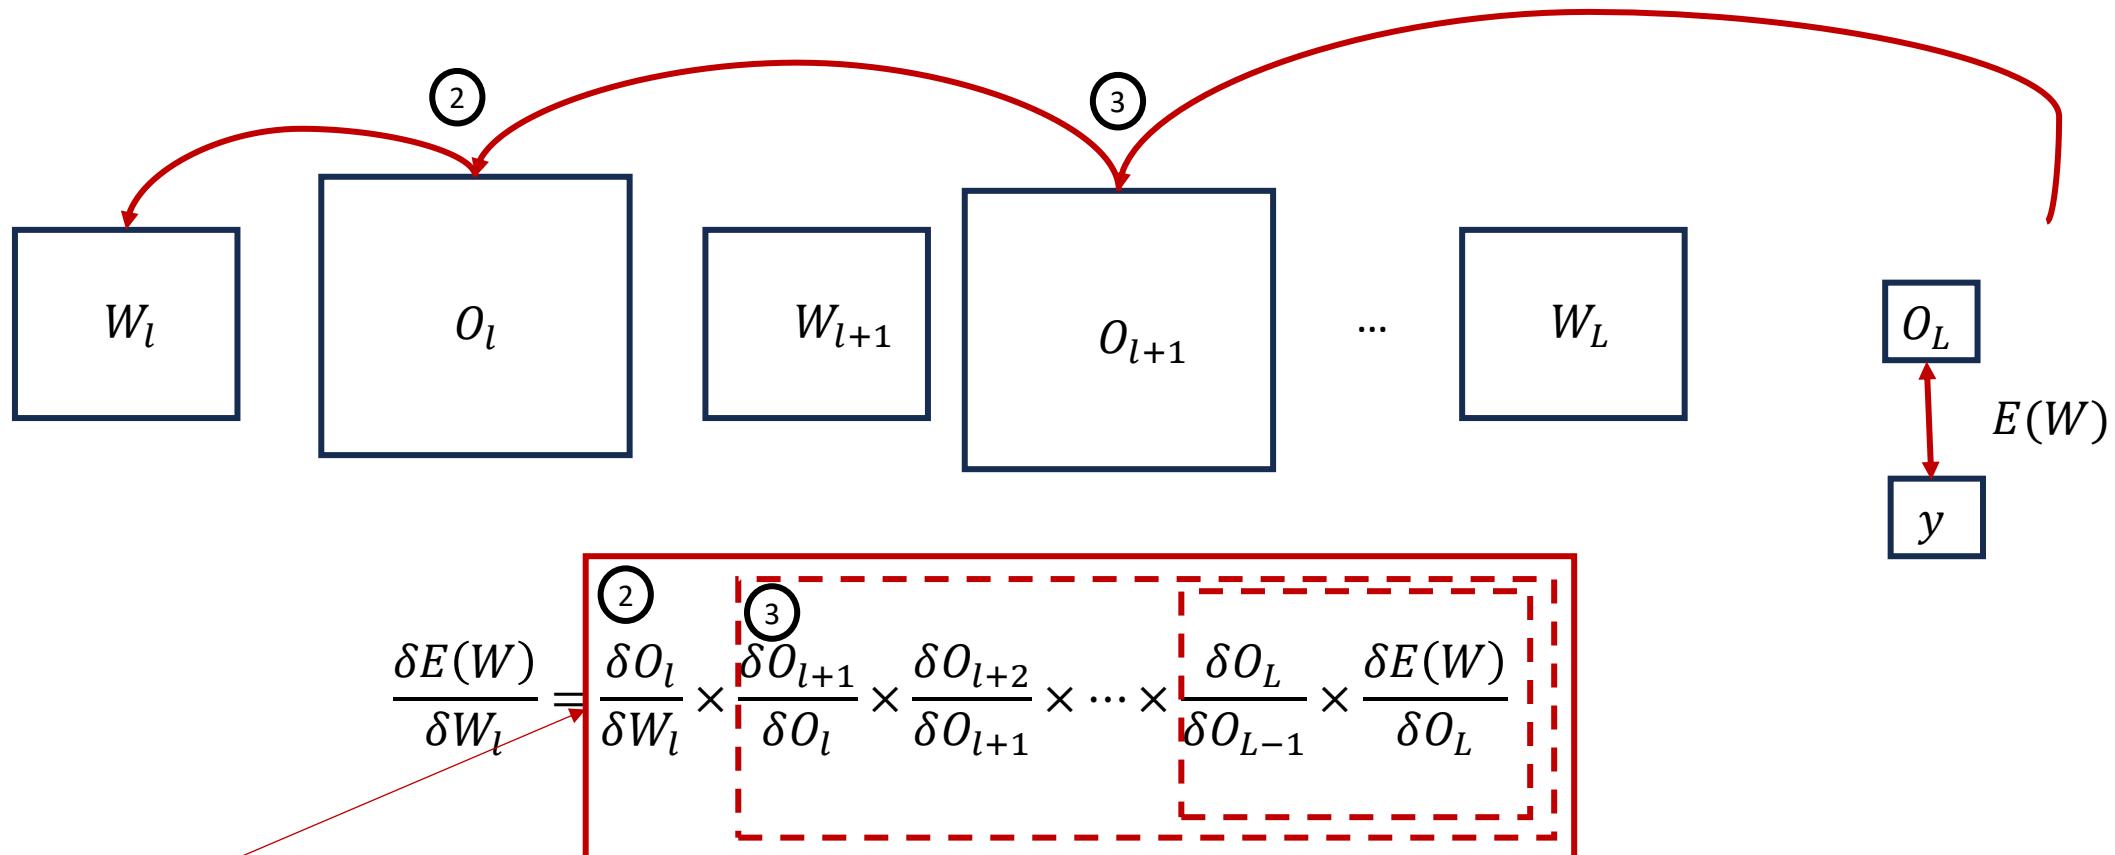
# CNN Training: Backward Propagation

- $\frac{\delta E(W)}{\delta o_L} = \frac{\delta}{\delta o_L} (y_i - o_L)^2 = ??$

# CNN Training: Backward Propagation

- $\frac{\delta E(W)}{\delta o_L} = \frac{\delta}{\delta o_L} (y_i - o_L)^2 = -2 * Error_i$
- For  $i$  image

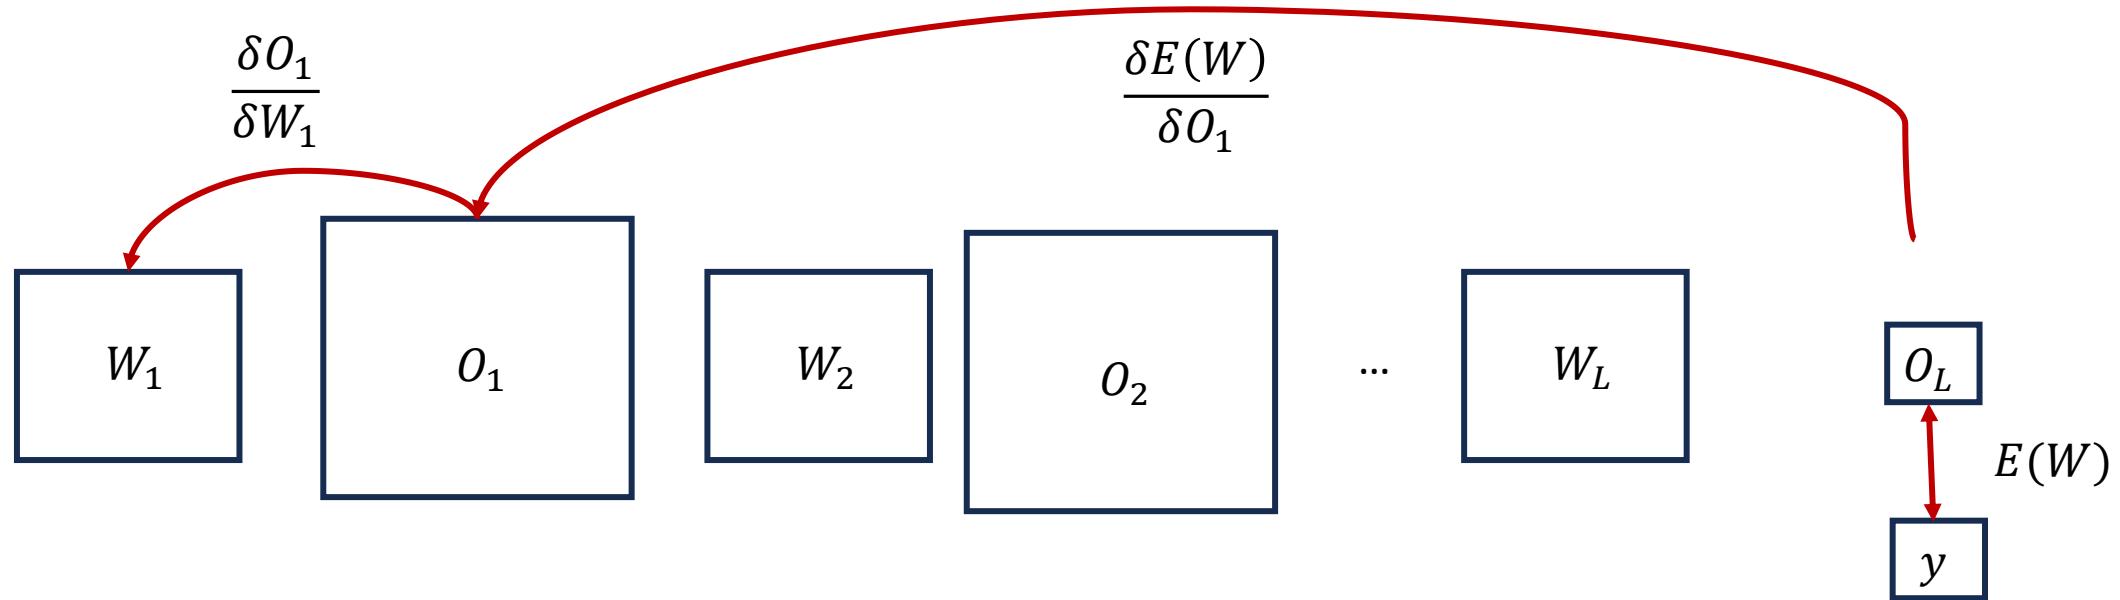
# CNN Training: Backward Propagation



We need to compute 3 types of terms

For a single image

# CNN Training: Gradient Calculations



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta E(W)}{\delta O_1}$$

# CNN Training: Gradient Calculations

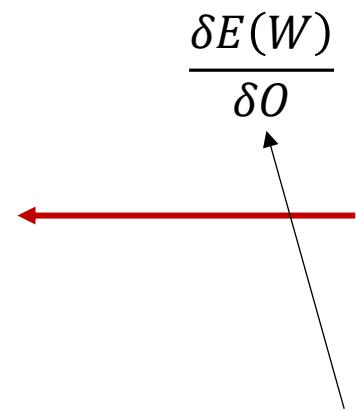
$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$

$W_1$	$W_2$
$W_3$	$W_4$

$$\frac{\delta O}{\delta W}$$

$O_1$	$O_2$	$O_3$
$O_4$	$O_5$	$O_6$
$O_7$	$O_8$	$O_9$

$$\frac{\delta E(W)}{\delta O}$$



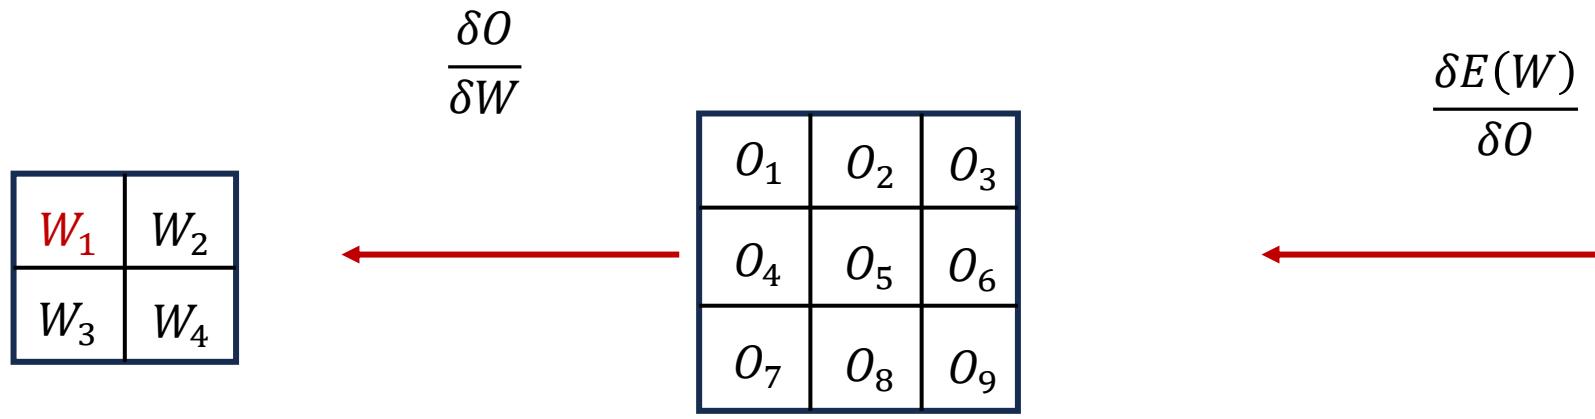
$$\frac{\delta E(W)}{\delta W} = \frac{\delta O}{\delta W} \times \frac{\delta E(W)}{\delta O}$$

For a layer  $l$ , removing the subscript  $l$  for clarity

Outputs of the previous layer that we produced in Forward Propagation

# CNN Training: Gradient Calculations

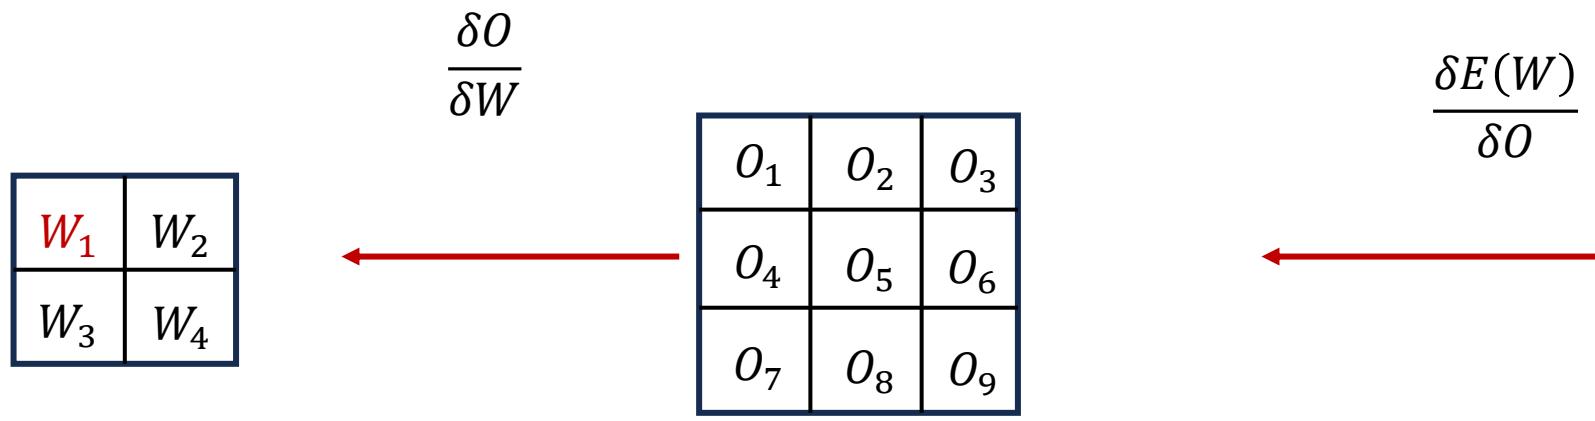
$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O}{\delta W_1} \times \frac{\delta E(W)}{\delta O}$$

# CNN Training: Gradient Calculations

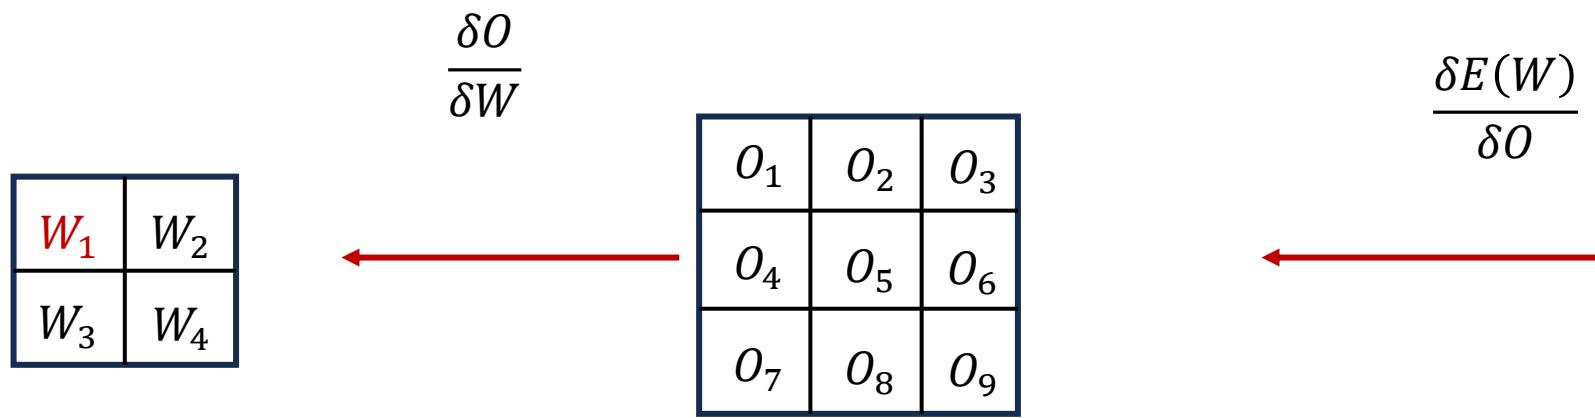
$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta E(W)}{\delta O_1} + \frac{\delta O_2}{\delta W_1} \times \frac{\delta E(W)}{\delta O_2} + \frac{\delta O_3}{\delta W_1} \times \frac{\delta E(W)}{\delta O_3} \dots + \frac{\delta O_9}{\delta W_1} \times \frac{\delta E(W)}{\delta O_9}$$

# CNN Training: Gradient Calculations

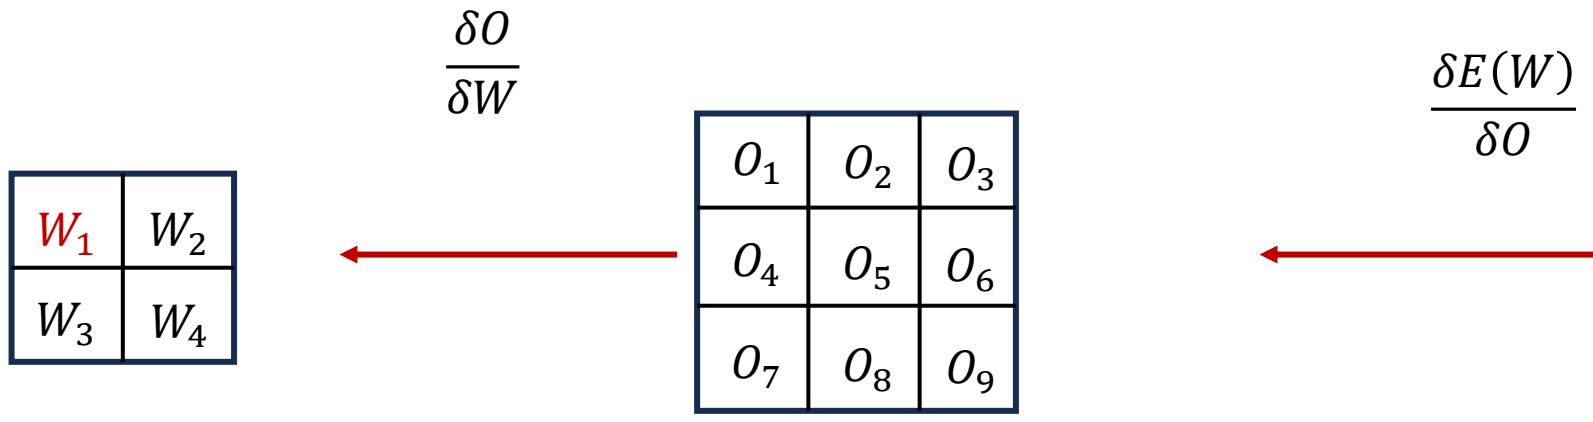
$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta E(W)}{\delta O_1} + \frac{\delta O_2}{\delta W_1} \times \frac{\delta E(W)}{\delta O_2} + \frac{\delta O_3}{\delta W_1} \times \frac{\delta E(W)}{\delta O_3} \dots + \frac{\delta O_9}{\delta W_1} \times \frac{\delta E(W)}{\delta O_9}$$

# CNN Training: Gradient Calculations

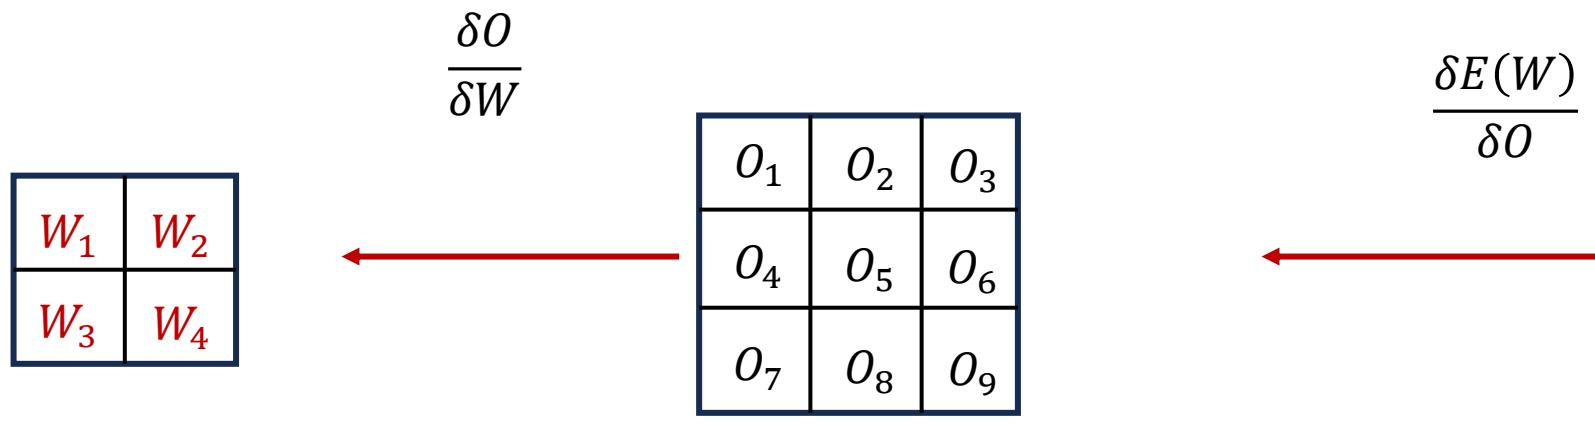
$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$



$$\frac{\delta E(W)}{\delta \textcolor{red}{W}_1} = I_1 \times \frac{\delta E(W)}{\delta O_1} + I_2 \times \frac{\delta E(W)}{\delta O_2} + I_3 \times \frac{\delta E(W)}{\delta O_3} \dots + I_{11} \times \frac{\delta E(W)}{\delta O_9}$$

# CNN Training: Gradient Calculations

$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$



$$\frac{\delta E(W)}{\delta \textcolor{red}{W}_1} = I_1 \times \frac{\delta E(W)}{\delta O_1} + I_2 \times \frac{\delta E(W)}{\delta O_2} + I_3 \times \frac{\delta E(W)}{\delta O_3} \dots + I_{11} \times \frac{\delta E(W)}{\delta O_9}$$

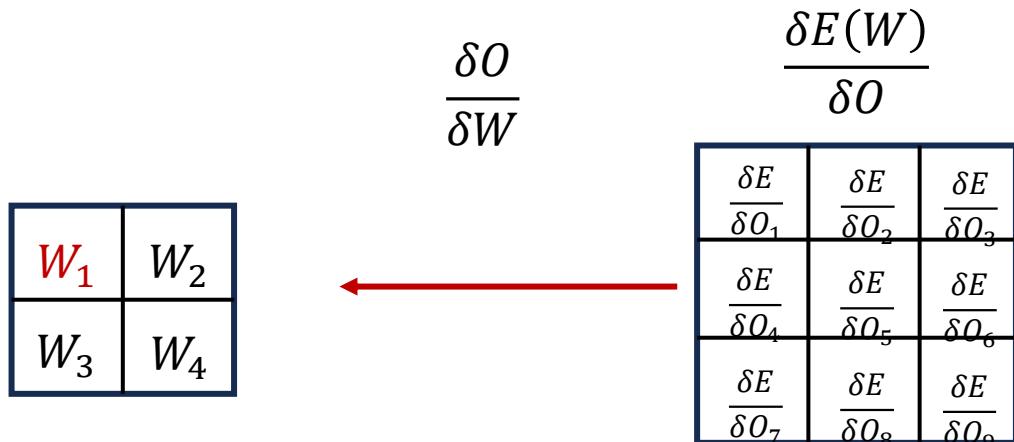
$$\frac{\delta E(W)}{\delta \textcolor{red}{W}_2} = I_2 \times \frac{\delta E(W)}{\delta O_1} + I_3 \times \frac{\delta E(W)}{\delta O_2} + I_6 \times \frac{\delta E(W)}{\delta O_3} \dots + I_{12} \times \frac{\delta E(W)}{\delta O_9}$$

$$\frac{\delta E(W)}{\delta \textcolor{red}{W}_3} = I_5 \times \frac{\delta E(W)}{\delta O_1} + I_6 \times \frac{\delta E(W)}{\delta O_2} + I_7 \times \frac{\delta E(W)}{\delta O_3} \dots + I_{13} \times \frac{\delta E(W)}{\delta O_9}$$

...

# CNN Training: Gradient Calculations

$I_1$	$I_2$	$I_3$	$I_4$
$I_5$	$I_6$	$I_7$	$I_8$
$I_9$	$I_{10}$	$I_{11}$	$I_{12}$
$I_{13}$	$I_{14}$	$I_{15}$	$I_{16}$

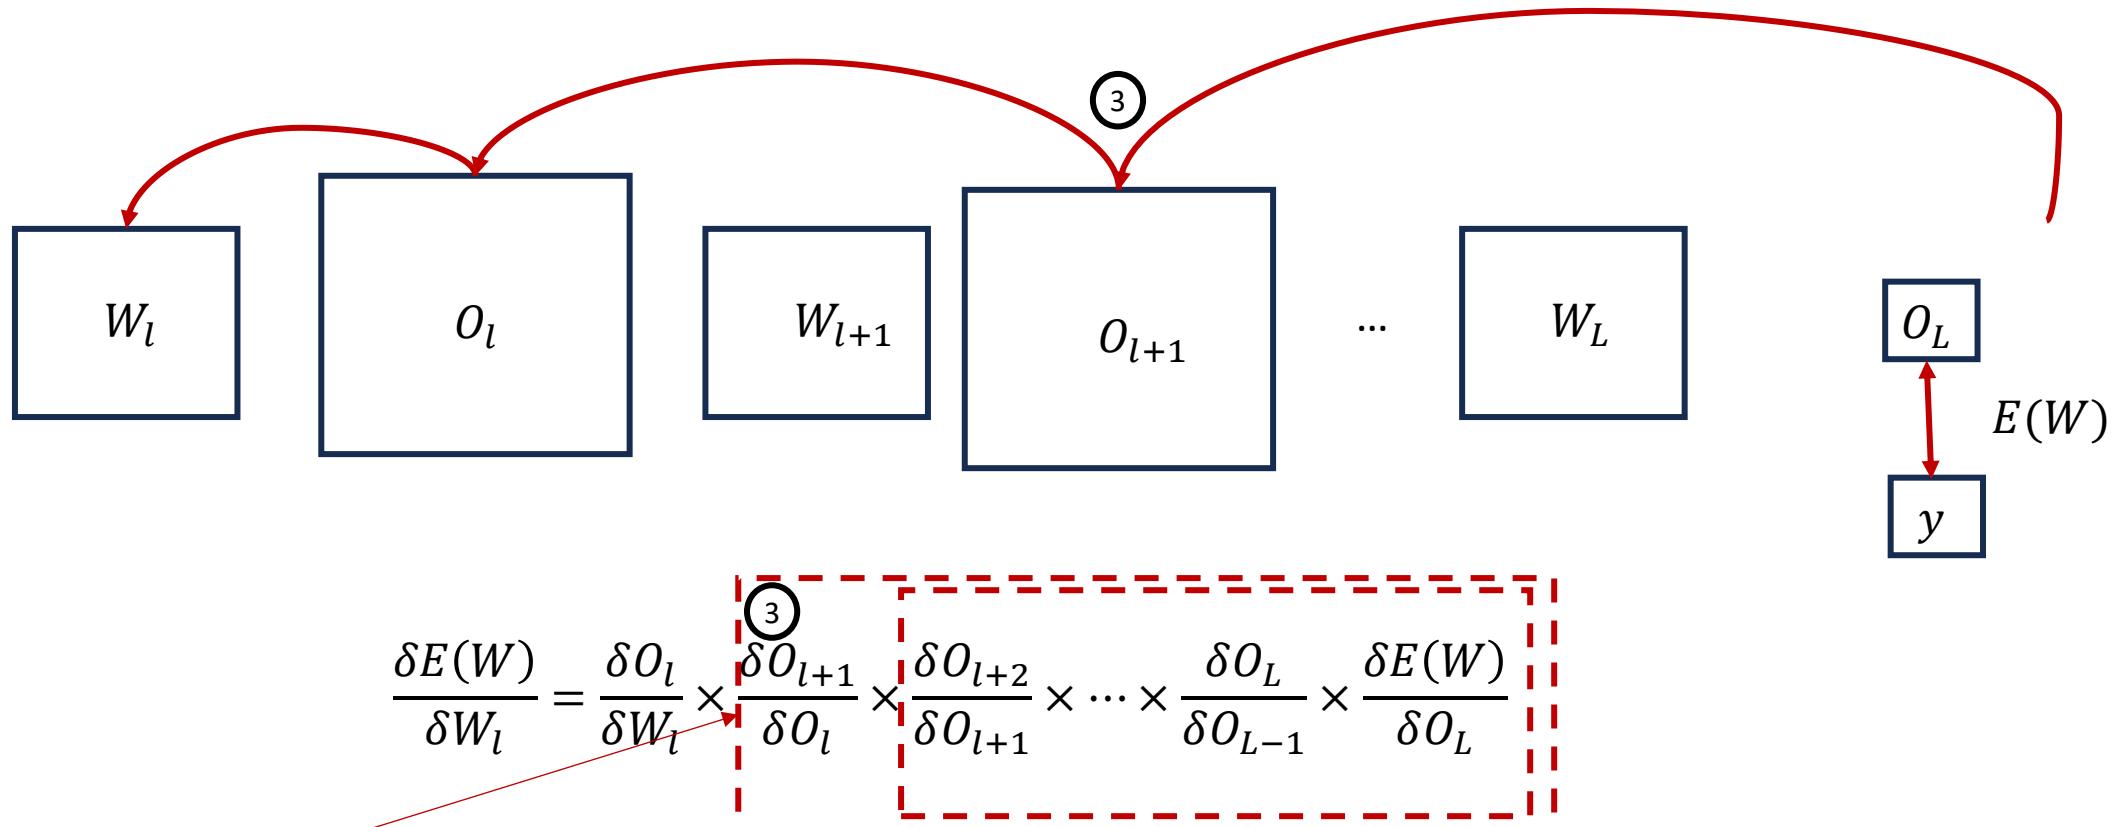


$$\frac{\delta E(W)}{\delta W} = \text{Conv}(I, \frac{\delta E(W)}{\delta O})$$

Note: Only inputs  $I$  and gradients  $\frac{\delta E(W)}{\delta O}$  are needed. We don't need the output  $O$  for this layer. Why?

Note: This operation where the output dimension is greater than input dimension is called deconvolution or transposed convolution

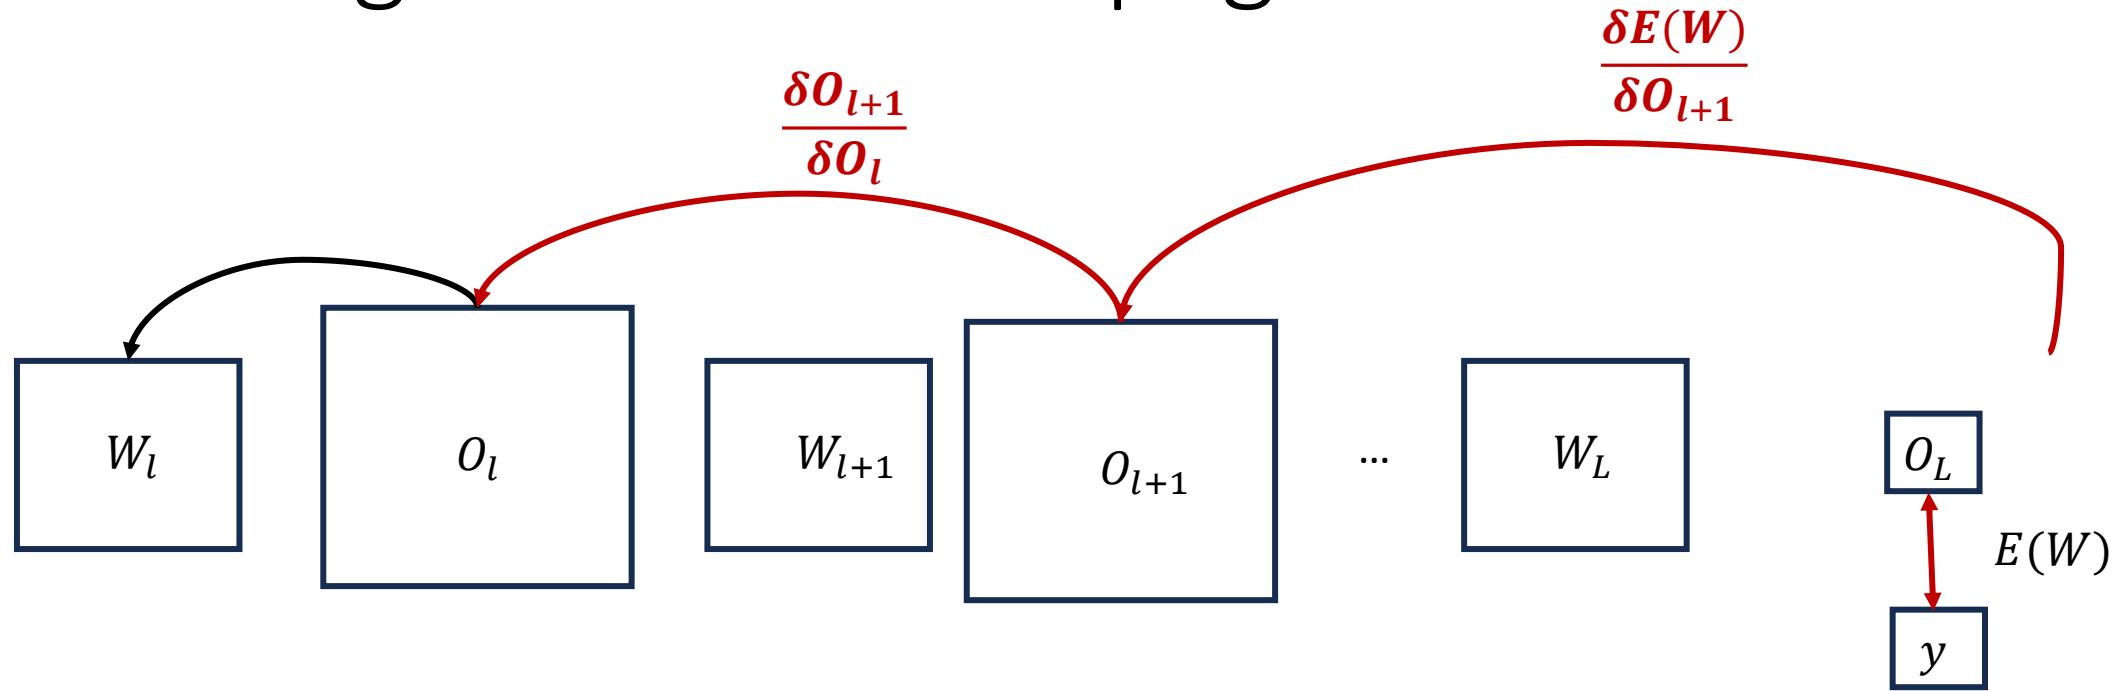
# CNN Training: Backward Propagation



We need to compute 3 types of terms

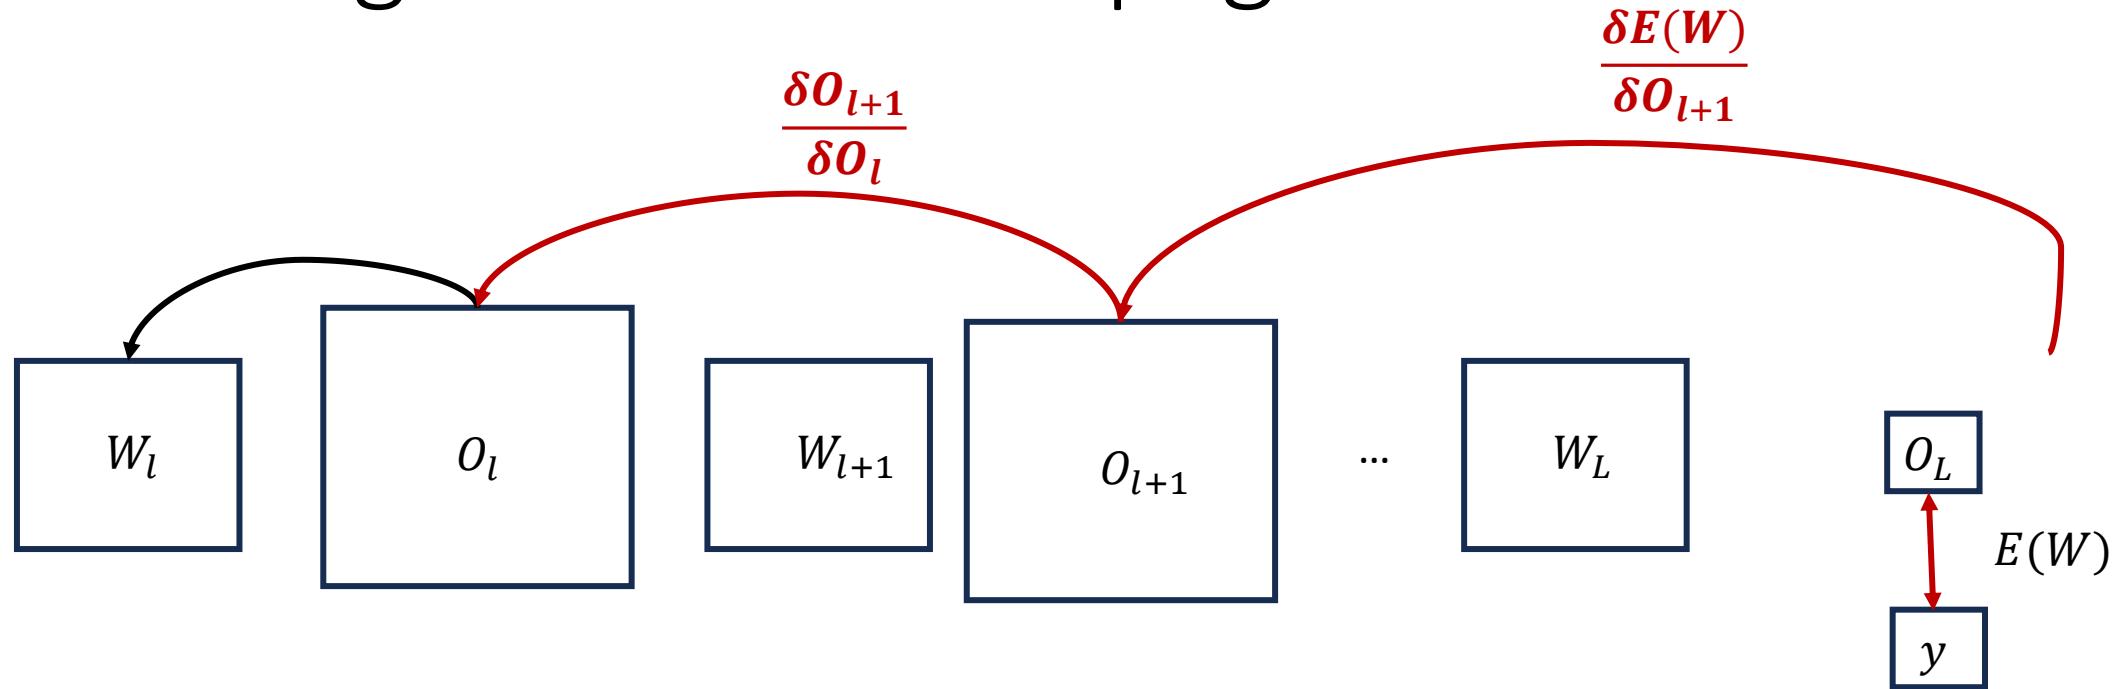
For a single image

# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta O_l} = \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta o_l} = \frac{\delta o_{l+1}}{\delta o_l} \times \frac{\delta E(W)}{\delta o_{l+1}}$$

# CNN Training: Backward Propagation

$O_l$	$O_1$	$O_2$	$O_3$	$O_4$
	$O_5$	$O_6$	$O_7$	$O_8$
	$O_9$	$O_{10}$	$O_{11}$	$O_{12}$
	$O_{13}$	$O_{14}$	$O_{15}$	$O_{16}$

$W_1$	$W_2$
$W_3$	$W_4$

$O_{l+1}$

$\frac{\delta E}{\delta O_1}$	$\frac{\delta E}{\delta O_2}$	$\frac{\delta E}{\delta O_3}$
$\frac{\delta E}{\delta O_4}$	$\frac{\delta E}{\delta O_5}$	$\frac{\delta E}{\delta O_6}$
$\frac{\delta E}{\delta O_7}$	$\frac{\delta E}{\delta O_8}$	$\frac{\delta E}{\delta O_9}$

What all do we need for the computation? Do we need inputs or outputs to the layer?

$$\frac{\delta E(W)}{\delta O_l} = \text{Conv}(W_{180}, \frac{\delta E(W)}{\delta O_{l+1}})$$

HW:  $W_{180}$  –  $W$  rotated by 180 degrees

# Key Operation in CNN

- Forward Propagation:

$$O_l = \textcolor{red}{Conv}(W_l, I_l)$$

- Backward Propagation

$$\frac{\delta E(W)}{\delta W_l} = \textcolor{red}{Conv}(I_l, \frac{\delta E(W)}{\delta O_l})$$

$$\frac{\delta E(W)}{\delta O_l} = \textcolor{red}{Conv}(W_{180}, \frac{\delta E(W)}{\delta O_{l+1}})$$

# Training

$$E(W): \min_W \sum (y_i - F(x_i: W))^2$$

$$W_{t+1} \leftarrow W_t - \alpha \frac{\delta E(W)}{\delta W}$$

Iteratively

$$\frac{\delta E(W)}{\delta W} = 2 \sum \underbrace{(y_i - F(x_i: W))}_{\text{Error}} \times \frac{\delta F(x_i: W)}{\delta W}$$

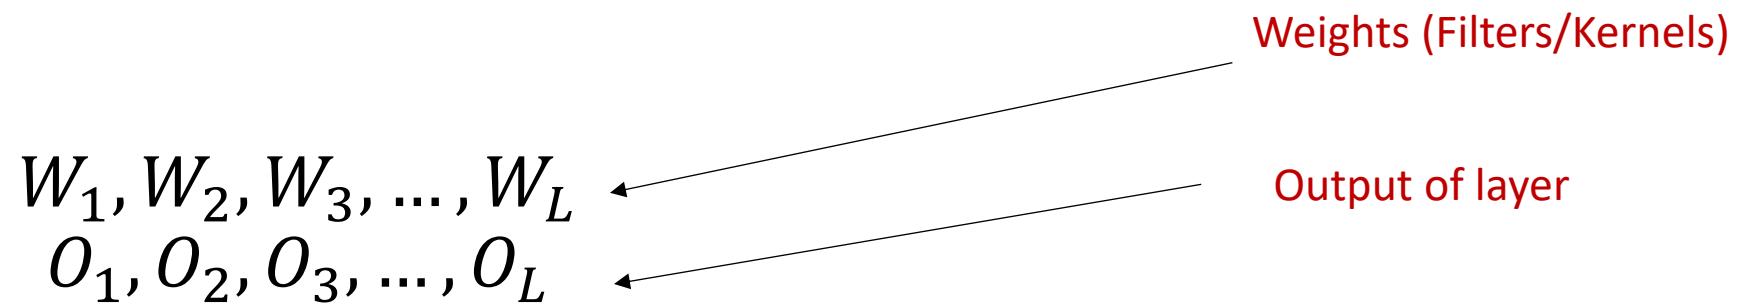
For all samples, in each iteration

For multiple images in a batch: Simply sum up the gradients

$$\frac{\delta E(W)}{\delta W_l} = \boxed{\frac{\delta O_l}{\delta W_l} \times \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}}} \times \boxed{\frac{\delta E(W)}{\delta O_L}}$$

# Training

$L$  layers in CNN



*2 \* Error \**   $\frac{\delta F(x_i: W_{1:L})}{\delta W_1}, \frac{\delta F(x_i: W_{1:L})}{\delta W_2}, \dots, \frac{\delta F(x_i: W_{1:L})}{\delta W_L}$

# Training – Key Performance Objective

- Accuracy – We want to train a model that can predict with low error on unseen data
- Throughput – Given a large dataset, we want to train as quickly as possible

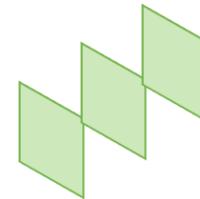
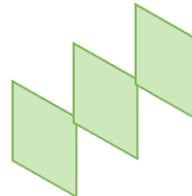
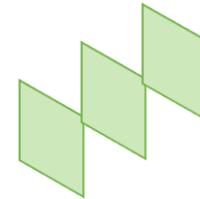
# Training – Key Requirement

- For each image,
- Need to store the inputs of each layer
  - Note: Input  $I_l$  of layer  $l$  is output (activation)  $O_{l-1}$  of layer  $l - 1$
- Need to have gradients  $\frac{\delta E(W)}{\delta W_l}, \frac{\delta E(W)}{\delta O_l}, \frac{\delta E(W)}{\delta O_{l+1}}$  for each layer  $l$
- Ungraded HW assignment: Can you calculate the difference in storage requirements between computing gradients from left to right vs right to left?

# Convolutional Neural Network

- **Training**

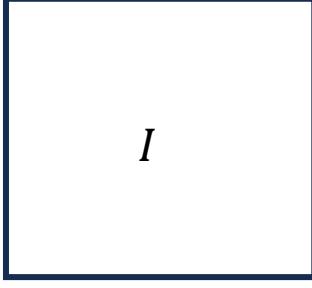
- Given  $x_i \in X$  (image from a database),  $y_i$ : Label for the image
- “Learn” a model  $F$ , defined by the filters  $W$  such that:
  - $\sum(y_i - F(x_i; W))^2$  is minimized



- **Inference**

- Given a new image  $x'$
- Predict a label  $y' = W(x')$
- No change in the filters

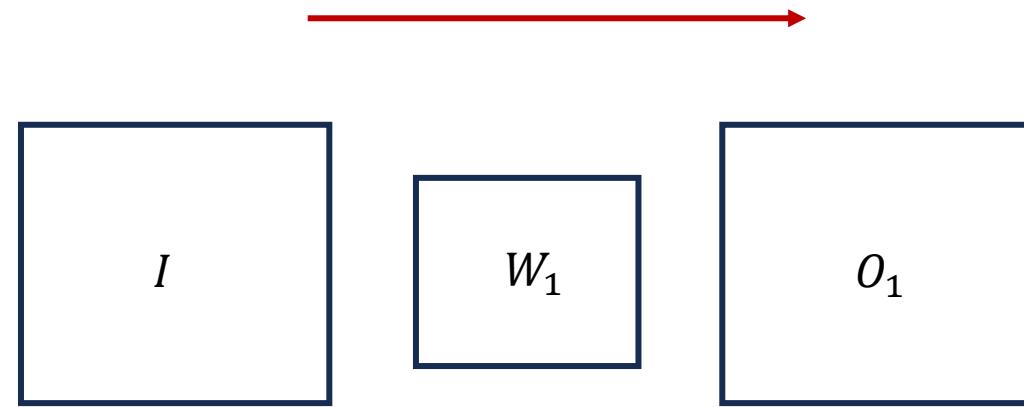
# CNN Training: Forward Propagation



*I*

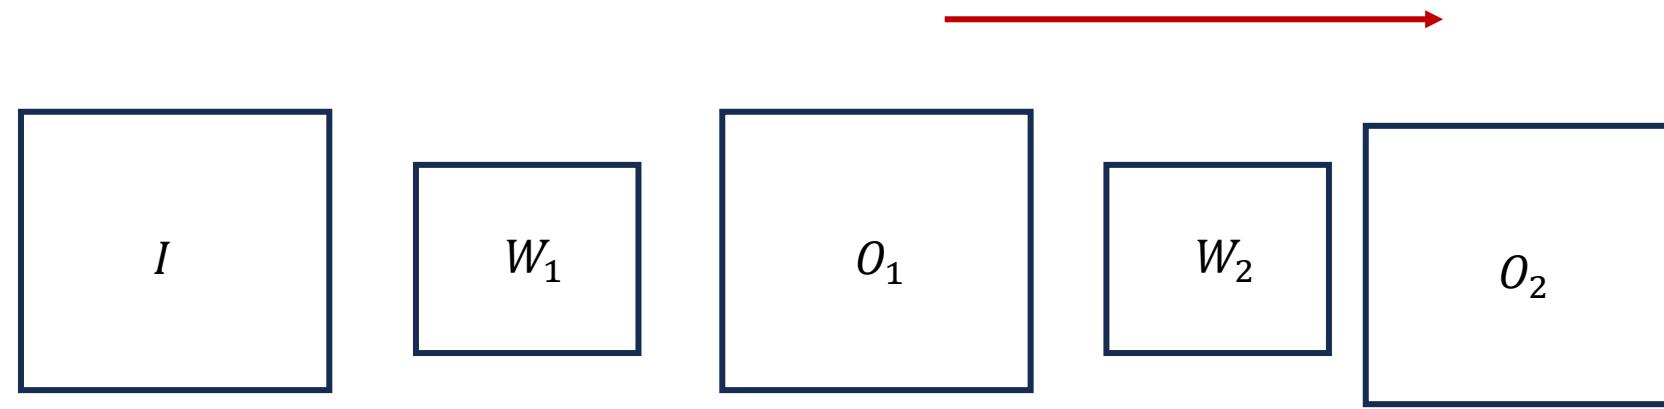
For a single image

# CNN Training: Forward Propagation



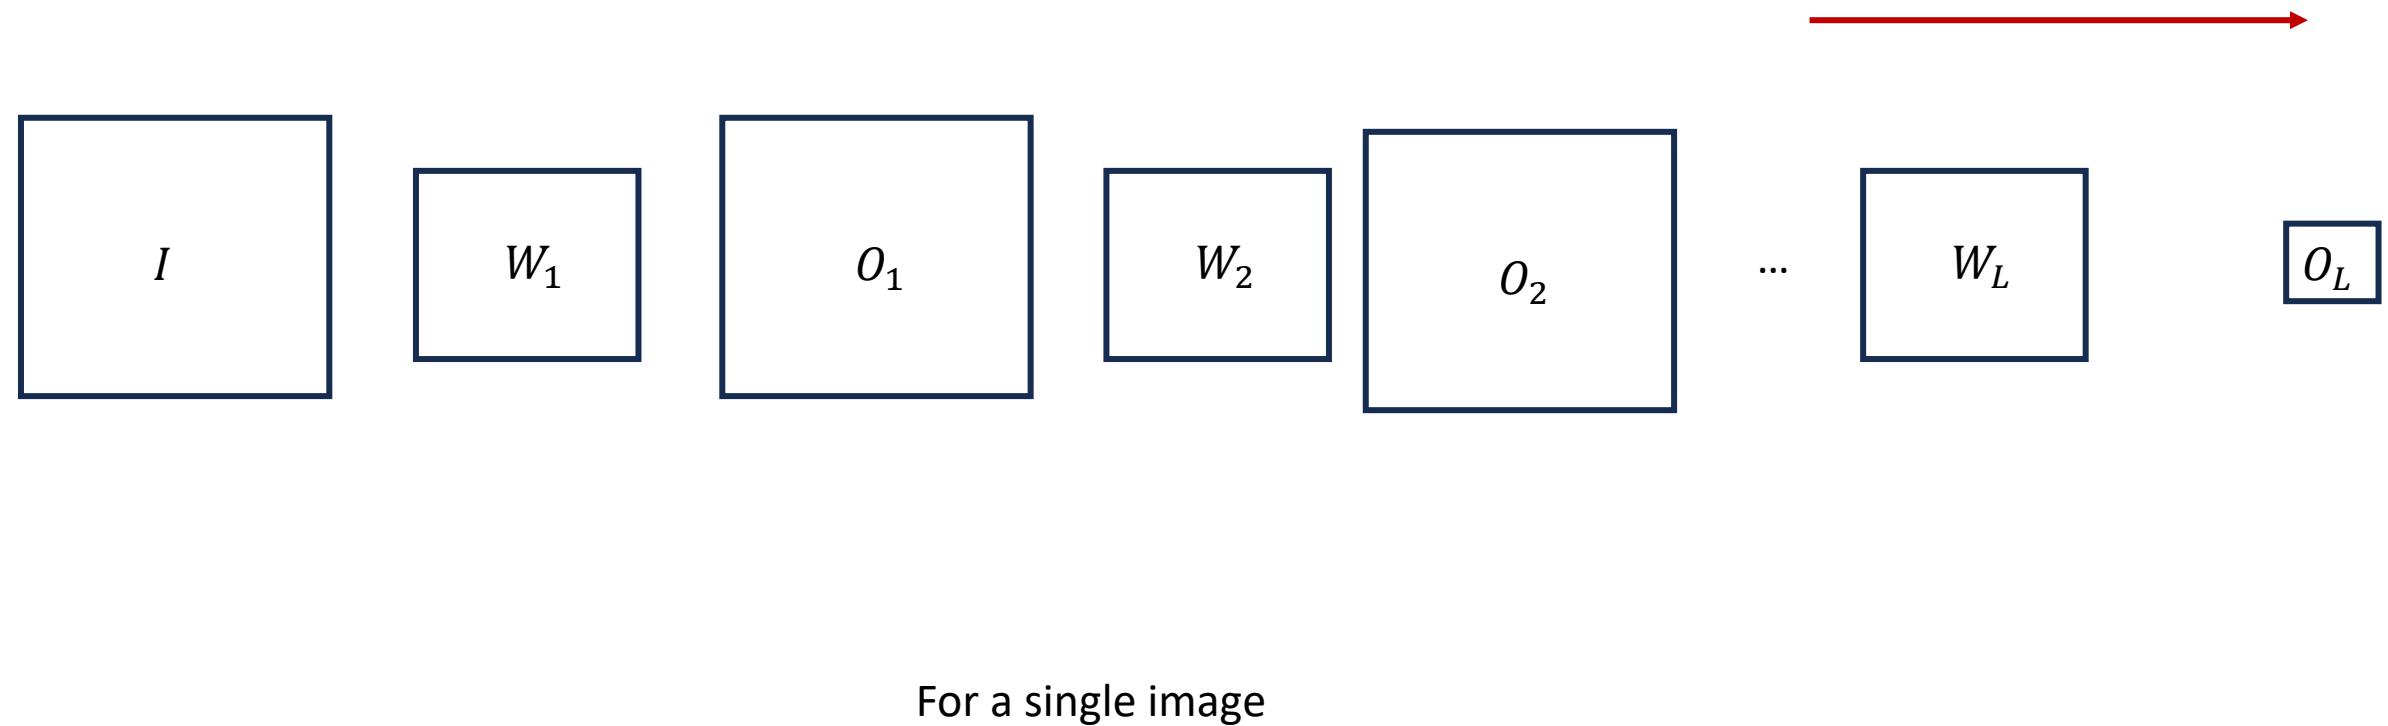
For a single image

# CNN Training: Forward Propagation



For a single image

# CNN Training: Forward Propagation



# Inference

- Simply the forward propagation portion
- No error calculation or backpropagation as ground truth doesn't exist
- Performance Objectives:
  - Accuracy – Although ground truth is not known
  - Latency – Perform inference as fast as possible on single or a batch of images

# Next Class

- 9/25 Lecture 10
  - Accelerating Convolutional Neural Networks: Convolution as Matrix Multiplication

# Thank You

- Questions?
- Email: [sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)