

CSDS 451: Designing High Performant Systems for AI

Lecture 12

10/7/2025

Sanmukh Kuppannagari

sanmukh.kuppannagari@case.edu

<https://sanmukh.research.st/>

Case Western Reserve University

Outline

- Pytorch Basics

Announcements

- Midterm this Thursday - 10/9
 - Will cover topics till the previous lecture
- I will upload WA1 and WA2 solutions by tonight

Outline

- Pytorch Basics

Training

$$E(W): \min_W \sum (y_i - F(x_i: W))^2$$

$$W_{t+1} \leftarrow W_t - \alpha \frac{\delta E(W)}{\delta W}$$

Iteratively

$$\frac{\delta E(W)}{\delta W} = 2 \sum_i (y_i - F(x_i: W)) \times \frac{\delta F(x_i: W)}{\delta W}$$

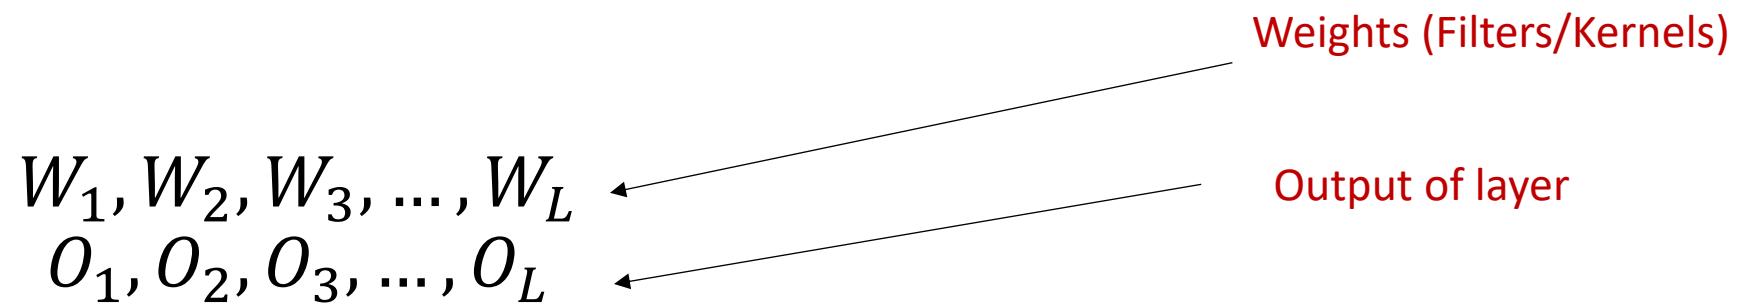
For all samples, in each iteration



Error

Training

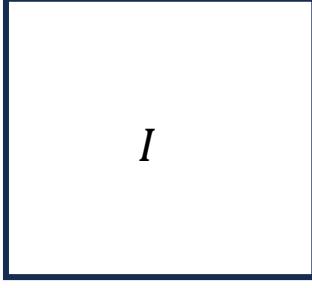
L layers in CNN



$$2 * \text{Error} * \frac{\delta F(x_i: W_{1:L})}{\delta W_1}, \frac{\delta F(x_i: W_{1:L})}{\delta W_2}, \dots, \frac{\delta F(x_i: W_{1:L})}{\delta W_L}$$

Compute these for
all the samples

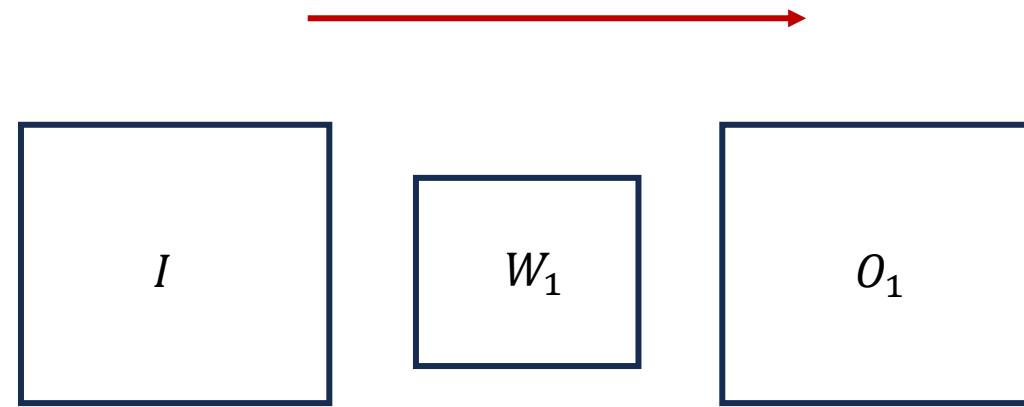
CNN Training: Forward Propagation



I

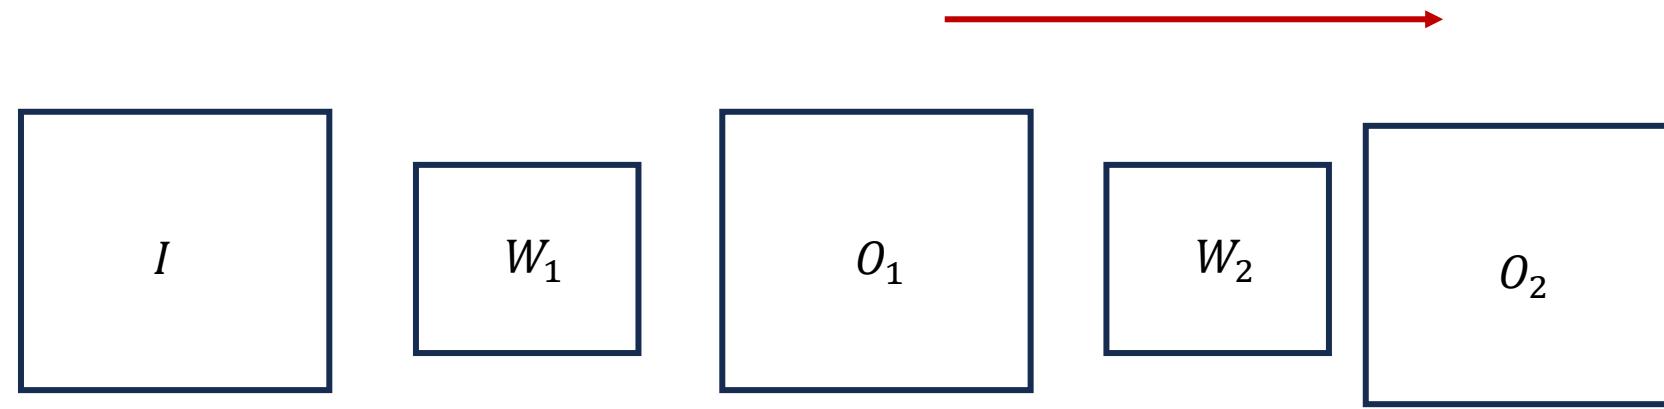
For a single image

CNN Training: Forward Propagation



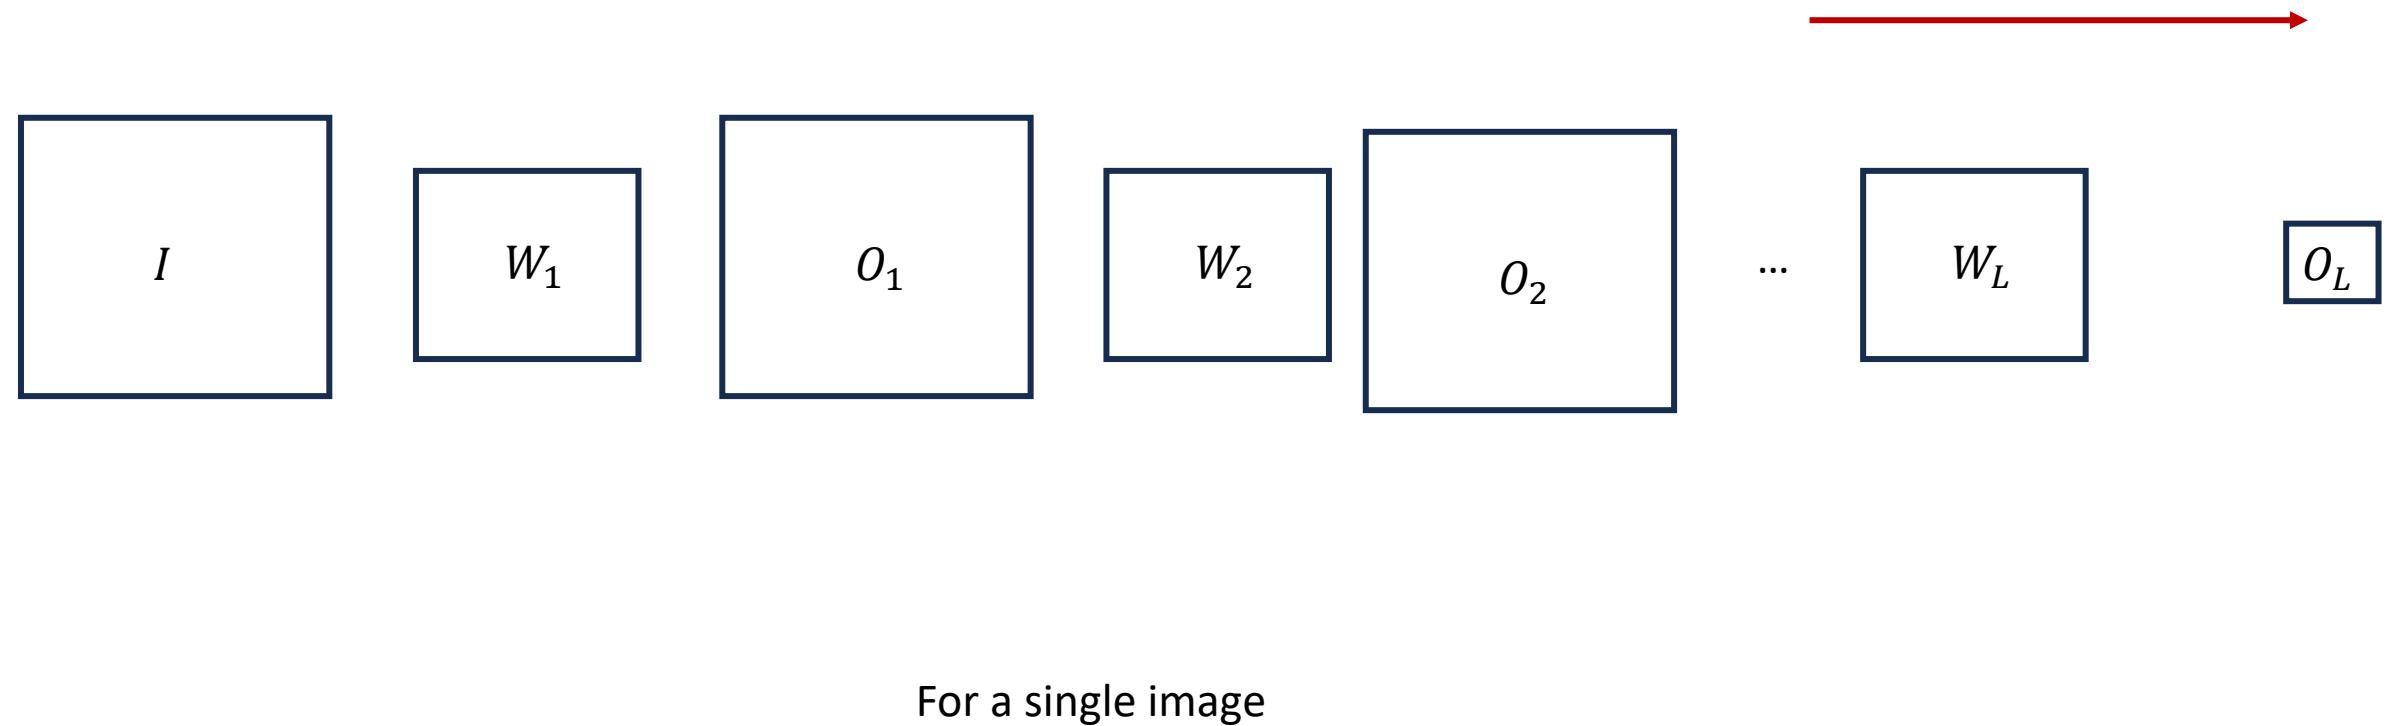
For a single image

CNN Training: Forward Propagation



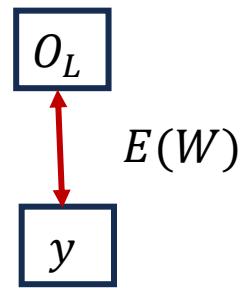
For a single image

CNN Training: Forward Propagation



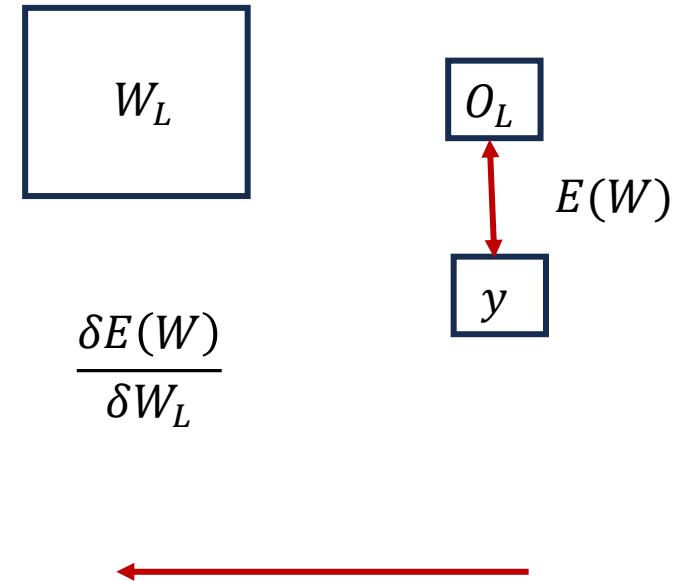
CNN Training: Error Calculation

For a single image

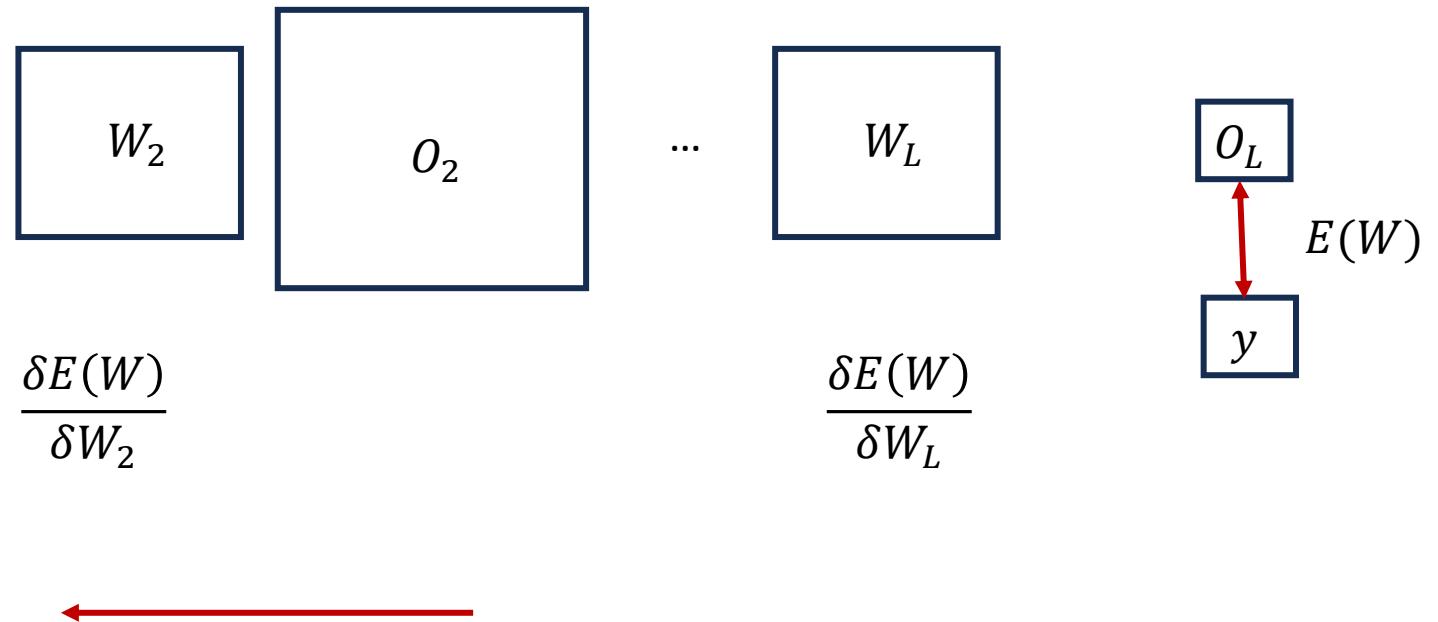


CNN Training: Backward Propagation

For a single image

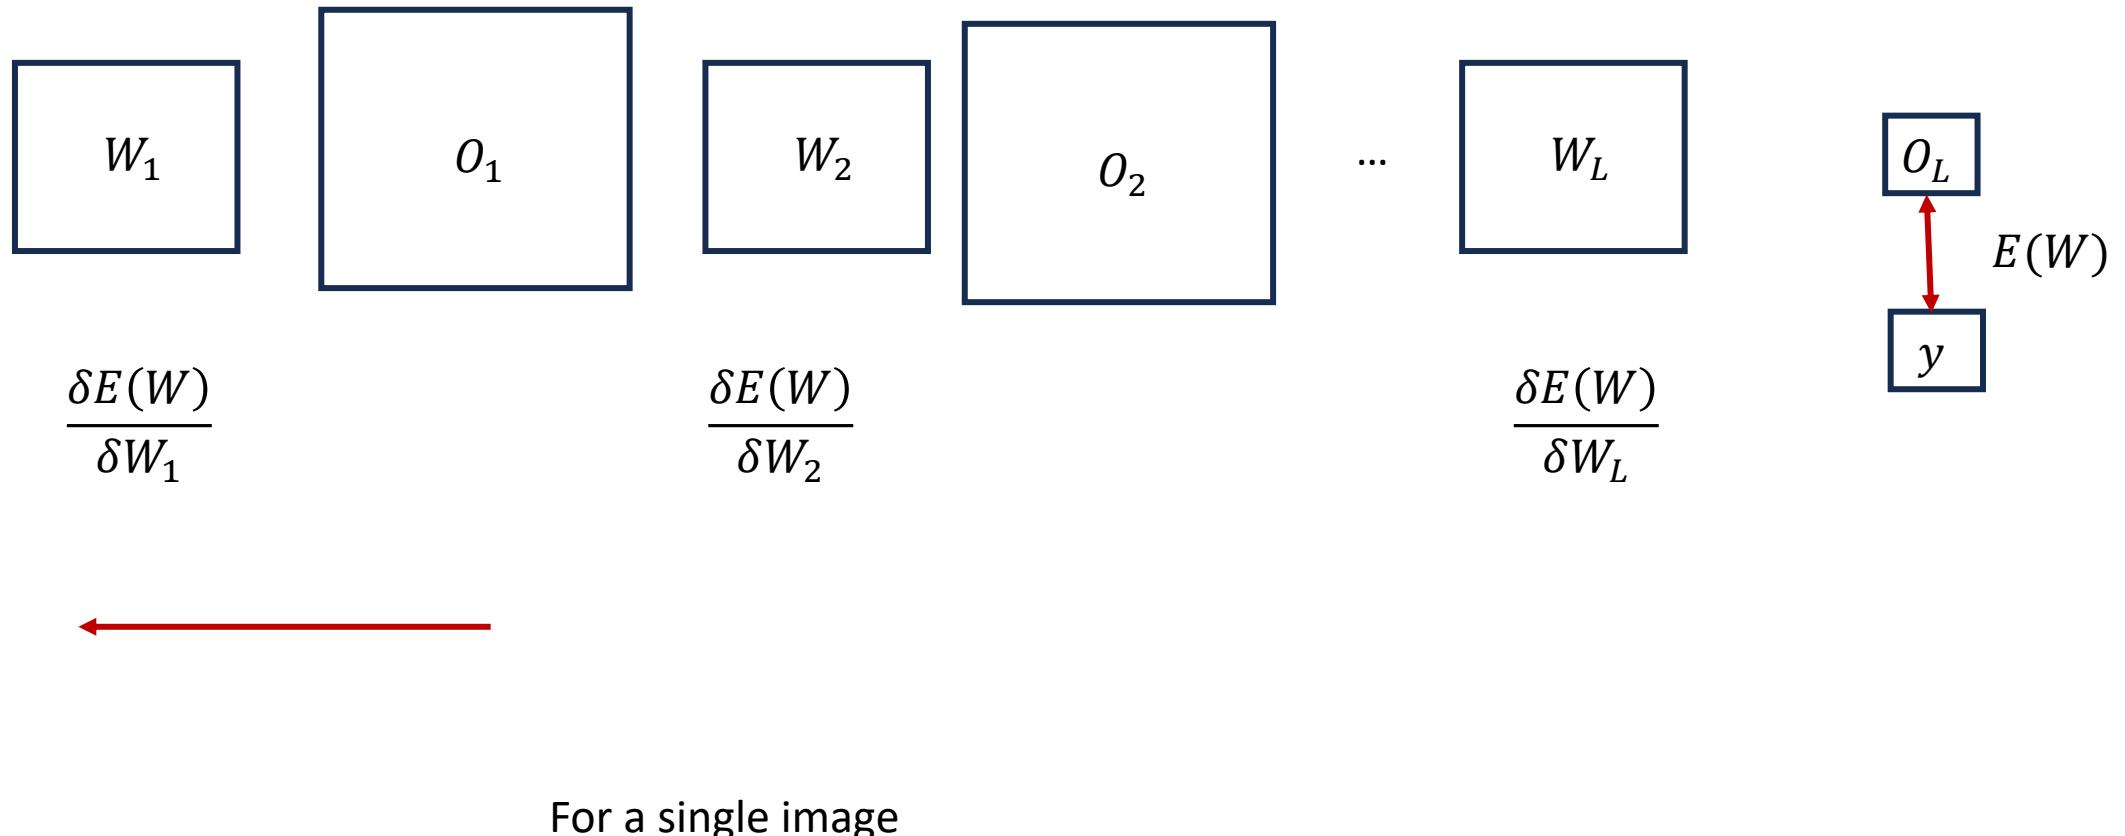


CNN Training: Backward Propagation

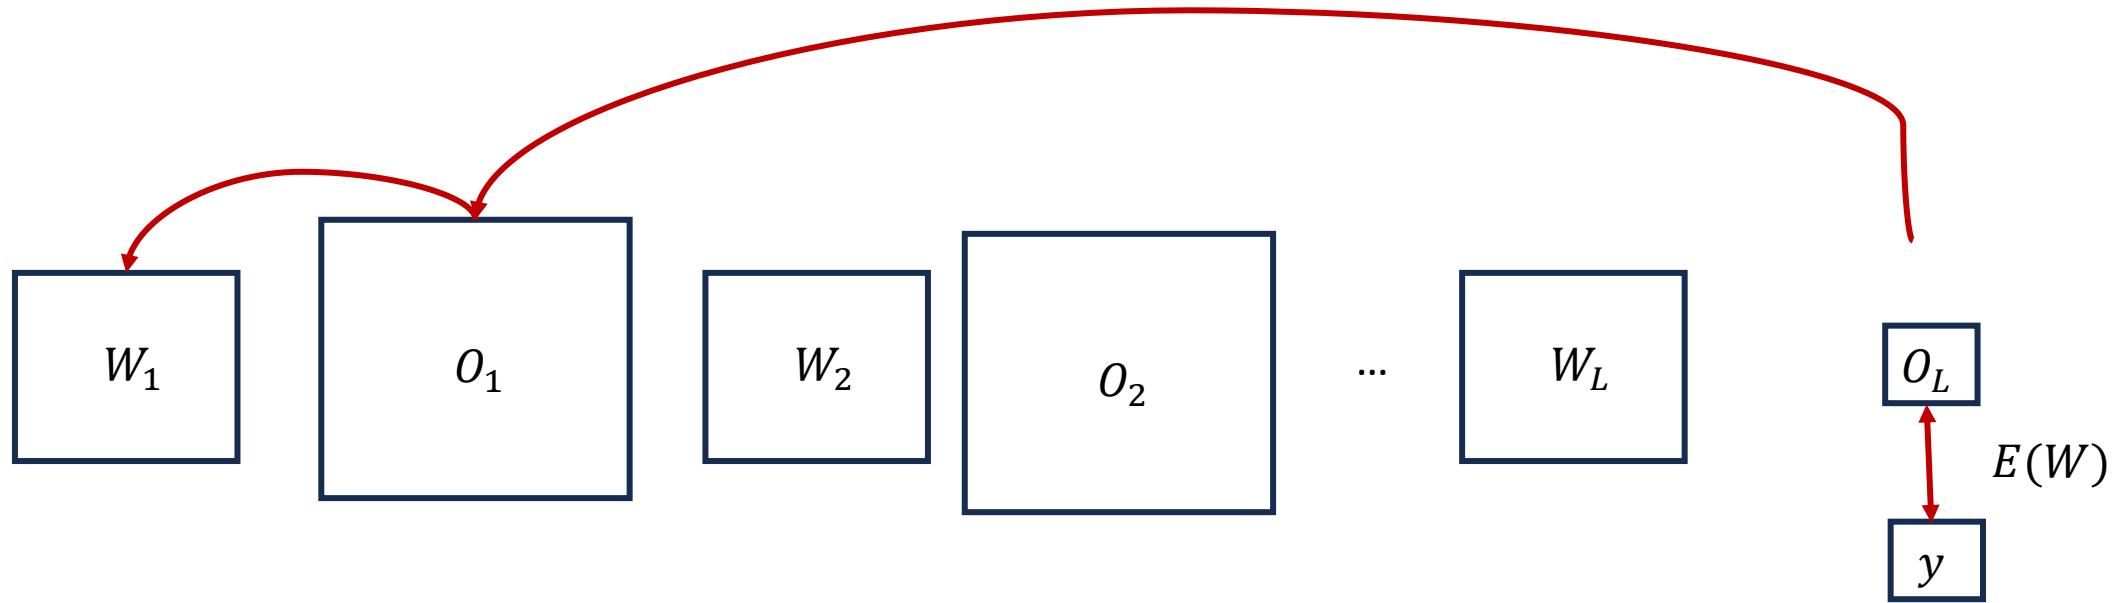


For a single image

CNN Training: Backward Propagation



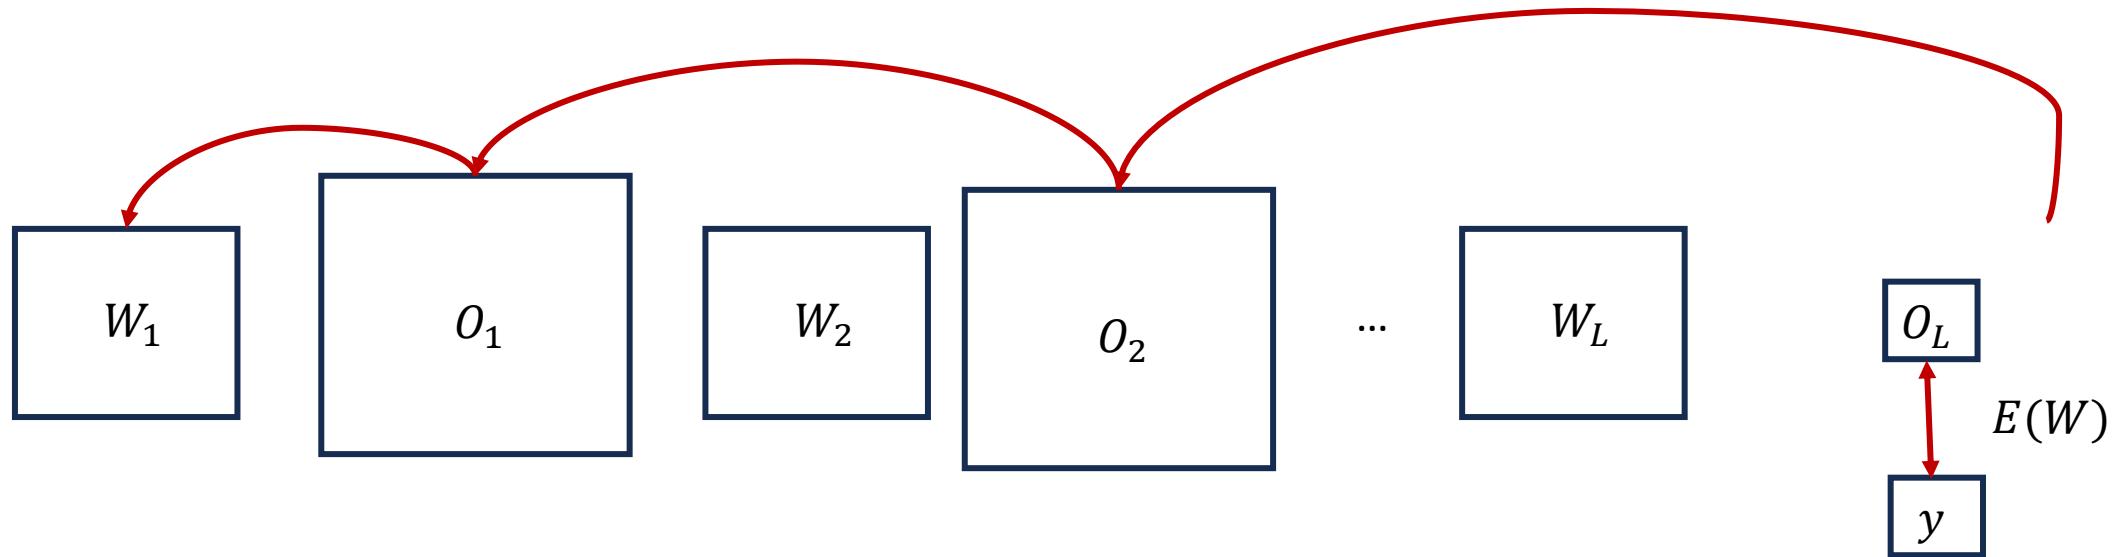
CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta E(W)}{\delta O_1}$$

For a single image

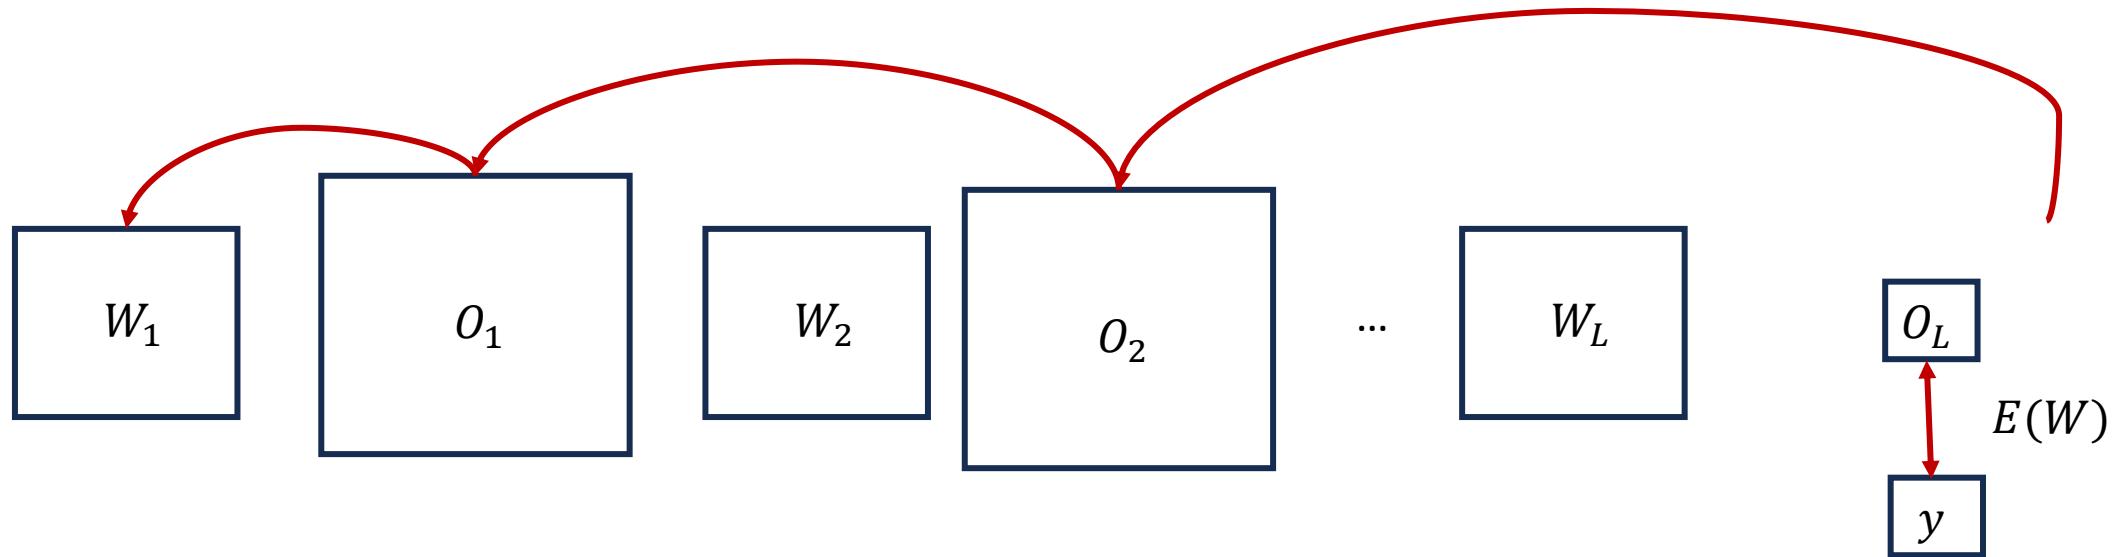
CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta O_2}{\delta O_1} \times \frac{\delta E(W)}{\delta O_2}$$

For a single image

CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_1} = \frac{\delta O_1}{\delta W_1} \times \frac{\delta O_2}{\delta O_1} \times \frac{\delta O_3}{\delta O_2} \times \cdots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

For a single image

Training

$$E(W): \min_W \sum (y_i - F(x_i: W))^2$$

$$W_{t+1} \leftarrow W_t - \alpha \frac{\delta E(W)}{\delta W}$$

Iteratively

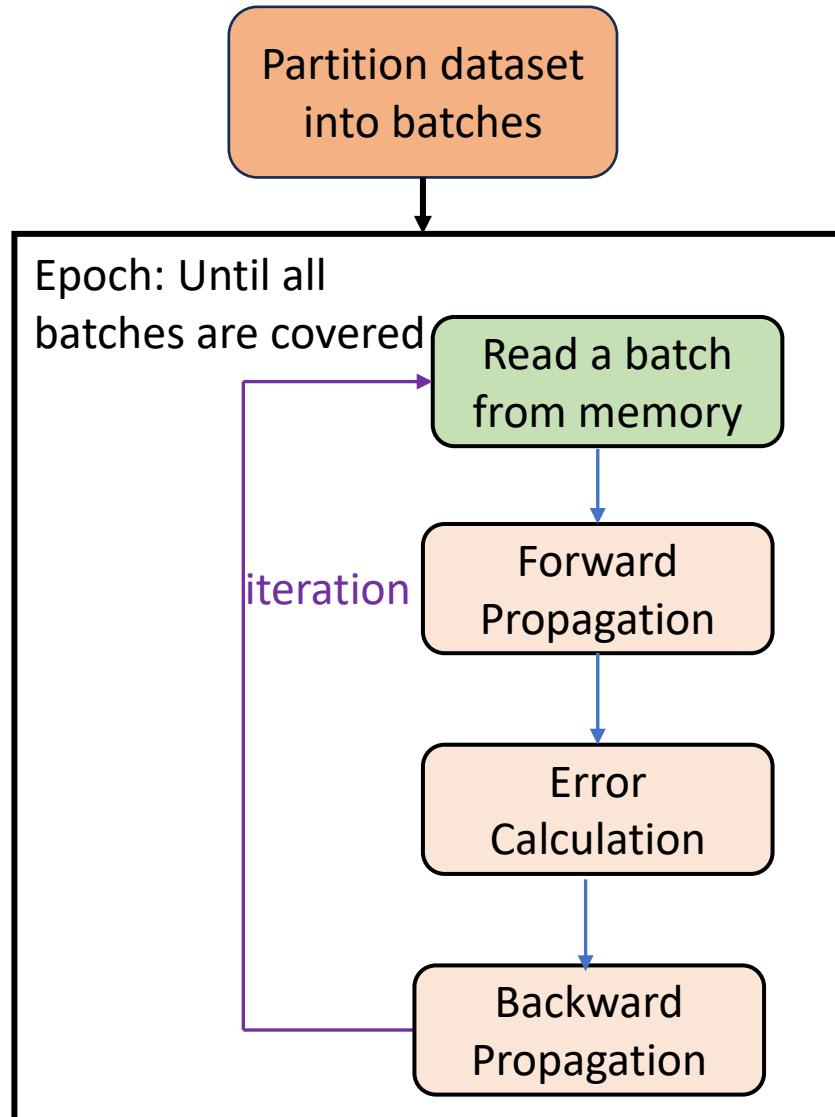
$$\frac{\delta E(W)}{\delta W} = 2 \sum \underbrace{(y_i - F(x_i: W))}_{\text{Error}} \times \frac{\delta F(x_i: W)}{\delta W}$$

For all samples, in each iteration

For multiple images in a batch: Simply sum up the gradients

$$\frac{\delta E(W)}{\delta W_l} = \boxed{\frac{\delta O_l}{\delta W_l} \times \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}}} \times \boxed{\frac{\delta E(W)}{\delta O_L}}$$

Training a Machine Learning Model



Pytorch

- Python wrapper for Torch library
 - Open source library for working with tensors
- Great accelerator support.
 - Automatically links to device specific accelerated libraries such as cuDNN, hipBLAS.
- Abstracts away the hardware/system complexity of building and training AI/ML models



Pytorch - Workflows

- Inference: Use a pretrained model to perform inference on data
- Training: Train a model using a dataset
 - Fine-tune a pre-trained model using a dataset
 - Train a new custom model from scratch

Pytorch - Inference

- Pre-trained models available at pytorch
- Vision Models - <https://docs.pytorch.org/vision/stable/models.html>
- Transformer Models (for NLP) -
<https://docs.pytorch.org/text/stable/models.html>

Pytorch - Inference

- Download the model weights from pytorch ‘cloud’
- Apply them to a model architecture
- Transform the image on which you need to perform inference based on the requirements of the model
- Perform inference using the model

Pytorch - Inference

- Download the model weights from pytorch ‘cloud’
- Apply them to a model architecture
- Transform the image on which you need to perform inference based on the requirements of the model
- Perform inference using the model

Pytorch - Inference

```
from torchvision.io import decode_image
from torchvision.models import resnet50, ResNet50_Weights

img = decode_image("test/assets/encode_jpeg/grace_hopper_517x606.jpg")

# Step 1: Initialize model with the best available weights
weights = ResNet50_Weights.DEFAULT
model = resnet50(weights=weights)
model.eval()

# Step 2: Initialize the inference transforms
preprocess = weights.transforms()

# Step 3: Apply inference preprocessing transforms
batch = preprocess(img).unsqueeze(0)

# Step 4: Use the model and print the predicted category
prediction = model(batch).squeeze(0).softmax(0)
class_id = prediction.argmax().item()
score = prediction[class_id].item()
category_name = weights.meta["categories"][class_id]
print(f"{category_name}: {100 * score:.1f}%")
```

Pytorch - Inference

- Download the model weights from pytorch ‘cloud’
- **Apply them to a model architecture**
- Transform the image on which you need to perform inference based on the requirements of the model
- Perform inference using the model

Pytorch - Inference

```
from torchvision.io import decode_image
from torchvision.models import resnet50, ResNet50_Weights

img = decode_image("test/assets/encode_jpeg/grace_hopper_517x606.jpg")

# Step 1: Initialize model with the best available weights
weights = ResNet50_Weights.DEFAULT
model = resnet50(weights=weights)
model.eval()

# Step 2: Initialize the inference transforms
preprocess = weights.transforms()

# Step 3: Apply inference preprocessing transforms
batch = preprocess(img).unsqueeze(0)

# Step 4: Use the model and print the predicted category
prediction = model(batch).squeeze(0).softmax(0)
class_id = prediction.argmax().item()
score = prediction[class_id].item()
category_name = weights.meta["categories"][class_id]
print(f"{category_name}: {100 * score:.1f}%")
```

Set model's model to evaluation – We are telling the model that this is not training to avoid some training specific operations

Pytorch - Inference

- Download the model weights from pytorch ‘cloud’
- Apply them to a model architecture
- Transform the image on which you need to perform inference based on the requirements of the model
- Perform inference using the model

Pytorch - Inference

```
from torchvision.io import decode_image
from torchvision.models import resnet50, ResNet50_Weights

img = decode_image("test/assets/encode_jpeg/grace_hopper_517x606.jpg")

# Step 1: Initialize model with the best available weights
weights = ResNet50_Weights.DEFAULT
model = resnet50(weights=weights)
model.eval()

# Step 2: Initialize the inference transforms
preprocess = weights.transforms()

# Step 3: Apply inference preprocessing transforms
batch = preprocess(img).unsqueeze(0)

# Step 4: Use the model and print the predicted category
prediction = model(batch).squeeze(0).softmax(0)
class_id = prediction.argmax().item()
score = prediction[class_id].item()
category_name = weights.meta["categories"][class_id]
print(f"{category_name}: {100 * score:.1f}%")
```

Transformations can be cropping or resizing the image so that it is similar in dimensions to what model is expecting

Pytorch - Inference

- Download the model weights from pytorch ‘cloud’
- Apply them to a model architecture
- Transform the image on which you need to perform inference based on the requirements of the model
- Perform inference using the model

Pytorch - Inference

```
from torchvision.io import decode_image
from torchvision.models import resnet50, ResNet50_Weights

img = decode_image("test/assets/encode_jpeg/grace_hopper_517x606.jpg")

# Step 1: Initialize model with the best available weights
weights = ResNet50_Weights.DEFAULT
model = resnet50(weights=weights)
model.eval()

# Step 2: Initialize the inference transforms
preprocess = weights.transforms()

# Step 3: Apply inference preprocessing transforms
batch = preprocess(img).unsqueeze(0)

# Step 4: Use the model and print the predicted category
prediction = model(batch).squeeze(0).softmax(0)
class_id = prediction.argmax().item()
score = prediction[class_id].item()
category_name = weights.meta["categories"][class_id]
print(f"{category_name}: {100 * score:.1f}%")
```

For a batch of images, the model produces the prediction output.

Pytorch – Training

- Decide what dataset to use
- Create a Dataloader object – responsible for creating batches
- Define the model architecture
- Choose a loss function
- Training loop – actual testing occurs here
- Testing – periodically test to ensure model training is proceeding on the right direction

Pytorch – Training

- Datasets – List of common datasets are available on pytorch
- Vision - <https://docs.pytorch.org/vision/stable/datasets.html>
- Text - <https://docs.pytorch.org/text/stable/datasets.html>
- You can also use your own custom dataset -
https://docs.pytorch.org/tutorials/beginner/data_loading_tutorial.html

Pytorch – Training

- Loading/Downloading a dataset from Pytorch

Train = True/False -> training/testing portion

Transform = Applies the transform that is needed for the model that you wish to use (such as resizing, cropping)

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False,  
download=True, transform=transform)
```

Pytorch – Training

- Creating a DataLoader object
 - Enables accessing data in batches
 - Performs random shuffling to improve AI/ML convergence

Pytorch – Training

- Defining the model architecture
- Layers (and structure) need to be outlined in the `__init__` method.
 - Specific attributes can be passed in here.
- Forward describes the data's movement through the model.
 - Activation functions tend to be applied here.
- Create an instance of the model.

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

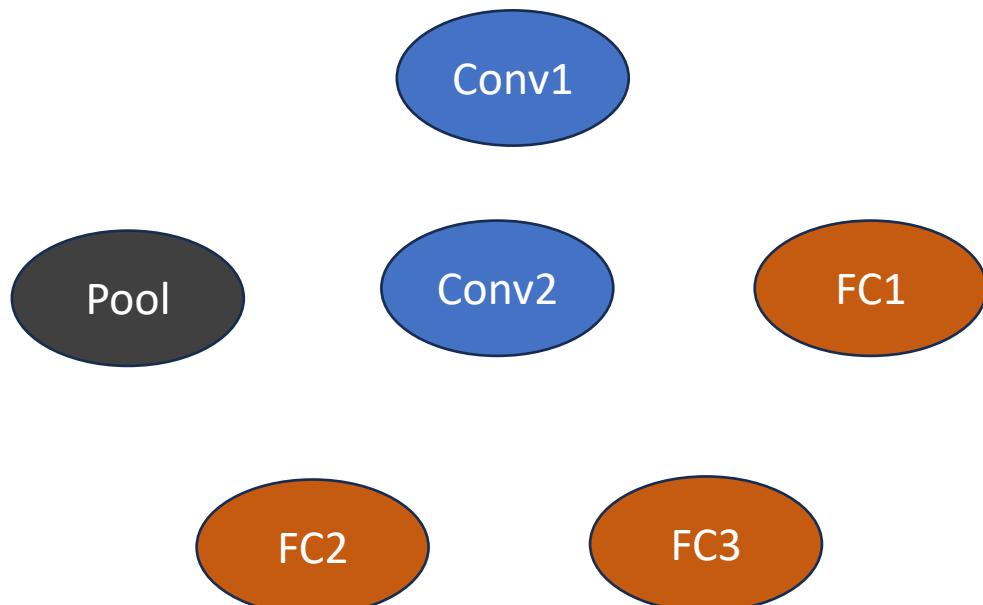
    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

Pytorch – Training

- The `__init__` method is analogous to **defining types** of certain vertices within a task graph.
 - These are layers, but they also perform operations.
- However, there is nothing specifying quantity or order (how data flows through these).

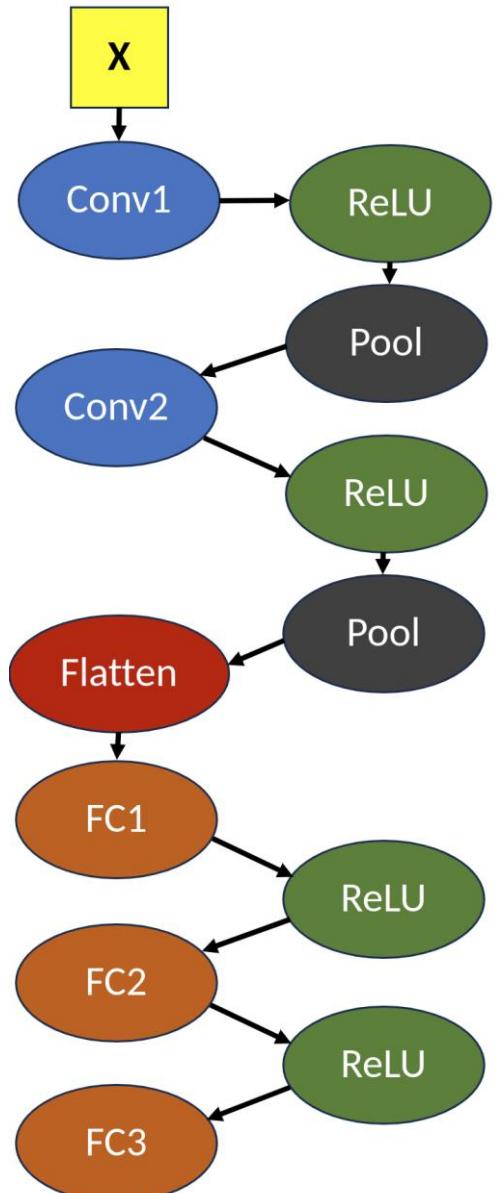
```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```



Pytorch – Training

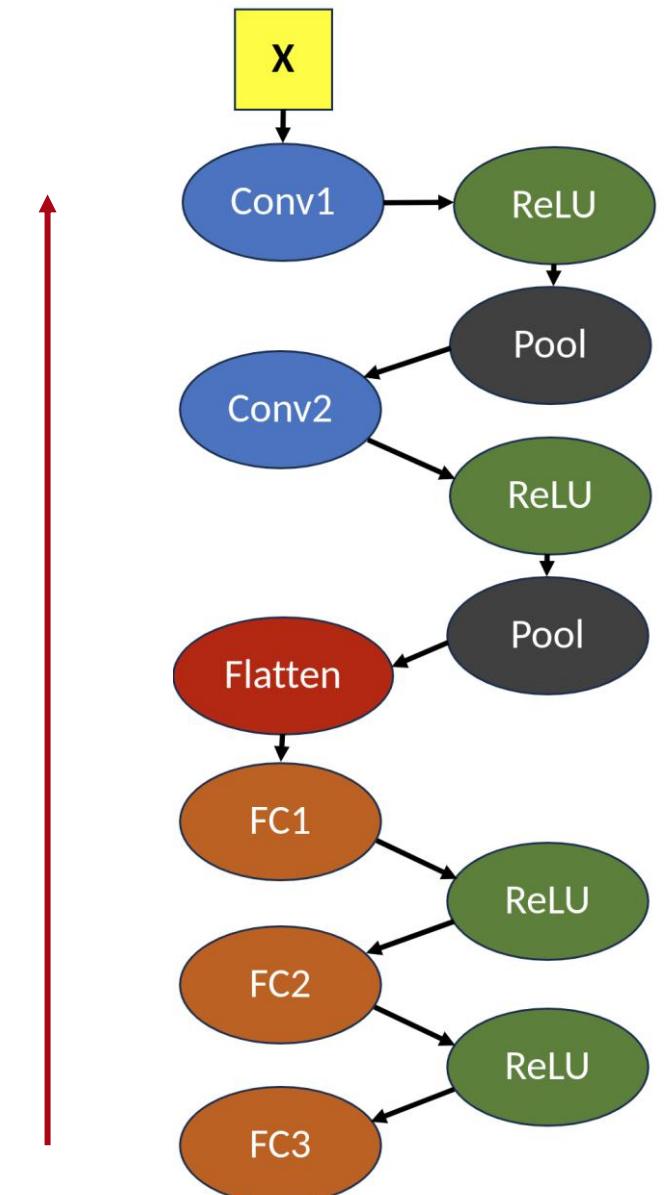
- `forward()` specifies connectivity and data flow
- During inference, when we call `model(img)`, the `forward` function gets called

```
def forward(self, x):
    x = self.pool(F.relu(self.conv1(x)))
    x = self.pool(F.relu(self.conv2(x)))
    x = torch.flatten(x, 1) # flatten all dimensions except batch
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
    return x
```



Pytorch – Training

- In backward propagation, the gradients flow in the reverse direction of the connectivity
- Do we explicitly define the backward() function?
- No, **Autograd** functionality takes care of this in most circumstances



Pytorch – Training

- Choosing a loss function
- Appropriate loss function for your task, you can also build custom loss functions
- We also need to choose an optimizer
 - It applies the weight updates based on the gradients calculated in backpropagation
 - SGD, ADAM are popular optimizers

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

Pytorch – Training

- Training Loop
- Epoch: Go over the entire dataset in batches to perform forward and backward propagation
- Stop training after a set number of epochs (or some other convergence criteria)

```
for epoch in range(2):    # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999:    # print every 2000 mini-batches
            print(f'[ {epoch + 1}, {i + 1:5d} ] loss: {running_loss / 2000:.3f}')
            running_loss = 0.0

print('Finished Training')
```

Pytorch – Training

- Testing Loop
- Test the performance of the model in an unknown data

```
correct = 0
total = 0
# since we're not training, we don't need to calculate the gradients for our outputs
with torch.no_grad():
    for data in testloader:
        images, labels = data
        # calculate outputs by running images through the network
        outputs = net(images)
        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the network on the 10000 test images: {100 * correct // total} %')
```

Pytorch – Execution on Accelerators

- PyTorch defaults to the CPU, even if there are other resources available.
- Specifying a GPU is easy.

- Indicate what device to use:

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

- Send your model to the device:

```
net.to(device)
```

- Send your data to the device:

```
inputs, labels = data[0].to(device), data[1].to(device)
```

Pytorch – Execution on Accelerators

- For simple networks, the speedup may not be noticeable.
 - However, on larger ones it can be quite substantial.
- Table is from the DarkNet repo (YOLO).
 - tkDNN is a neural network library built using cuDNN and tensorRT.

GeForce RTX 2080 Ti:

Network Size	Darknet, FPS (avg)	tkDNN TensorRT FP32, FPS	tkDNN TensorRT FP16, FPS	OpenCV FP16, FPS	tkDNN TensorRT FP16 batch=4, FPS	OpenCV FP16 batch=4, FPS	tkDNN Speedup
320	100	116	202	183	423	430	4.3x
416	82	103	162	159	284	294	3.6x
512	69	91	134	138	206	216	3.1x
608	53	62	103	115	150	150	2.8x
Tiny 416	443	609	790	773	1774	1353	3.5x
Tiny 416 CPU Core i7 7700HQ	3.4	-	-	42	-	39	12x

Next Class

- 10/9 Midterm
 - Best of luck!
- 10/14 Lecture 14
 - Pytorch-Lightning

Thank You

- Questions?
- Email: sanmukh.kuppannagari@case.edu