

# CSDS 451: Designing High Performant Systems for AI

Lecture 16

10/23/2024

Sanmukh Kuppannagari

[sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)

<https://sanmukh.research.st/>

Case Western Reserve University

# Outline

- Flash Attention - Memory Access Aware Attention Computation


# Outline

- Flash Attention - Memory Access Aware Attention Computation

# Transformer Models


- Encoder-Decoder Architecture
- Encoder:
  - Map a sequence of input symbols  $(x_1, x_2, \dots, x_n)$ , to
  - A sequence of continuous representation  $(z_1, z_2, \dots, z_n)$
- Decoder:
  - Given  $(z_1, z_2, \dots, z_n)$  and output from previous iteration  $(y_1^{t-1}, y_2^{t-1}, \dots, y_m^{t-1})$
  - Output  $(y_1^t, y_2^t, \dots, y_m^t)$

# Transformer Models



# Recall: Multi-headed Attention


- Perform  $h$  attention computations
- Assuming dimension of embeddings is  $d_m$  (also called **model dimension**)
- Project  $Q, K, V$  using  $W_i^Q \in R^{d_m \times d_k}, W_i^K \in R^{d_m \times d_k}, W_i^V \in R^{d_m \times d_v} \forall k \in \{1, 2, \dots, h\}$
- $head_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$
- $\text{Concat}(head_1, head_2, \dots, head_h)W^O$  where  $W^O \in R^{hd_v \times d_m}$




# Recall: Multi-headed Attention

- Perform  $h$  attention computations
  - Assuming dimension of embeddings is  $d_m$  (also called **model dimension**)
  - Project  $X$  into  $Q, K, V$  using  $W_i^Q \in R^{d_m \times d_k}, W_i^K \in R^{d_m \times d_k}, W_i^V \in R^{d_m \times d_v} \forall k \in \{1, 2, \dots, h\}$
  - $head_i = \text{Attention}(Q, K, V)$
  - $\text{Concat}(head_1, head_2, \dots, head_h)W^O$  where  $W^O \in R^{hd_v \times d_m}$
- Most computation and memory intensive process.
- We will focus on attention calculation for a single head. Multiple heads are processed in parallel using the same techniques


# Recall: Multi-headed attention




# Recall: Multi-headed attention



# Recall: Multi-headed attention



# Zooming into Attention




Q Matrix: Each row is a query

$K^T$  Matrix: Each column is a key


V Matrix: Each row is a value

X: Matrix M

# Zooming into Attention




# Zooming into Attention



$$[w_{i1}, w_{i2}, \dots, w_{iL}] = \text{softmax}([q_1 \cdot k_1, q_1 \cdot k_2, \dots, q_1 \cdot k_L])$$

# Zooming into Attention




$$[w_{i1}, w_{i2}, \dots, w_{iL}] = \text{softmax}([q_i \cdot k_1, q_i \cdot k_2, \dots, q_i \cdot k_L])$$

Exponent of dot product of query  
 $i$  with key  $j$


$$w_{ij} = \frac{e^{q_i \cdot k_j}}{\sum_l e^{q_i \cdot k_l}}$$

Sum of exponents of dot product of  
query  $i$  with ALL keys


# Zooming into Attention



# Zooming into Attention




# Problem: Lets Track the Memory Requirements




Memory Needed: ??

# Problem: Lets Track the Memory Requirements



Memory Needed:  $2Ld_k + L^2$


# Problem: Lets Track the Memory Requirements



$$[w_{i1}, w_{i2}, \dots, w_{iL}] = \text{softmax}([q_1 \cdot k_1, q_1 \cdot k_2, \dots, q_1 \cdot k_L])$$

Memory Needed:  $L^2$

# Problem: Lets Track the Memory Requirements



Memory Needed:  $L^2 + 2Ld_v$

# Problem: Lets Track the Memory Requirements


- Atleast  $L^2$  memory needed at any given time
- Assume, data type = 16 bits (2 Bytes)
- $L = 1024$
- Memory Needed = ??

# Problem: Lets Track the Memory Requirements


- Atleast  $L^2$  memory needed at any given time
- Assume, data type = 16 bits (2 Bytes)
- $L = 1024$
- Memory Needed = **2 MB**

# Recall: Compute Unit

- Each VALU consists of Vector Register File
  - 64KB - 256 total registers – each register is 64 4-byte-wide entries
  - Private to each compute core
- Each CU consists of Local Data Share (LDS)
  - **64 KB** – Shared Cache in our GPU model
  - Can be used to share data between all threads mapped to the same CU
- A smaller Vector Cache/L1 Cache is also available, not user controllable




=



# Problem


- Naïve Matrix Multiplication based implementation requires materialization of  $L^2$  sized temporary softmax matrix
- Compute Units quickly run out of memory, leading to transfers between global memory and shared cache of CU
- This was the standard implementation in Pytorch before **FlashAttention**

# Recall: GPU Modeling



- Excessive memory transfers on the red arrow
- From our GPU Model:
  - Time taken to transfer data from processor to cache = ??
  - Time taken to transfer data from cache to DRAM = ??

# Recall: GPU Modeling



- Excessive memory transfers on the red arrow
- From our GPU Model:
  - Time taken to transfer data from processor to cache = 1 cycle
  - Time taken to transfer data from cache to DRAM = latency + data rate cycles

# FlashAttention: Resources

- **Original FlashAttention:** Dao, T., Fu, D., Ermon, S., Rudra, A., & Ré, C. (2022). Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems*, 35, 16344-16359.
- **FlashAttention2:** Dao, Tri. "Flashattention-2: Faster attention with better parallelism and work partitioning." *arXiv preprint arXiv:2307.08691* (2023).
- **FlashAttention3:** Shah, J., Bikshandi, G., Zhang, Y., Thakkar, V., Ramani, P., & Dao, T. (2024). Flashattention-3: Fast and accurate attention with asynchrony and low-precision. *Advances in Neural Information Processing Systems*, 37, 68658-68685.


# FlashAttention

- Key Idea: Memory Access Pattern Aware Attention
- (They mention I/O aware attention, but it maybe a misnomer. I/O refers to storage (SSD, HDD, etc.) in systems community)

# FlashAttention


- Key Idea: Memory Access Pattern Aware Attention
- Specifically, perform computations similar to blocked matrix multiplications
  - Do not need to store the entire temporary  $O(L \times L)$  matrix
  - Generate smaller blocks, use them for multiplication with values, discard them

# FlashAttention




Memory Needed:  $L^2$

# FlashAttention



Memory Needed: ??

# FlashAttention






Memory Needed:  $B_r \times B_c$

# FlashAttention

Ignore softmax for now

What rows/Columns of Q and K will we load onto the shared memory?




Memory Needed:  $B_r \times B_c$

# FlashAttention

Ignore softmax for now

What rows/Columns of Q and K will we load onto the shared memory?




Memory Needed:  $B_r \times B_c + ??$



# FlashAttention

Ignore softmax for now

What rows/Columns of Q  
and K will we load onto the  
shared memory?





Memory Needed:  $B_r \times B_c + d_k B_r + d_k B_c$

# FlashAttention

Ignore softmax for now

We can bring portion of V too to calculate the final output

What rows/Columns of V will we load onto the shared memory?




# FlashAttention

Ignore softmax for now

We can bring portion of V too to calculate the final output


What rows/Columns of V will we load onto the shared memory?




# FlashAttention

Ignore softmax for now

What rows/Columns of O will this produce (that will again be temporarily stored in the shared memory)?




# FlashAttention



# FlashAttention

Ignore softmax for now

What rows/Columns of O will this produce (that will again be temporarily stored in the shared memory)?



Memory Needed:  $B_r \times B_c + d_k B_r + d_k B_c + d_v B_c + d_v B_r$

$L$

$d_v$


$L$

$d_v$

# FlashAttention

Ignore softmax for now






Notice: For outputs  $(o_1, o_2)$  only partial outputs have been produced



# FlashAttention

Ignore softmax for now

How to produce the full output  
for  $B_r$  rows of O?




Memory Needed:  $B_r \times B_c + d_k B_r + d_k B_c + d_v B_c + d_v B_r$

# FlashAttention

Ignore softmax for now

How to produce the full output  
for  $B_r$  rows of O? **Iterate over  
columns of K**



Memory Needed:  $B_r \times B_c + d_k B_r + d_k B_c + d_v B_c + d_v B_r$

$L$


$d_v$

$B_c$

# FlashAttention

Ignore softmax for now


How to produce the full output  
for  $B_r$  rows of O? **Iterate over  
columns of K**



# FlashAttention

Ignore softmax for now


How to produce the full output  
for  $B_r$  rows of O? **Iterate over  
columns of K**



# FlashAttention

Ignore softmax for now

Continue this for  $L/B_r$  times to produce all  $L$  rows of the output



# FlashAttention - PseudoCode

```

5: for  $1 \leq j \leq T_c$  do
6:   Load  $\mathbf{K}_j, \mathbf{V}_j$  from HBM to on-chip SRAM.
7:   for  $1 \leq i \leq T_r$  do
8:     Load  $\mathbf{Q}_i, \mathbf{O}_i, \ell_i, m_i$  from HBM to on-chip SRAM.
9:     On chip, compute  $\mathbf{S}_{ij} = \mathbf{Q}_i \mathbf{K}_j^T \in \mathbb{R}^{B_r \times B_c}$ .
10:    On chip, compute  $\tilde{m}_{ij} = \text{rowmax}(\mathbf{S}_{ij}) \in \mathbb{R}^{B_r}$ ,  $\tilde{\mathbf{P}}_{ij} = \exp(\mathbf{S}_{ij} - \tilde{m}_{ij}) \in \mathbb{R}^{B_r \times B_c}$  (pointwise),  $\tilde{\ell}_{ij} = \text{rowsum}(\tilde{\mathbf{P}}_{ij}) \in \mathbb{R}^{B_r}$ .
11:    On chip, compute  $m_i^{\text{new}} = \max(m_i, \tilde{m}_{ij}) \in \mathbb{R}^{B_r}$ ,  $\ell_i^{\text{new}} = e^{m_i - m_i^{\text{new}}} \ell_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\ell}_{ij} \in \mathbb{R}^{B_r}$ .
12:    Write  $\mathbf{O}_i \leftarrow \text{diag}(\ell_i^{\text{new}})^{-1} (\text{diag}(\ell_i) e^{m_i - m_i^{\text{new}}} \mathbf{O}_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\mathbf{P}}_{ij} \mathbf{V}_j)$  to HBM.
13:    Write  $\ell_i \leftarrow \ell_i^{\text{new}}$ ,  $m_i \leftarrow m_i^{\text{new}}$  to HBM.
14:  end for
15: end for

```

For each column block of  $K, V$ :  $T_c = L/B_c$

Iterate through row blocks of  $Q$  to produce  $O$ :  $T_r = L/B_r$

Additional bookkeeping to compute softmax that we ignored till now.

# FlashAttention

- How do we determine  $B_r, B_c$ ?
- $B_r \times B_c + d_k B_r + d_k B_c + d_v B_c + d_v B_r \leq S$
- Typically sizes of {64,128} used

# How do we compute “partial” softmaxes

- Concept similar to running averages

$$[w_{i1}, w_{i2}, w_{i3}, w_{i4}, w_{i5}, w_{i6}] = \text{softmax}([q_i \cdot k_1, q_i \cdot k_2, q_i \cdot k_3, q_i \cdot k_4, q_i \cdot k_5, q_i \cdot k_6])$$

- Assuming  $B_c = 2$

# How do we compute “partial” softmaxes

- Concept similar to running averages

$$[w_{i1}, w_{i2}, w_{i3}, w_{i4}, w_{i5}, w_{i6}] = [e^{q_i \cdot k_1}, e^{q_i \cdot k_2}, e^{q_i \cdot k_3}, e^{q_i \cdot k_4}, e^{q_i \cdot k_5}, e^{q_i \cdot k_6}] / \sum_j e^{q_i \cdot k_j}$$

- Assuming  $B_c = 2$

# How do we compute “partial” softmaxes

- Concept similar to running averages

$$[w_{i1}, w_{i2}, w_{i3}] = [e^{q_i \cdot k_1}, e^{q_i \cdot k_2}, e^{q_i \cdot k_3}] / \sum_{j=\{1,2,3\}} e^{q_i \cdot k_j} \rightarrow l$$

- Assuming  $B_c = 2$

# How do we compute “partial” softmaxes

- Concept similar to running averages

$$[w_{i4}, w_{i5}, w_{i6}] = [e^{q_i \cdot k_4}, e^{q_i \cdot k_5}, e^{q_i \cdot k_6}] / (\sum_{j=\{4,5,6\}} e^{q_i \cdot k_j} + l)$$

- Assuming  $B_c = 2$

# How do we compute “partial” softmaxes

- Concept similar to running averages

$$[w_{i4}, w_{i5}, w_{i6}] = [e^{q_i \cdot k_4}, e^{q_i \cdot k_5}, e^{q_i \cdot k_6}] / (\sum_{j=\{4,5,6\}} e^{q_i \cdot k_j} + l)$$

- Assuming  $B_c = 2$
- We need to recompute  $[w_{i1}, w_{i2}, w_{i3}]$  Why???

# How do we compute “partial” softmaxes

- Concept similar to running averages

$$[w_{i4}, w_{i5}, w_{i6}] = [e^{q_i \cdot k_4}, e^{q_i \cdot k_5}, e^{q_i \cdot k_6}] / (\sum_{j=\{4,5,6\}} e^{q_i \cdot k_j} + l) \rightarrow l_2$$

- Assuming  $B_c = 2$
- We need to recompute  $[w_{i1}, w_{i2}, w_{i3}]$  Why??? The sum in denominator has changed.

# How do we compute “partial” softmaxes

- $[w_{i1}, w_{i2}, w_{i3}] = [w_{i1}, w_{i2}, w_{i3}] \times l/l_2$
- In practice,
  - Output O is rescaled
  - In softmax computation, scaling by the maximum value of row is performed to avoid numerical stability issues
- Ungraded HW Assignment: Try to understand this. Will not ask in exam.

On chip, compute  $\tilde{m}_{ij} = \text{rowmax}(\mathbf{S}_{ij}) \in \mathbb{R}^{B_r}$ ,  $\tilde{\mathbf{P}}_{ij} = \exp(\mathbf{S}_{ij} - \tilde{m}_{ij}) \in \mathbb{R}^{B_r \times B_c}$  (pointwise),  $\tilde{\ell}_{ij} = \text{rowsum}(\tilde{\mathbf{P}}_{ij}) \in \mathbb{R}^{B_r}$ .

On chip, compute  $m_i^{\text{new}} = \max(m_i, \tilde{m}_{ij}) \in \mathbb{R}^{B_r}$ ,  $\ell_i^{\text{new}} = e^{m_i - m_i^{\text{new}}} \ell_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\ell}_{ij} \in \mathbb{R}^{B_r}$ .

Write  $\mathbf{O}_i \leftarrow \text{diag}(\ell_i^{\text{new}})^{-1} (\text{diag}(\ell_i) e^{m_i - m_i^{\text{new}}} \mathbf{O}_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\mathbf{P}}_{ij} \mathbf{V}_j)$  to HBM.

Write  $\ell_i \leftarrow \ell_i^{\text{new}}$ ,  $m_i \leftarrow m_i^{\text{new}}$  to HBM.

# Flashattention – Scheduling on GPUs

- Each block  $B_r \times B_c$  is mapped to a SMP
- Flashattention 2 optimizations
  - Improve partitioning of blocks into warps
  - Improve online softmax/output calculation
- Flashattention 3 optimizations
  - Nvidia Hopper specific optimizations – asynchronous scheduling of online softmax/output calculation and matrix multiplications

# Next Class

- 10/28 Lecture 17
  - Accelerating Transformer Model: Sparsity

# Thank You

- Questions?
- Email: [sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)