

CSDS 451: Designing High Performant Systems for AI

Lecture 17

10/28/2025

Sanmukh Kuppannagari

sanmukh.kuppannagari@case.edu

<https://sanmukh.research.st/>

Case Western Reserve University

Outline

- Sparse Transformers Basics
- Sparse Masks


Outline

- Sparse Transformers Basics
- Sparse Masks

Attention Mechanism – Fast Facts

- Three Key Operations
 - Operation #1: $Y = QK^T$: Product of Q and K^T
 - Operation #2: $Z = \text{Softmax}(Y)$
 - Operation #3: $O = ZV$: Product of Z and V matrices
 - Q, K^T, V, Z : Dense matrices


Attention




Q Matrix: Each row is a query
 K^T Matrix: Each column is a key
 V Matrix: Each row is a value

X : Matrix M

Zooming into Attention




Zooming into Attention



$$[w_{i1}, w_{i2}, \dots, w_{iL}] = \text{softmax}([q_1 \cdot k_1, q_1 \cdot k_2, \dots, q_1 \cdot k_L])$$

Zooming into Attention




$$[w_{i1}, w_{i2}, \dots, w_{iL}] = \text{softmax}([q_i \cdot k_1, q_i \cdot k_2, \dots, q_i \cdot k_L])$$

Exponent of dot product of query
 i with key j


$$w_{ij} = \frac{e^{q_i \cdot k_j}}{\sum_l e^{q_i \cdot k_l}}$$

Sum of exponents of dot product of
query i with ALL keys

Zooming into Attention



Zooming into Attention




Attention Mechanism – Fast Facts

- Naïve Matrix Multiplication based implementation requires materialization of L^2 sized temporary softmax matrix
- Compute Units quickly run out of memory, leading to transfers between global memory and shared cache of CU
- This was the standard implementation in Pytorch before **FlashAttention**
- Flashattention approach – Partition Z into independent tiles of size $B_r \times B_c$, and figure out the inputs tiles (Q, K, V) and output tiles (O) to produce

FlashAttention

Ignore softmax for now

How to produce the full output
for B_r rows of O? **Iterate over
columns of K**



Memory Needed: $B_r \times B_c + d_k B_r + d_k B_c + d_v B_c + d_v B_r$

L

d_v

B_c

B_c

B_r

d_v

II


L

X

FlashAttention

Ignore softmax for now


How to produce the full output
for B_r rows of O? **Iterate over
columns of K**



FlashAttention

Ignore softmax for now


How to produce the full output
for B_r rows of O? **Iterate over
columns of K**



FlashAttention

Ignore softmax for now


How to produce the full output
for B_r rows of O? **Iterate over**
columns of K



FlashAttention

Ignore softmax for now


How to produce the full output
for B_r rows of O? **Iterate over
columns of K**



FlashAttention

Ignore softmax for now

How to produce the full output
for B_r rows of O? **Iterate over
columns of K**



How do we compute “partial” softmaxes

- $[w_{i1}, w_{i2}, w_{i3}] = [w_{i1}, w_{i2}, w_{i3}] \times l/l_2$
- In practice,
 - Output O is rescaled
 - In softmax computation, scaling by the maximum value of row is performed to avoid numerical stability issues
- Ungraded HW Assignment: Try to understand this. Will not ask in exam.

On chip, compute $\tilde{m}_{ij} = \text{rowmax}(\mathbf{S}_{ij}) \in \mathbb{R}^{B_r}$, $\tilde{\mathbf{P}}_{ij} = \exp(\mathbf{S}_{ij} - \tilde{m}_{ij}) \in \mathbb{R}^{B_r \times B_c}$ (pointwise), $\tilde{\ell}_{ij} = \text{rowsum}(\tilde{\mathbf{P}}_{ij}) \in \mathbb{R}^{B_r}$.

On chip, compute $m_i^{\text{new}} = \max(m_i, \tilde{m}_{ij}) \in \mathbb{R}^{B_r}$, $\ell_i^{\text{new}} = e^{m_i - m_i^{\text{new}}} \ell_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\ell}_{ij} \in \mathbb{R}^{B_r}$.

Write $\mathbf{O}_i \leftarrow \text{diag}(\ell_i^{\text{new}})^{-1} (\text{diag}(\ell_i) e^{m_i - m_i^{\text{new}}} \mathbf{O}_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\mathbf{P}}_{ij} \mathbf{V}_j)$ to HBM.

Write $\ell_i \leftarrow \ell_i^{\text{new}}$, $m_i \leftarrow m_i^{\text{new}}$ to HBM.

Flash Attention

- Fuse and tile the three Key Operations
 - Operation #1: $Y = QK^T$: Product of Q and K^T
 - Operation #2: $Z = \text{Softmax}(Y)$
 - Operation #3: $O = ZV$: Product of Z and V matrices
 - Q, K^T, V, Z : Dense matrices
- Notice how you are taking a tile of Q, K, V and producing a partial sum for a tile of the output O
- This is also known as Fused attention

Flash attention


- Reduces the memory requirements from $O(L^2)$ to ??

Sparse Transformers

- Flashattention reduces the memory requirements from $O(L^2)$ to $B_r \times B_c$
- However, no change in computations – still require $O(L^2)$ computations
- For longer values of L this can be exorbitant
- Solution: Sparse Transformers


Key Idea

- Considering soft-lookup view of attention
- Not every query-key pairs need to interact
- As keys and queries are projection of tokens -> Not every pair of token need to interact



Introducing Sparsity in Attention

- Key Idea: Not every pair of query and key/value need to interact
- If we represent the query-key interactions as an adjacency matrix, we want to prune the entries of the matrix




Attention Mask

Self-Attention in
Transformer Blocks:



All 1s adjacency matrix

(colored cell is 1)

Each row corresponds to a query





Attention Mask





The L^2 entries of M represent the L^2 dot-products needed for the matrix multiplication QK^T

QK Product with Attention Mask





If some of the L^2 entries of M are 0, the corresponding dot-products should not be computed

QK Product with Attention Mask





Do you recall anything from the exam?

QK Product with Attention Mask





Last Question: Product of two matrices with mask

QK Product with Attention Mask




Each entry in the attention mask M corresponds to 1 dot-product
 l^2 entries in total
 Sparsity factor s determines how many dot-products actually need to
 be computed

QK Product with Attention Mask




Product will lead to a l^2 sized matrix with 0 and non-zero elements similar to mask:
 How? Will ask in WA 3.
 In practice, elements corresponding to 0 mask are set to $-\infty$: Why? Will ask in WA 3

Softmax with Attention Mask




M

$$Z = \text{softmax}(QK^T)$$




Softmax will replace all $-\infty$ elements with 0: How? Will ask in WA 3.
The pattern of 0 and non-zero still remains the same

Product with V Matrix



$$Z = \text{softmax}(QK^T)$$



$$V$$

Product of a sparse matrix ($Z = \text{softmax}(QK^T)$) and dense V matrix

Attention with Sparse Attention Mask

- Three Key Operations
- Operation #1: $Y = QK^T | M$: Product of Q and K^T matrices under mask M
- Operation #2: $Z = \text{Softmax}(Y)$
- Operation #3: $O = ZV$: Product of Z and V matrices
- Q, K^T, V : Dense matrices
- Z : Sparse matrix

Theoretical Reduction in Computation

- For sparsity factor $0 \leq s \leq 1$ defined as the ratio of the number of non-zero elements to the number of zero elements
- $KQ^T : O(sL^2 \times d_k)$: How? Will be a question in WA3
- $ZV : O(sL^2 \times d_v)$: How? Will be a question in WA3
- For very small values of s , significant reduction in computations

Reduction in Computations

- In practice, pytorch algorithms do not fully exploit this sparsity.
 - Recall our implementation that in the last question of Exam 1, it was not work-optimal
- We will discuss this issue and potential solutions (some developed by my students) in the next class (or two)

Research on Sparse Transformers

- Two main directions
- #1 Build masks that reduce the computations but still preserve accuracy
 - Today's class
- #2 Exploit Sparsity on Hardware such as GPUs
 - Next one or two classes

Outline

- Sparse Transformers Basics
- Sparse Masks

Types of Attention Masks


- Two main types
- Dynamic Attention: Only consider query-key pairs for attending that are the most similar
 - No explicit attention mask
 - Pairs to attend determined at runtime
- Static Attention: Fixed attention mask used to determine attending query-key pairs

Types of Attention Masks

- Two main types
- **Dynamic Attention: Only consider query-key pairs for attending that are the most similar**
 - No explicit attention mask
 - Pairs to attend determined at runtime
- Static Attention: Fixed attention mask used to determine attending query-key pairs

Reformer

- Key Idea: Reduce the $O(L^2)$ complexity of attention calculation by only considering key values pairs that are more similar
- Less similar key-value pairs -> small weights -> less information gain



Kitaev, N., Kaiser, L., & Levskaya, A. (2019, September). Reformer: The Efficient Transformer. In *International Conference on Learning Representations*.

Reformer

- How do we know what key-value pairs are more similar without computing the values???

Reformer

- How do we know what key-value pairs are more similar without computing the values
- Locality Sensitive Hashing

Locality Sensitive Hashing

- A technique to perform efficient nearest neighbor search
- Hashes “similar” items into same “buckets” with “high” probability


Locality Sensitive Hashing

- Consider a set of hash functions $h \in \mathcal{F}$, where $h: M \rightarrow S$
 - M : some metric space with distance metric d
 - S : buckets
- Consider the following additional parameters
 - $r > 0$: a threshold value
 - $\epsilon > 0$: an approximation factor
 - p_1, p_2 : probability values

Locality Sensitive Hashing

- The set of hash functions \mathcal{F} is called an LSH family if,
- For any two points $a, b \in M$ and a randomly selected hash function $h \in \mathcal{F}$
- If $d(a, b) \leq r$, then $h(a) = h(b)$ with probability $\geq p_1$
- If $d(a, b) > \epsilon r$, then $h(a) = h(b)$ with probability $\leq p_2$

Locality Sensitive Hashing



M

High Probability $> p_1$
of conflict

Low Probability $< p_2$
of conflict

Locality Sensitive Hashing

- Key Design Parameters:
- Determining the hash functions that can achieve this property
- Look at this survey for more information: Jafari, Omid, Preeti Maurya, Parth Nagarkar, Khandker Mushfiqul Islam, and Chidambaram Crushev. "A survey on locality sensitive hashing algorithms and their applications." *arXiv preprint arXiv:2102.08942* (2021).

Reformer


- Hashing Scheme: angular locality sensitive hashing
- For a key with dimension d^k , obtain b values as follows:
- Fix a random matrix R of size $d^k \times b/2$
- $h(x) = \arg \max\{xR; -xR\}$; -> concatenation

Reformer

- Intuition: The dimension with largest value for points close to each other should be same even after projection.
- For formal proof, and details on how to create the random matrix, and other parameters, see:
 - Andoni, A., Indyk, P., Laarhoven, T., Razenshteyn, I., & Schmidt, L. (2015). Practical and optimal LSH for angular distance. *Advances in neural information processing systems*, 28.


Reformer

- Attention with LSH
- Sparse adjacency matrix
- However, inefficient for computation
 - Why???




Reformer

- Attention with LSH
- Sparse adjacency matrix
- However, inefficient for computation
 - Too many random accesses
- How do we fix???




Reformer

- Attention with LSH
- Sparse adjacency matrix
- However, inefficient for computation
 - Too many random accesses
- Sort the sequence to cluster buckets together



Reformer

- Sort the sequence to cluster buckets together
- Perform attention mechanism on blocks
 - Smaller matrix multiplications
- (Make a note of this sorting technique, we will discuss more sophisticated methods in the next one/two classes)



Reformer

- Other Contributions:
- Multi-round LSH attention
 - Single hash computation may skip some connections due to probabilistic nature
 - Do n different (and parallel) LSH computations to find out the neighbors for populating the adjacency matrix
- $N_i = \bigcup_{r=1}^n N_i^r$
- N_i^r : neighbors of query i obtained in i th LSH computation
- N_i : all neighbors of query i used to populate the adjacency matrix

Types of Attention Masks


- Two main types
- Dynamic Attention: Only consider query-key pairs for attending that are the most similar
 - No explicit attention mask
 - Pairs to attend determined at runtime
- **Static Attention: Fixed attention mask used to determine attending query-key pairs**

Big Bird

- Key Idea:
 - Think of a Transformer model as information flows between tokens of the input sequence
 - A transformer block represents single edge connections in a graph with L nodes
 - A Transformer model with L_m layers can be thought of as information flow using paths of length L_m


Zaheer, M., Guruganesh, G., Dubey, K. A., Ainslie, J., Alberti, C., Ontanon, S.,
... & Ahmed, A. (2020). Big bird: Transformers for longer sequences.
Advances in neural information processing systems, 33, 17283-17297.

Big Bird



Transformer: Information flow between all tokens

Big Bird




Fully Connected Attention: Information flow between all tokens in Single Layer

Big Bird

- Consider the following adjacency matrix for attention mechanism
- Which two tokens have no information flow in a single layer?


x_1	x_2	x_3	x_4



Big Bird

- Consider the following adjacency matrix for attention mechanism
- $x_1:x_3, x_2:x_4$
- How many layers will be needed to ensure information flows between them???


x_1	x_2	x_3	x_4



Big Bird

- Consider the following adjacency matrix for attention mechanism
- $x_1: x_3, x_2: x_4$
- ≥ 2


x_1	x_2	x_3	x_4



Big Bird

- Consider the following adjacency matrix for attention mechanism
- $x_1: x_3, x_2: x_4$
- ≥ 2

x_1	x_2	x_3	x_4




Big Bird


- Key Idea:
 - Think of a Transformer model as information flows between tokens of the input sequence
 - A transformer block represents single edge connections in a graph with L nodes
 - A Transformer model with L_m layers can be thought of as information flow using paths of length L_m
- Objective: Find graphs (adjacency matrix) with $O(L)$ nodes and small path lengths

Big Bird


- Consists of three types of connections




(a) Random attention



(b) Window attention



(c) Global Attention




(d) BIGBIRD

Big Bird

- Random Attention
 - Inspired by graph theory (graph sparsification problem)
 - Random connections between tokens
- Erdos-Renyi model for random graph construction
 - Each edge randomly chosen with fixed probability
- Property: Shortest path between any two nodes is logarithmic in the number of nodes.
- So, for L tokens, only $L_m = O(\log L)$ layers needed for full information flow

Big Bird


- Random Attention: Select edges with a fixed probability



(a) Random attention

Big Bird


- Windowed Attention:
 - Inspired from linguistics/computational biology
 - Sequences have locality of reference
 - w : window size is the key parameter



(b) Window attention

Big Bird


- Global Attention: Marks some token as global (or add new tokens) and connects them to all the tokens in the sequence
 - Obtained using the theoretical analysis in the Reformer paper
 - g : number of global tokens is the key parameter



(c) Global Attention

Big Bird

- Combining everything, we get
- $O(L)$ connections – proof in the paper. We will not discuss here.




(d) BIGBIRD

Other Attention Patterns

- 1D Dilated Attention: Create uniform gaps of a set size between attended tokens across a row
- Window size: w
- Dilation factor: r

```
if ((abs(i - j) < w) && (abs(i - j) % r == 0)) {  
    return 1;  
} else {  
    return 0;  
}
```



I. Beltagy, M. E. Peters, and A. Cohan, “Longformer: The longdocument transformer,” arXiv preprint arXiv:2004.05150, 2020


Other Attention Patterns

- 2D Dilated Attention: Dilation along both dimensions
 - Probability of query and key attending exponentially reduces with distance between the corresponding tokens


```
if (floor(i / (L/w)) == floor(j / (L/w))) {  
    i_w = i % w;  
    j_w = j % w;  
    if ((i_w % r == 0) && (j_w % r == 0)) {  
        return 1;  
    } else {  
        return 0;  
    }  
} else {  
    return 0;  
}
```

Other Attention Patterns


- 2D Dilated Attention
- Understand the equation
- May ask a question in WA 3




Segment Length: 4
Dilated Rate: 1



Segment Length: 8
Dilated Rate: 2



Segment Length: 16
Dilated Rate: 4



Next Class

- 10/30 Lecture 18
 - Accelerating Transformer Model: Sparse Transformers II
 - Hardware Acceleration of Sparse Transformers

Thank You

- Questions?
- Email: sanmukh.kuppannagari@case.edu