

CSDS 451: Designing High Performant Systems for AI

Lecture 19

11/4/2025

Sanmukh Kuppannagari

sanmukh.kuppannagari@case.edu

<https://sanmukh.research.st/>

Case Western Reserve University

Outline

- Accelerating Sparse Transformers – Part II

Announcements

- WA 3 was out
- Due by next Saturday

Next Few Classes and Deadlines

- 11/6 – Modeling Cluster of Accelerators
- 11/11 – Distributed Training Techniques
- 11/13 – Distributed Training Techniques (WA 3 due)
- 11/18 – No Class (Traveling for a conference)
- 11/20 – Distributed Training Techniques

Next Few Classes and Deadlines

- 11/25 – Exam
- 11/27 – Thanksgiving Holiday
- 12/2 – Guest Lecture
- 12/4 – Guest Lecture (WA4 due)
- 12/10 – NO final. Projects Due



Outline

- Accelerating Sparse Transformers – Part II

Attention with Sparse Attention Mask

- Three Key Operations
- Operation #1: $Y = QK^T | M$: Product of Q and K^T matrices under mask M
- Operation #2: $Z = \text{Softmax}(Y)$
- Operation #3: $O = ZV$: Product of Z and V matrices
- Q, K^T, V : Dense matrices
- Z : Sparse matrix



QK Product with Attention Mask



Each entry in the attention mask M corresponds to 1 dot-product
 l^2 entries in total
 Sparsity factor s determines how many dot-products actually need to
 be computed


QK Product with Attention Mask

$$Y = QK^T$$




Product will lead to a l^2 sized matrix with 0 and non-zero elements similar to mask:
 How? Will ask in WA 3.
 In practice, elements corresponding to 0 mask are set to $-\infty$: Why? Will ask in WA 3

Softmax with Attention Mask




M




Softmax will replace all $-\infty$ elements with 0: How? Will ask in WA 3.
 The pattern of 0 and non-zero still remains the same

Product with V Matrix



$$Z = \text{softmax}(QK^T)$$



V

Product of a sparse matrix ($Z = \text{softmax}(QK^T)$) and dense V matrix

Attention with Sparse Attention Mask

- Three Key Operations
- Operation #1: $Y = QK^T | M$: Product of Q and K^T matrices under mask M
- Operation #2: $Z = \text{Softmax}(Y)$
- Operation #3: $O = ZV$: Product of Z and V matrices
- Q, K^T, V : Dense matrices
- Z : Sparse matrix

Attention with Sparse Attention Mask

- Three Key Operations
- Operation #1: $Y = QK^T \mid M$: **Product of two dense matrices under a mask**
- Operation #2: $Z = \text{Softmax}(Y)$
- Operation #3: $O = ZV$: Product of Z and V matrices **Product of a sparse and dense matrix**
- Q, K^T, V : Dense matrices
- Z : Sparse matrix

Attention with Sparse Attention Mask

- Three Key Operations

- Operation #1: $Y = QK^T$ | **M: Product of two dense matrices under a mask**
- Operation #2: $Z = \text{Softmax}(Y)$
- Operation #3: $O = ZV$: Product of Z and V matrices **Product of a sparse and dense matrix**

- Q, K^T, V : Dense matrices
- Z : Sparse matrix

$Y = QK^T \mid M$: Product of two dense matrices under a mask

- Known as Sampled Dense Dense Matrix Multiplication
- Just a handful of papers on accelerating this kernel – we will discuss in next class

Attention with Sparse Attention Mask

- Three Key Operations
 - Operation #1: $Y = QK^T$ | **M: Product of dense matrices under a sparse mask**
 - Operation #2: $Z = \text{Softmax}(Y)$
 - Operation #3: $O = ZV$: Product of Z and V matrices: **Product of a sparse and dense matrix**
-
- Q, K^T, V : Dense matrices
 - Z : Sparse matrix

Sparse-Dense Matrix Multiplication (SpMM)


- An important kernel in scientific computing
- In machine learning, an important kernel in Graph Neural Networks (GNNs), pruned CNNs, and Sparse Transformers
- Extensive Research has been performed (and still continuing) on accelerating SpMM

Sparse Dense Matrix Multiplication (SpMM)

- Let nnz be the number of non-zero elements in Z
- Key Challenge: Cannot use dense-dense matrix multiplication techniques
 - Inefficient Storage
 - Non-useful computations
- Optimizations
 - Efficient Storage of Sparse Matrices
 - Work optimal task scheduling

SpMM Storage Format - COO

- Coordinate format
- Each non-zero entry is assigned a coordinate
 - RowID,ColumnID
- Three arrays, each of size nnz are used to store the matrix
 - Row: Array of RowIDs
 - Column: Array of ColumnIDs
 - Value: Array of Values



SpMM Storage Format - COO

DENSE MATRIX

	0	1	2	3
0	1.0		2.0	
1		3.0		
2				
3	4.0	5.0		
4		6.0	7.0	8.0

COORDINATE FORMAT - COO
(ZERO-BASE INDEX)

Row
INDICES

0	1	2	3	4	5	6	7
0	0	1	4	4	5	5	5

COLUMN
INDICES


0	1	2	3	4	5	6	7
0	2	1	0	1	1	2	3

VALUES


0	1	2	3	4	5	6	7
1.0	2.0	3.0	4.0	5.0	6.0	7.0	8.0

SpMM Storage Format - CSR

- Compressed Sparse Row Format
- Row array is compressed so that its entries point to offsets in column whenever new row starts
- Row: Array of offsets into Column – size: L (number of rows)
- Column: Array of ColumnIDs – size: nnz
- Value: Array of Values – size: nnz




SpMM Storage Format - CSR




Sparse Dense Matrix Multiplication (SpMM)

- Approach #1: Assign each row to a SMP




Sparse Dense Matrix Multiplication (SpMM)

- Approach #1: Assign each row to a SMP
- Processors execute the non-zero elements of the row.



Sparse Dense Matrix Multiplication (SpMM)



Step 1: Scale the rows of the second matrix
with assigned non-zero values

$$\text{Processor 1: } O_{p1} = 1.0 \times v_1$$


$$\text{Processor 2: } O_{p2} = 2.0 \times v_3$$

Note: Assuming processors
are 1 indexed

Step 2: All-reduce to produce a single output

$$O_1 = O_{p1} + O_{p2} + \dots$$



Sparse Dense Matrix Multiplication (SpMM)



Question in WA 3


Given BlockID, ThreadID, can you write a GPU code for SpMM?

- Divide rows evenly among the blocks of GPU
- Within each block, divide column indices of rows evenly among the threads
- Each thread scales rows that it owns (based on column indices) of the second matrix
- Syncthreads
- All-reduce of the scaled outputs to produce final output (Recall Parallel Sum? In practice, GPUs have APIs to do this)



Sparse Dense Matrix Multiplication (SpMM)

- Approach #1: Assign each row to a SMP
- Processors execute the non-zero elements of the row.



Pros: CSR format naturally leads to the assignment

Cons: Load imbalance across SMPs

Attention with Sparse Attention Mask

- Three Key Operations
- Operation #1: $Y = QK^T | M$: **Product of dense matrices under a sparse mask**
- Operation #2: $Z = \text{Softmax}(Y)$
- Operation #3: $O = ZV$: Product of Z and V matrices: **Product of a sparse and dense matrix**
- Q, K^T, V : Dense matrices
- Z : Sparse matrix

$$QK^T \mid M$$

- Pytorch Way
- Dense-dense matrix multiplication of $Y = QK^T$
 - Computation Complexity: $O(L^2 d_k)$
 - Storage Requirements: $O(L^2)$ (or $B_r \times B_C + d (B_r + B_c)$ if using flashattention)
- Invalidate entries (Set them to $-\infty$) in Y using mask M
 - Alternatively, initialize Y to $-\infty$ and only update the valid entries

$$QK^T \mid M$$

- Can we directly use SpMM kernel?
- No. But similar ideas can be utilized – Sampled Dense Dense Matrix Multiplication

Today's Lecture

- SDDMM Algorithm Design Idea
 - You will implement the algorithm in WA3
- Using Flash attention on sparse masks
 - Block Sparsity
- Algorithms to reduce the number of blocks in block sparsity

Today's Lecture

- SDDMM Algorithm Design Idea
 - You will implement the algorithm in WA3
- Using Flash attention on sparse masks
 - Block Sparsity
- Algorithms to reduce the number of blocks in block sparsity

Sampled Dense Dense Matrix Multiplication (SDDMM)

- Key Idea/Similarity with SpMM
- SpMM: $O = ZV$, Z : sparse matrix, V : dense matrix
 - For each non-zero element of Z at location Z_{ij}
 - Fetch row j of V
 - Multiply with the element
 - Add the result to the output O_i
- SDDMM: $O = AB | M$: A, B : dense matrices, M : sparse mask
 - For each non-zero element of M at location M_{ij}
 - Fetch row i of A , Fetch column j of B ,
 - Perform dot product between the fetched row and column
 - Write the result to the output O_{ij}






Sampled Dense Dense Matrix Multiplication (SDDMM)

- Key Idea/Similarity with SpMM
- For each non-zero, instead of a vector scaling operation
 - You perform a dot product
- WA3, second part of Q5
 - I will make that extra credits. (so 10 points extra credits)
 - Will not ask in exam
- (This lecture is already getting too big :'-))


Today's Lecture

- SDDMM Algorithm Design Idea
 - You will implement the algorithm in WA3
- Using Flash attention on sparse masks
 - Block Sparsity
- Algorithms to reduce the number of blocks in block sparsity


Recall: Flashattention



With Mask (so crowded)




Flash Attention on Attention Mask




$$B_c = B_r = 2$$

$$QK^T/Y$$




Flash Attention on Attention Mask

M




$$B_c = B_r = 2$$

QK^T/Y




Flash Attention on Attention Mask

M




$$B_c = B_r = 2$$

QK^T/Y




Flash Attention on Attention Mask

M



$$B_c = B_r = 2$$

QK^T/Y




Flash Attention on Sparse Attention Mask


- Pretty much the same thing happens
- If a block has at least one non-zero element, Flashattention ignores the mask and computes the entire block

Flash Attention on Attention Mask

M



QK^T




$$B_c = B_r = 2$$

- Flashattention will ignore the mask
- Notice how all the elements are getting calculated


Flash Attention on Attention Mask

M



QK^T


$B_c = B_r = 2$



- Flashattention will ignore the mask
- Notice how all the elements are getting calculated


Flash Attention on Attention Mask

M



QK^T


$B_c = B_r = 2$




- Flashattention will ignore the mask
- Notice how all the elements are getting calculated

Flash Attention on Attention Mask

M



QK^T




$$B_c = B_r = 2$$

- Flashattention will ignore the mask
- Notice how all the elements are getting calculated
- **If a block is fully empty, Flashattention will ignore it**


Flash Attention on Attention Mask

M



QK^T


$B_c = B_r = 2$




- Flashattention will ignore the mask
- Notice how all the elements are getting calculated

Flash Attention on Attention Mask

M



QK^T



$$B_c = B_r = 2$$

- Flashattention will ignore the mask
- Notice how all the elements are getting calculated

Flash Attention on Attention Mask

- Total Memory Needed?

Flash Attention on Attention Mask

- Total Memory Needed? Still $O(M)$, cache size
- Total Computations?

Flash Attention on Attention Mask

- Total Memory Needed? Still $O(M)$, cache size
- Total Computations? Still $O(L^2d)$ in worst case
- Performs a lot of non-useful computations

How to Improve Performance?

- ???


How to Improve Performance?

- Idea #1: Design masks that have block like structure
 - Non-zeros are packed together to get dense blocks on which Flashattention performs the computations
- Idea #2: Reorder the attention mask to achieve better block like structure


Today's Lecture

- SDDMM Algorithm Design Idea
 - You will implement the algorithm in WA3
- Using Flash attention on sparse masks
 - Block Sparsity
- Algorithms to reduce the number of blocks in block sparsity


#1 Design masks that have block like structure




Dense Mask



Streaming Mask
(Token Granularity)



Streaming Mask
(Block Granularity)



Block Sparse Mask


One Example: <https://github.com/mit-han-lab/Block-Sparse-Attention>
We will not discuss how to build these.

#2: Reorder Attention Masks


- Resources:
 - Sharma, Agniv, and Jonas Geiping. "Efficiently Dispatching Flash Attention For Partially Filled Attention Masks." *arXiv preprint arXiv:2409.15097* (2024).
- Other Resource: “Exploring Binary Block Masking with Hypergraph-Based Token Reordering for GPU-Efficient Attention,” Kutay Tasci, Nathaniel Tomczak, Sanmukh Kuppannagari (Under review)
 - (Our technique is much better than the above, but still under review so unavailable for you to read. ;-))

Introducing Sparsity in Flashattention

M



QK^T




$$B_c = B_r = 2$$

$$\#B = 0$$

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



QK^T

$B_c = B_r = 2$




#B = 1

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



QK^T

$B_c = B_r = 2$




#B = 2

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



$$B_c = B_r = 2$$

QK^T




Do not process

#B = 2

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



$$B_c = B_r = 2$$

QK^T




Do not process

$$\#B = 2$$

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



QK^T

$B_c = B_r = 2$




$\#B = 3$


Lets calculate the number of blocks computed by
FlashAttention

Introducing Sparsity in Flashattention

M



QK^T




$$B_c = B_r = 2$$

$$\#B = 5$$

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



$$B_c = B_r = 2$$

QK^T




#B = ??

Lets calculate the number of blocks computed by
FlashAttention


Introducing Sparsity in Flashattention

M



$$B_c = B_r = 2$$

QK^T




$\#B = 12$

- Significant reduction in computations can be obtained if the number of blocks with non-zeros are less


Introducing Sparsity in Flashattention

M



$B_c = B_r = 2$

QK^T




#B = 12

- What could have led to an even better reduction in computation?


Introducing Sparsity in Flashattention

M



$B_c = B_r = 2$


QK^T



- What could have led to an even better reduction in computation?


Introducing Sparsity in Flashattention

M



$B_c = B_r = 2$


QK^T



- What could have led to an even better reduction in computation? If non-zero elements are packed closer to the diagonal


Introducing Sparsity in Flashattention

M



$B_c = B_r = 2$


QK^T



- However, we cannot simply take a non-zero and move it.


Introducing Sparsity in Flashattention

M



$B_c = B_r = 2$

QK^T



#B = 9

- However, we cannot simply take a non-zero and move it.
- **Solution: Matrix Reordering**

Matrix Reordering


- A key preprocessing step in scientific computations
- Given a sparse matrix, reorder rows/columns with the following objective:
 - Minimize bandwidth: Longest distance across row from the diagonal elements
 - Fill in: Number of entries that change from 0 to non-zero, typically after matrix factorization
- This link provides a good visualization:
<https://www.mathworks.com/help/matlab/math/sparse-matrix-reordering.html>

Matrix Reordering

- In our case, the sparse matrix is an adjacency matrix of a graph
 - For undirected graphs: adjacency matrix is symmetric
 - For masks that are non-causal, attention masks are symmetric
- And we do not need to do any factorization, so we don't care about Fill in
- So, the problem boils down to: Relabel the vertices of the graph to minimize bandwidth


Bandwidth

- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0$, if $|i - j| > K$



Bandwidth


- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0$, if $|i - j| > K$
- Procedure:
 - For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
 - Maximum across all the rows is the bandwidth



Bandwidth

- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0$, if $|i - j| > K$
- Procedure:
 - For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
 - Maximum across all the rows is the bandwidth


$$\begin{aligned} i &= 0 \\ j &= 1 \\ K &= 1 \end{aligned}$$



Bandwidth

- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0$, if $|i - j| > K$
- Procedure:
 - For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
 - Maximum across all the rows is the bandwidth


$$\begin{aligned} i &= 2 \\ j &= ?? \\ K &= ?? \end{aligned}$$



Bandwidth


- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0$, if $|i - j| > K$
- Procedure:
 - For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
 - Maximum across all the rows is the bandwidth

$$\begin{aligned} i &= 2 \\ j &= 7 \\ K &= 5 \end{aligned}$$



Bandwidth

- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0$, if $|i - j| > K$
- Procedure:
 - For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
 - Maximum across all the rows is the bandwidth



Bandwidth


- Definition: For a matrix M with entries denoted using i, j

- Bandwidth is the number K s.t. :

- $M_{ij} = 0$, if $|i - j| > K$

- Procedure:

- For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
- Maximum across all the rows is the bandwidth



Bandwidth

- Definition: For a matrix M with entries denoted using i, j


- Bandwidth is the number K s.t. :

- $M_{ij} = 0$, if $|i - j| > K$

- Bandwidth = 7

- **Procedure:**

- For each row i , find j such that all elements outside the band $i - j, i + j$ are 0.
 - Maximum across all the rows is the bandwidth.




Bandwidth – After Reordering




- Definition: For a matrix M with entries denoted using i, j
- Bandwidth is the number K s.t. :
- $M_{ij} = 0 , \text{ if } |i - j| > K$
- Bandwidth = 3

A lower bandwidth is being used as a proxy for better packing


$$K = 3$$







Graph Vertex Relabelling



Notice how the adjacency matrix changes with the reordering of the vertices






Graph Vertex Relabelling




Graph Bandwidth Minimization:
Finding a re-ordering of graph to
minimize the bandwidth of its
corresponding adjacency matrix

Graph Vertex Relabelling






https://en.wikipedia.org/wiki/Graph_bandwidth

This problem is NP-hard




Graph Vertex Relabelling



https://en.wikipedia.org/wiki/Graph_bandwidth

This problem is NP-hard

Fast heuristics exists




Cuthill–McKee algorithm (CM)

- Given a square symmetric matrix that is an adjacency matrix of a graph, relabel the vertices to obtain low bandwidth in the matrix
- Cuthill, Elizabeth, and James McKee. "Reducing the bandwidth of sparse symmetric matrices." *Proceedings of the 1969 24th national conference*. 1969.
<https://dl.acm.org/doi/pdf/10.1145/800195.805928>
- Type written, uses fortran to discuss the algorithm. It is a fun read.

Cuthill–McKee algorithm (CM)

- Algorithm:
([https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee algorithm](https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee_algorithm))
- Let $R = \{\}$
- Choose a vertex x with lowest degree (randomly select if multiple) and set $R_0 = \{x\}$. Append R_0 to R
- While $|R| < n$, for $i = 1, 2, \dots$
 - Identify all the neighbors of R_{i-1} and put them into a set A_i (remove any vertices already in R)
 - Sort A_i in ascending order of degree, call this sorted set R_i
 - Append R_i to R


Cuthill–McKee algorithm (CM)



Cuthill–McKee algorithm (CM)

- Algorithm:
([https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee algorithm](https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee_algorithm))
- Let $R = \{\}$
- Choose a vertex x with lowest degree (randomly select if multiple) and set $R_0 = \{x\}$. Append R_0 to R
- While $|R| < n$, for $i = 1, 2, \dots$
 - Identify all the neighbors of R_{i-1} and put them into a set A_i (remove any vertices already in R)
 - Sort A_i in ascending order of degree, call this sorted set R_i
 - Append R_i to R



Cuthill–McKee algorithm (CM)



Cuthill–McKee algorithm (CM)

- Algorithm:
([https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee algorithm](https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee_algorithm))
- Let $R = \{\}$
- Choose a vertex x with lowest degree (randomly select if multiple) and set $R_0 = \{x\}$. Append R_0 to R
- While $|R| < n$, for $i = 1, 2, \dots$
 - Identify all the neighbors of R_{i-1} and put them into a set A_i (remove any vertices already in R)
 - Sort A_i in ascending order of degree, call this sorted set R_i
 - Append R_i to R



Cuthill–McKee algorithm (CM)





Cuthill–McKee algorithm (CM)

- Algorithm:
([https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee algorithm](https://en.wikipedia.org/wiki/Cuthill%E2%80%93McKee_algorithm))
- Let $R = \{\}$
- Choose a vertex x with lowest degree (randomly select if multiple) and set $R_0 = \{x\}$. Append R_0 to R
- **While $|R| < n$, for $i = 1, 2, \dots$**
 - Identify all the neighbors of R_{i-1} and put them into a set A_i (remove any vertices already in R)
 - Sort A_i in ascending order of degree, call this sorted set R_i
 - **Append R_i to R**

Cuthill–McKee algorithm (CM)



Cuthill–McKee algorithm (CM)




Cuthill–McKee algorithm (CM)

- Ungraded HW assignment: Try to run the algorithm on a graph of your choice and see if you understand.
- Reverse Cuthill–McKee algorithm (RCM): Same as CM, but reverse the node numbering
 - Leads to lower fill-in after factorization (beyond the scope of this class)


Going back: Introducing Sparsity in Flashattention

M



QK^T


$B_c = B_r = 2$



- What could have led to an even better reduction in computation?


Going back: Introducing Sparsity in Flashattention

M



QK^T

$B_c = B_r = 2$



- What could have led to an even better reduction in computation? **RCM algorithm before calling the modified flashattention**

Sparsity in Attention

- Reorder attention masks to reduce the number of blocks
- Call Flashattention, to only operate on non-zero blocks
- Caveats:
 - Reordering takes time, so not applicable to all use cases
 - Sometimes the block structure may become worse
- Our proposed technique: Using bandwidth as a proxy for reducing the number of blocks works well for scientific computing, but not for attention masks
 - Use hypergraph partitioning techniques that directly optimize for reducing the number of blocks

Next Class

- 11/6 Lecture 20
 - Modeling Cluster of Accelerators

Thank You

- Questions?
- Email: sanmukh.kuppannagari@case.edu