

# CSDS 451: Designing High Performant Systems for AI

Lecture 20

11/6/2025

Sanmukh Kuppannagari

[sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)

<https://sanmukh.research.st/>

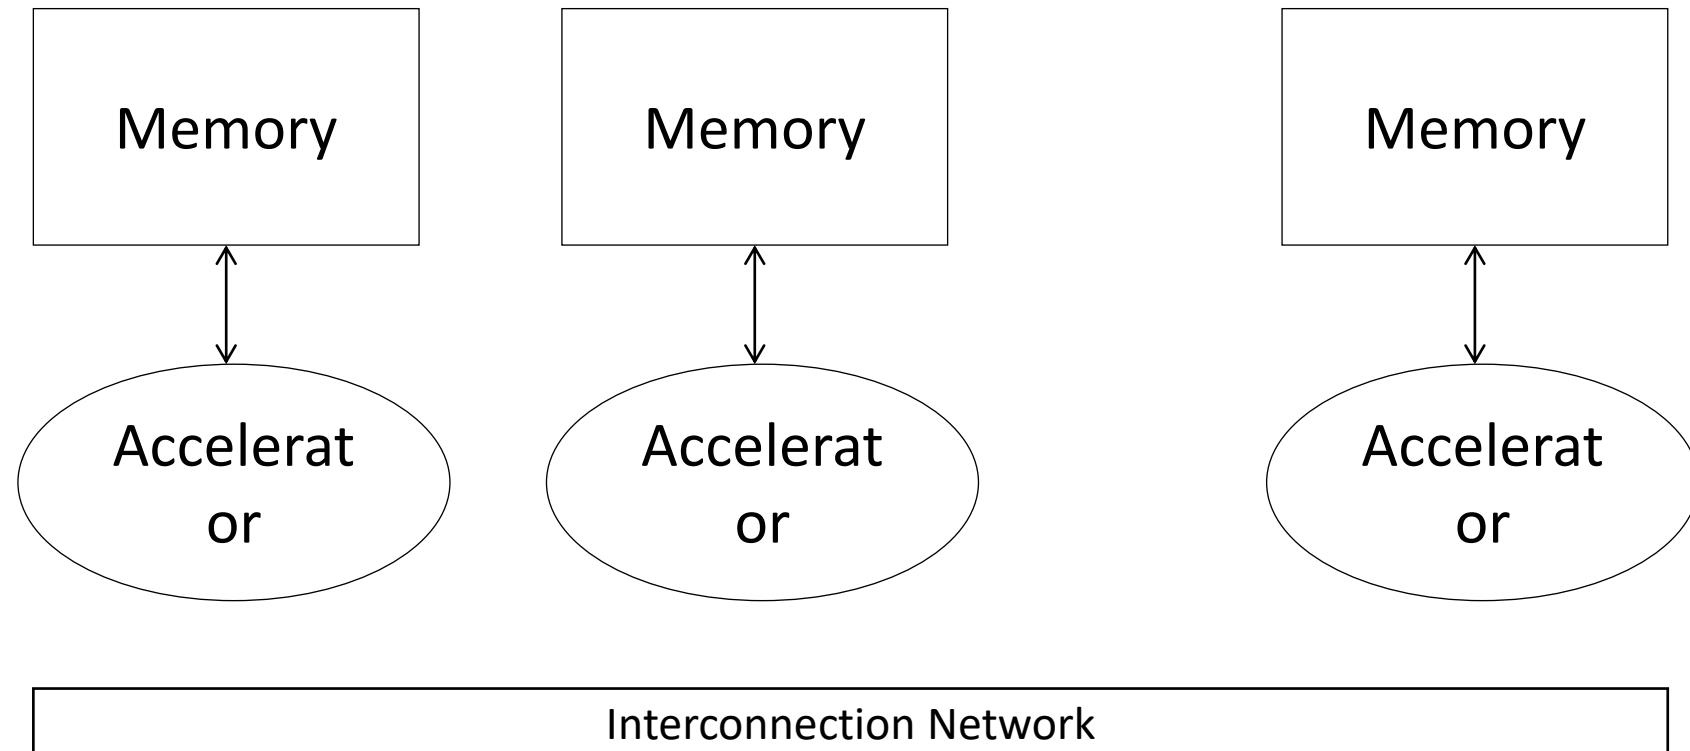
Case Western Reserve University

# Outline

- Cluster of Accelerators – Basics

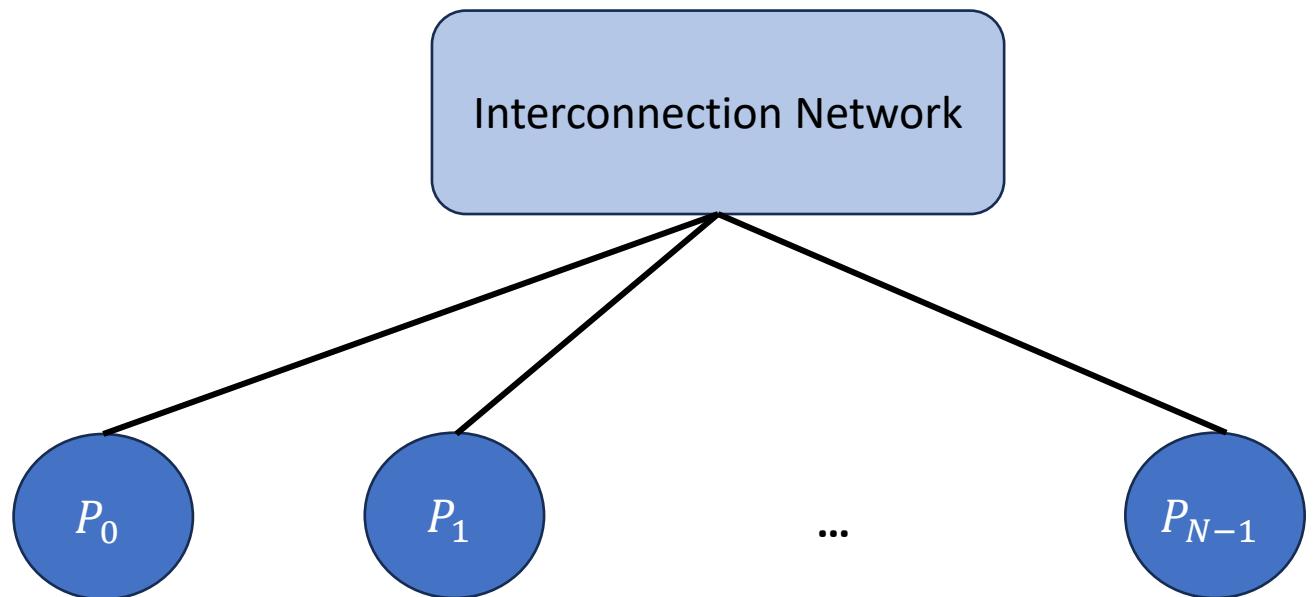
# Cluster of Accelerators

- Multiple Accelerator-Memory devices connected through an interconnection network
- Accelerator – CPU, GPU, Systolic Array, ...



# Cluster of Accelerators

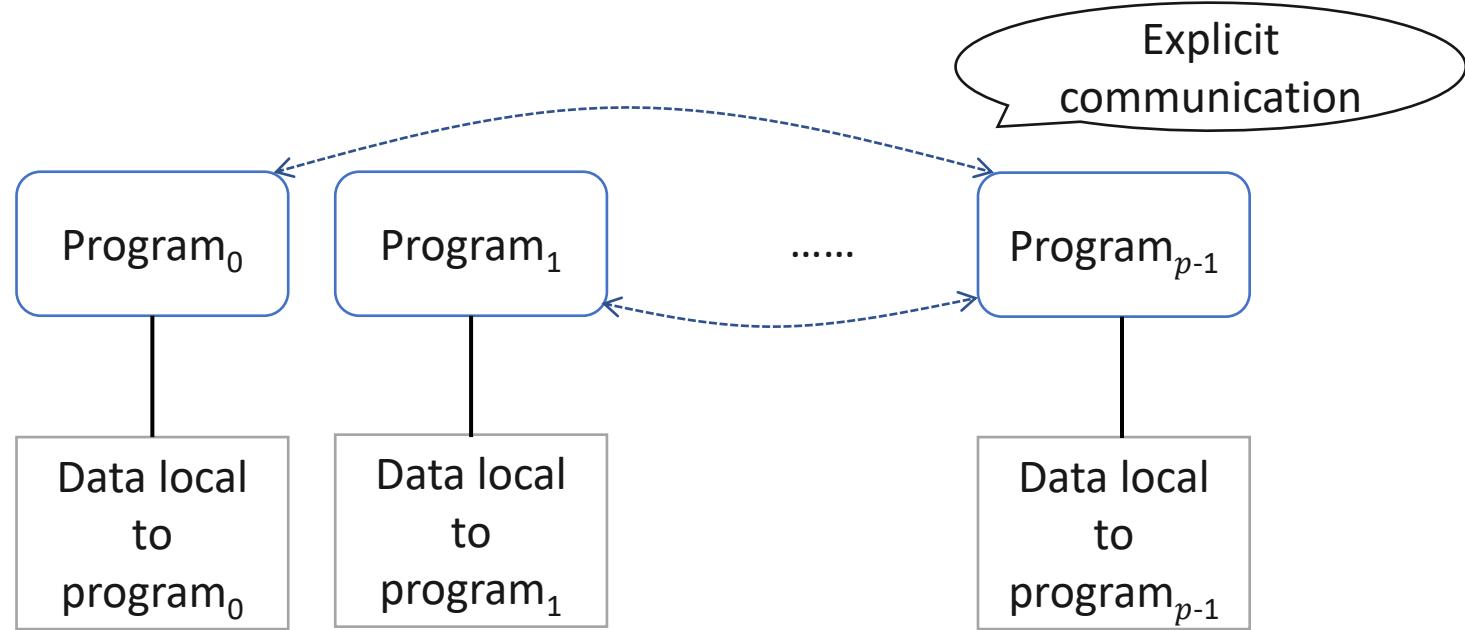
- $N$  processors (memory + accelerator)
  - Local compute
  - Local memory
- Connected using an Interconnection Network
- Communication through Message Passing



# Message Passing Programming Model (1)

- Message passing
  - One of the oldest parallel programming paradigms
  - Widely used
  - Key features
    - Partition address space
      - local data, remote data
    - Explicit parallelization
      - user is responsible to specify and manage concurrency

# Message Passing Programming Model (2)



Communication - needs coordination among the communicating processes

Most Popular Framework – Message Passing Interface (MPI)

# Message Passing Programming

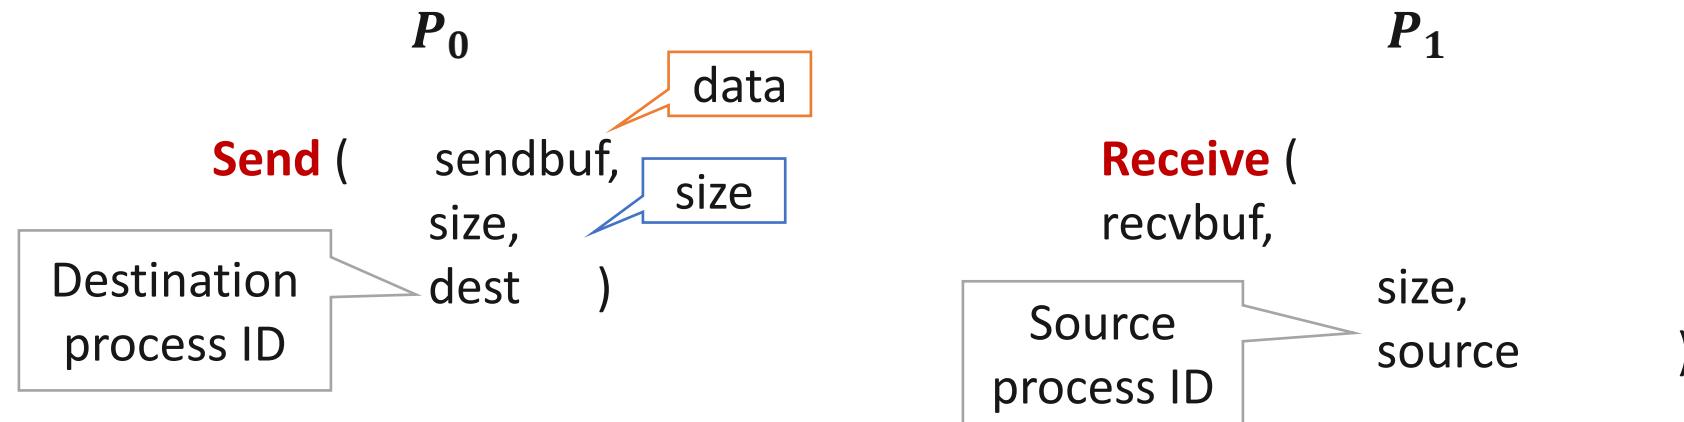
- How do we program clusters using Message Passing?
- We need to specify the local computations that each processor will execute
- We also need to specify at what time which pairs of processors will communicate and what data they will transfer to each other

# Resources

- Slides from University of Stuttgart -  
[https://fs.hlrs.de/projects/par/par\\_prog\\_ws/pdf/mpi\\_3.1\\_rab.pdf](https://fs.hlrs.de/projects/par/par_prog_ws/pdf/mpi_3.1_rab.pdf)
- Covers all the important details of MPI (744 slides)
- We will cover only enough for its application to AI

# Communication Operations (1)

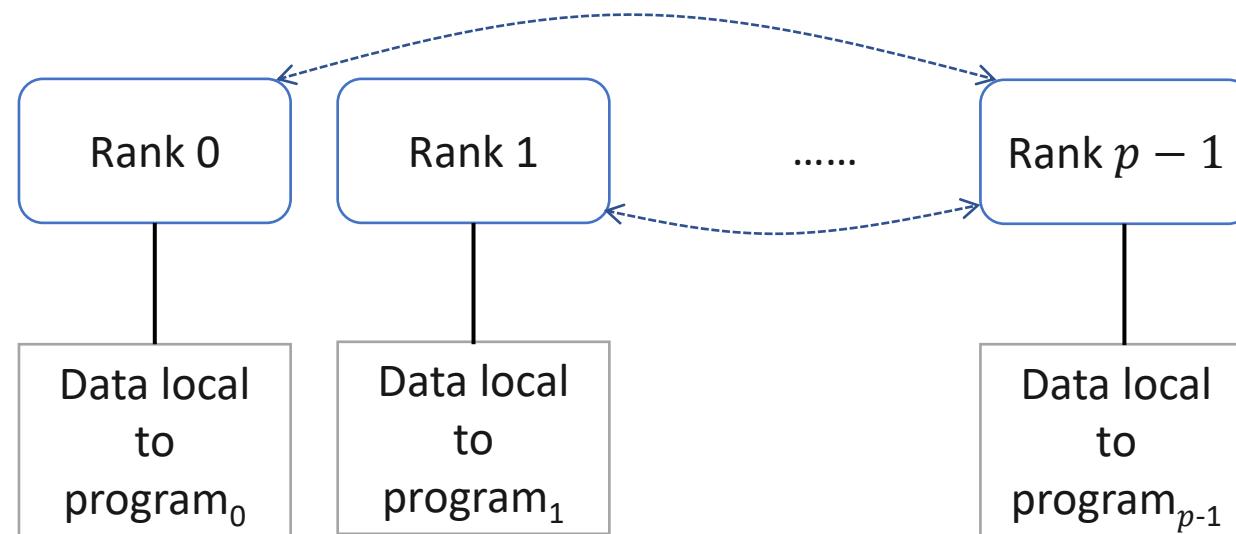
- A pair of Send and Receive is used to implement a communication step



- Processor  $P_0$  *sends* size amount of data to Processor  $P_1$
- Processor  $P_1$  *receives* size amount of data from Processor  $P_0$
- Needs to be symmetric: If not, will lead to hard to debug errors

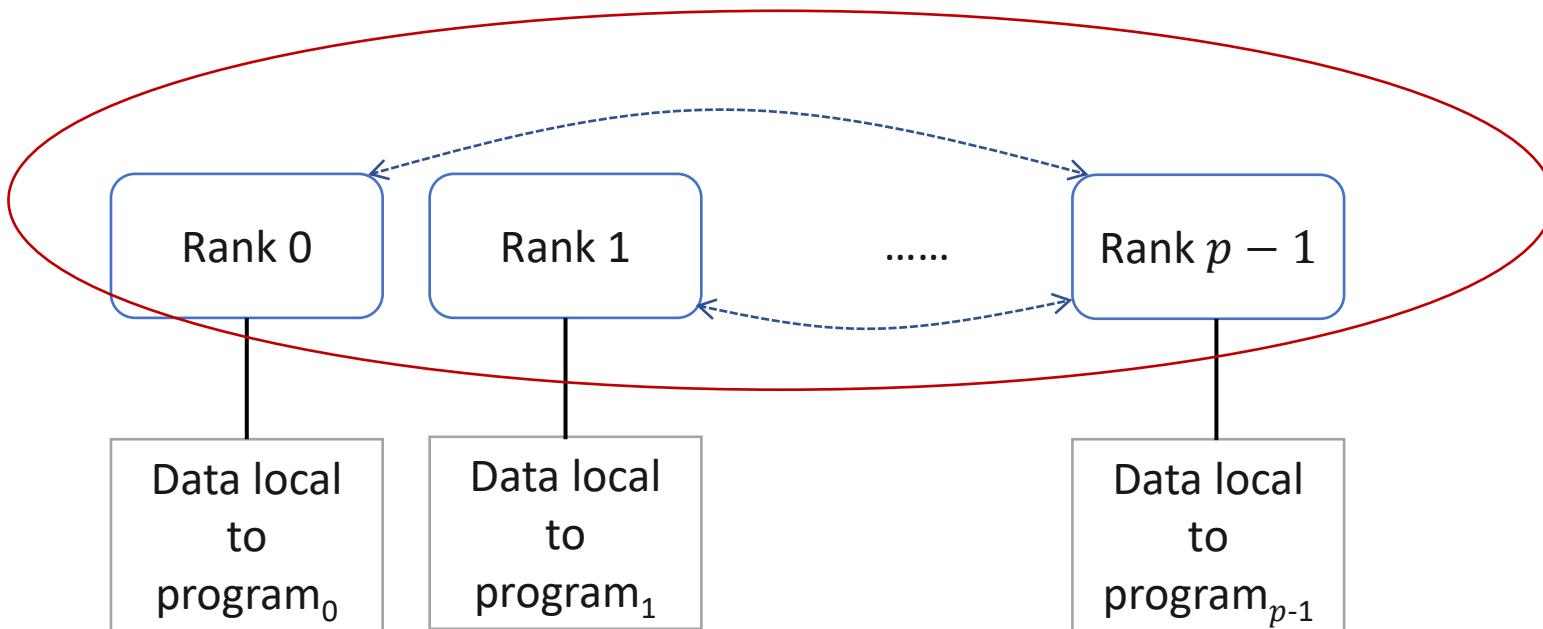
# Communication Operations (2)

- How do we identify processor IDs? **Ranks**



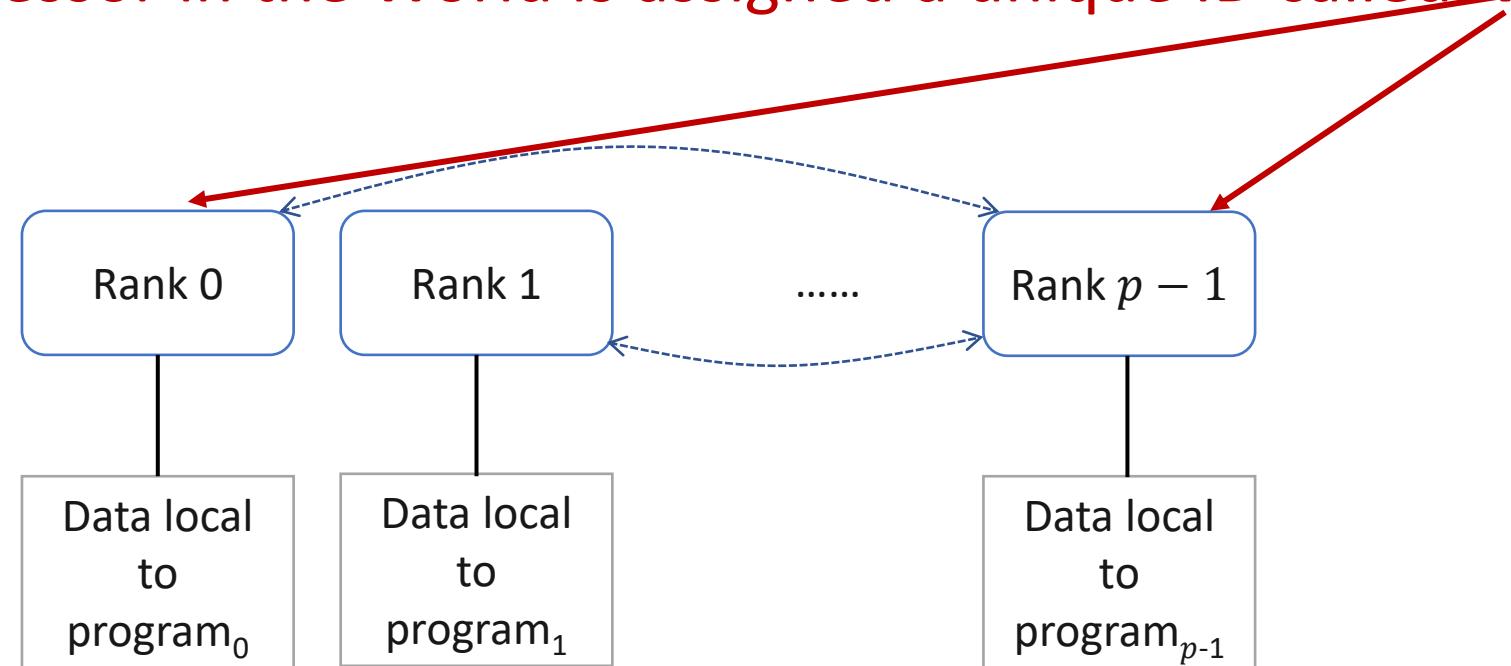
# Communication Operations (3)

- How do we identify processor IDs? **Ranks**
- Collection of all the processors is called **World**



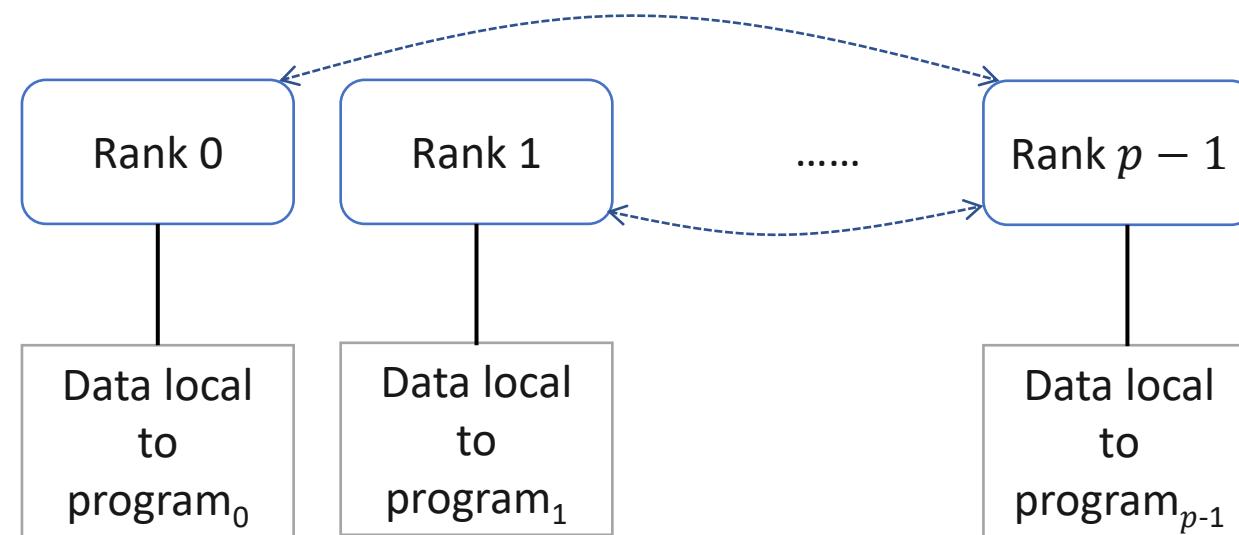
# Communication Operations (4)

- How do we identify processor IDs? **Ranks**
- Each Processor in the World is assigned a unique ID called **Rank**



# Communication Operations (5)

- How do we identify processor IDs? **Ranks**
- It is possible to create partitions of the world. We will not discuss in this class



# Anatomy of an MPI Program

- `MPI_Init(...)`
- `MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)`
- `MPI_Comm_size(MPI_COMM_WORLD, &num_procs);`
- Do rank specific work

Each processor will run the exact same program

Actual instructions that get executed may vary depending upon the rank

Mpiprogram.c

# Anatomy of an MPI Program

- **MPI\_Init(...)**
- **MPI\_Comm\_rank(MPI\_COMM\_WORLD, &my\_rank)**      API to initialize the MPI framework
- **MPI\_Comm\_size(MPI\_COMM\_WORLD, &num\_procs);**
- Do rank specific work

# Anatomy of an MPI Program

- `MPI_Init(...)`
- `MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)`
- `MPI_Comm_size(MPI_COMM_WORLD,  
&num_procs);`
- Do rank specific work

Obtain the rank of  
the process and the  
size of the world

# Anatomy of an MPI Program

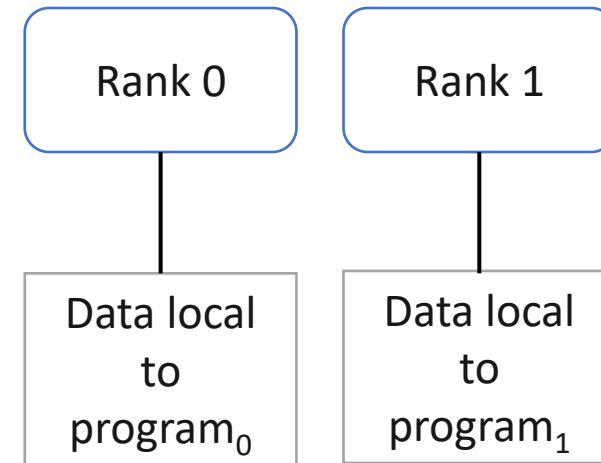
- `MPI_Init(...)`
- `MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)`
- `MPI_Comm_size(MPI_COMM_WORLD,  
&num_procs);`
- **Do rank specific work**

Do work which is  
specific to the rank

# Inter-Process Communication (1)

```
MPI_Init(...)  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)  
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

```
If (rank == 0) {  
    D0 = C00;  
    Send(D0, size, P1);  
    C01 ; }  
Else {  
    C11;  
    Receive(D1, size, P1);  
    C12(D1); }
```



Two processor world

# Inter-Process Communication (2)

```
MPI_Init(...)  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)  
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

```
If (rank == 0) {  
    D0 = C00;  
    Send(D0, size, P1);  
    C01 ; }  
Else {  
    C11;  
    Receive(D1, size, P1);  
    C12(D1); }
```

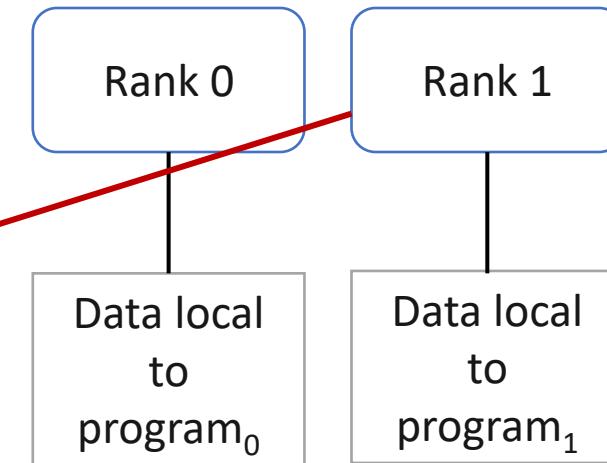


$P_0$ : Executes If portion

# Inter-Process Communication (3)

```
MPI_Init(...)  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)  
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

```
If (rank == 0) {  
    D0 = C00;  
    Send(D0, size, P1);  
    C01 ; }  
Else {  
    C11;  
    Receive(D1, size, P1);  
    C12(D1); }
```



$P_1$ : Executes Else portion

# Inter-Process Communication (4)

P0:

$D_0 = C_{00};$   
`Send(D0, size, P1);`  
 $C_{01} ;$

P1:

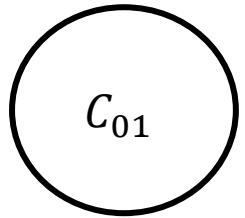
$C_{11};$   
`Receive(D1, size, P1);`  
 $C_{12}(D_1)$

Compute operations

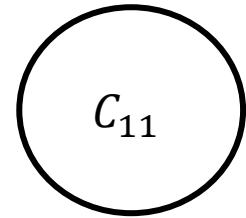
For better visualization, we will represent programs like this. This is equivalent to the previous program

# Inter-Process Communication (5)

P0

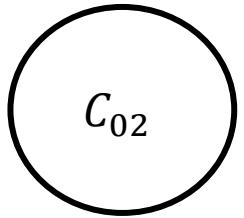


P1

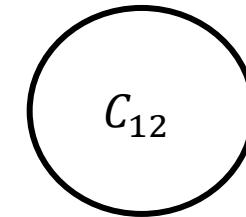


Can you show the  
dependencies between  
these tasks?

$C_{02}$



$C_{12}$

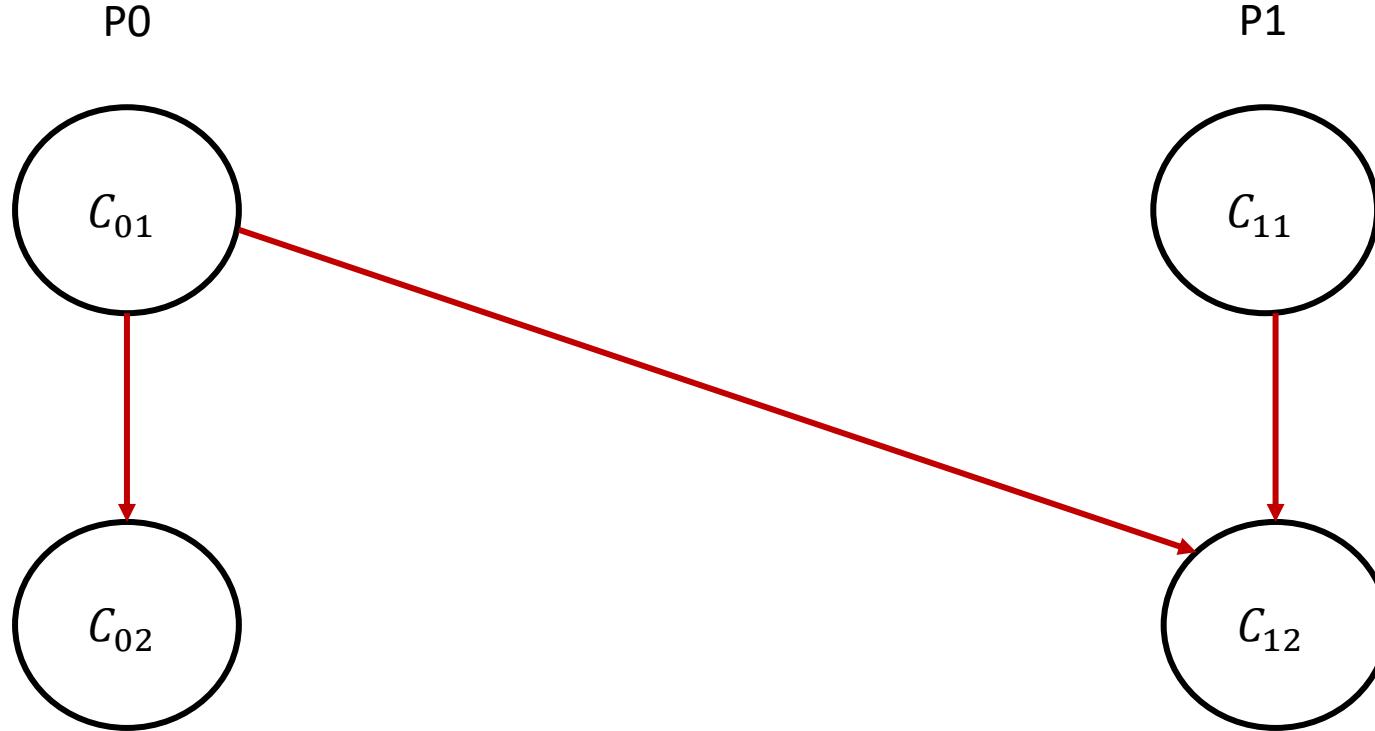


# Inter-Process Communication (6)



Sequential Dependencies due  
to same processor

# Inter-Process Communication (7)



Dependency across processors due to  
communication operation

# Blocking Semantics

P0:

$D_0 = C_{00};$

`Send( $D_0$ , size, P1);`

$C_{01} ;$

- Does  $P0$  wait for the send operation to complete?  
**Blocking/Non-Blocking**
- Does MPI directly transfer the data stored on  $D_0$  or does it make another copy? **Buffered/Non-Buffered**

P1:

$C_{11};$

`Receive( $D_1$ , size, P1);`

$C_{12}(D_1)$

# Blocking Semantics

P0:

$D_0 = C_{00};$

$\text{Send}(D_0, \text{size}, \text{P1});$

$C_{01} ;$

- Blocking Non-Buffered Send
  - Block sending process
  - Send request to receiving process
  - Wait for receiving process to acknowledge (matched receive operation)
  - Upon receiving acknowledgement, start the transfer
  - No buffers are used for data to be sent

P1:

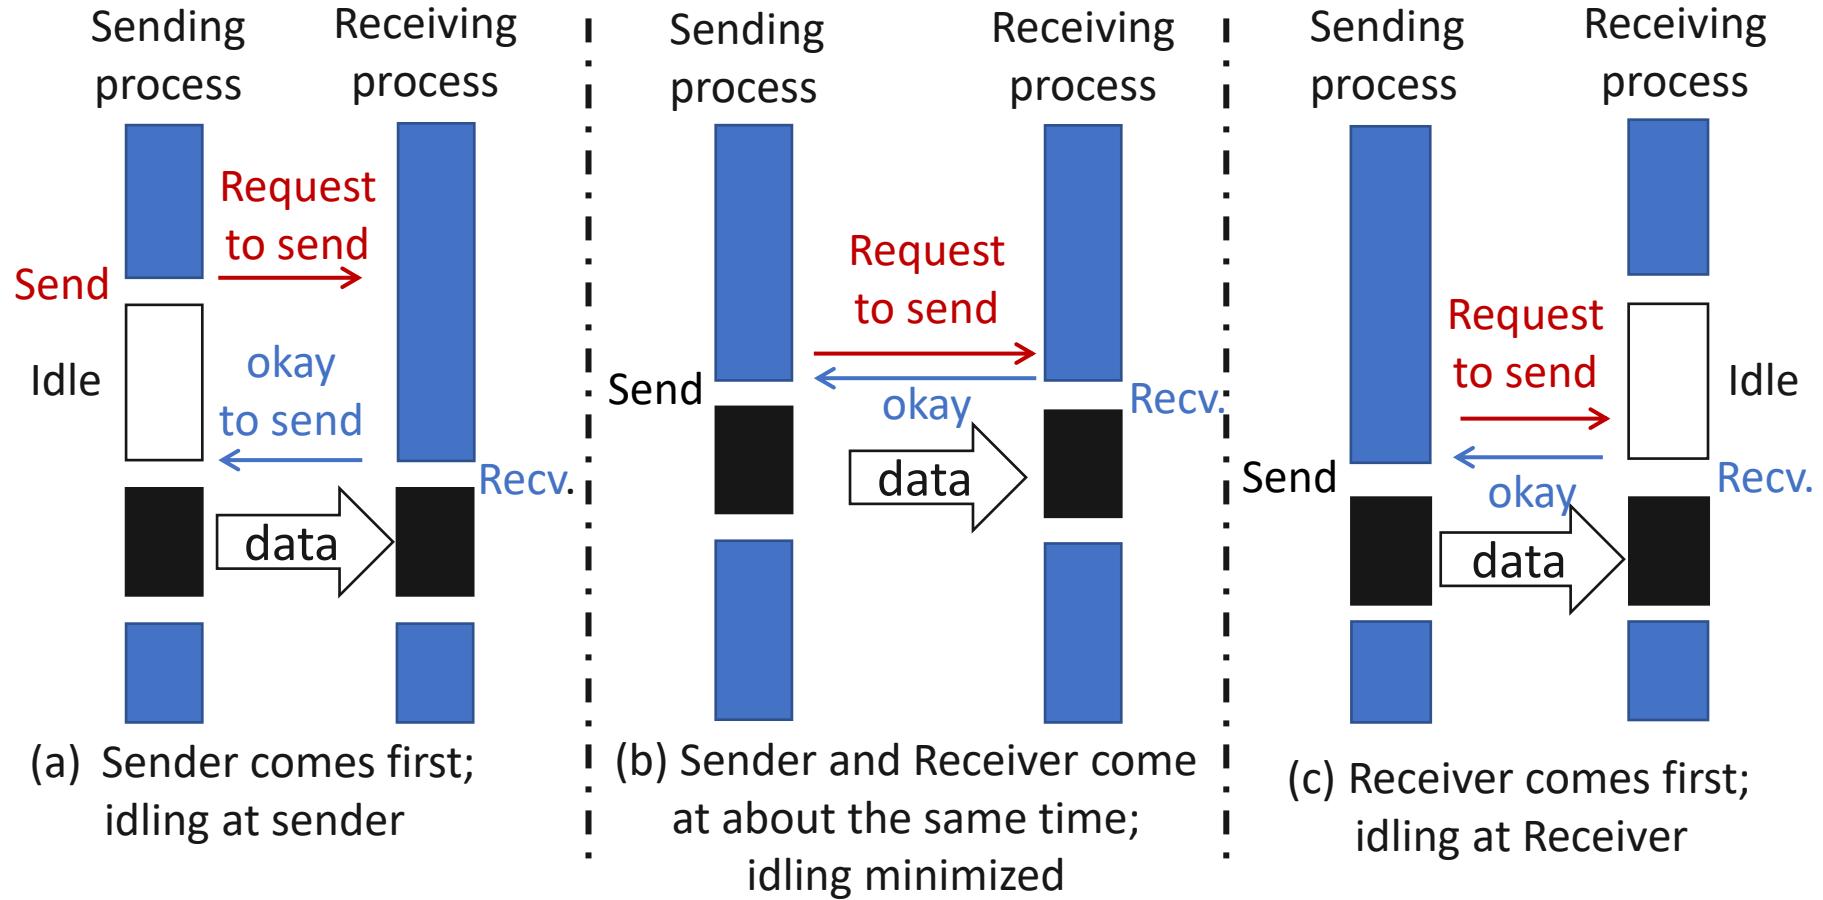
$C_{11};$

$\text{Receive}(D_1, \text{size}, \text{P1});$

$C_{12}(D_1)$

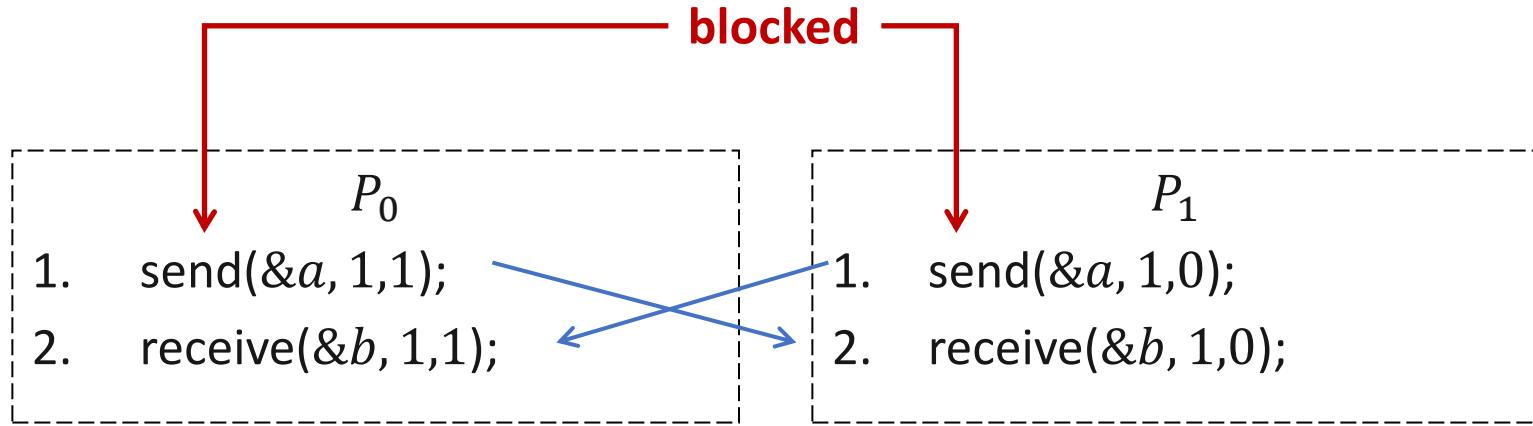
Same happens at the receiver end

# Blocking Send/Receive (1)



# Blocking Send/Receive (2)

## Issue #2: Deadlocks



Deadlocks are very easy in blocking protocols

# Blocking Send/Receive (3)

- Non-Block Buffered/Non-buffered addresses these issues, however, it complicates the implementation.
- They are widely used, but we will not discuss here
- You can refer to the Slides from University of Stuttgart -  
[https://fs.hlrs.de/projects/par/par\\_prog\\_ws/pdf/mpi\\_3.1\\_rab.pdf](https://fs.hlrs.de/projects/par/par_prog_ws/pdf/mpi_3.1_rab.pdf)

# Message Passing Programming

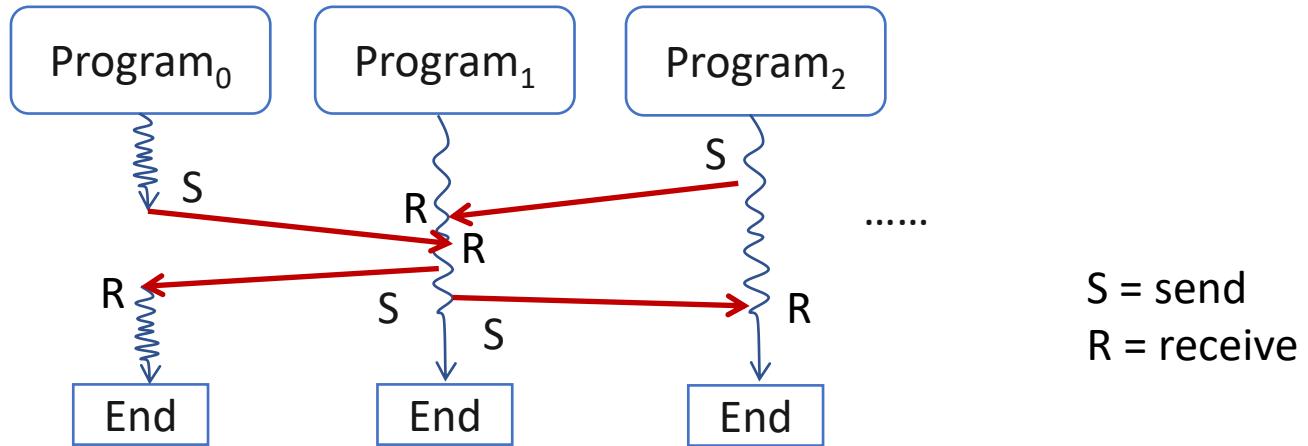
- How do we program clusters using Message Passing?
- We need to specify the local computations that each processor will execute
- We also need to specify at what time which pairs of processors will communicate and what data they will transfer to each other
  - $P^2$  pairs of processors that can communicate at any point
  - Seems too chaotic

# Concurrency Models in Message Passing Programming

- Bring some structure to the programming
- Decide points in the program where processes will not communicate at all or when all processes communicate

# Message Passing Program (1)

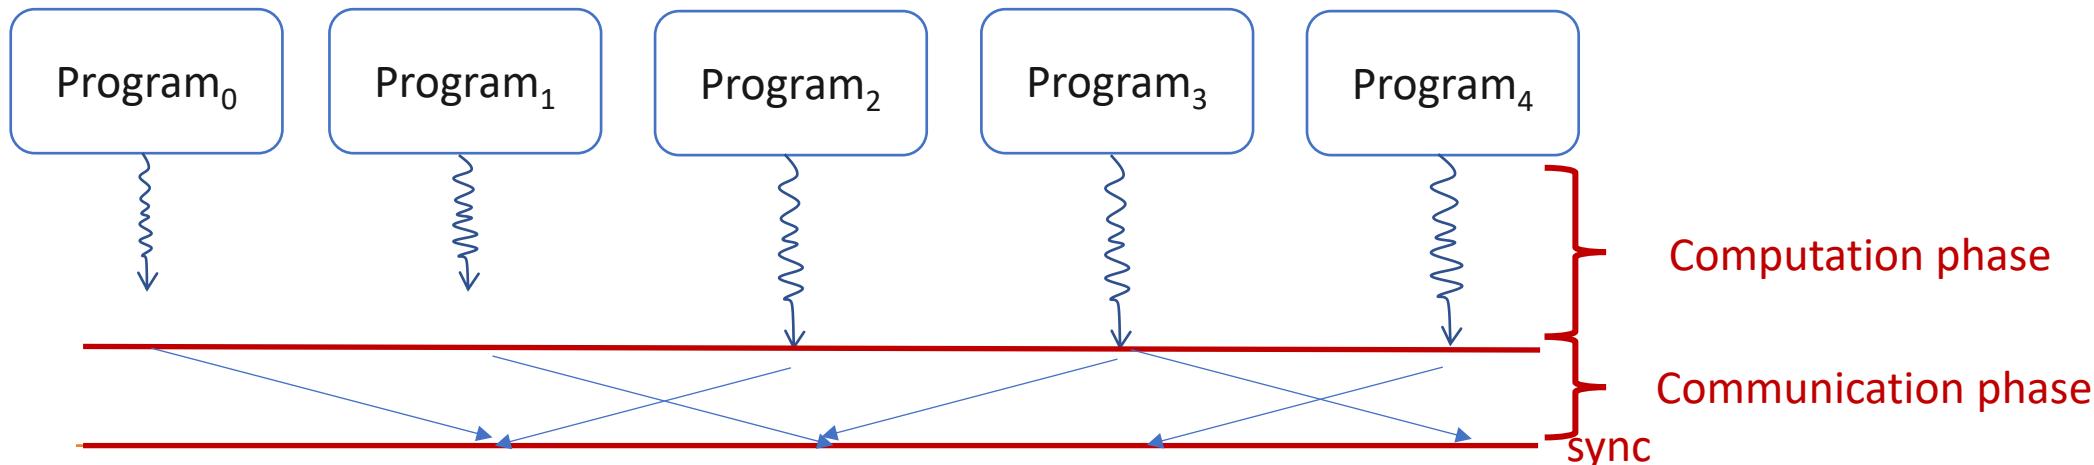
## Most General Model: Asynchronous



- No structure with respect to instructions, interactions
- No global clock
- Execution is asynchronous
- Programs  $0, 1, \dots, p - 1$  can be all distinct
- Hard to write/debug

# Message Passing Program (2)

## Bulk synchronous

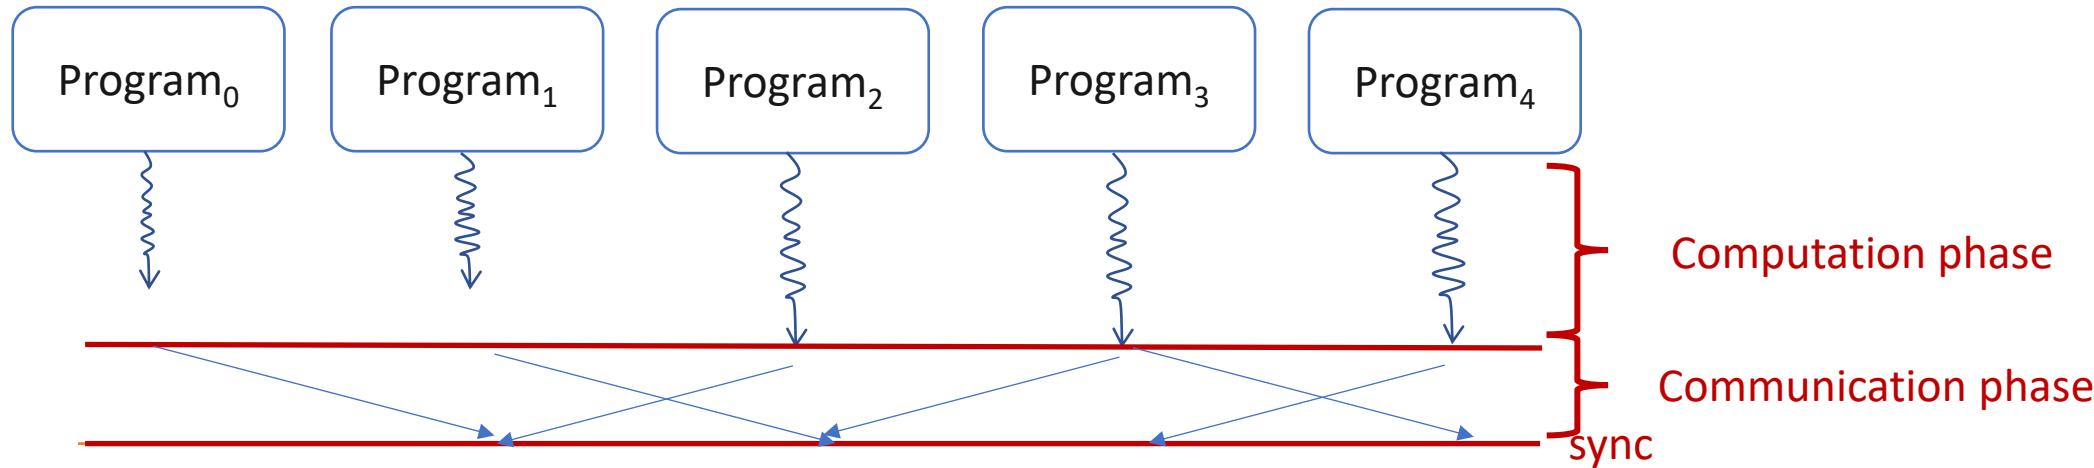


Two phases to the Program

- Computation Phase: Each process executes independently.  
No communication
- Communication Phase: Processes communicate with each other (usually using group communication primitives – we will discuss in the next class)

# Message Passing Program (2)

## Bulk synchronous

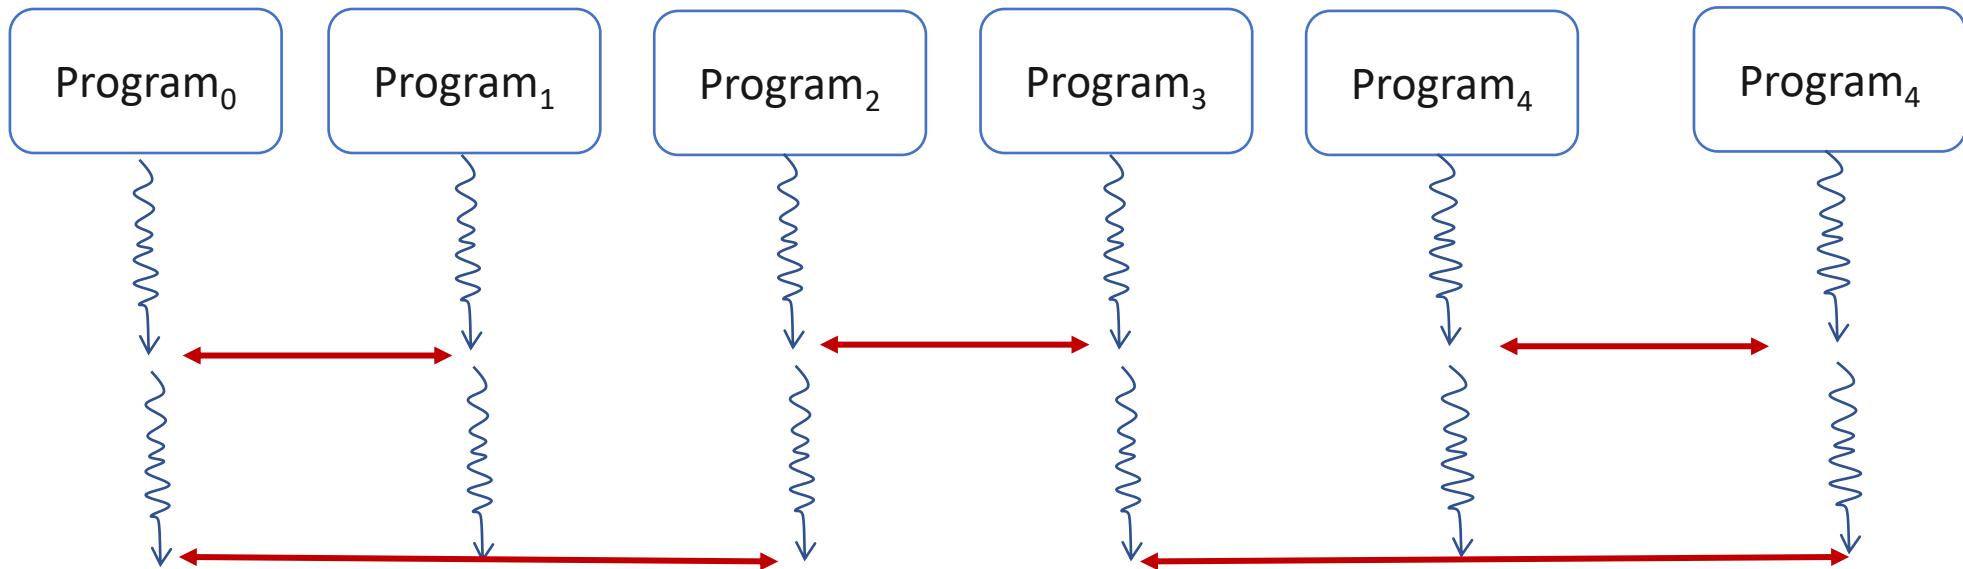


- Simplifies implementation by bringing in structure to the program
- Communication doesn't happen at random times, easier to debug issues
- Using Group communication primitives further simplifies implementations
- Widely use in Machine Learning Training

# Message Passing Program (3)

## SPMD (Single Program Multiple Data)

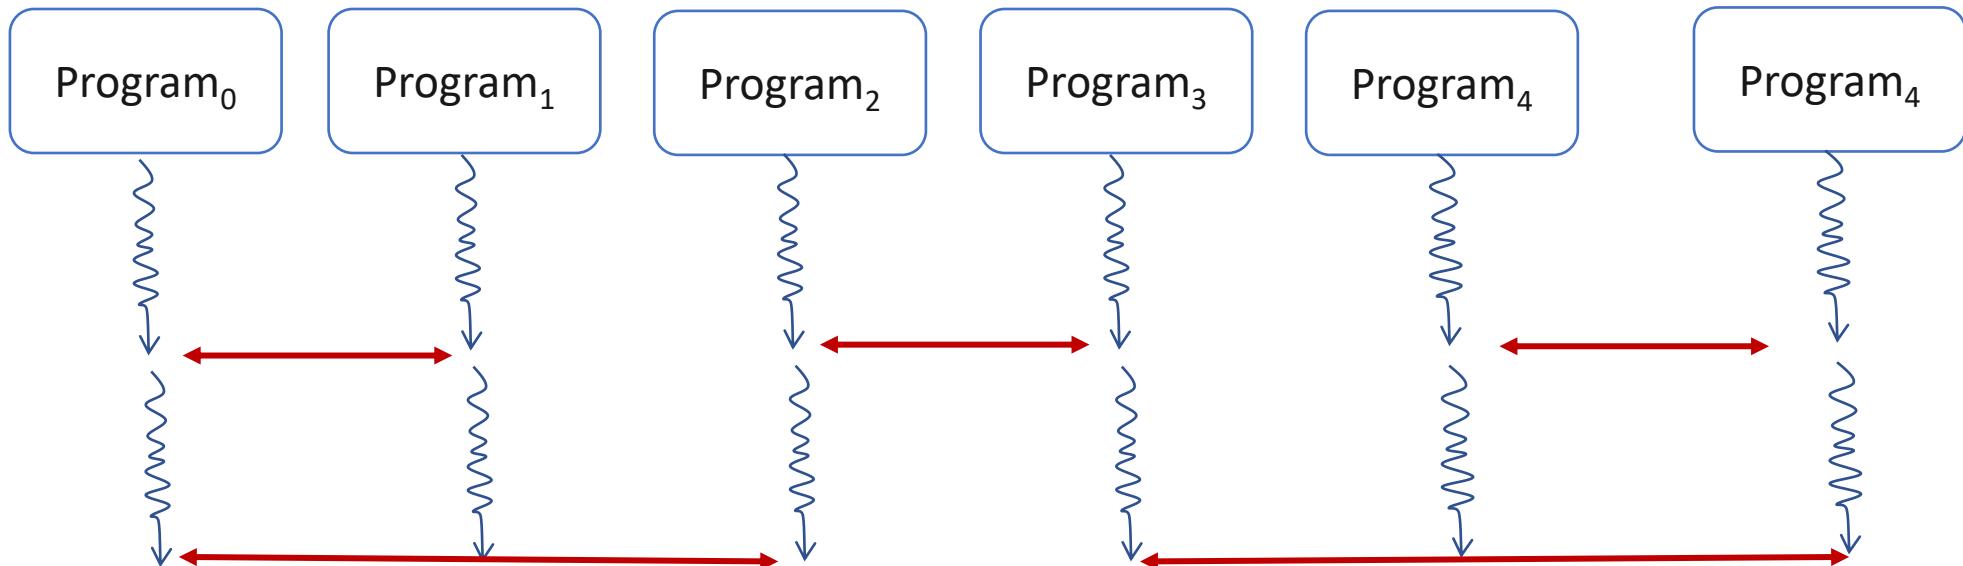
- Code is same in all the processes except for initialization
- Restrictive model, easy to write and debug
- Easy to do performance analysis
- Widely used in Machine Learning Training



# Message Passing Program (3)

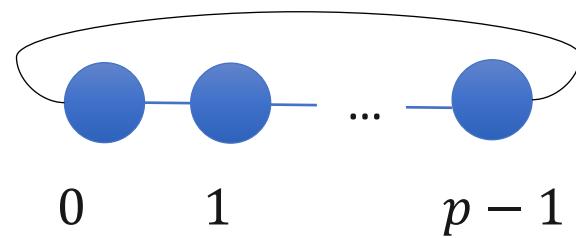
## SPMD (Single Program Multiple Data)

- BSP
  - Programs can be different
  - Explicit Barrier (synchronization) between Computation and Communication Phases
- SPMD
  - Programs have to be same – different data
  - No explicit barrier needed (but maybe implied if using blocking send/receive)

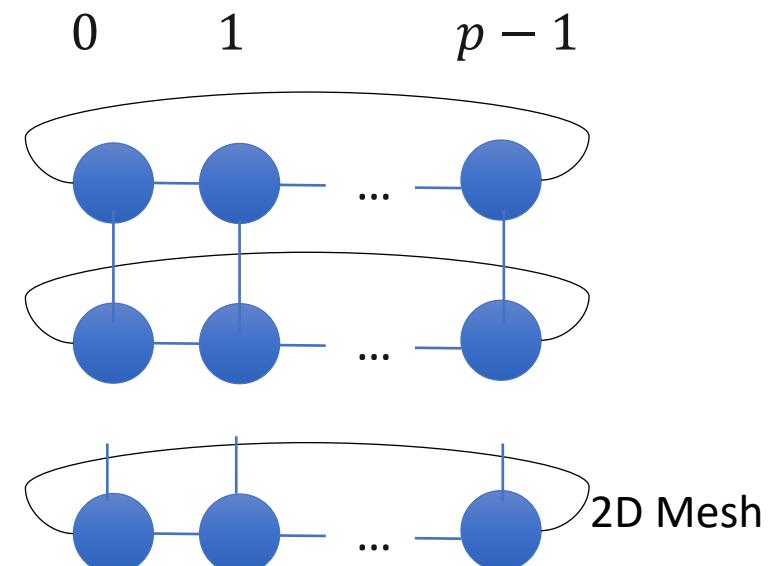


# Virtual Topology

- Define the “connectivity pattern” of the processors
- Helps us in developing more intuitive notions of message passing algorithms
- Think of it as 1D versus 2D arrays. It will be hard to visualize matrix multiplication if we write algorithms using 1D arrays



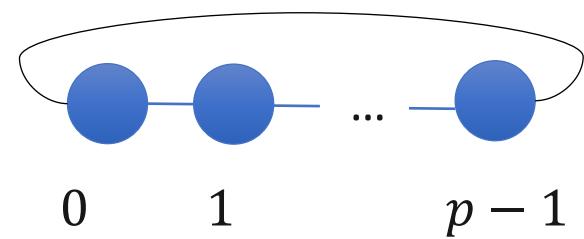
1D Mesh



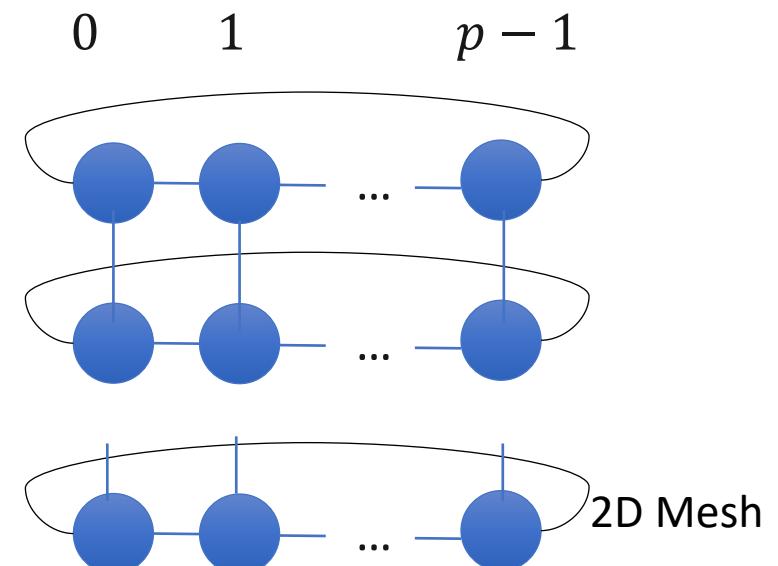
2D Mesh

# Virtual Topology

- Define the “connectivity pattern” of the processors
- Helps us in developing more intuitive notions of message passing algorithms
- (also has uses in optimizing the mapping of message passing algorithms onto real interconnection networks)



1D Mesh



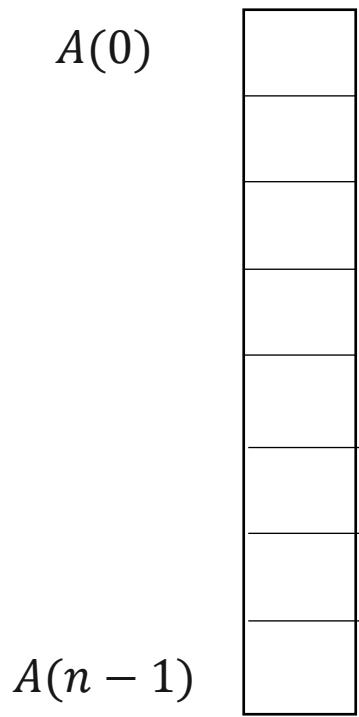
2D Mesh

# Message Passing Programming Paradigm

- User Specifies the following
  - Concurrency Model (BSP, SPMD, None/Asynchronous)
    - Processes: The number of processes and the work performed for each process
    - Send, receive that enable data (we will only use blocking non-buffered semantics)
  - A virtual topology of the processes
- We will discuss it at algorithmic level.
  - Skip initialization, rank calculation, etc.

# Adding Using Message Passing on 1D Mesh (1)

$$\text{Output} = \sum_{i=0}^{n-1} A(i) \text{ in } A(0)$$



# Adding Using Message Passing on 1D Mesh (1)

- Concurrency Model? **SPMD**
- How many Processors?  **$n$**
- What does each processor do?
- Send/receive commands?
- Virtual topology – **1D mesh**

# Adding Using Message Passing on 1D Mesh (2)

- Key Idea: Recursive doubling

$A(0)$



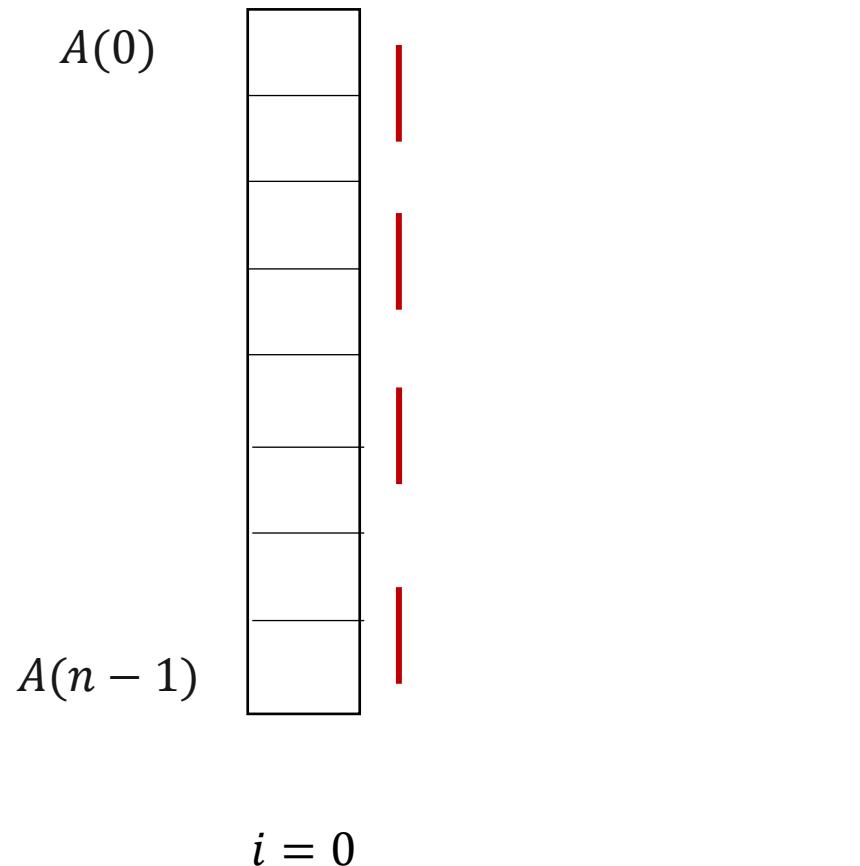
- In iteration  $i$ :

- Processors  $j$  and  $j + 2^i$  communicate, where  $j = k \cdot 2^{i+1}$
- Processor  $j$  computes
- $j$  – Active Processors

$A(n - 1)$

# Adding Using Message Passing on 1D Mesh (3)

- Key Idea: Recursive doubling
- In iteration  $i$ :
  - Processors  $j$  and  $j + 2^i$  communicate, where  $j = k \cdot 2^{i+1}$
  - Processor  $j$  computes
  - 0, 2, 4, 6 – Active Processors

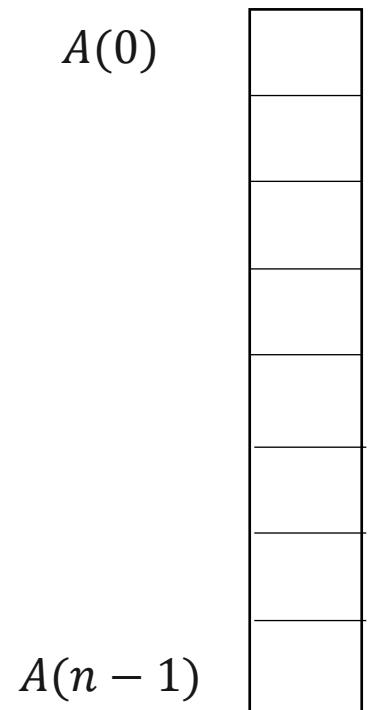


# Adding Using Message Passing on 1D Mesh (4)

- Key Idea: Recursive doubling

- In iteration  $i$ :

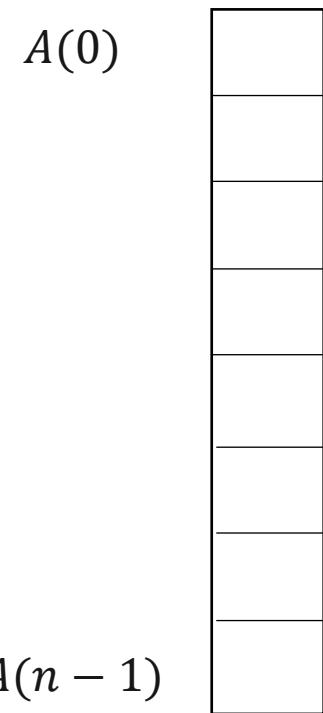
- Processors  $j$  and  $j + 2^i$  communicate, where  $j = k \cdot 2^{i+1}$
- Processor  $j$  computes
- 0, 4 – Active Processors



$i = 1$

# Adding Using Message Passing on 1D Mesh (5)

- Key Idea: Recursive doubling



- In iteration  $i$ :

- Processors  $j$  and  $j + 2^i$  communicate, where  $j = k \cdot 2^{i+1}$
- Processor  $j$  computes
- 0 – Active Processors

$i = 2$

# Adding Using Message Passing on 1D Mesh (6)

- Message Passing Algorithm (SPMD model)

Program in process  $j$ ,  $0 \leq j \leq n - 1$

1. Do  $i = 0$  to  $\text{??}$
2. If  $j = k \cdot 2^{i+1} + 2^i$ , for some  $k \in N$   
**Send  $A(j)$  to process  $j - 2^i$**   $2^i$  distance communication
3. Else if  $j = k \cdot 2^{i+1}$ , for some  $k \in N$   
**Receive  $A(j + 2^i)$  from process  $j + 2^i$**
4.  $A(j) \leftarrow A(j) + A(j + 2^i)$
5. End
6. **Barrier**
7. Note:  
 $A(j)$  is local to process  $j$
8.  $N = \text{set of natural numbers} = \{0, 1, \dots\}$
9. End

# Adding Using Message Passing on 1D Mesh (6)

- Message Passing Algorithm (SPMD model)

Program in process  $j$ ,  $0 \leq j \leq n - 1$

1. Do  $i = 0$  to  $\log_2 n - 1$
2. If  $j = k \cdot 2^{i+1} + 2^i$ , for some  $k \in N$   
**Send  $A(j)$  to process  $j - 2^i$**   $2^i$  distance communication
3. Else if  $j = k \cdot 2^{i+1}$ , for some  $k \in N$   
**Receive  $A(j + 2^i)$  from process  $j + 2^i$**
4.  $A(j) \leftarrow A(j) + A(j + 2^i)$
5. End
6. **Barrier**
7. Note:  
 $A(j)$  is local to process  $j$
8.  $N = \text{set of natural numbers} = \{0, 1, \dots\}$
9. End

# Performance Analysis (1)

- Total Computation time: Number of rounds  $\times$  computation time per round
  - Note each processor is doing the same computation in each round
- Computation time per round – Calculate using accelerator model – GPU/Systolic array
  - Computation time per round:  $O(1)$
- Number of rounds -  $\log N$
- Total Computation time -  $O(\log N)$

# Performance Analysis (2)

- Total Communication time ???
- Depends upon the underlying interconnection topology
- We will assume fully connected in this class
  - A transfer of  $k$  data items from any processor to any processor takes  $O(k)$  amount of time.
  - Assuming  $t_w$  as the per word transfer time, this is equal to  $k \times t_w$
  - Note: Actual Interconnection modeling is much more complicated and could be a lecture or 2 in itself

# Performance Analysis (3)

- Total Communication time - Number of rounds  $\times$  communication time per round
- Number of rounds -  $\log N$
- Communication time per round -  $1 \times t_w$
- Total Communication time:  $\log N \times t_w$

# Performance Analysis (4)

- Challenge: Does the system have enough communication bandwidth to support the communication requirement of the algorithm?
- Maximum Communication Requirement – ??
- Note: Iteration 0 has the most ( $N/2$ ) processors active
- Maximum communication requirement =  $\frac{N}{2} \cdot 1 < B$
- $B$ : Maximum bandwidth supported by the system.

# Next Class

- 11/13 Lecture 21
  - Distributed Matrix Multiplication on a Cluster of Accelerators
  - Communication Primitives

# Thank You

- Questions?
- Email: [sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)