

# CSDS 451: Designing High Performant Systems for AI

Lecture 23

11/20/2025

Sanmukh Kuppannagari

[sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)

<https://sanmukh.research.st/>

Case Western Reserve University

# Outline

- Data Parallel Distributed Training
- Model Parallel Distributed Training

# Announcements

- WA 4 was out
  - Due December 4
- Finals next Tuesday
  - Transformer Models
  - Clusters
    - How to program clusters with MPI
    - Communication collectives with complexity
    - Distributed training – very simple question to check knowledge, no application/analysis questions

# Midterm Logistics

- 60 minute exam, on paper
- 15~20 Fill in the blank/True-False questions
  - With space for explanation
  - 1 point for answer
  - 2 for explanation - wrong answer with a reasonable explanation will receive partial credits
- 2 big questions

# Midterm Logistics

- Open Notes
  - You can use slides
  - I suggest you create a cheat sheet: You won't have enough time to search for concepts.
- No questions during exam (Gets too distractive)
  - If a question is ambiguous, write your assumption and solve
  - The assumption will be taken into consideration when grading
- You should not need a calculator
  - You can use your phone if you do need it

# Midterm Logistics


- Exam may be lengthy (intentionally)
  - If you know the concepts well, you should be able to solve in time
  - Searching your notes or slides will incur additional time
- Plan the questions you want to solve first accordingly
  - Suggestion: Get the big questions out of the way.
  - Don't get stuck on a question. If you can't solve go ahead and come back
  - If you don't know an answer, write your explanation. We will consider that while grading

# Outline

- Data Parallel Distributed Training
- Model Parallel Distributed Training

# Cluster of Accelerators

- $N$  processors (memory + accelerator)
  - Local compute
  - Local memory
- Connected using an Interconnection Network
- Communication through Message Passing




# Anatomy of an MPI Program

- `MPI_Init(...)`
- `MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)`
- `MPI_Comm_size(MPI_COMM_WORLD,  
&num_procs);`
- Do rank specific work

# Inter-Process Communication


```
MPI_Init(...)  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank)  
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

```
If (rank == 0) {  
    D0 = C00;  
    Send(D0, size, P1);  
    C01 ; }  
Else {  
    C11;  
    Receive(D1, size, P1);  
    C12(D1); }
```



Two processor world


# Inter-Process Communication



Dependency across processors due to  
communication operation

# Message Passing Program


## Most General Model: Asynchronous



- No structure with respect to instructions, interactions
- No global clock
- Execution is asynchronous
- Programs  $0, 1, \dots, p - 1$  can be all distinct
- Hard to write/debug

# Message Passing Program

## Bulk synchronous




Two phases to the Program

- Computation Phase: Each process executes independently.  
No communication
- Communication Phase: Processes communicate with each other


# Message Passing Program

## SPMD (Single Program Multiple Data)

- Code is same in all the processes except for initialization
- Restrictive model, easy to write and debug
- Easy to do performance analysis
- Widely used in Machine Learning Training




# Blocking Non-Buffered Send/Receive




# Virtual Topology

- Define the “connectivity pattern” of the processors
- Helps us in developing more intuitive notions of message passing algorithms
- Think of it as 1D versus 2D arrays. It will be hard to visualize matrix multiplication if we write algorithms using 1D arrays



1D Mesh




2D Mesh

# Message Passing Programming Paradigm


- User Specifies the following
  - Concurrency Model (BSP, SPMD, None/Asynchronous)
    - Processes: The number of processes and the work performed for each process
    - Send, receive that enable data (we will only use blocking non-buffered semantics)
  - A virtual topology of the processes
- We will discuss it at algorithmic level.
  - Skip initialization, rank calculation, etc.

# MPI\_Broadcast

- MPI\_Broadcast(buffer, size, root)
- Send the data stored in **buffer** at the **root** processor to all the processors




# MPI\_Scatter





# MPI\_Gather







# MPI\_Reduce

**before MPI\_Reduce**

- inbuf
- result



**after**



AoDoGoJoM

# MPI\_AllReduce

- `MPI_AllReduce(inbuf, resultbuf, size)`
- Same as `MPI_Reduce()` but stores the results in all the processors
- Equivalent to:
  - `MPI_Reduce(inbuf, ibsize, resultbuf, rbsize, Aggr_op, root)`
  - `MPI_Broadcast(resultbuf, rbsize, root )`

# Training

$$E(W): \min_W \sum (y_i - F(x_i: W))^2$$

$$W_{t+1} \leftarrow W_t - \alpha \frac{\delta E(W)}{\delta W}$$

Iteratively

$$\frac{\delta E(W)}{\delta W} = 2 \sum (y_i - F(x_i: W)) \times \frac{\delta F(x_i: W)}{\delta W}$$


For all samples, in each iteration



Error


# Training

$L$  layers in CNN



*2 \* Error \**   $\frac{\delta F(x_i: W_{1:L})}{\delta W_1}, \frac{\delta F(x_i: W_{1:L})}{\delta W_2}, \dots, \frac{\delta F(x_i: W_{1:L})}{\delta W_L}$


# CNN Training: Forward Propagation



*I*


For a single image

# CNN Training: Forward Propagation




For a single image

# CNN Training: Forward Propagation




For a single image

# CNN Training: Forward Propagation




# CNN Training: Error Calculation

For a single image




# CNN Training: Backward Propagation

For a single image




# CNN Training: Backward Propagation




For a single image

# CNN Training: Backward Propagation



# CNN Training: Backward Propagation



$$\frac{\delta E(W)}{\delta W_l} = \frac{\delta O_l}{\delta W_l} \times \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}} \times \frac{\delta E(W)}{\delta O_L}$$

For a single image

# Training

$$E(W): \min_W \sum (y_i - F(x_i: W))^2$$

$$W_{t+1} \leftarrow W_t - \alpha \frac{\delta E(W)}{\delta W}$$

Iteratively

$$\frac{\delta E(W)}{\delta W} = 2 \sum \underbrace{(y_i - F(x_i: W))}_{\text{Error}} \times \frac{\delta F(x_i: W)}{\delta W}$$

For all samples, in each iteration

**For multiple images in a batch: Simply sum up the gradients**

$$\frac{\delta E(W)}{\delta W_l} = \boxed{\frac{\delta O_l}{\delta W_l} \times \frac{\delta O_{l+1}}{\delta O_l} \times \frac{\delta O_{l+2}}{\delta O_{l+1}} \times \dots \times \frac{\delta O_L}{\delta O_{L-1}}} \times \boxed{\frac{\delta E(W)}{\delta O_L}}$$

# Gradient Update


$$W \leftarrow W - \alpha \frac{\delta E(W)}{\delta W}$$

- Lets dig deeper into Gradient Update

# Gradient Update

- For a mini-batch with  $M$  samples and mean squared error
- $E(W) = \sum_{i=1}^M E_i(W)$
- $E_i(w) = (O_{Li} - y_i)^2$
- $\frac{\delta E(W)}{\delta w} = \sum_{i=1}^M \frac{\delta E_i(W)}{\delta w}$

# Gradient Update


$$\frac{\delta E(W)}{\delta W} = \frac{\delta E_1(W)}{\delta W} + \frac{\delta E_2(W)}{\delta W} + \dots + \frac{\delta E_M(W)}{\delta W}$$
$$= 2(O_{L1} - y_1) \frac{\delta O_{L1}}{\delta W} + 2(O_{L2} - y_2) \frac{\delta O_{L2}}{\delta W} + \dots + 2(O_{LM} - y_M) \frac{\delta O_{LM}}{\delta W}$$


Can be calculated in parallel

# Gradient Update

$$\frac{\delta E(W)}{\delta W} = \frac{\delta E_1(W)}{\delta W} + \frac{\delta E_2(W)}{\delta W} + \dots + \frac{\delta E_M(W)}{\delta W}$$

$$= 2(O_{L1} - y_1) \frac{\delta O_{L1}}{\delta W} + 2(O_{L2} - y_2) \frac{\delta O_{L2}}{\delta W} + \dots + 2(O_{LM} - y_M) \frac{\delta O_{LM}}{\delta W}$$




Forward propagation in parallel – Need  
multiple copies of model

# Gradient Update

$$\frac{\delta E(W)}{\delta W} = \frac{\delta E_1(W)}{\delta W} + \frac{\delta E_2(W)}{\delta W} + \dots + \frac{\delta E_M(W)}{\delta W}$$

$$= 2(O_{L1} - y_1) \frac{\delta O_{L1}}{\delta W} + 2(O_{L2} - y_2) \frac{\delta O_{L2}}{\delta W} + \dots + 2(O_{LM} - y_M) \frac{\delta O_{LM}}{\delta W}$$




Error Calculation in parallel

# Gradient Update

$$\frac{\delta E(W)}{\delta W} = \frac{\delta E_1(W)}{\delta W} + \frac{\delta E_2(W)}{\delta W} + \dots + \frac{\delta E_M(W)}{\delta W}$$

$$= 2(O_{L1} - y_1) \frac{\delta O_{L1}}{\delta W} + 2(O_{L2} - y_2) \frac{\delta O_{L2}}{\delta W} + \dots + 2(O_{LM} - y_M) \frac{\delta O_{LM}}{\delta W}$$



# Gradient Update

- If using  $p$  processes (or GPUs) for a mini-batch of size  $M$
- Each process (or GPU) keeps a copy of the model  $W$
- Each process performs the computations of  $M/p$  samples independently using its own model
- Note: No synchronization was needed till now

# Gradient Update

- Weight Update Step:

$$W \leftarrow W - \alpha \frac{\delta E(W)}{\delta W}$$

- Do we need synchronization now???

# Gradient Update


- Weight Update Step:

$$W \leftarrow W - \alpha \frac{\delta E(W)}{\delta W}$$

- Do we need synchronization now??? Yes.
- Need to make sure that all GPUs have the same weights before they start processing the next mini-batch




# Data Parallel Training of Neural Networks



# Data Parallel Training of Neural Networks


- How do we synchronize?
- Model #1: Parameter Server Model
- A dedicated server performs the weight updates and distributes the models to all the workers
- Resource: Li, Mu, et al. "Scaling distributed machine learning with the parameter server." *11th USENIX Symposium on operating systems design and implementation (OSDI 14)*. 2014.






Cui, H., Zhang, H., Ganger, G. R., Gibbons, P. B., & Xing, E. P. (2016, April). Geeps: Scalable deep learning on distributed gpus with a gpu-specialized parameter server. In *Proceedings of the eleventh european conference on computer systems* (pp. 1-16).

# Parameter Server Model

Parameter Server




Data Parallel GPUs



# Parameter Server Model

Parameter Server

Data Parallel GPUs





# Parameter Server Model




Parameter Server



Data Parallel GPUs




Forward/Backward/  
Gradient Calculation









# Parameter Server Model

Parameter Server




Data Parallel GPUs



# Parameter Server Model


Parameter Server







Data Parallel GPUs

# Parameter Server Model

Parameter Server




Data Parallel GPUs



# Parameter Server Model

Parameter Server

Data Parallel GPUs






# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: Parameter server “sends” the weights to each processor
  - #2: Forward/Backward/gradient computation in each processor
  - #3: Each processor sends the gradient back to the parameter server which takes an average of the gradients
  - #4: Parameter server performs the gradient update step
- Repeat




w




# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: Parameter server “sends” the weights to each processor
  - #2: Forward/Backward/gradient computation in each processor
  - #3: Each processor sends the gradient back to the parameter server which takes an average of the gradients
  - #4: Parameter server performs the gradient update step
- Repeat




# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: `MPI_Broadcast(W)`
  - #2: Forward/Backward/gradient computation in each processor
  - #3: Each processor sends the gradient back to the parameter server which takes an average of the gradients
  - #4: Parameter server performs the gradient update step
- Repeat




# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: MPI\_Broadcast( $W$ )
  - #2: Forward/Backward/gradient computation in each processor
  - **#3: Each processor sends the gradient back to the parameter server which takes an average of the gradients**
  - #4: Parameter server performs the gradient update step
- Repeat




# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: MPI\_Broadcast( $W$ )
  - #2: Forward/Backward/gradient computation in each processor
  - **#3: MPI\_REDUCE( $\delta W$ ,AVG)**
  - #4: Parameter server performs the gradient update step
- Repeat



# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: MPI\_Broadcast( $W$ )
  - #2: Forward/Backward/gradient computation in each processor
  - #3: MPI\_REDUCE( $\delta W$ ,AVG)
  - #4: Parameter server performs the gradient update step
- Repeat
- Communication Complexity per Iteration (assume  $|W| = N$ ): ??











# MPI\_Broadcast

- MPI\_Broadcast:  $O(t_w \times \log P)$ 
  - $t_w$ : size of data to broadcast
  - $P$ : Number of processors
  - Algorithm proceeds in  $\log P$  steps, each step communicates  $t_w$  amount of data

# MPI\_Reduce


Can you calculate the communication complexity?  
 $1 + 1 + \dots 1 \text{ (log } P\text{)times} = O(\log P)$



Step 3


# Parameter Server Using Communication Collectives

- In each Iteration
  - #1: MPI\_Broadcast( $W$ )
  - #2: Forward/Backward/gradient computation in each processor
  - #3: MPI\_REDUCE( $\delta W$ ,AVG)
  - #4: Parameter server performs the gradient update step
- Repeat
- Communication Complexity per Iteration (assume  $|W| = N$ ):  
 $O(N \log(p + 1))$



# Data Parallel Training of Neural Networks


- How do we synchronize?
- Model #2: All-reduce based Synchronization
- No dedicated parameter server. Each worker calculates the gradients and broadcasts it to all other workers
- Each worker aggregates the gradients collected by all other workers and performs the update step



Iandola, F. N., Moskewicz, M. W., Ashraf, K., & Keutzer, K. (2016). Firecaffe: near-linear acceleration of deep neural network training on compute clusters. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 2592-2600).

# All Reduce Model


- In each Iteration:
  - Forward/Backward Propagation



# All Reduce Model


- In each Iteration:
  - Forward/Backward Propagation
  - Gradient Calculation

Data Parallel GPUs




# All Reduce Model

- In each Iteration:
  - Forward/Backward Propagation
  - Gradient Calculation
  - Each processor collects gradients from all the other processors.




# All Reduce Model

- In each Iteration:
  - Forward/Backward Propagation
  - Gradient Calculation
  - Each processor collects gradients from all the other processors.



# All Reduce Model

- In each Iteration:
  - Forward/Backward Propagation
  - Gradient Calculation
  - **MPI\_ALLREDUCE( $\delta W$ ,AVG)**




# MPI\_AllReduce

- `MPI_AllReduce(inbuf, resultbuf, size)`
- Same as `MPI_Reduce()` but stores the results in all the processors
- Equivalent to:
  - `MPI_Reduce(inbuf, ibsize, resultbuf, rbsize, Aggr_op, root)`
  - `MPI_Broadcast(resultbuf, rbsize, root )`


# All Reduce Model

- In each Iteration:
  - Forward/Backward Propagation
  - Gradient Calculation
  - **MPI\_ALLREDUCE( $\delta W$ ,AVG)**
- Complexity: ??



# All Reduce Model

- In each Iteration:
  - Forward/Backward Propagation
  - Gradient Calculation
  - $\text{MPI\_ALLREDUCE}(\delta W, \text{AVG})$
- Complexity:  $O(N \log p)$



# 3D Parallelism

- In practice, all three parallelisms are combined
- Data + Pipeline + Tensor = 3D parallelism
- DeepSpeed:  
<https://github.com/microsoft/DeepSpeed/tree/master?tab=readme-ov-file>

# Next Class

- 11/25 Final exam

# Thank You

- Questions?
- Email: [sanmukh.kuppannagari@case.edu](mailto:sanmukh.kuppannagari@case.edu)